

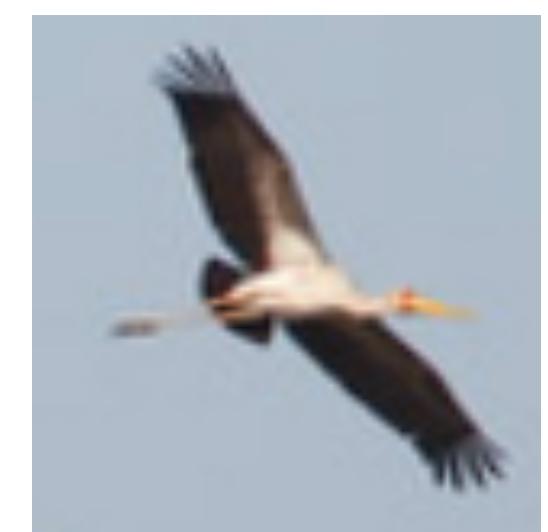
# Lecture 5

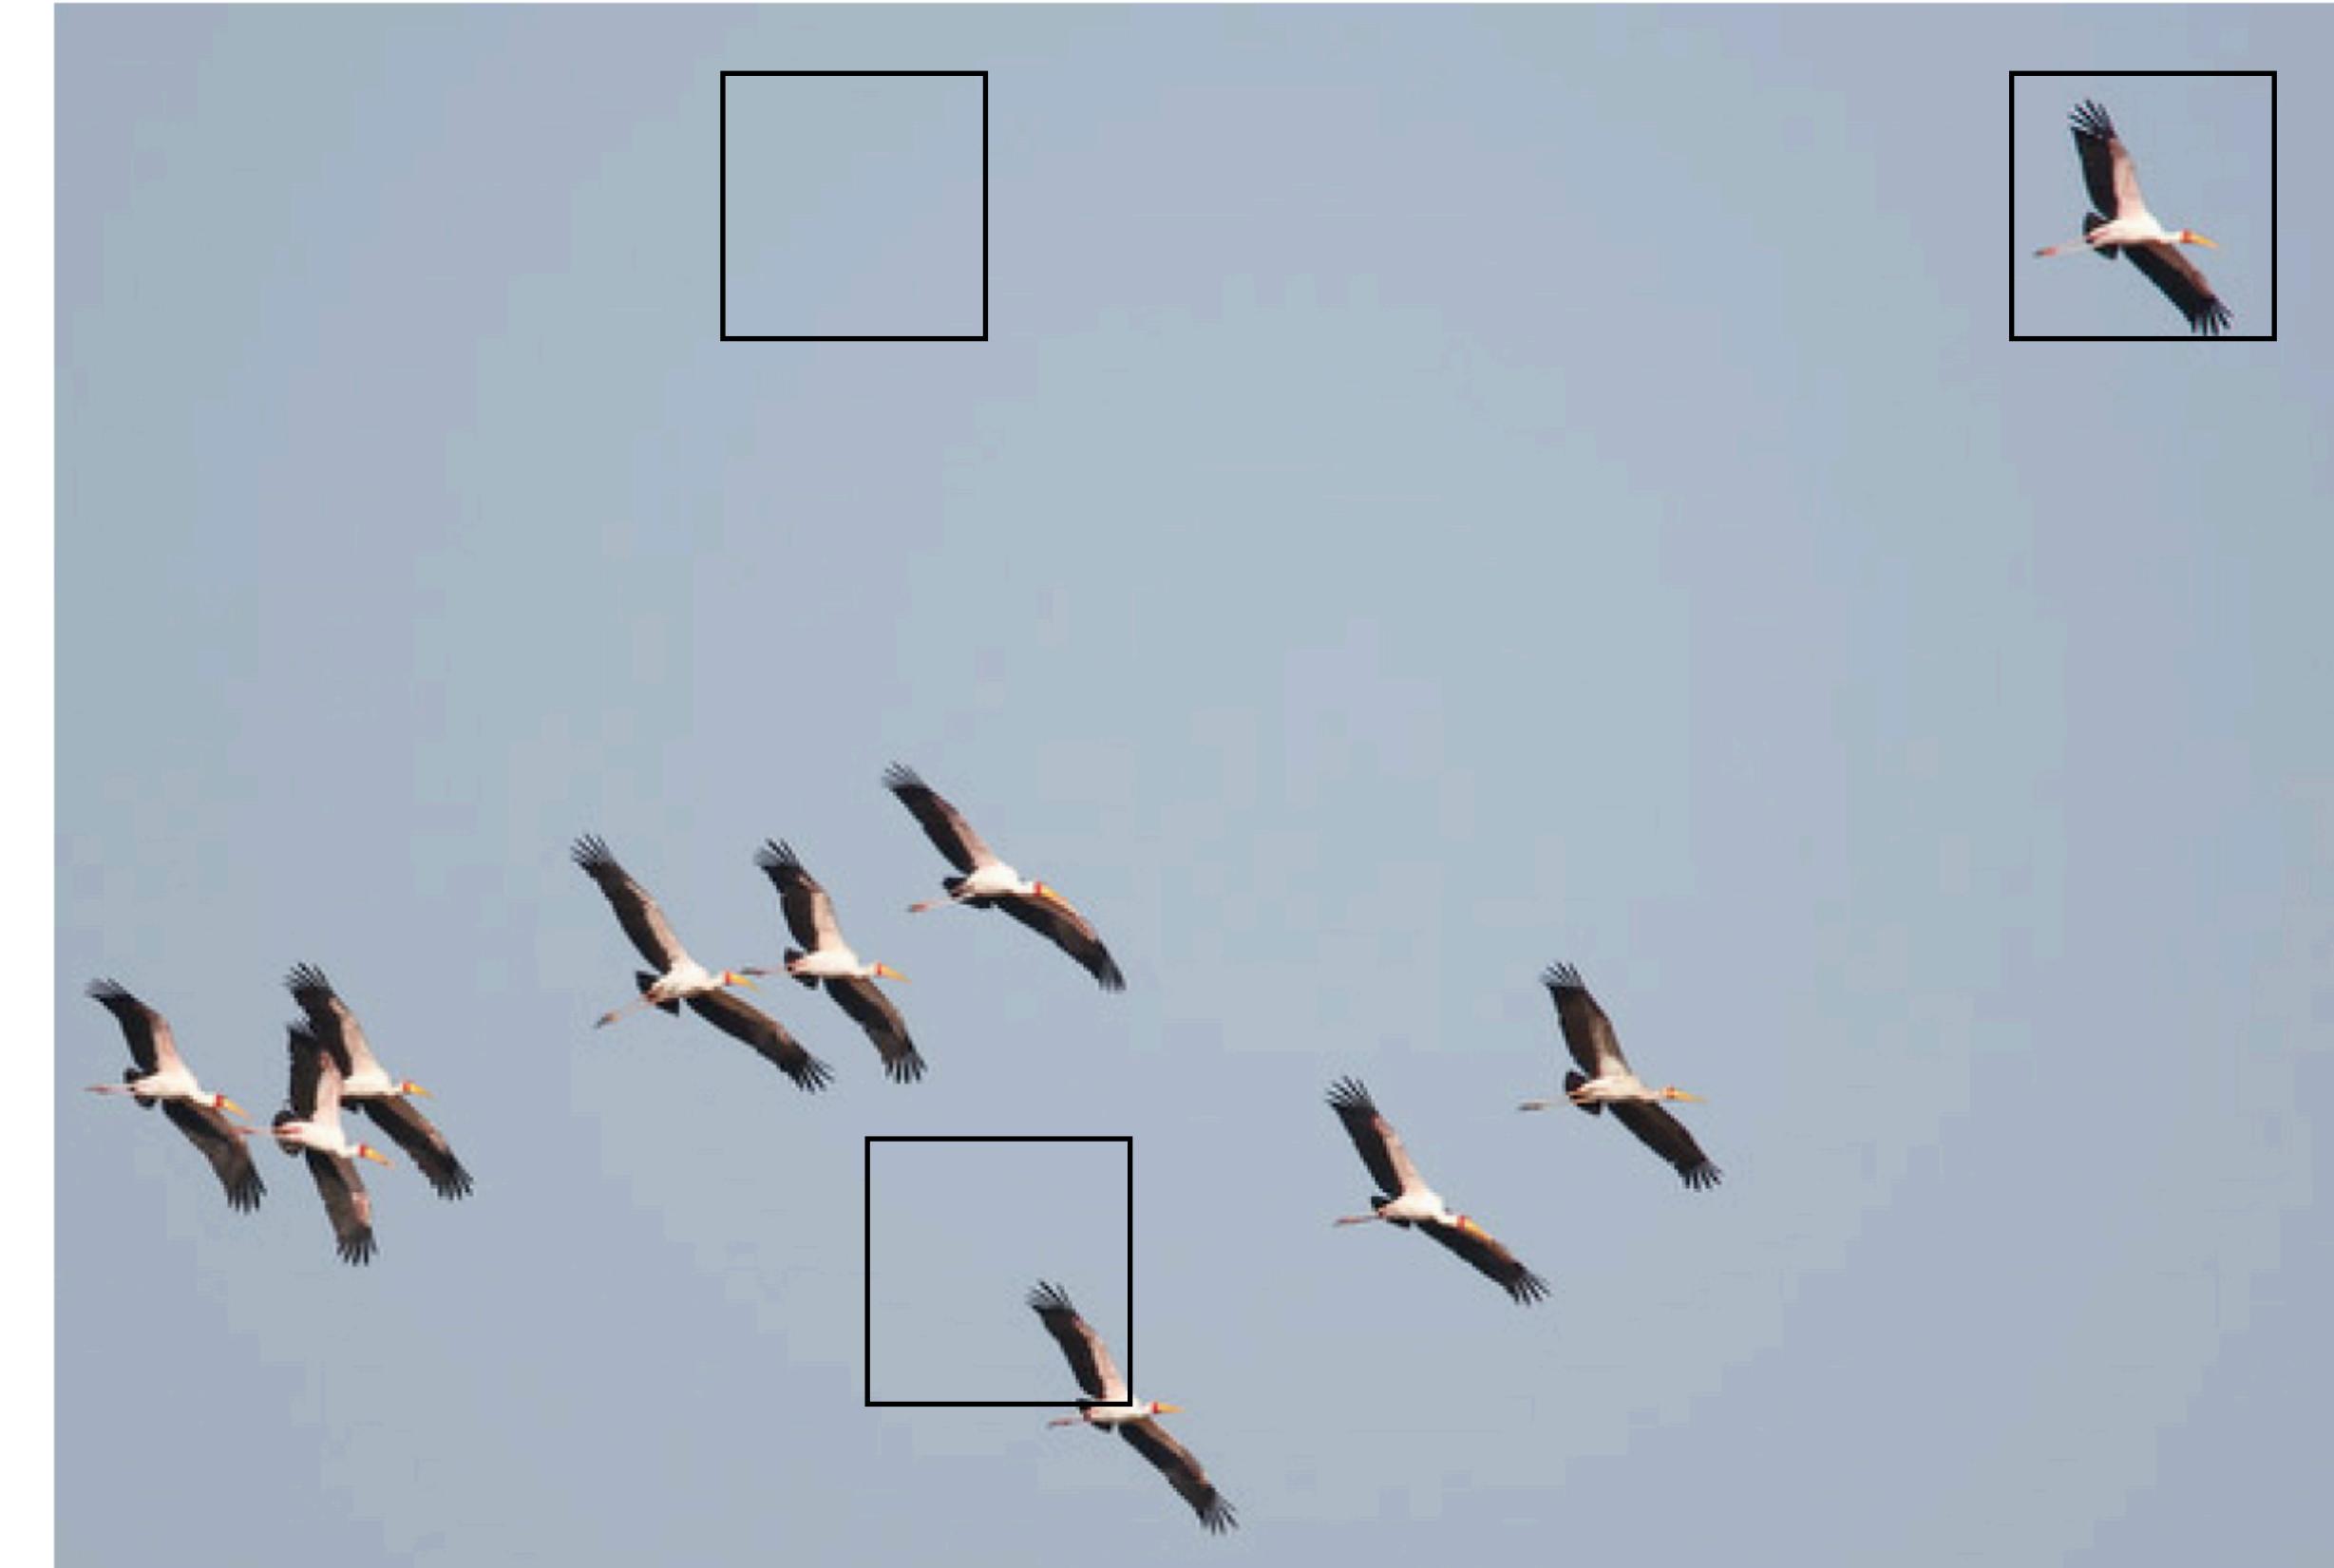
# Multi-Scale Pyramids

# Detecting a pattern in an image

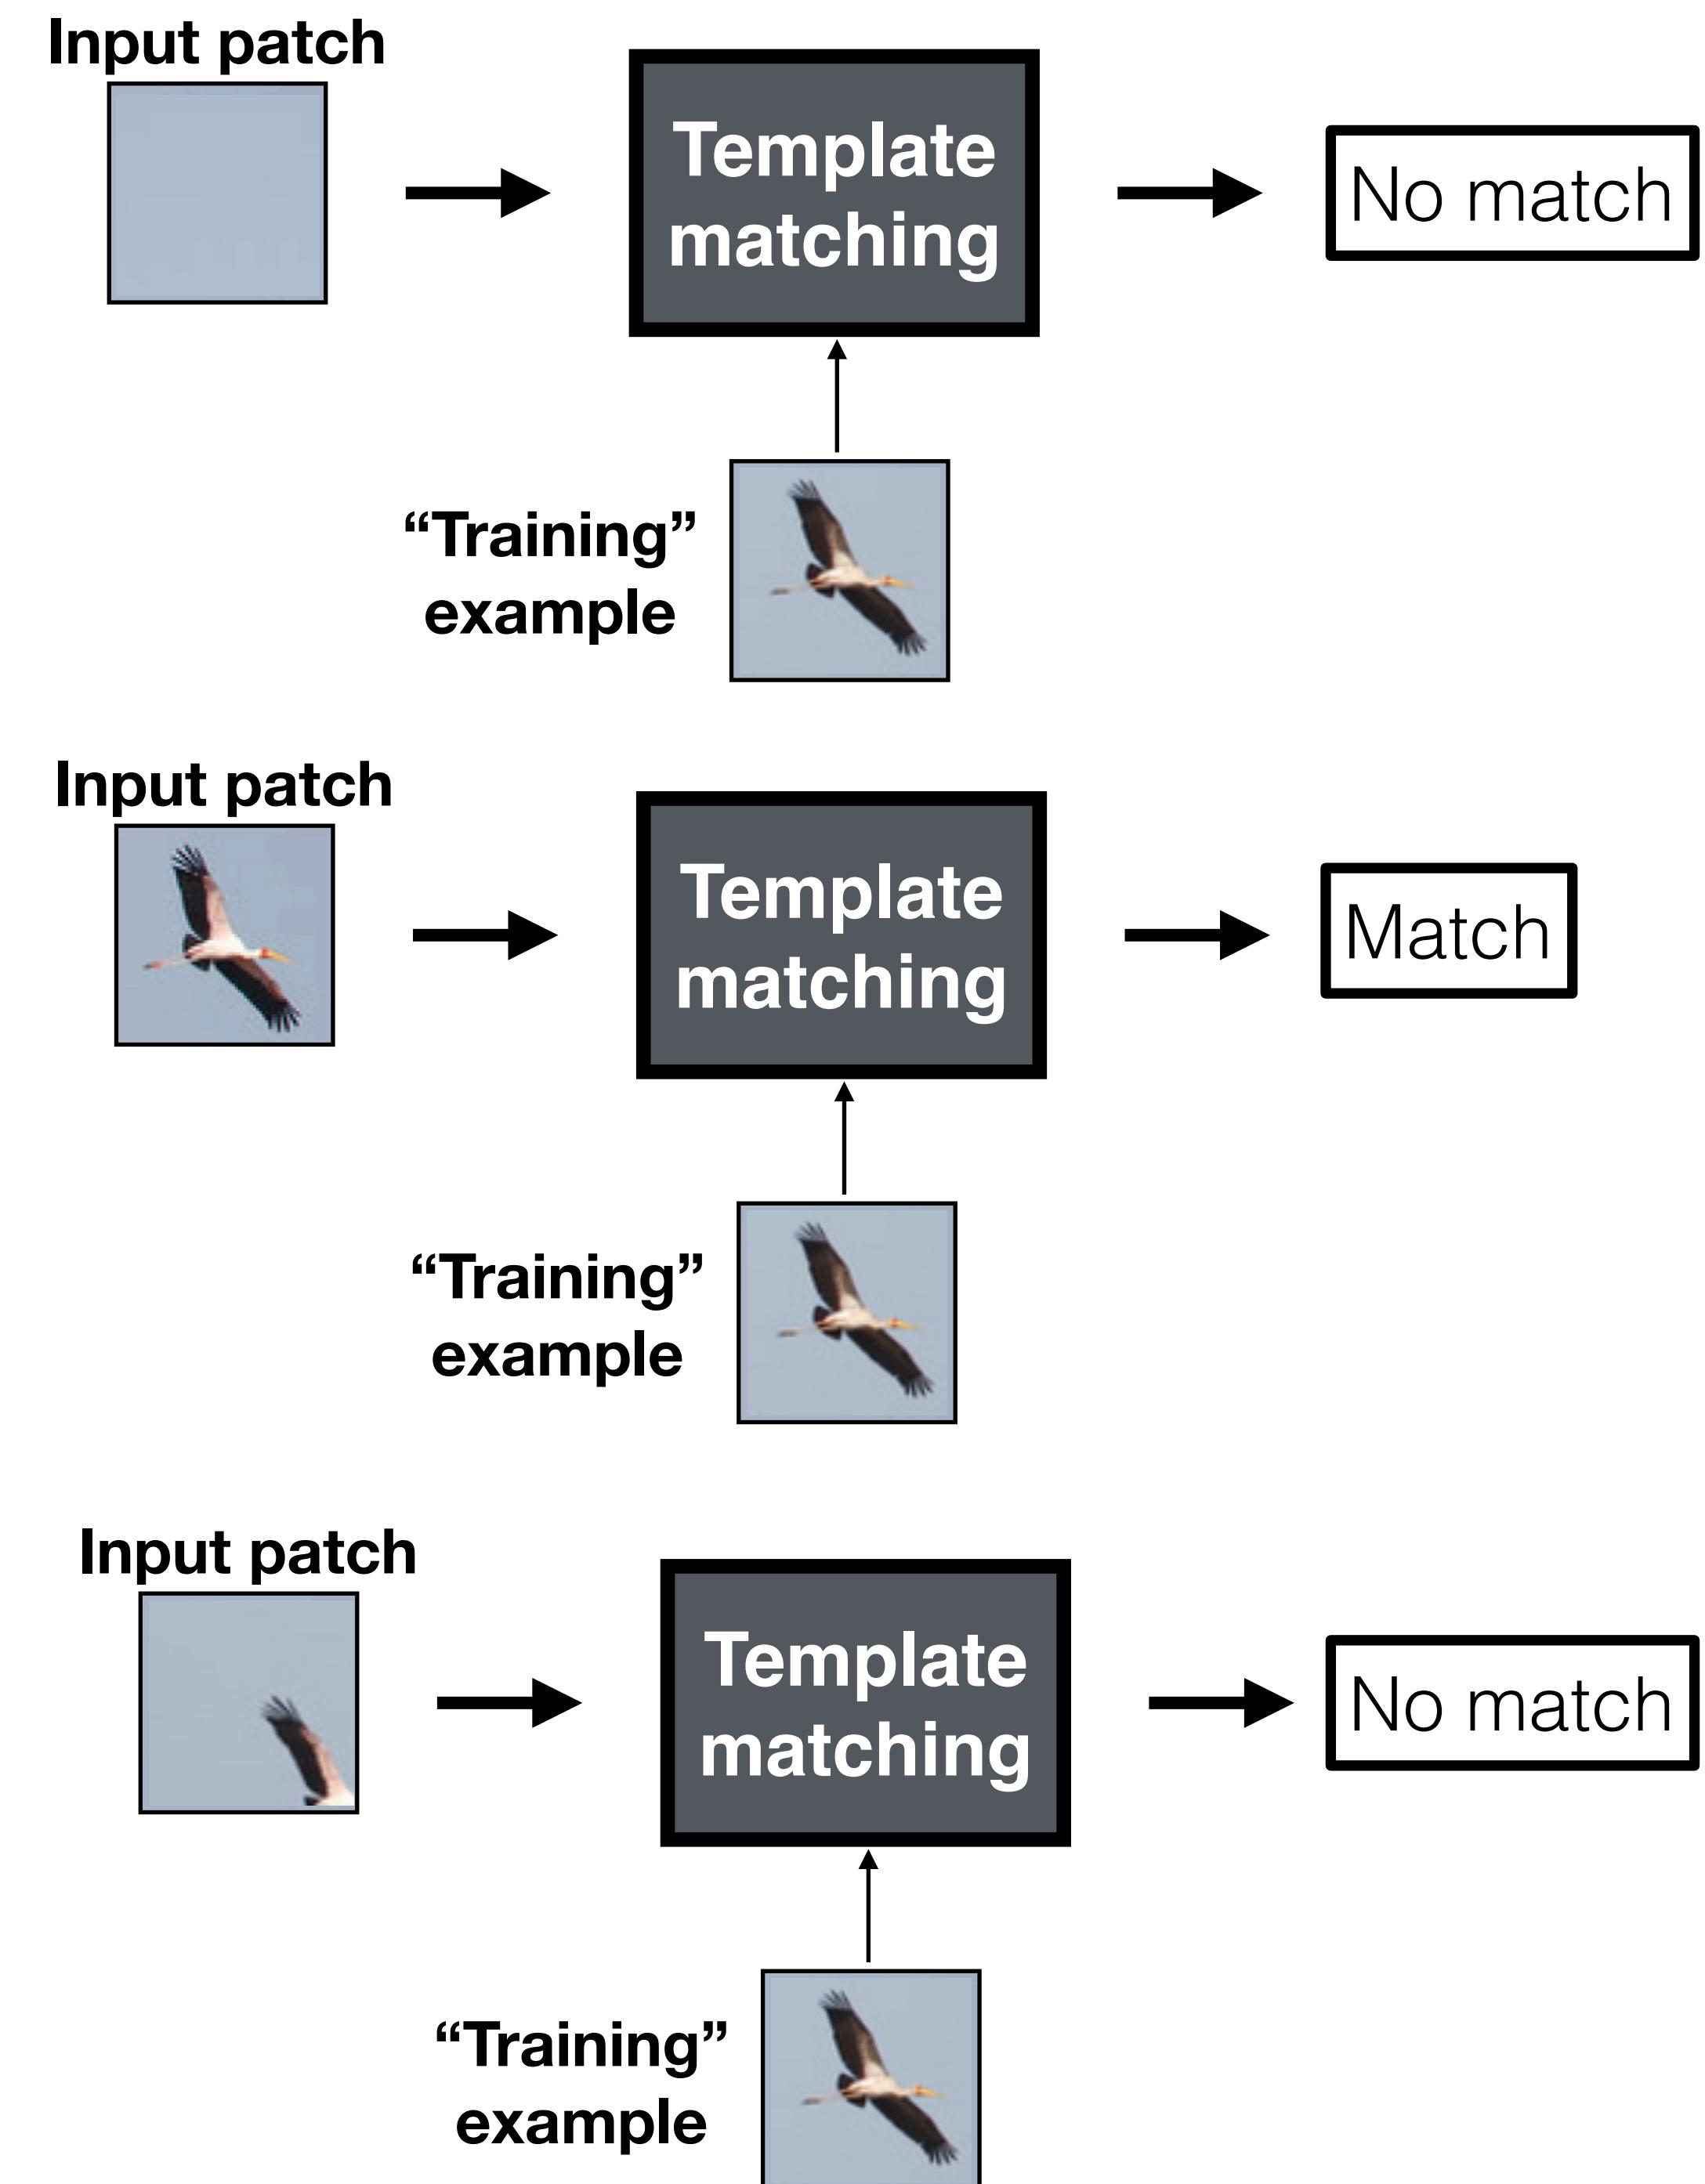


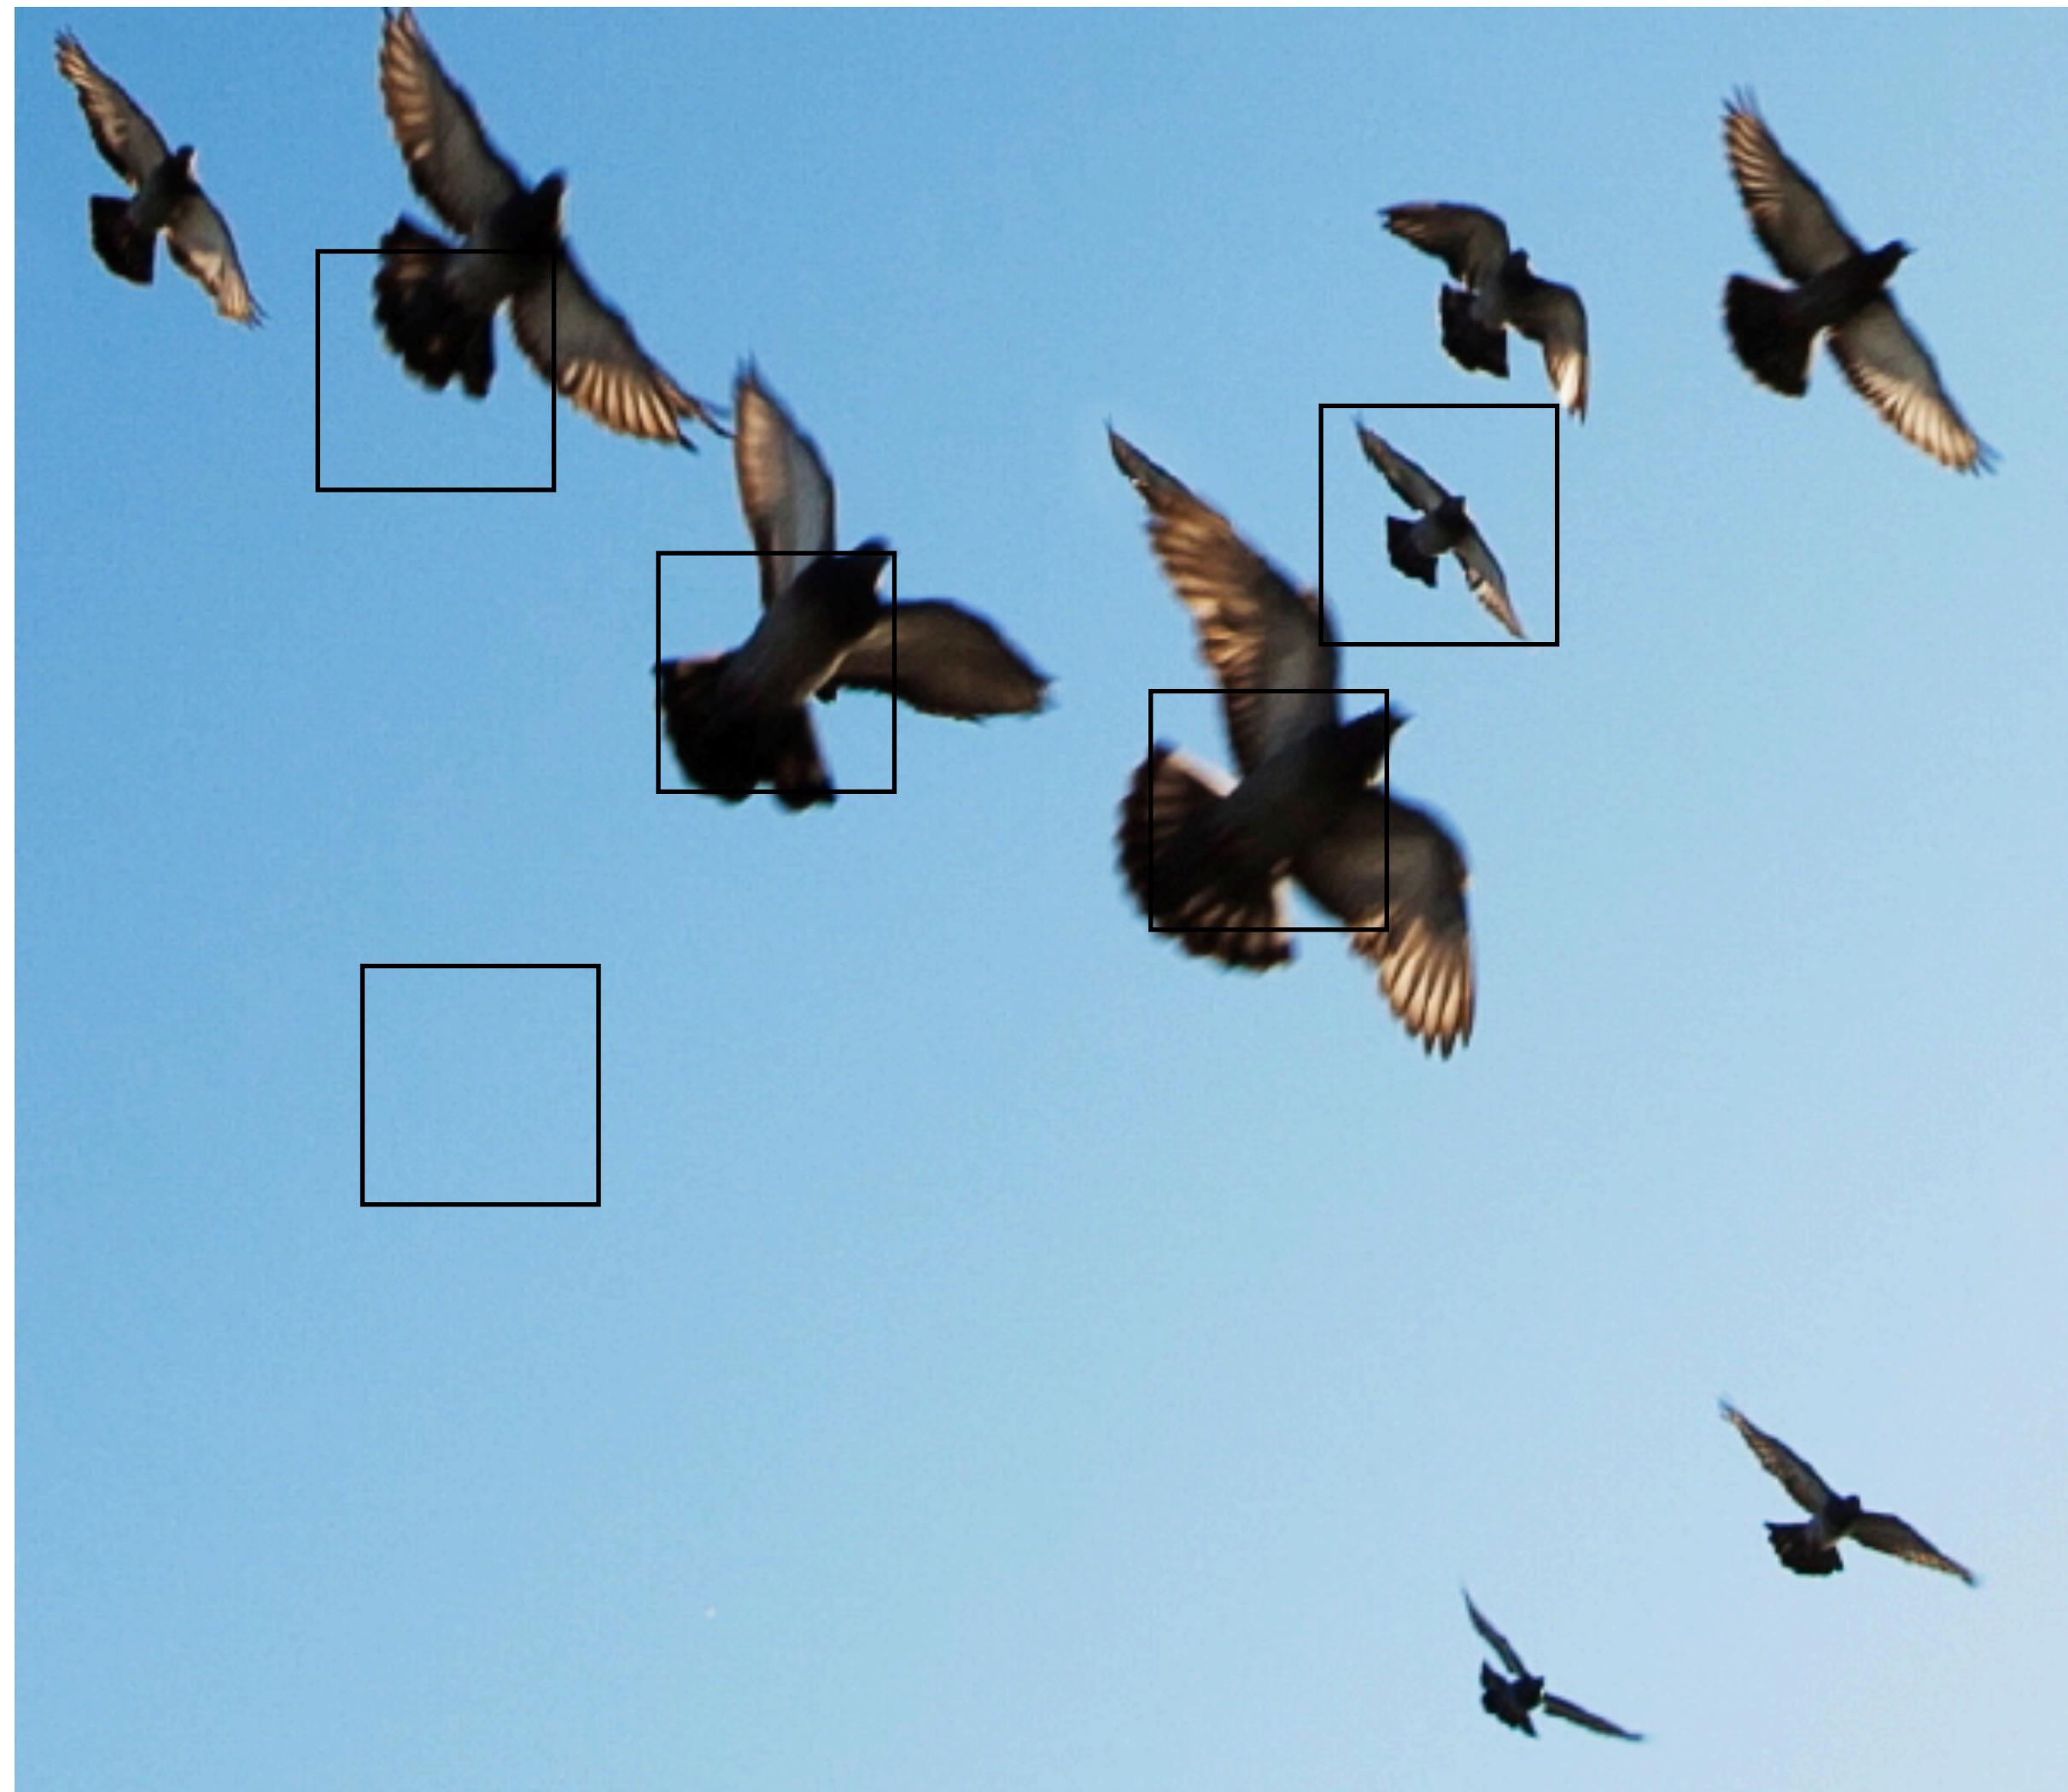
Pattern to detect



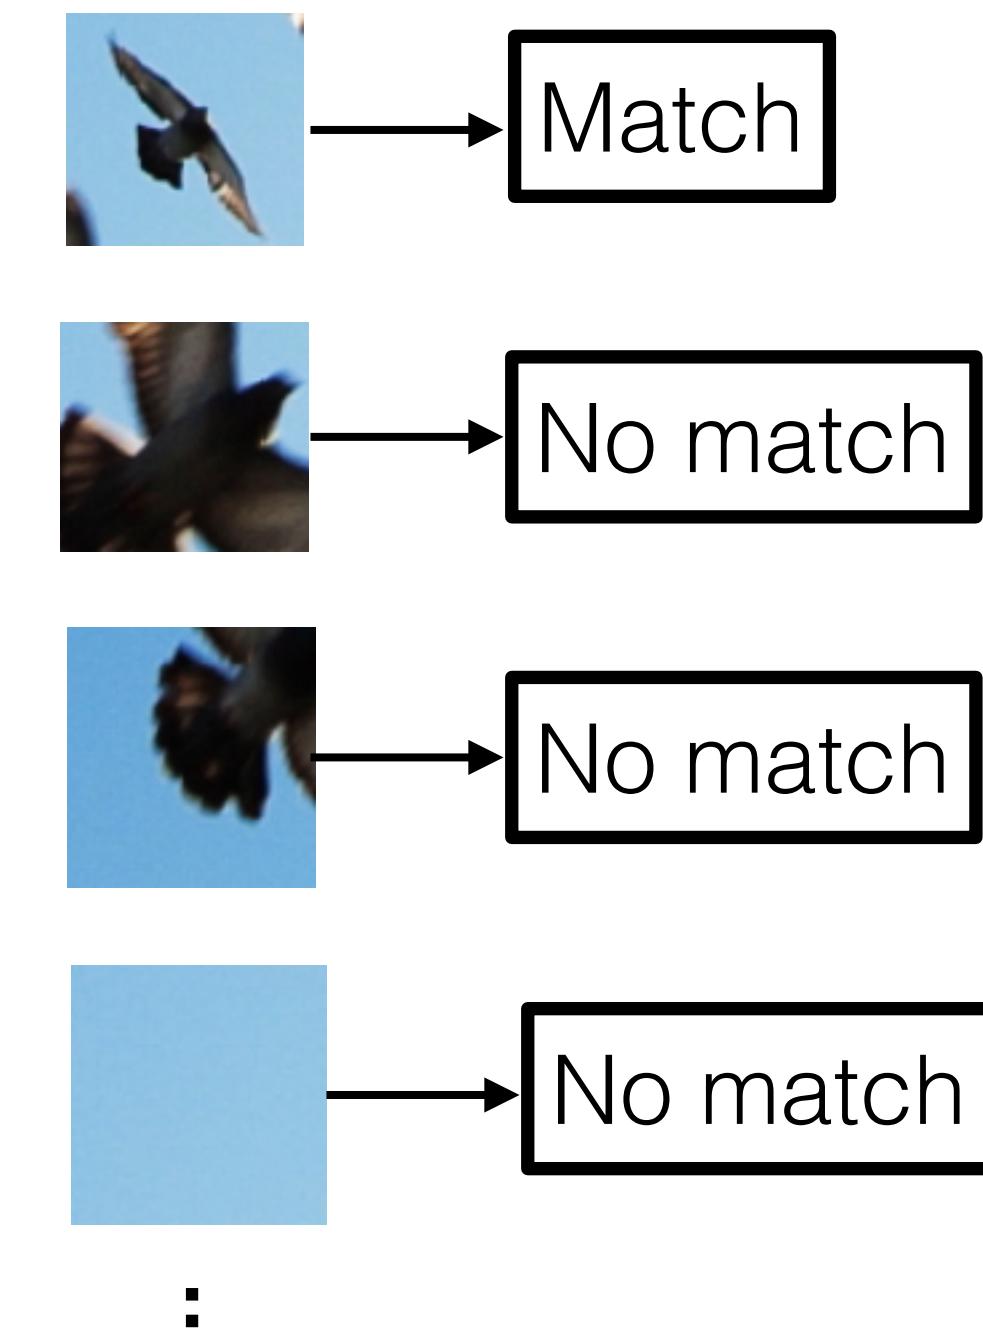


Motivation for translation invariance!

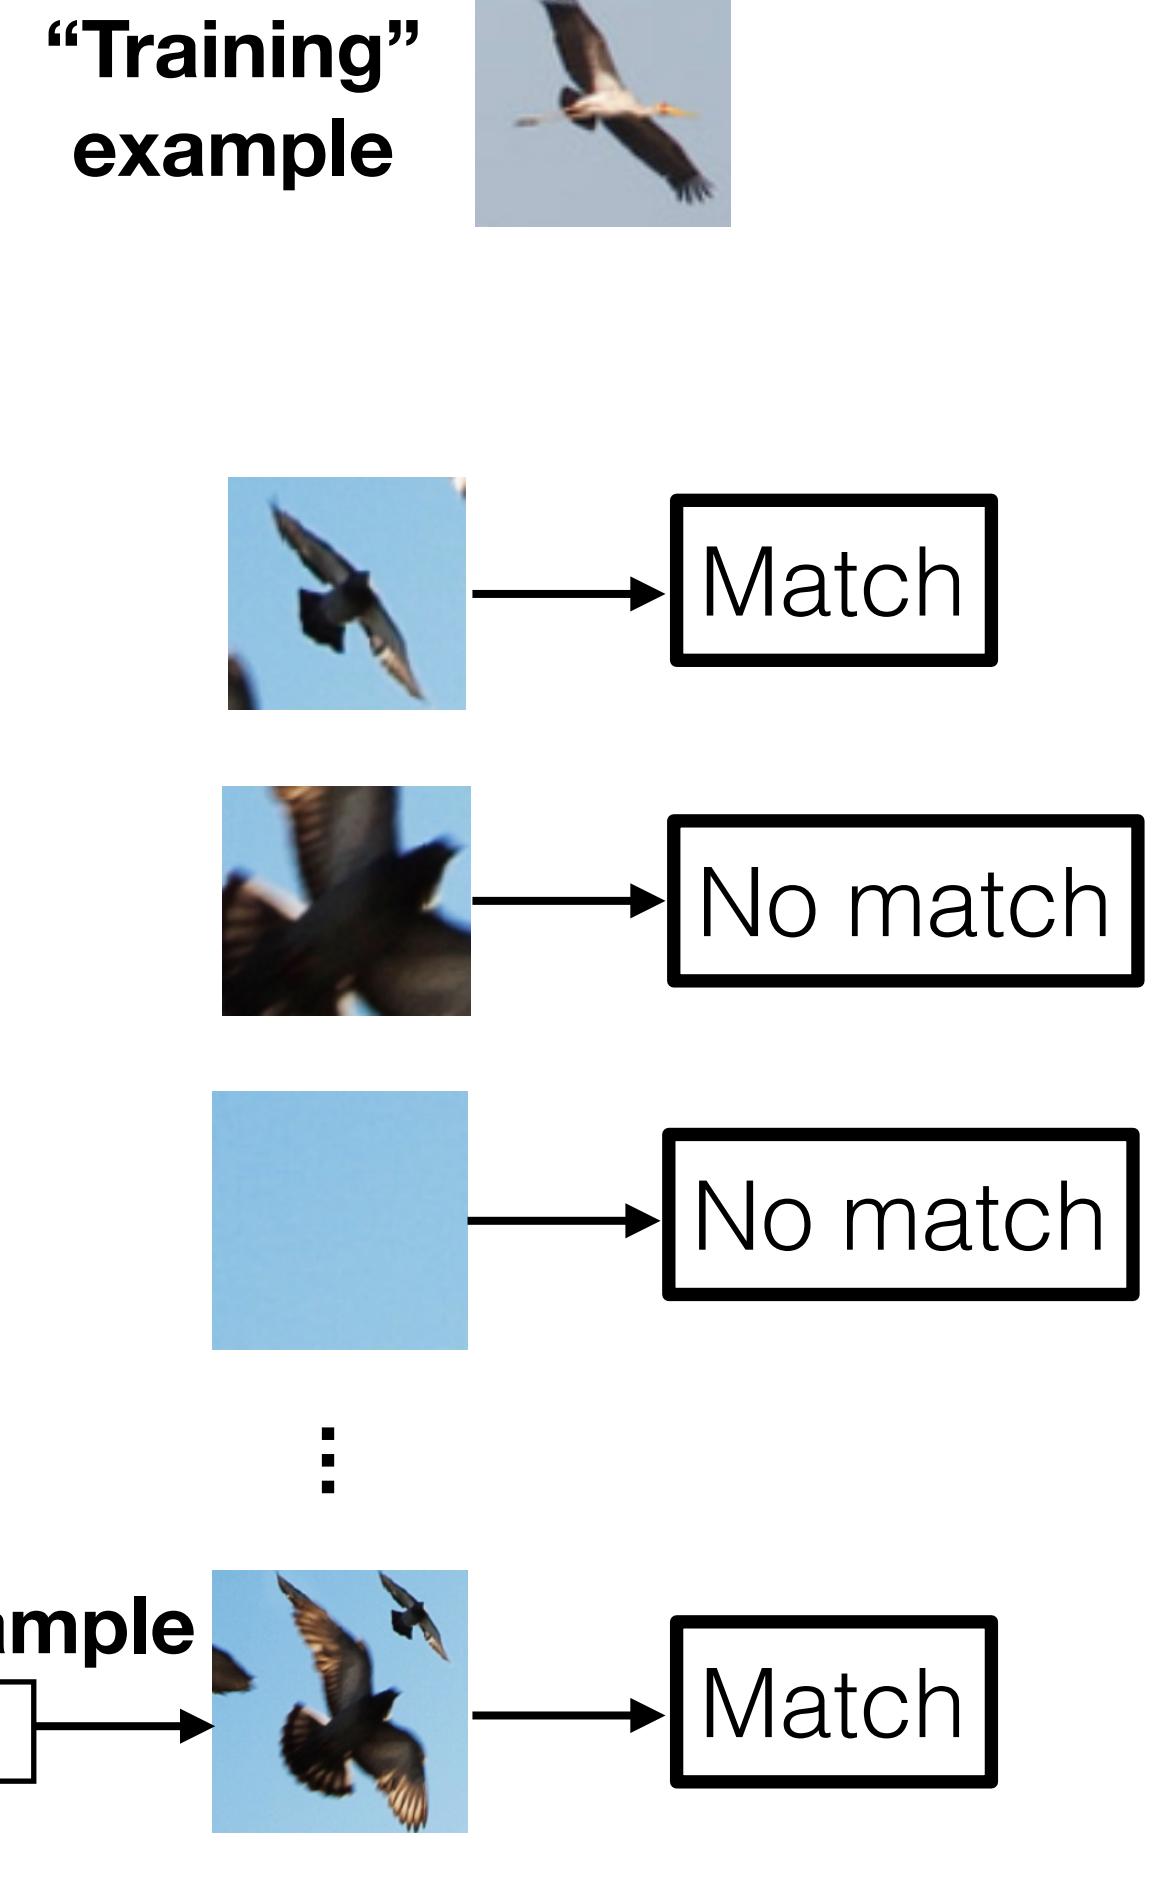
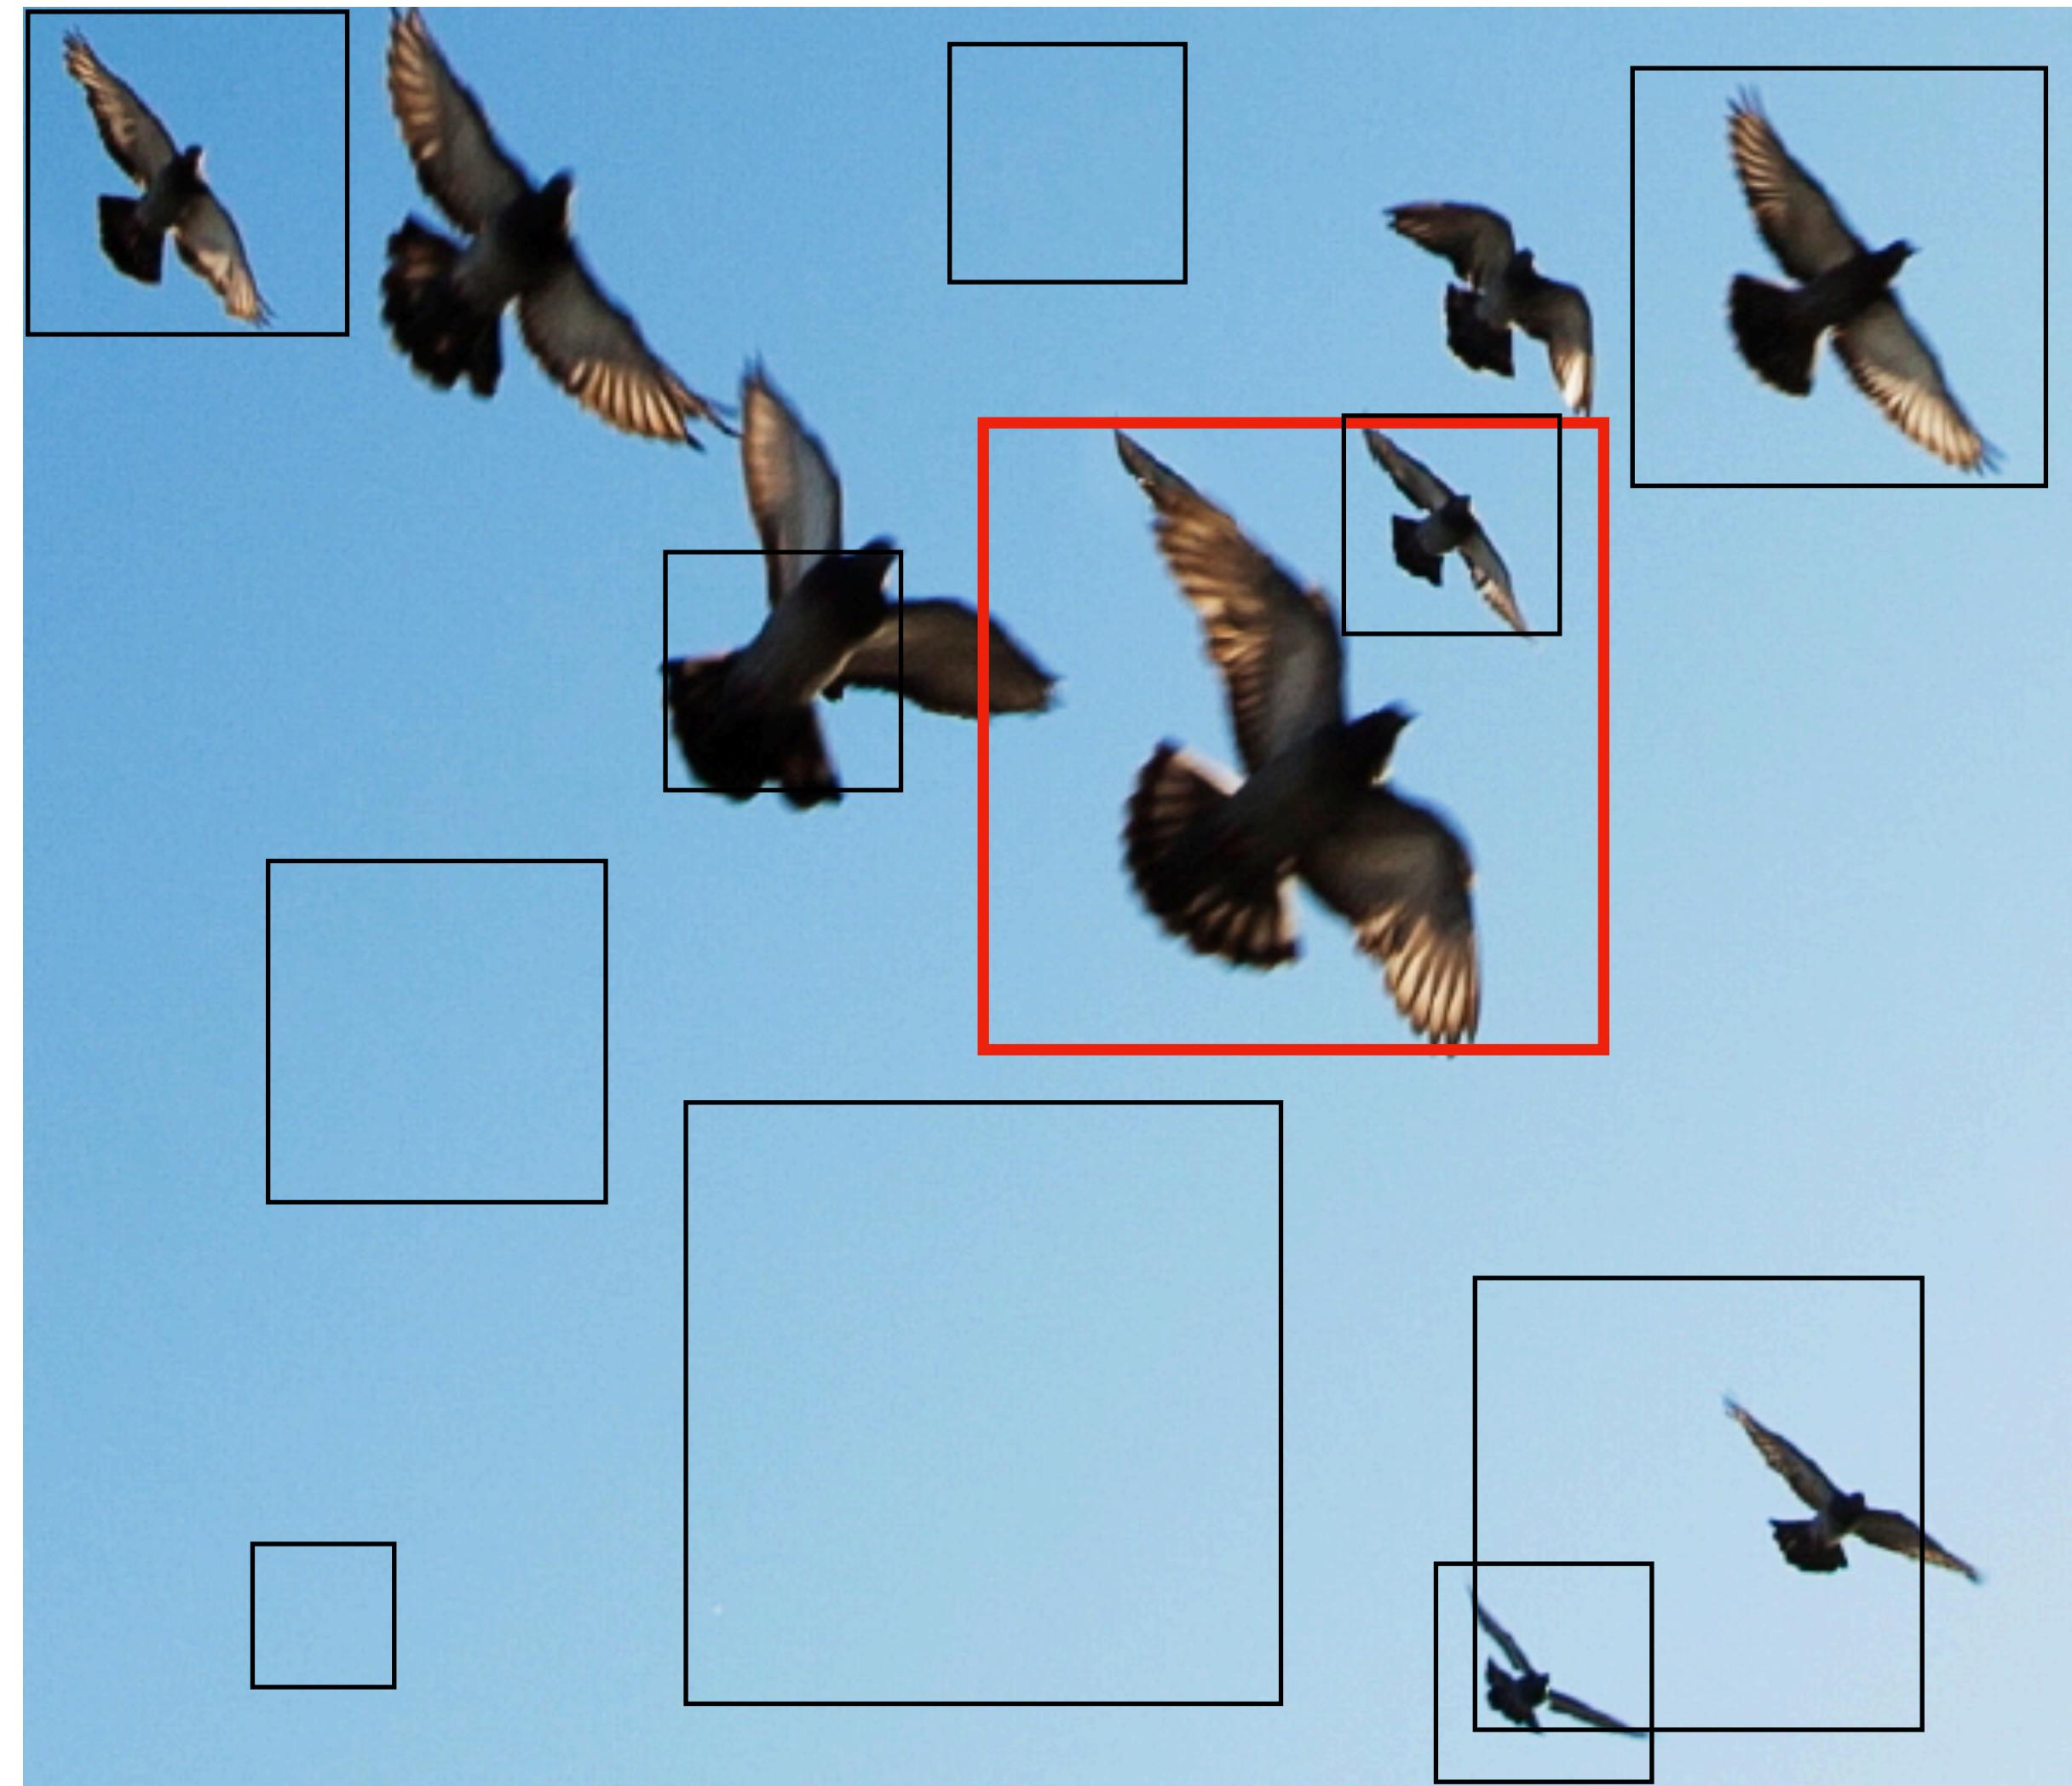




**“Training” example**



We need translation and **scale** invariance



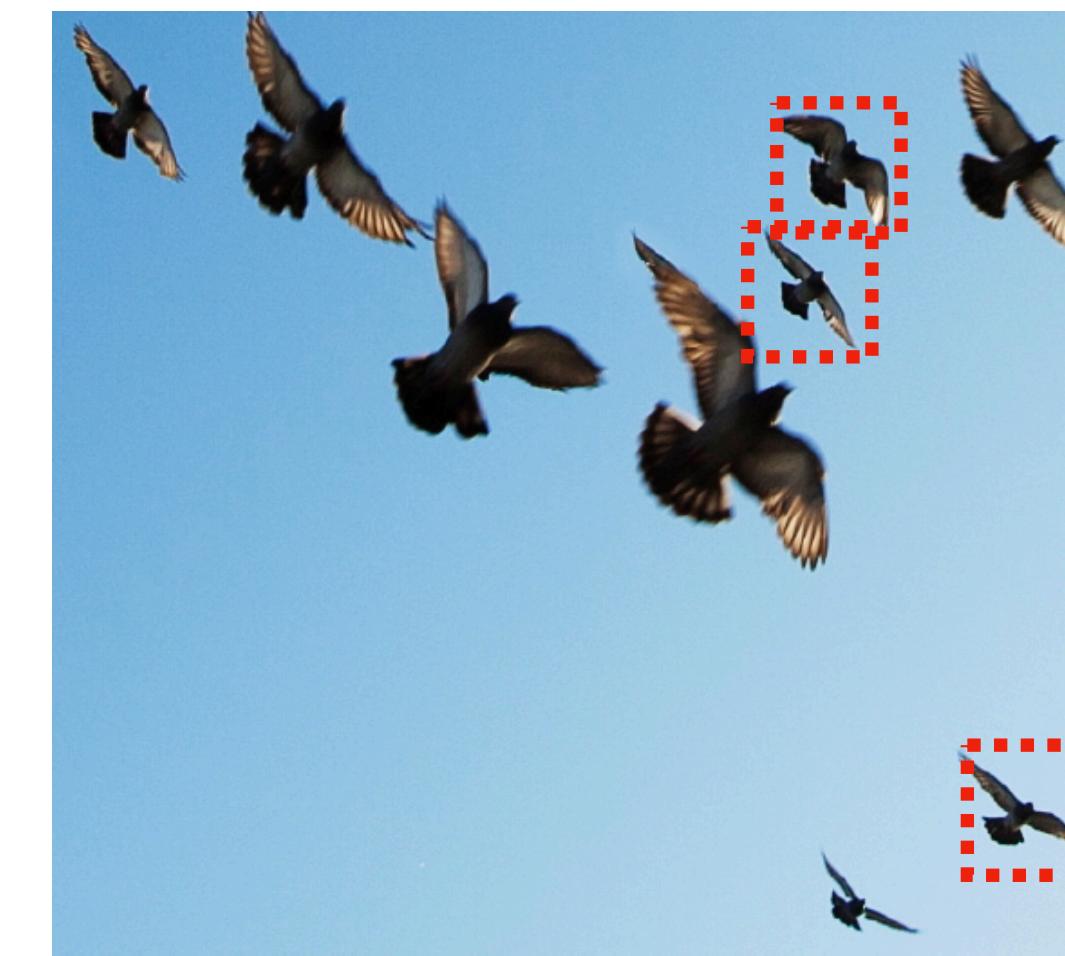
Consider all patch location and sizes

# Idea: Build Multi-Scale Pyramids

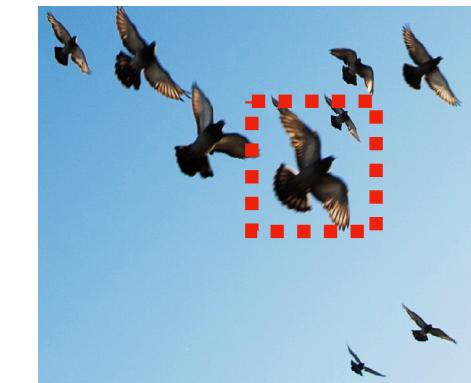
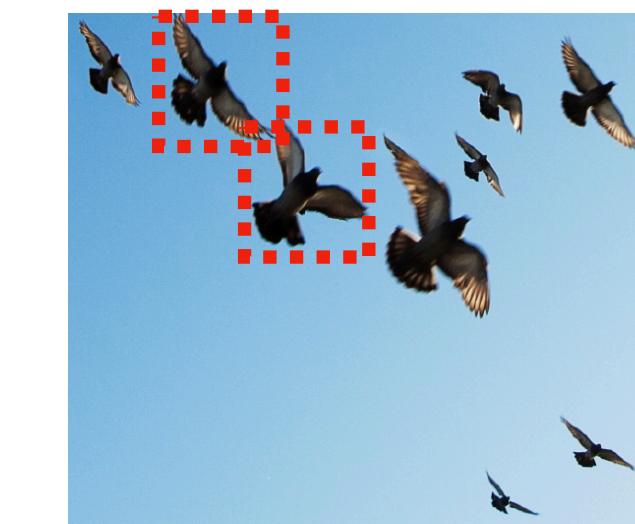
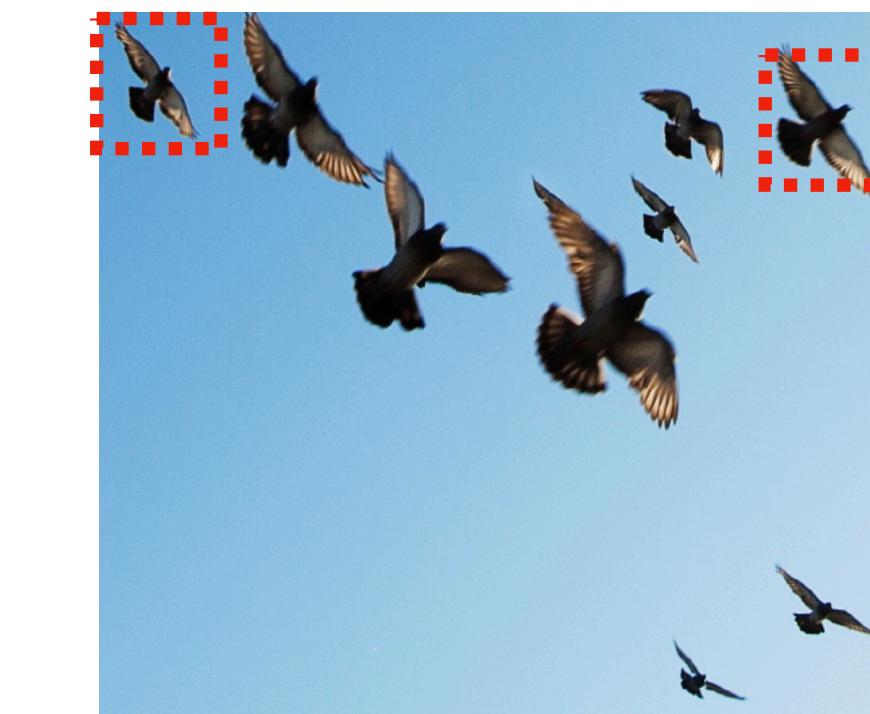


**Template**

# Idea: Build Multi-Scale Pyramids



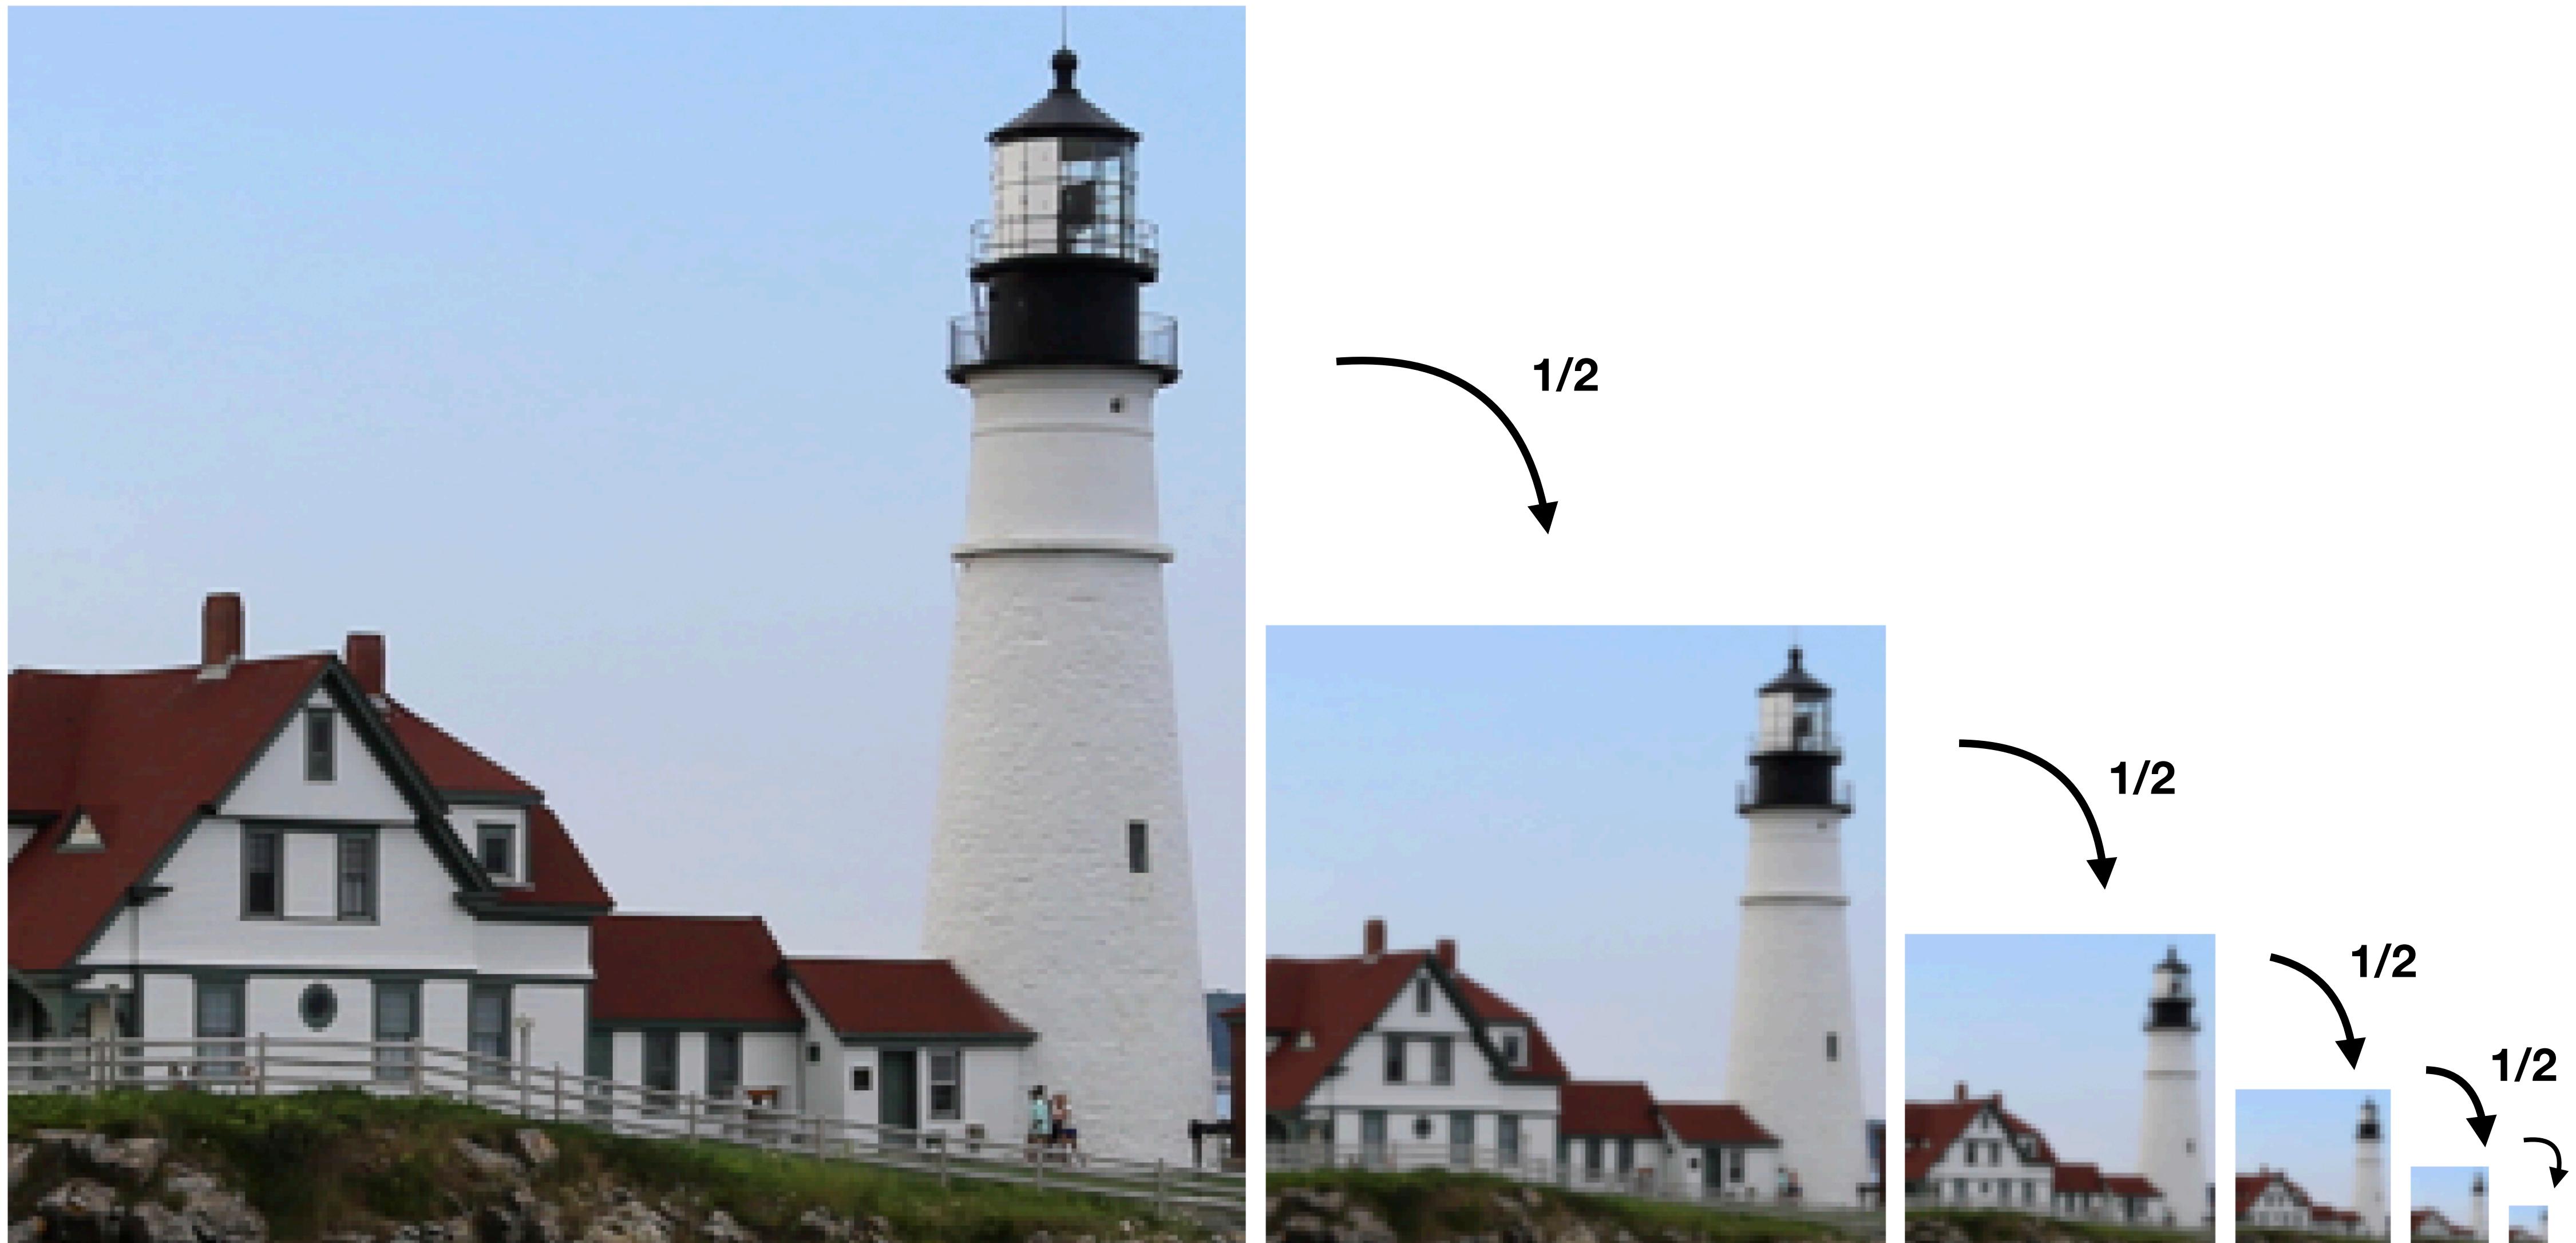
Multiscale image pyramid



Template

A fast & efficient way to provide scale & translation invariance!

# Gaussian Pyramid

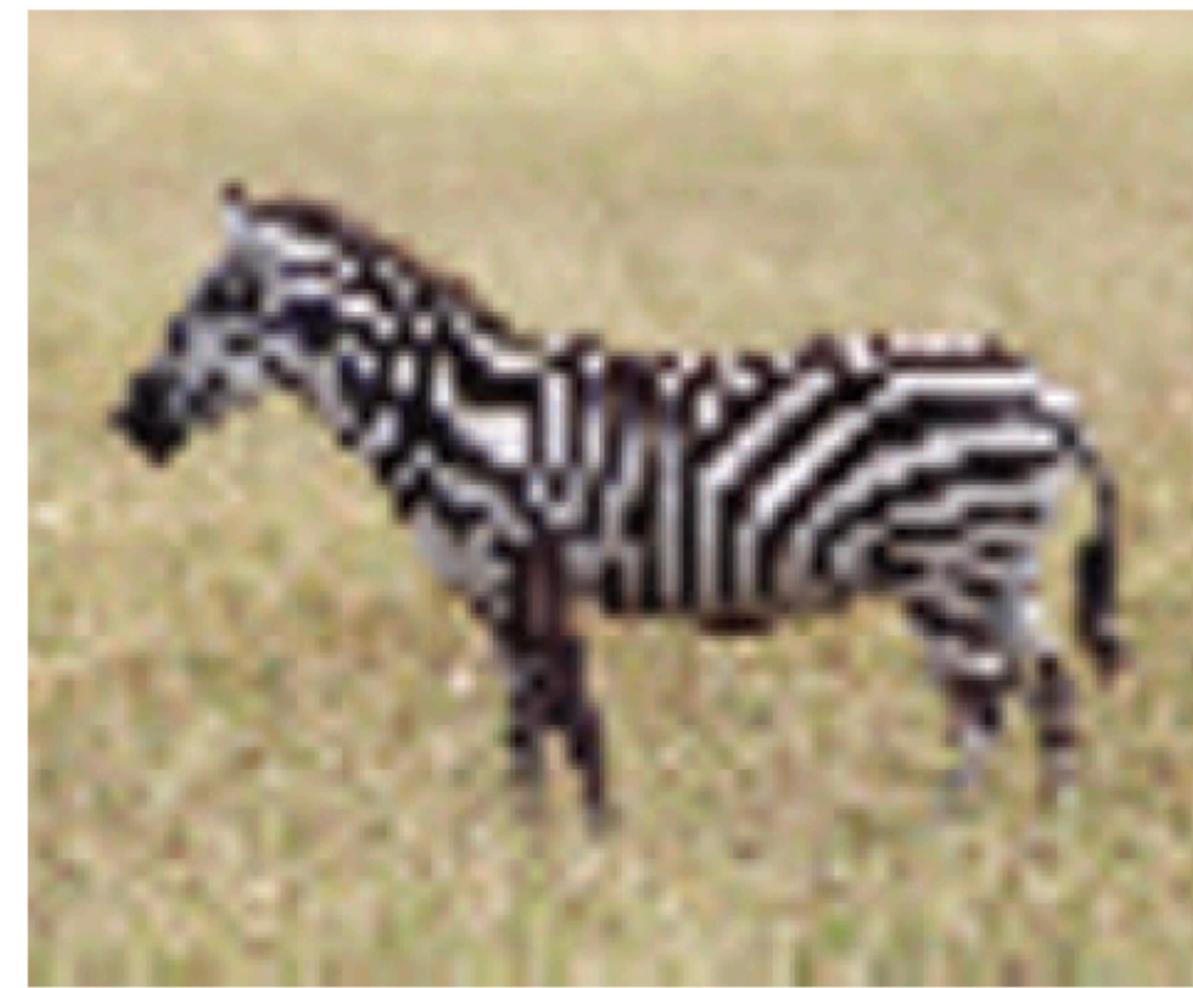


# Subsampling and aliasing

103×128



52×64



26×32



1/2

1/2

# The Gaussian pyramid

For each level

1. Blur input image with a Gaussian filter

$$[1, 4, 6, 4, 1]$$



$$\bigcirc [1, 4, 6, 4, 1] \bigcirc \begin{matrix} 6 \rightarrow \\ [1 \\ 4 \\ 4 \\ 1] \end{matrix}$$



(The Gaussian filter is approximated with a binomial filter)

# The Gaussian pyramid

For each level

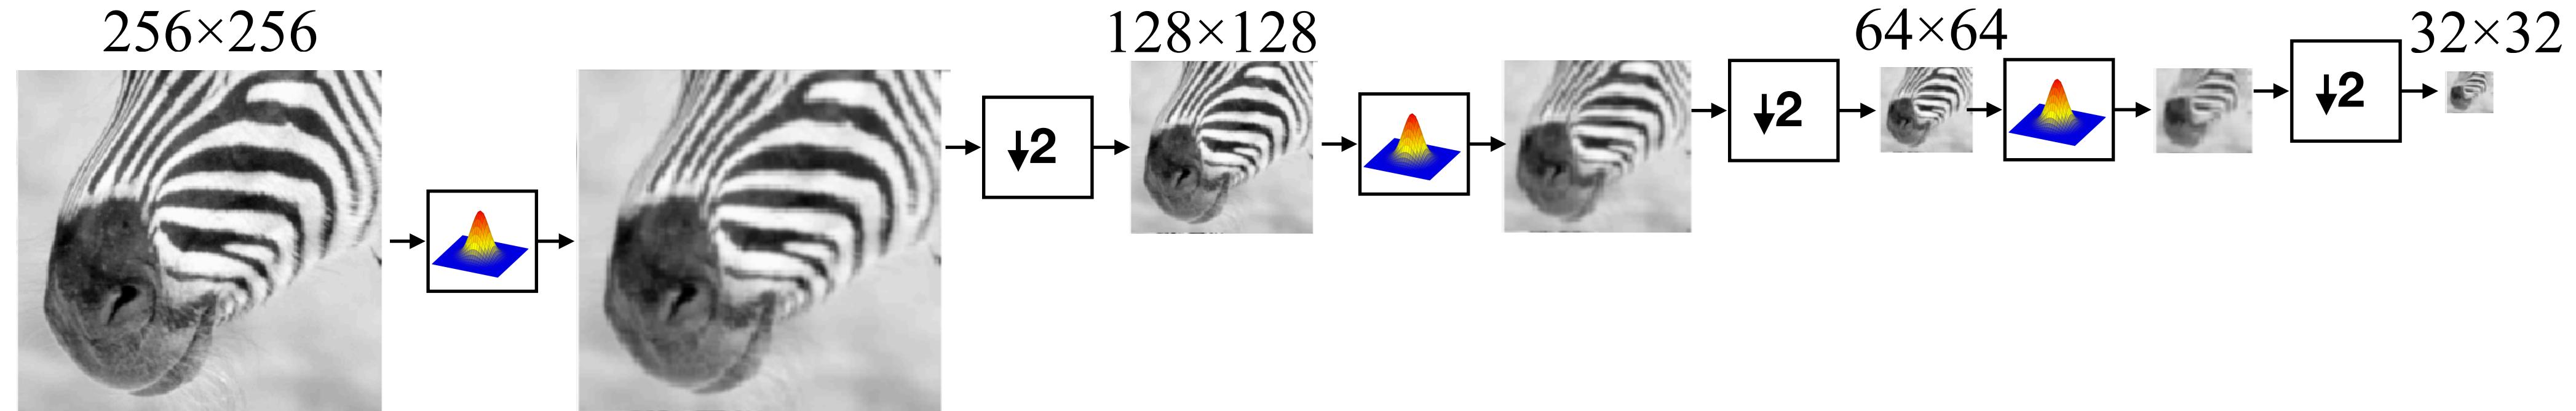
1. Blur input image with a Gaussian filter
2. Downsample image



$$\circ [1, 4, 6, 4, 1] \circ \begin{bmatrix} 1 \\ 4 \\ 6 \\ 4 \\ 1 \end{bmatrix}$$



# The Gaussian pyramid



# The Gaussian pyramid

512×512



(original image)

256×256



128×128



64×64



32×32

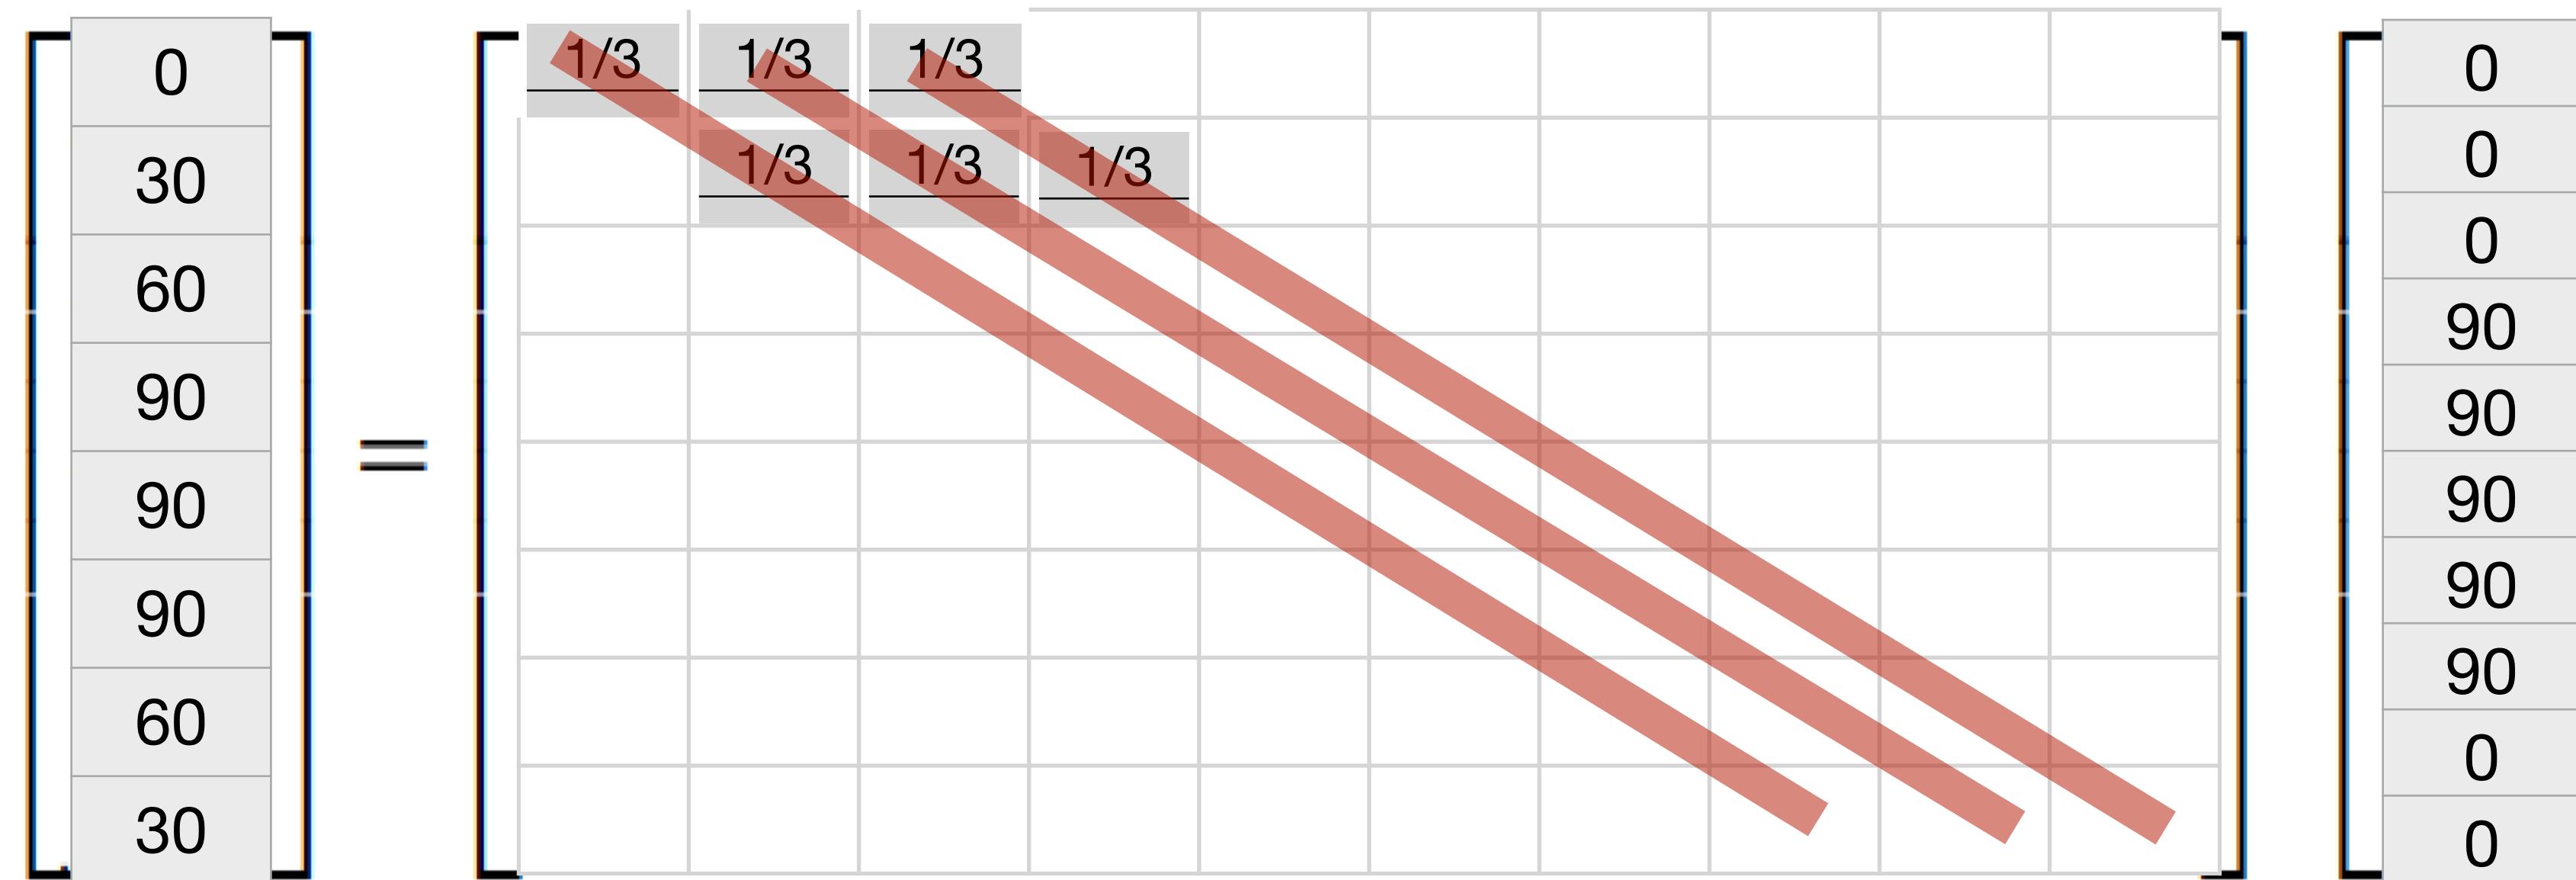


# Convolution as matrix multiplication

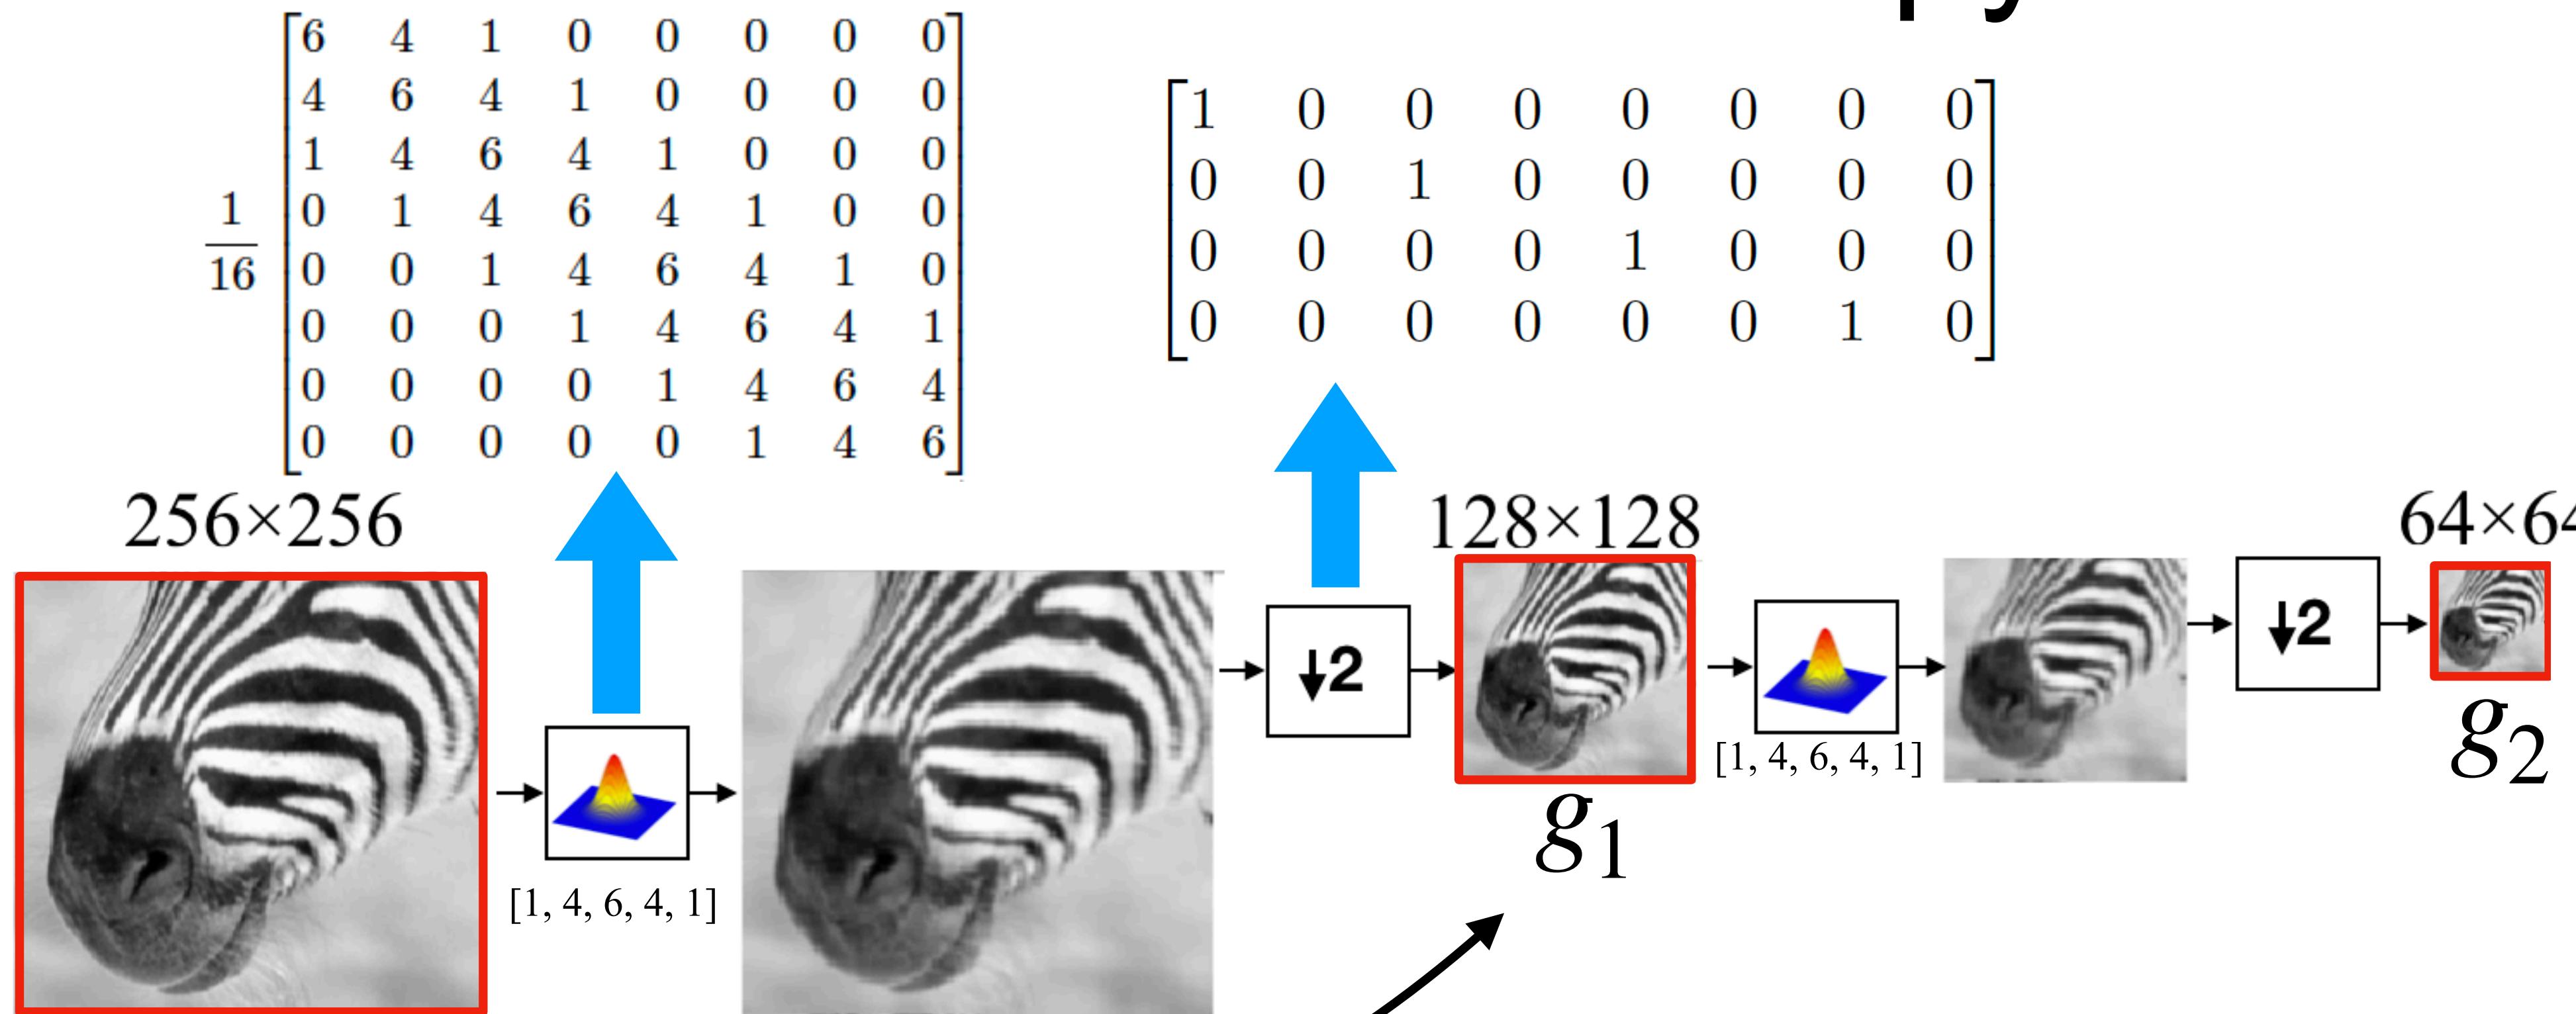
In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \end{bmatrix} \circ \begin{bmatrix} 1/3 & 1/3 & 1/3 \end{bmatrix} = \begin{bmatrix} \text{ } & 0 & 30 & 60 & 90 & 90 & 60 & 30 & \text{ } \end{bmatrix}$$

In the 1D case, it helps to make explicit the structure of the matrix:



# The Gaussian pyramid



$g_0$

$$g_1 = G_0 g_0$$

$$G_0 = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix} \frac{1}{16} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

(The arrays shown here are for 1D signals)

# The Gaussian pyramid

$$\frac{1}{16} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

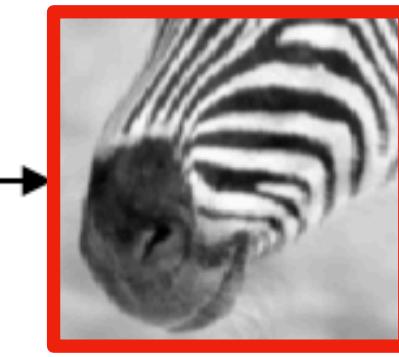
$256 \times 256$



$[1, 4, 6, 4, 1]$

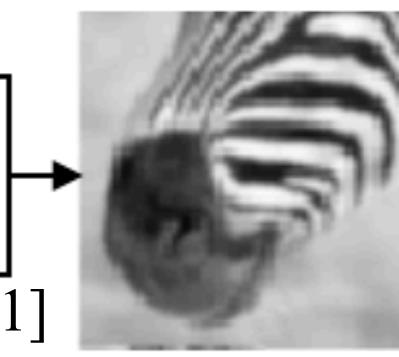
$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

$128 \times 128$



$g_1$

$$\downarrow 2$$



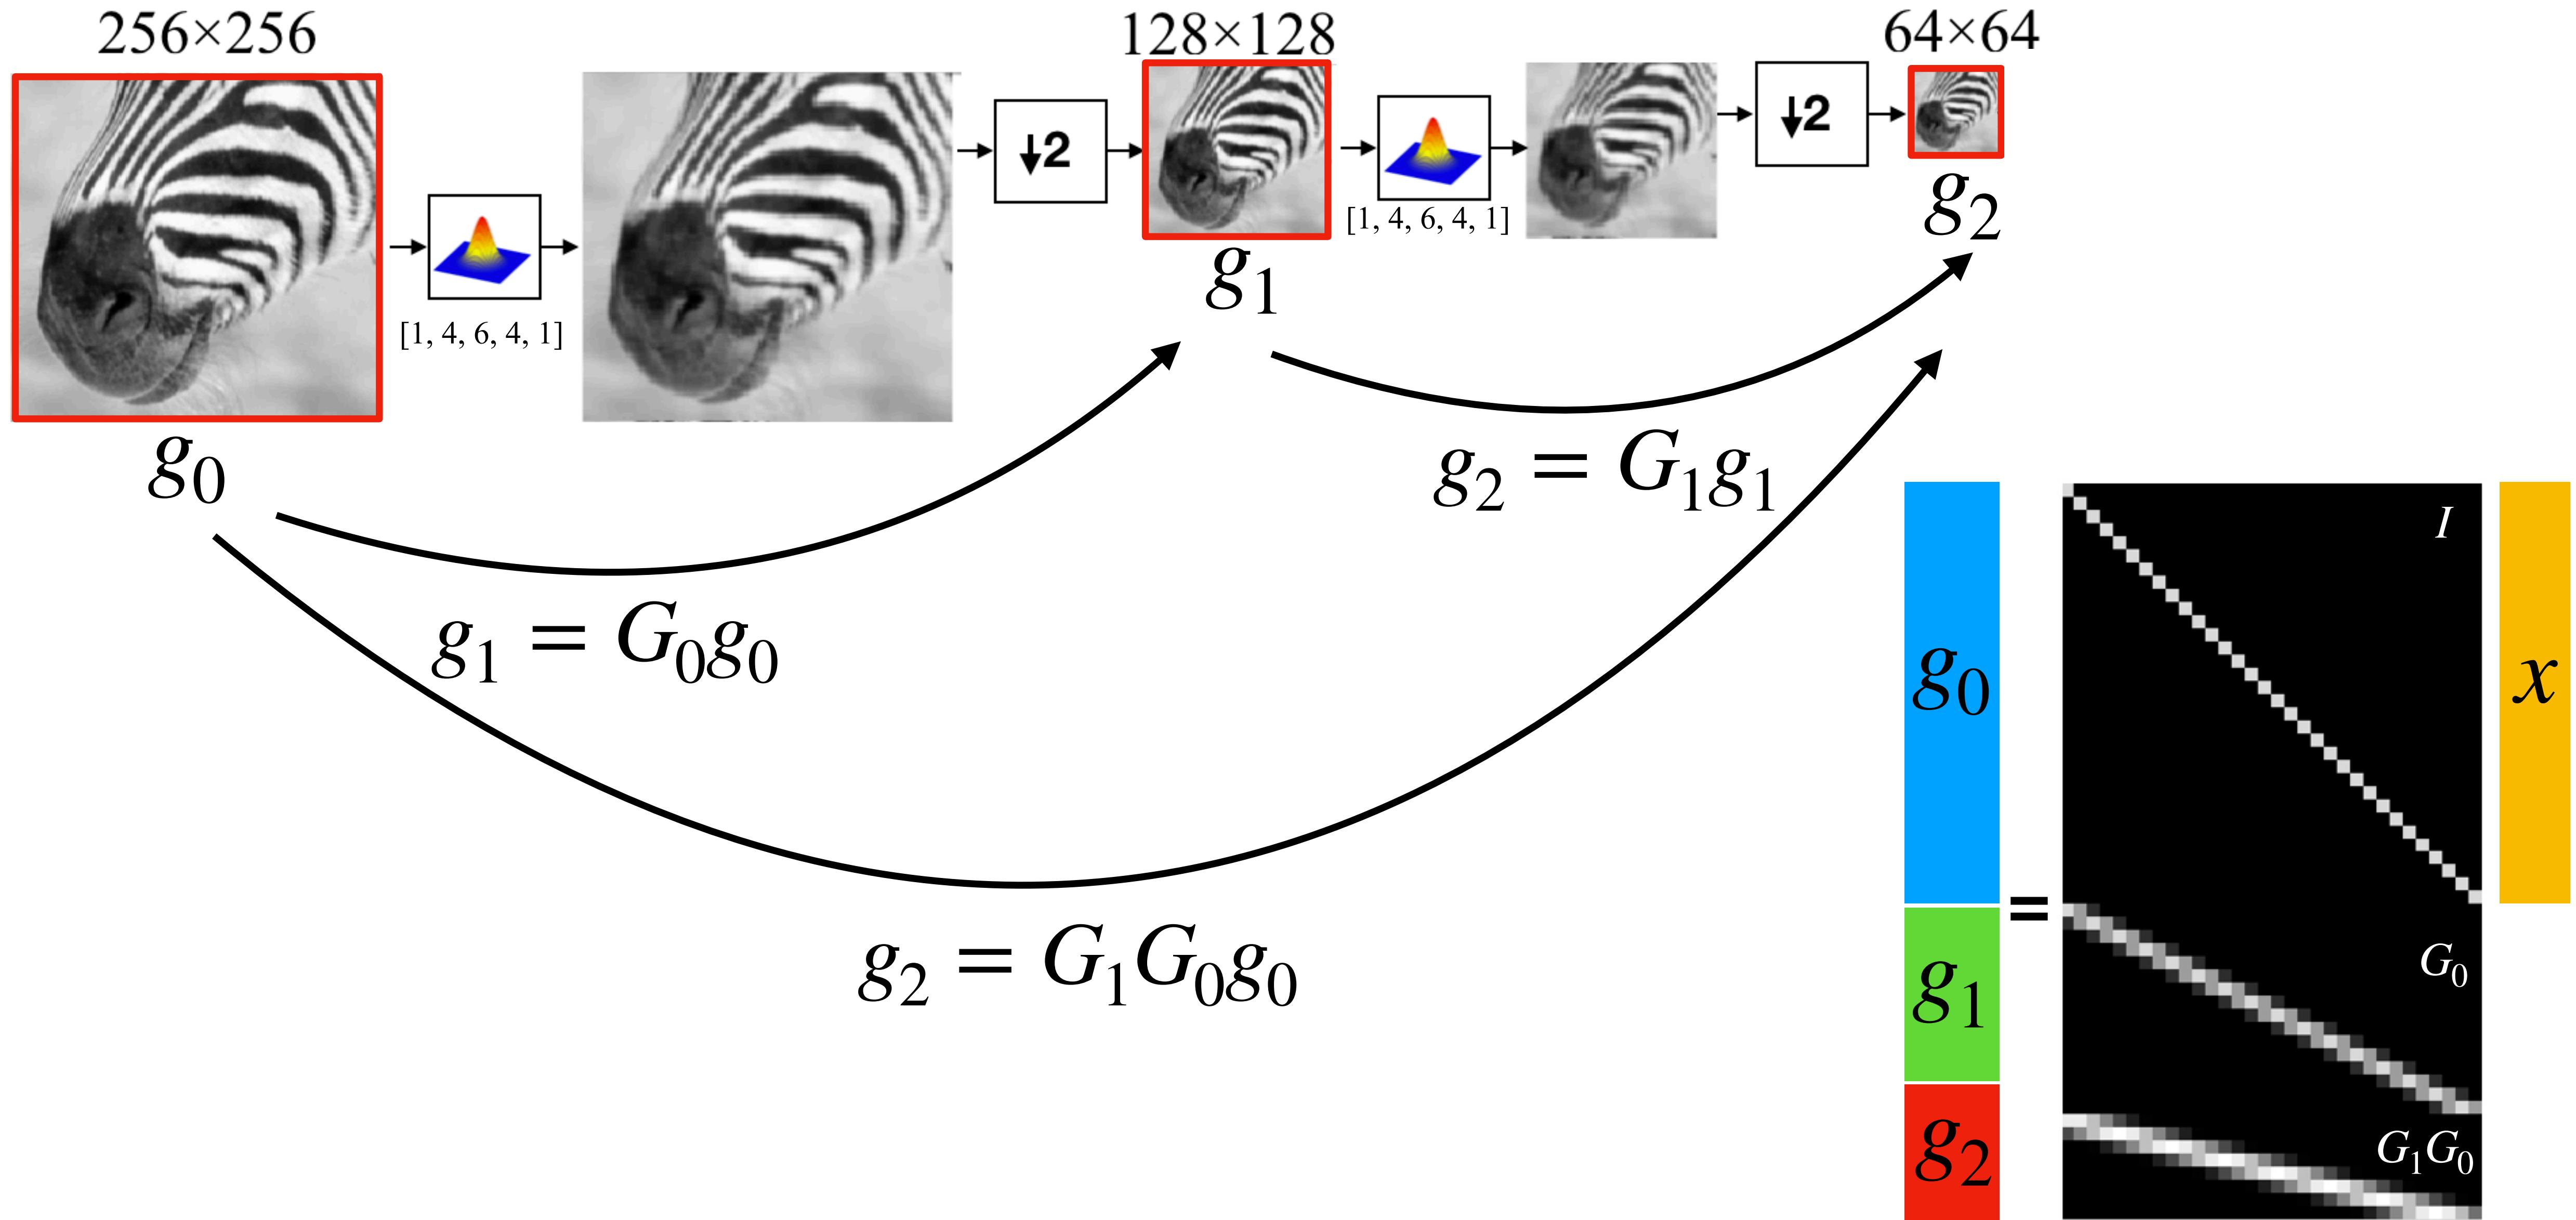
$64 \times 64$   
 $g_2$

$g_0$

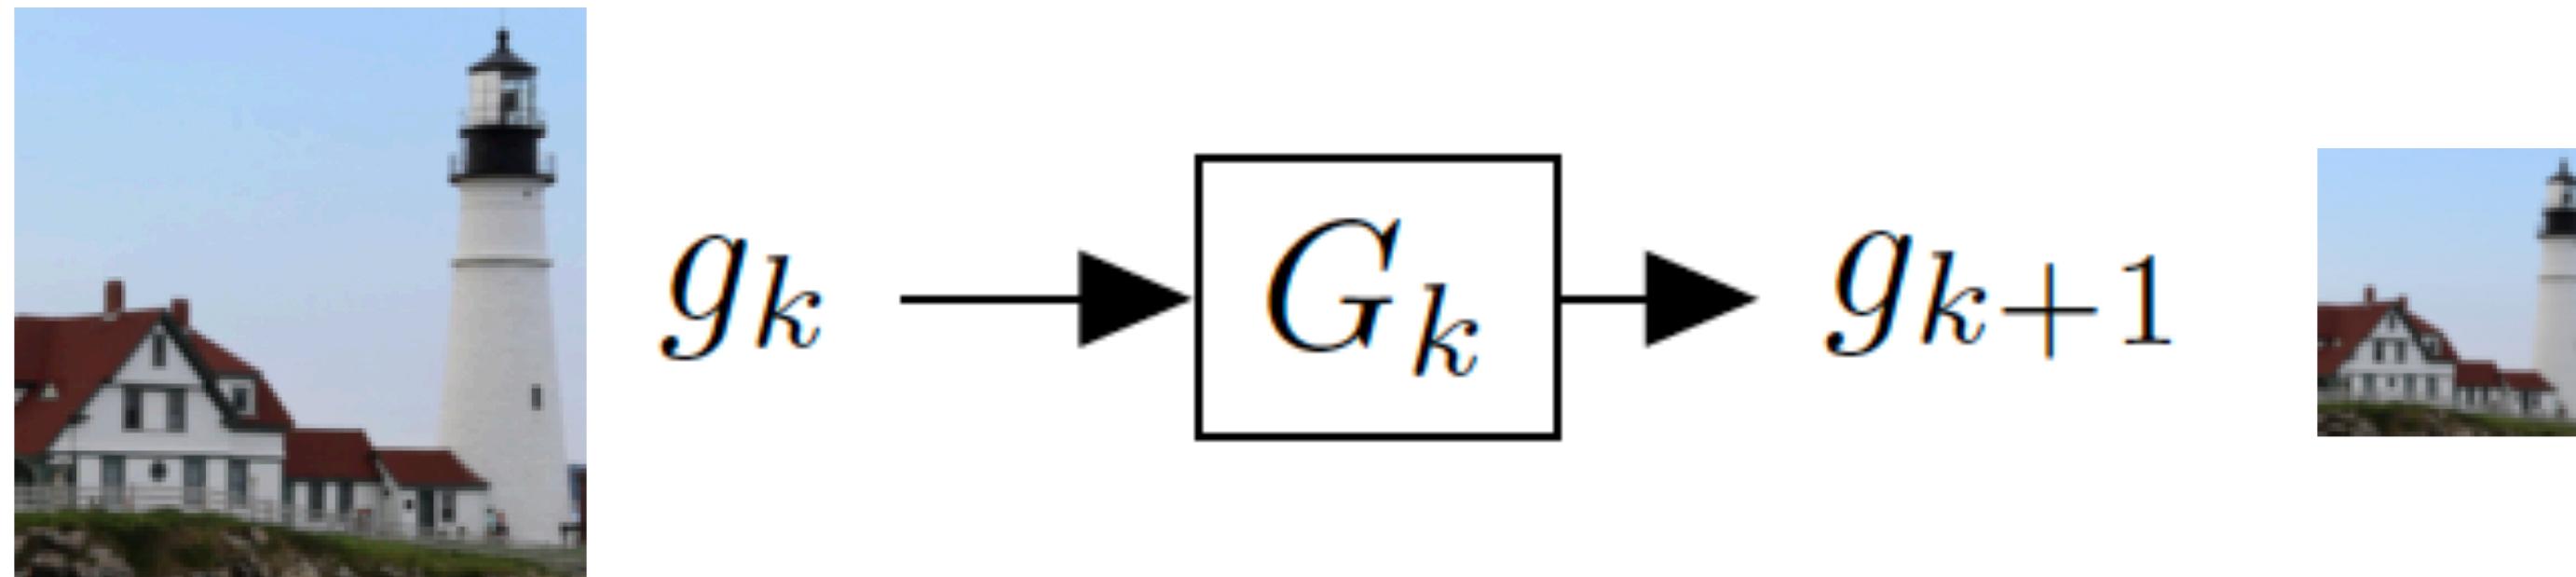
$$g_1 = G_0 g_0$$

$$G_0 = \frac{1}{16} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \end{bmatrix}$$

# The Gaussian pyramid



# The Gaussian pyramid



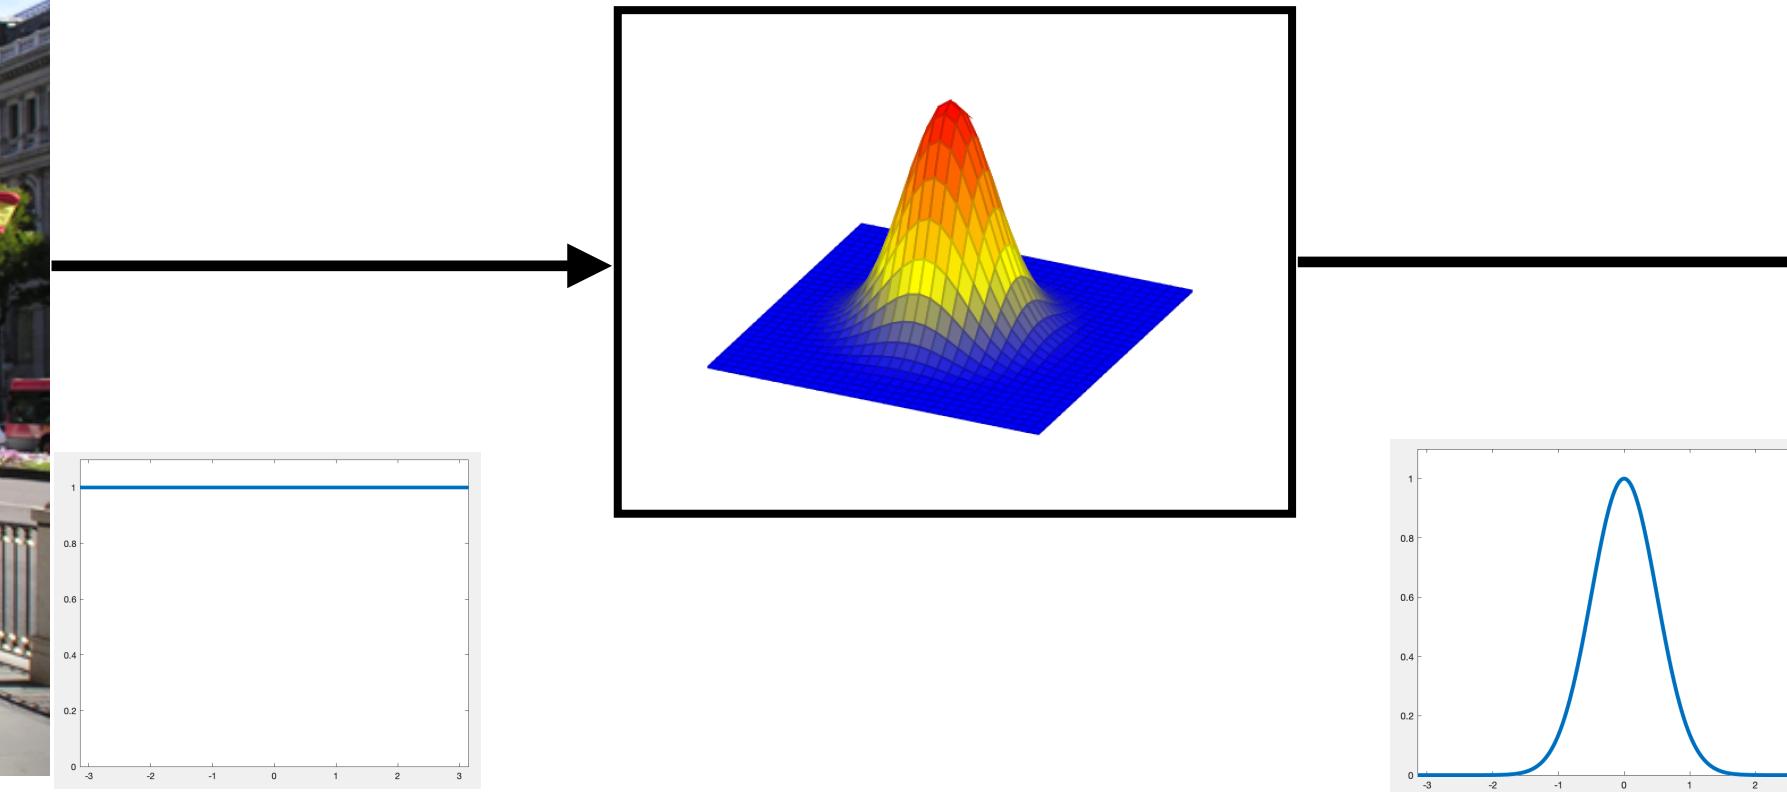
For each level

1. Blur input image with a Gaussian filter
2. Downsample image

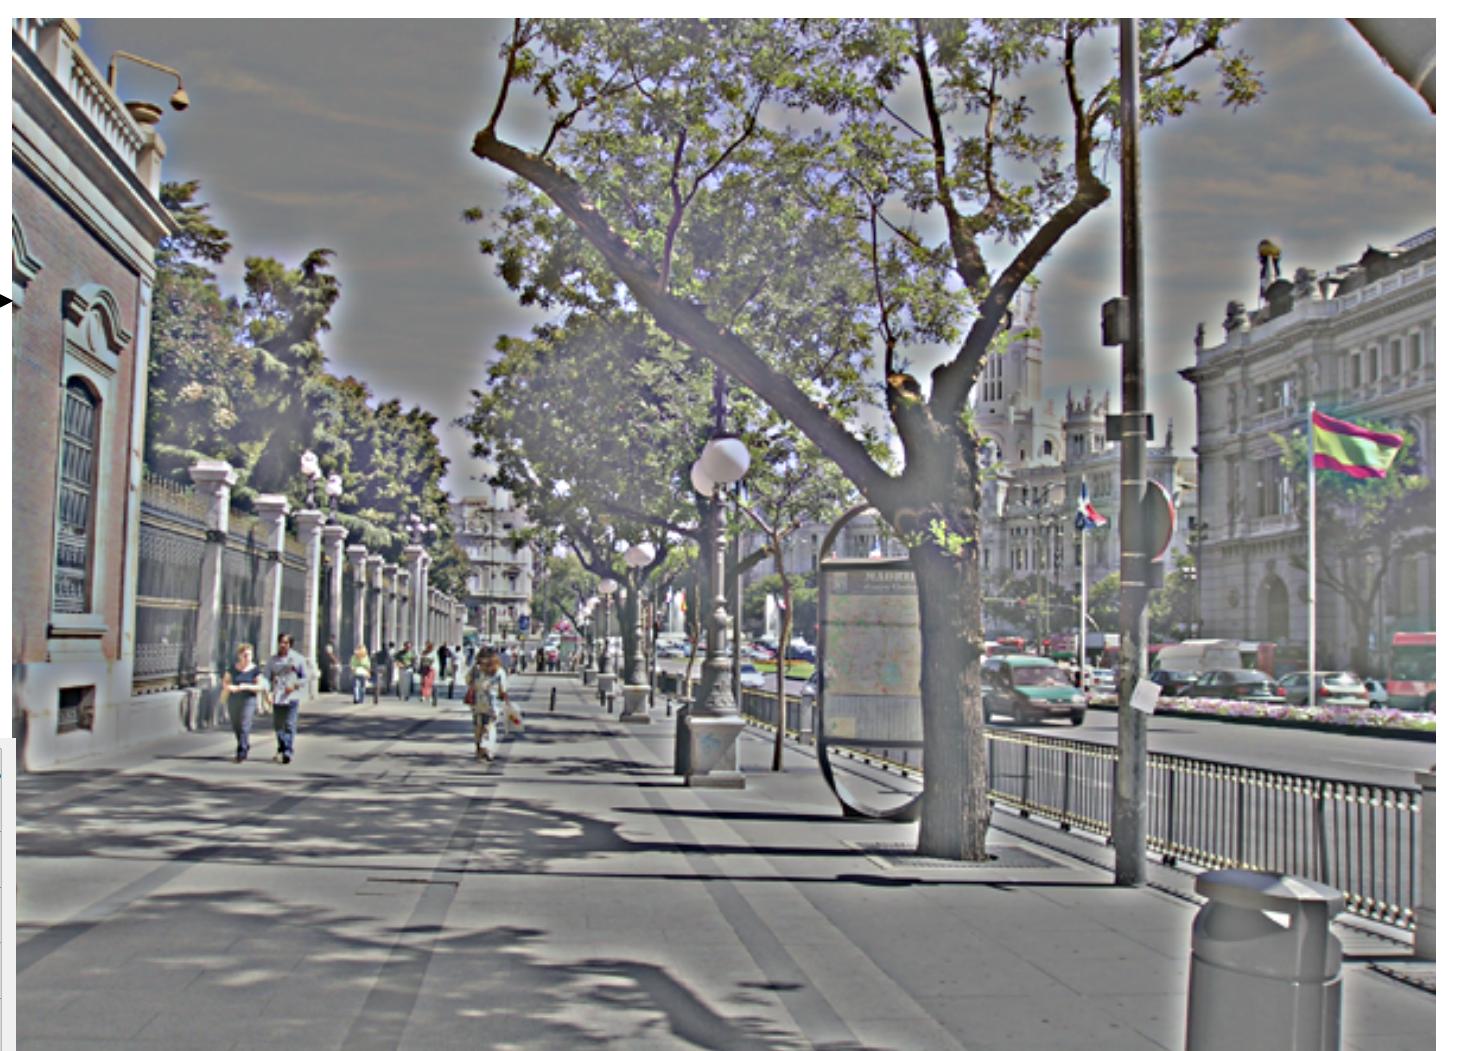
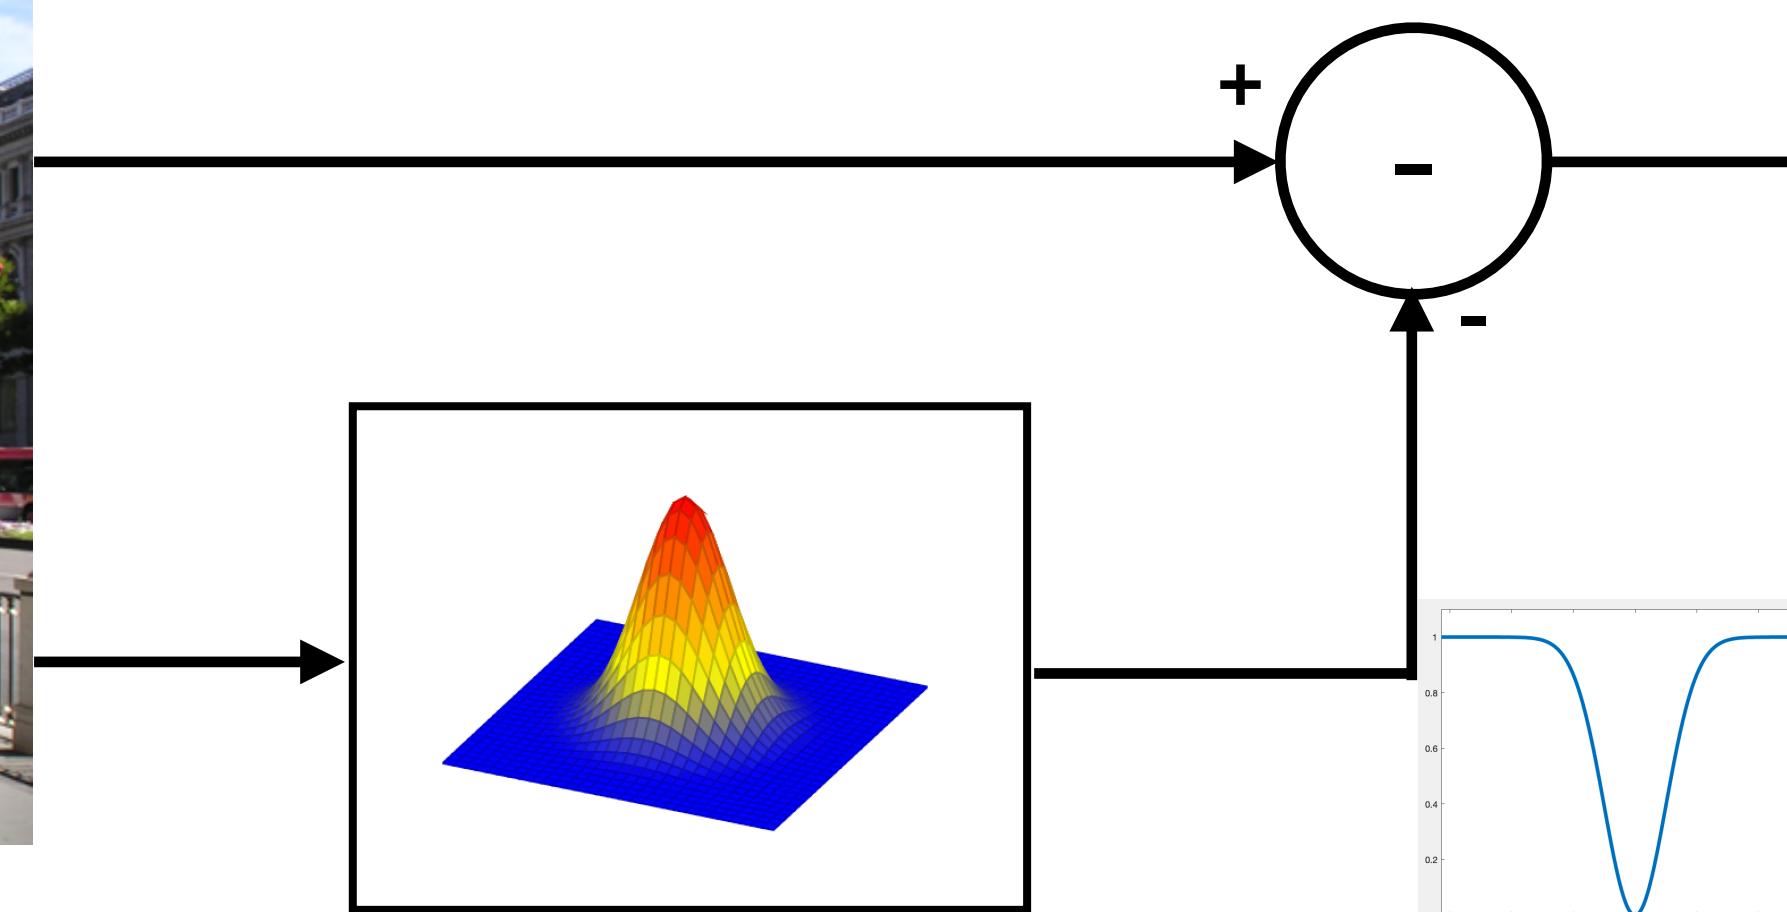
# What about the opposite of blurring?



Gaussian filter

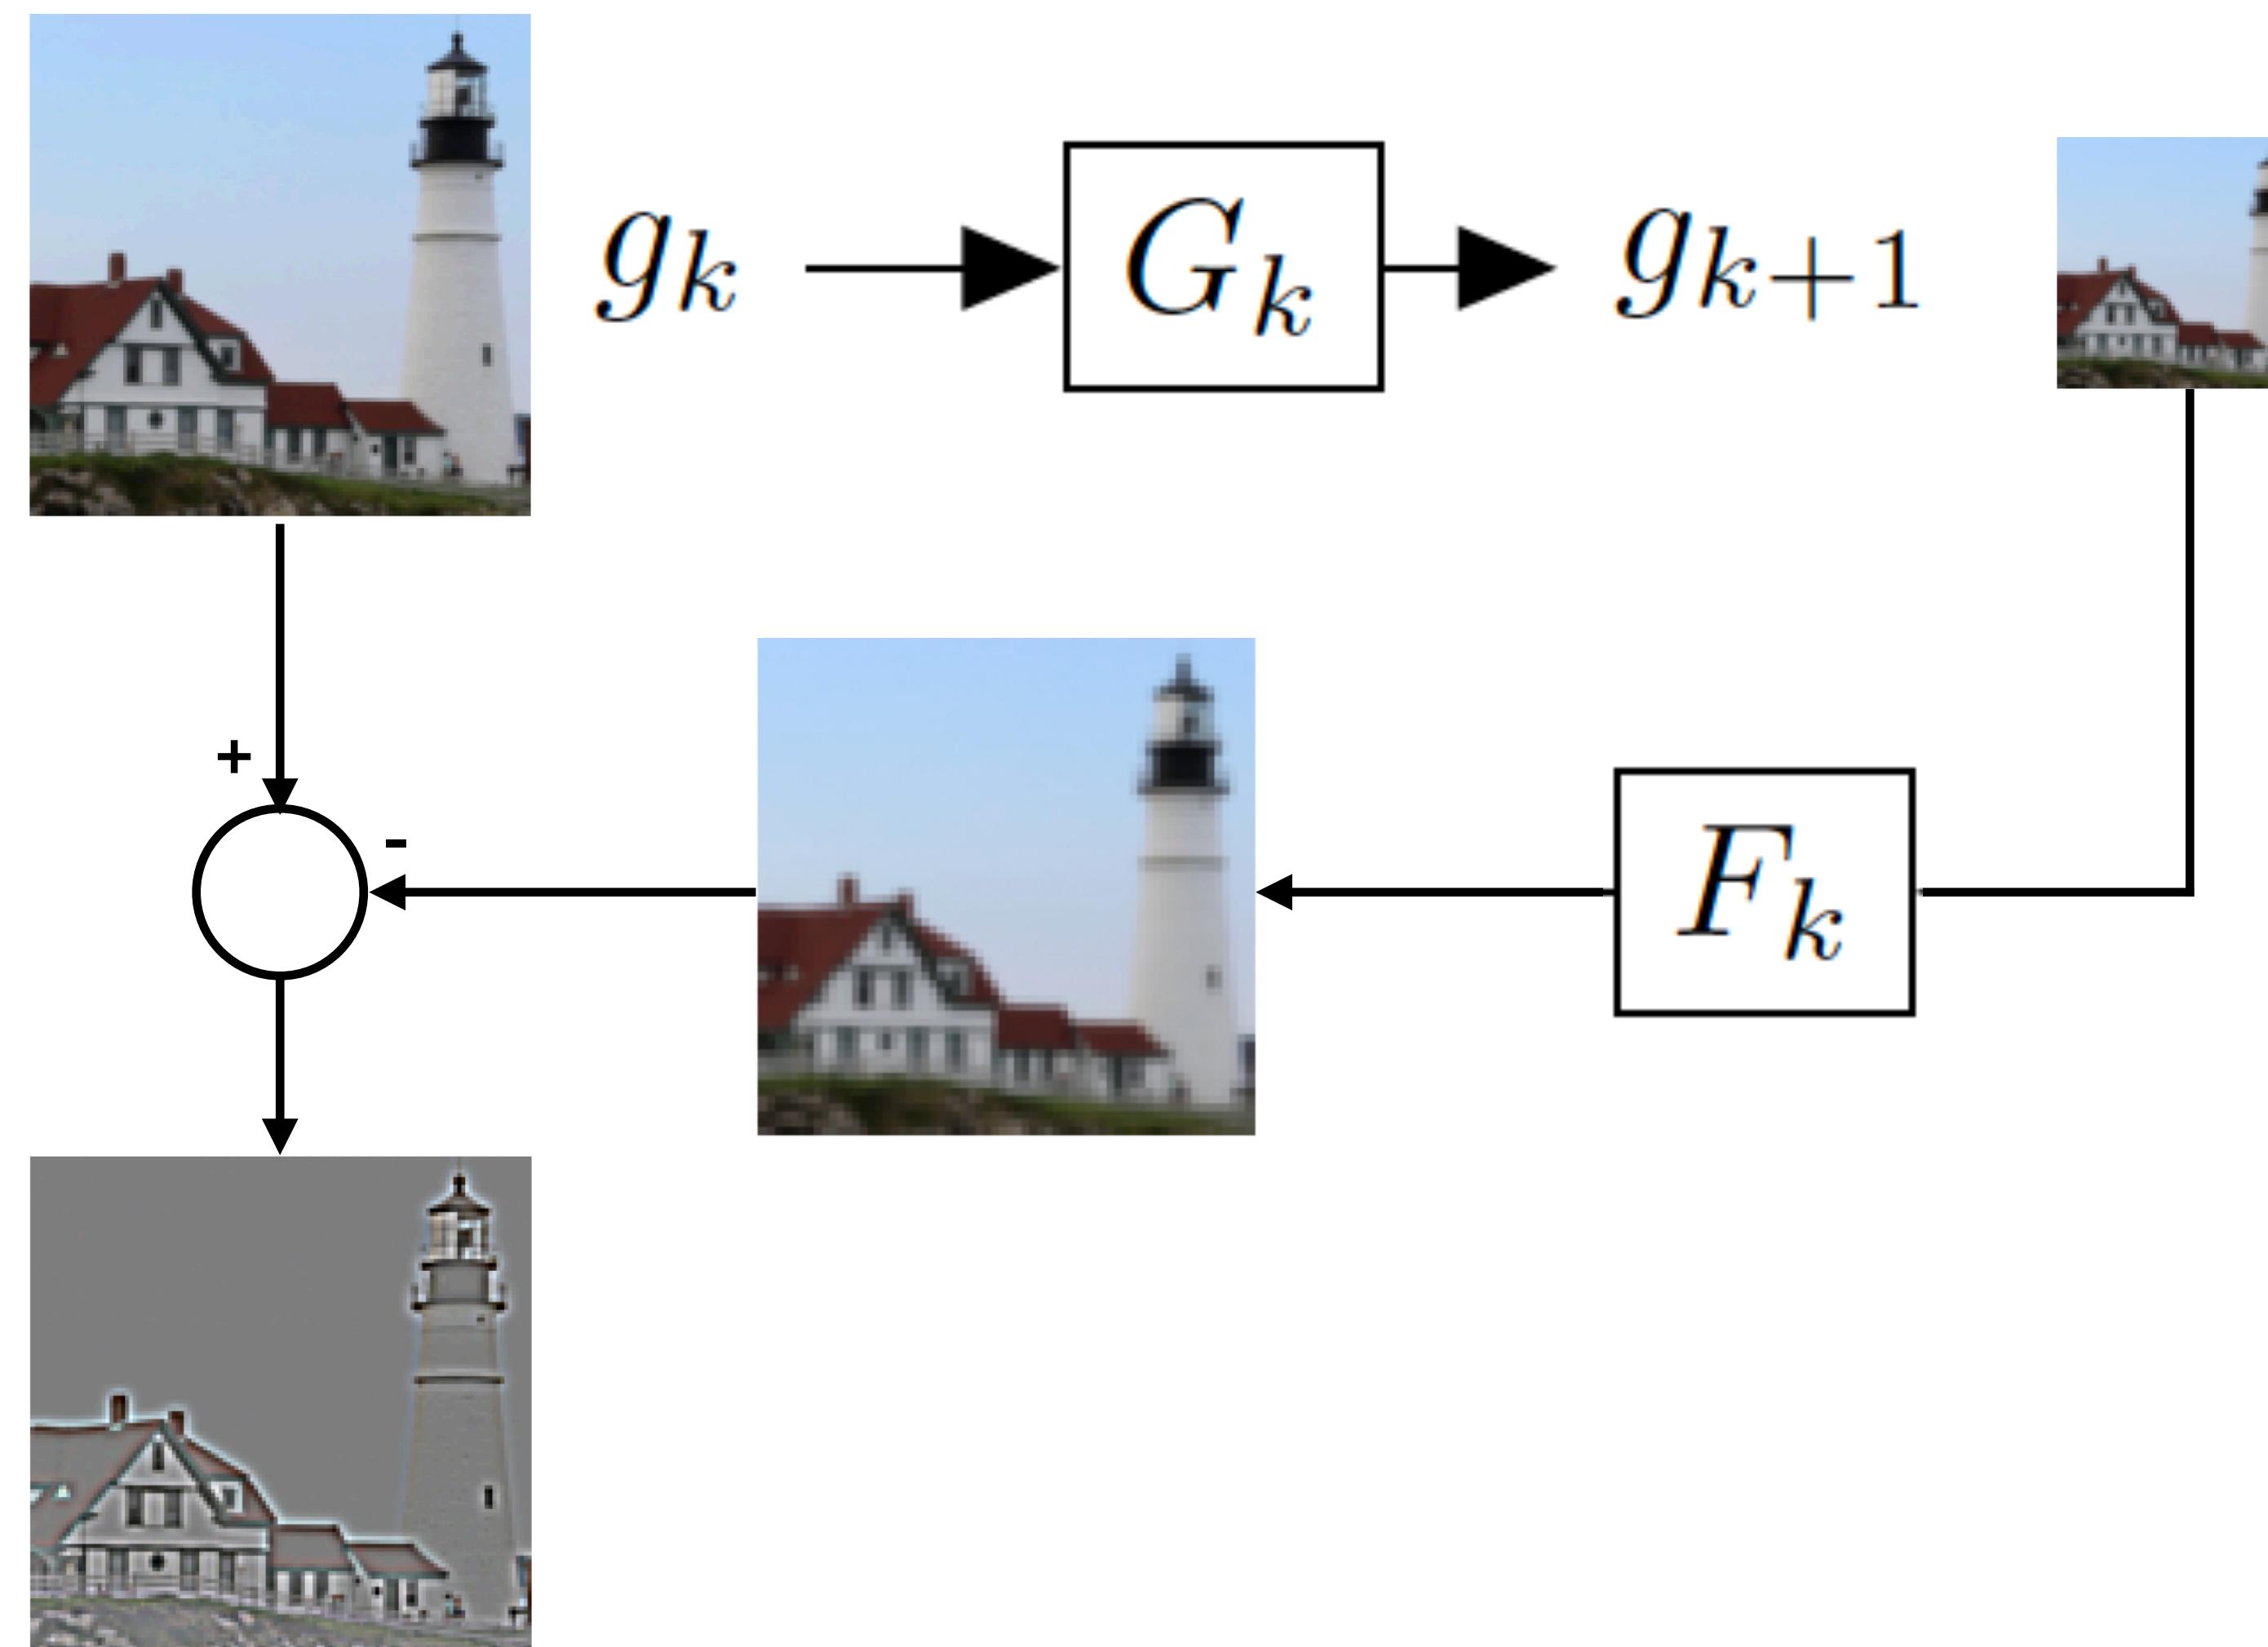


Laplacian filter

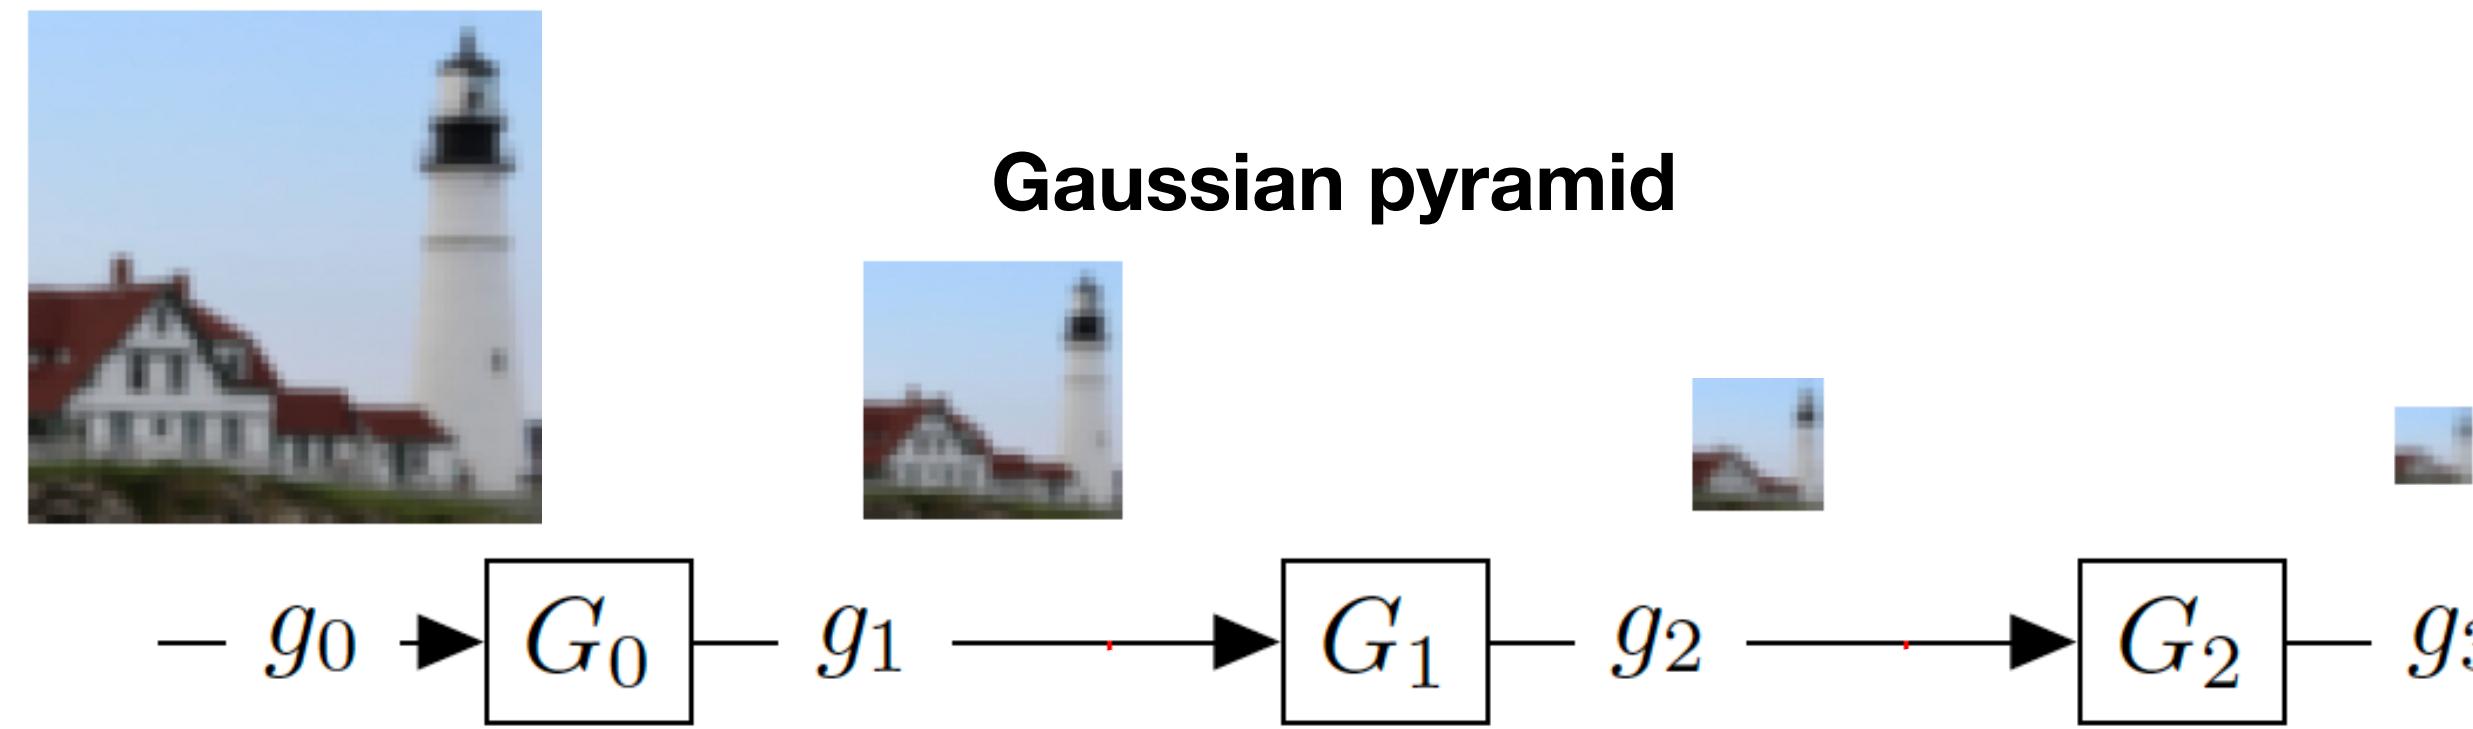


# The Laplacian Pyramid

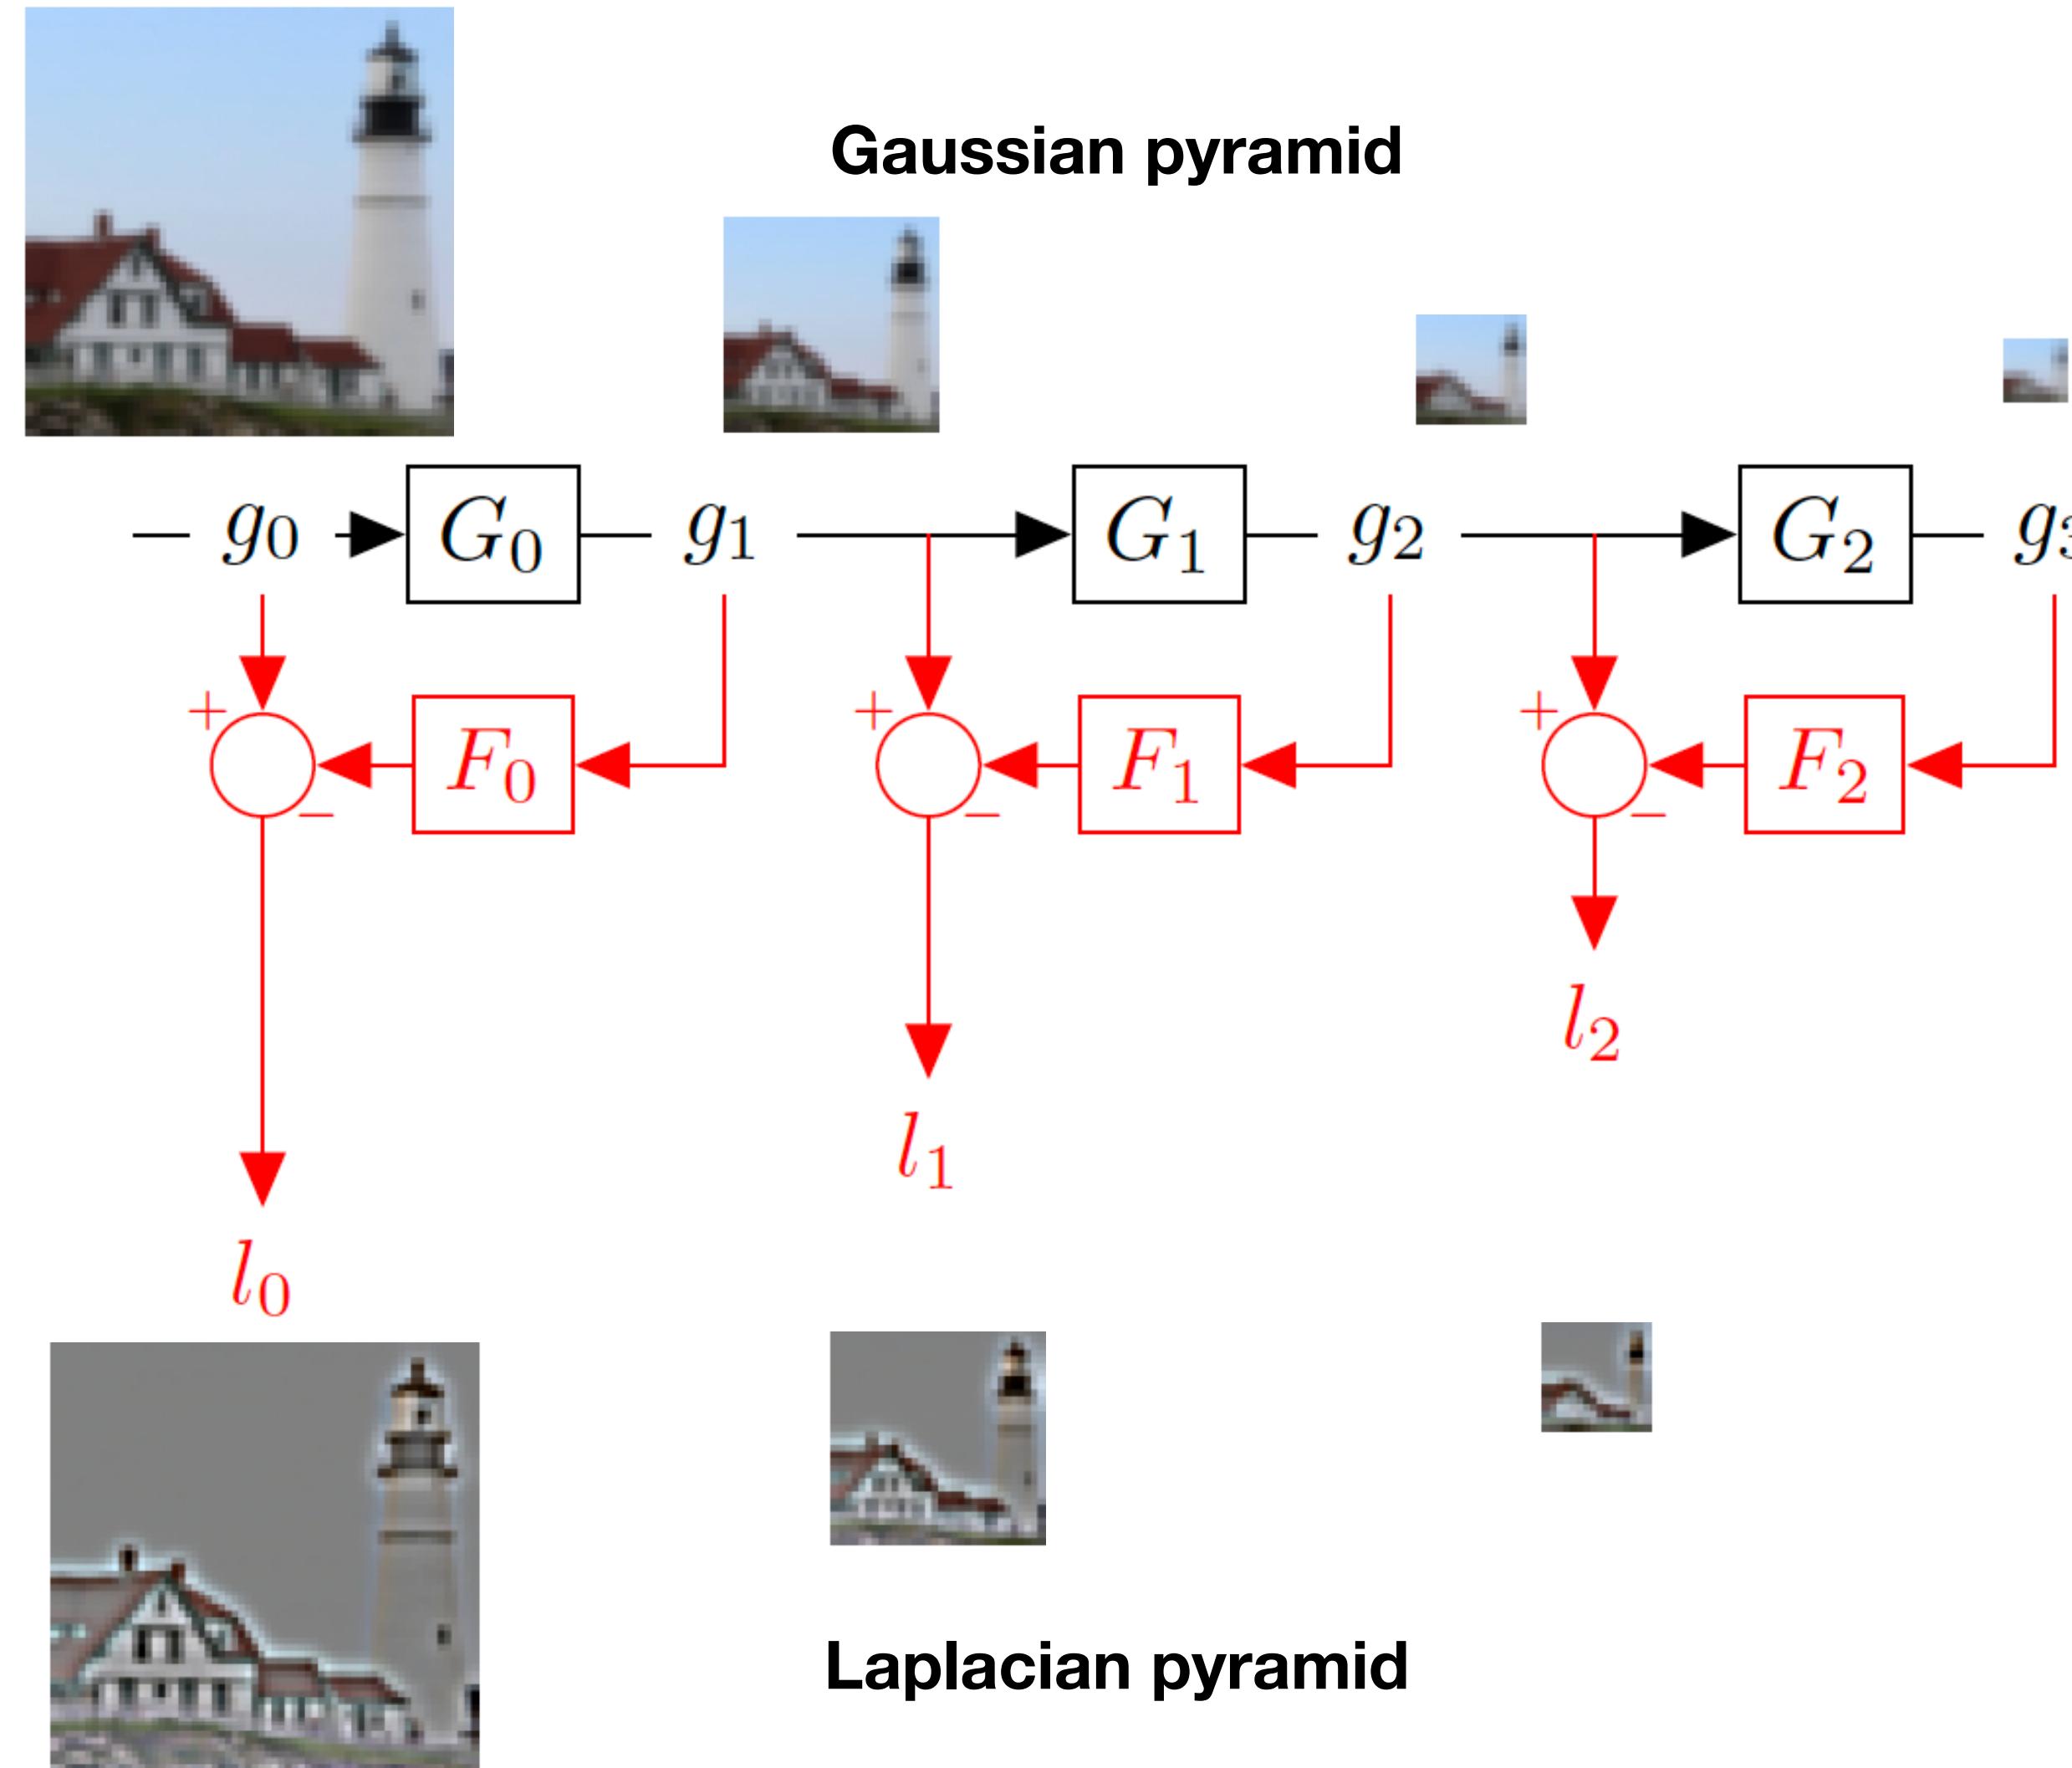
Compute the difference between upsampled Gaussian pyramid level  $k+1$  and Gaussian pyramid level  $k$ .



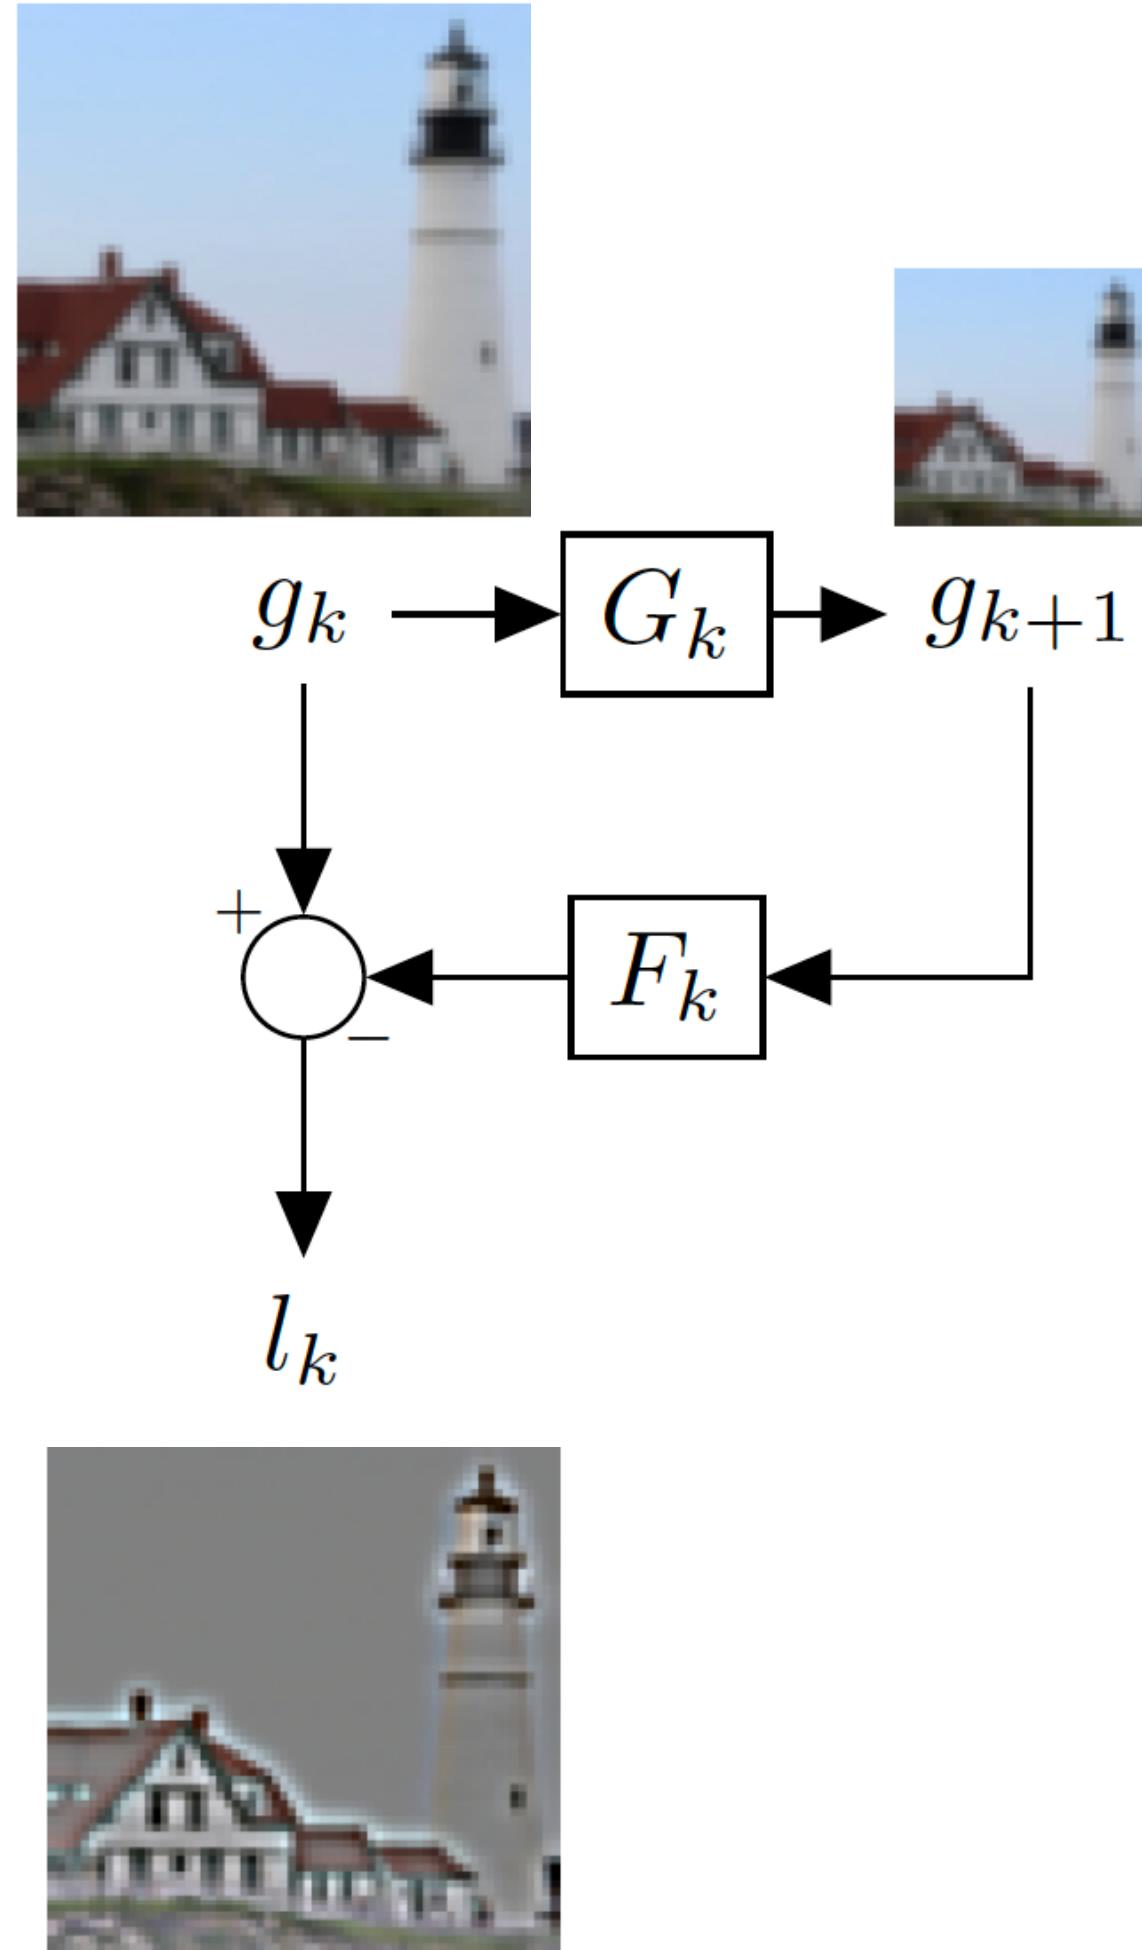
# The Laplacian Pyramid



# The Laplacian Pyramid



# The Laplacian Pyramid



**Blurring and downsampling:**

$$G_0 = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix} \frac{1}{16}$$

(Downsampling by 2)

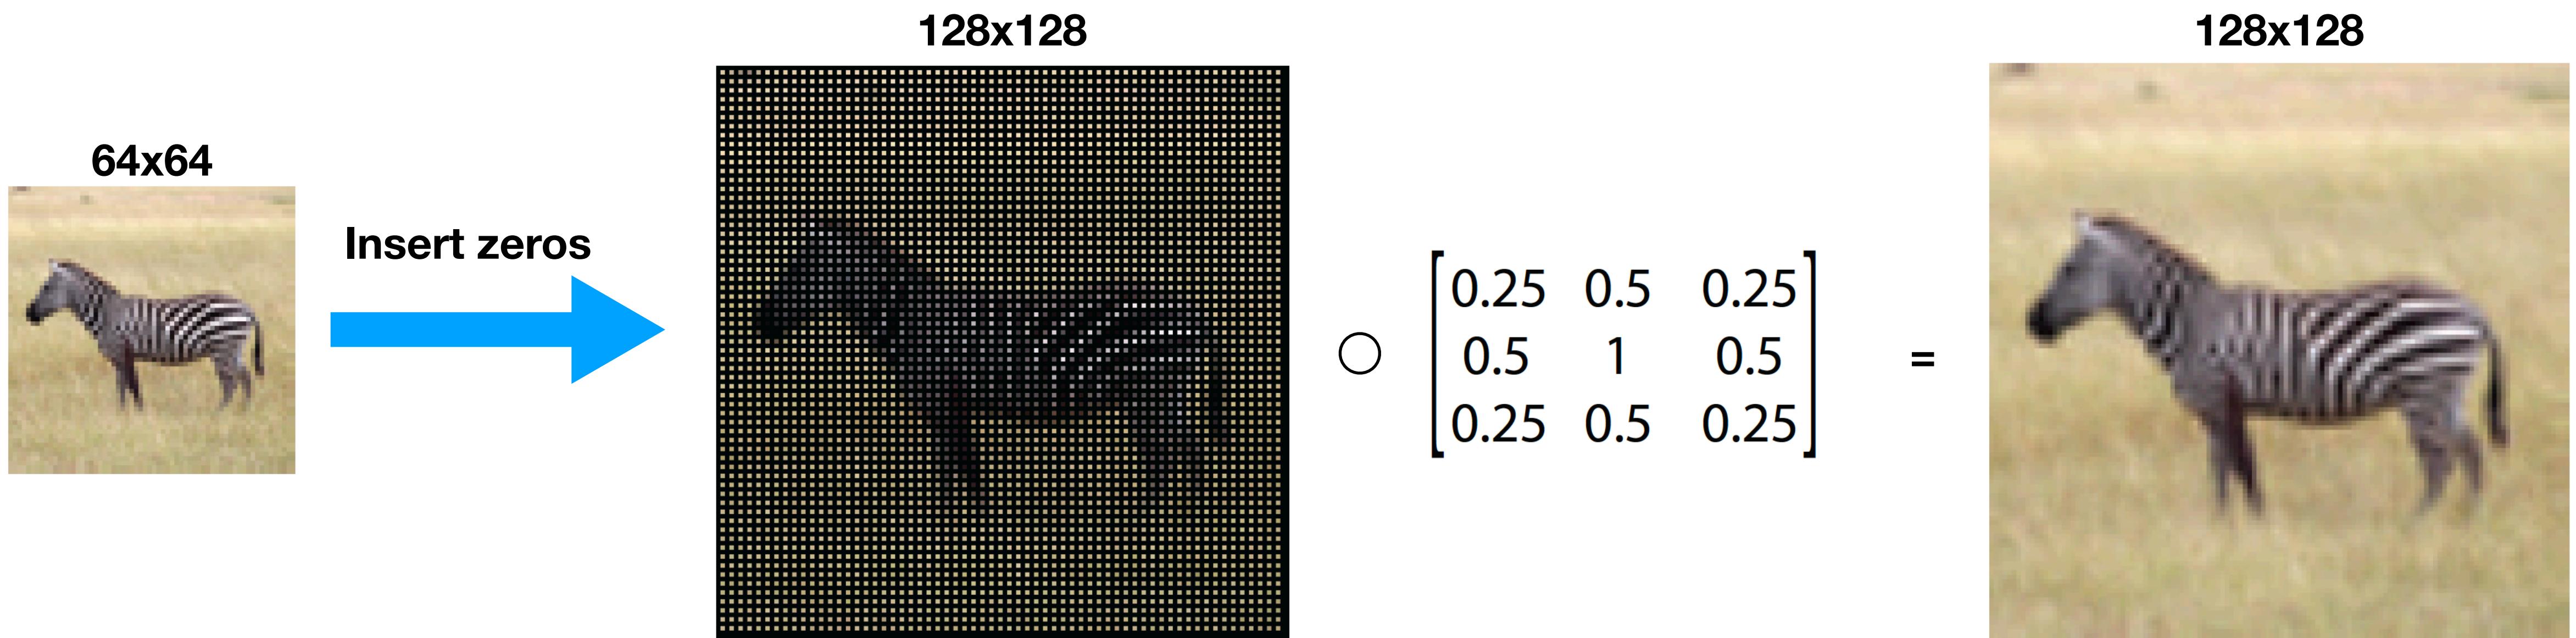
$$\begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

(blur)

**Upsampling and blurring:**

$$F_0 =$$

# Upsampling



# Upsampling

1	0	1	0	1
0	0	0	0	0
1	0	1	0	1
0	0	0	0	0
1	0	1	0	1

$$\circ \begin{bmatrix} 0.25 & 0.5 & 0.25 \\ 0.5 & 1 & 0.5 \\ 0.25 & 0.5 & 0.25 \end{bmatrix} = ?$$

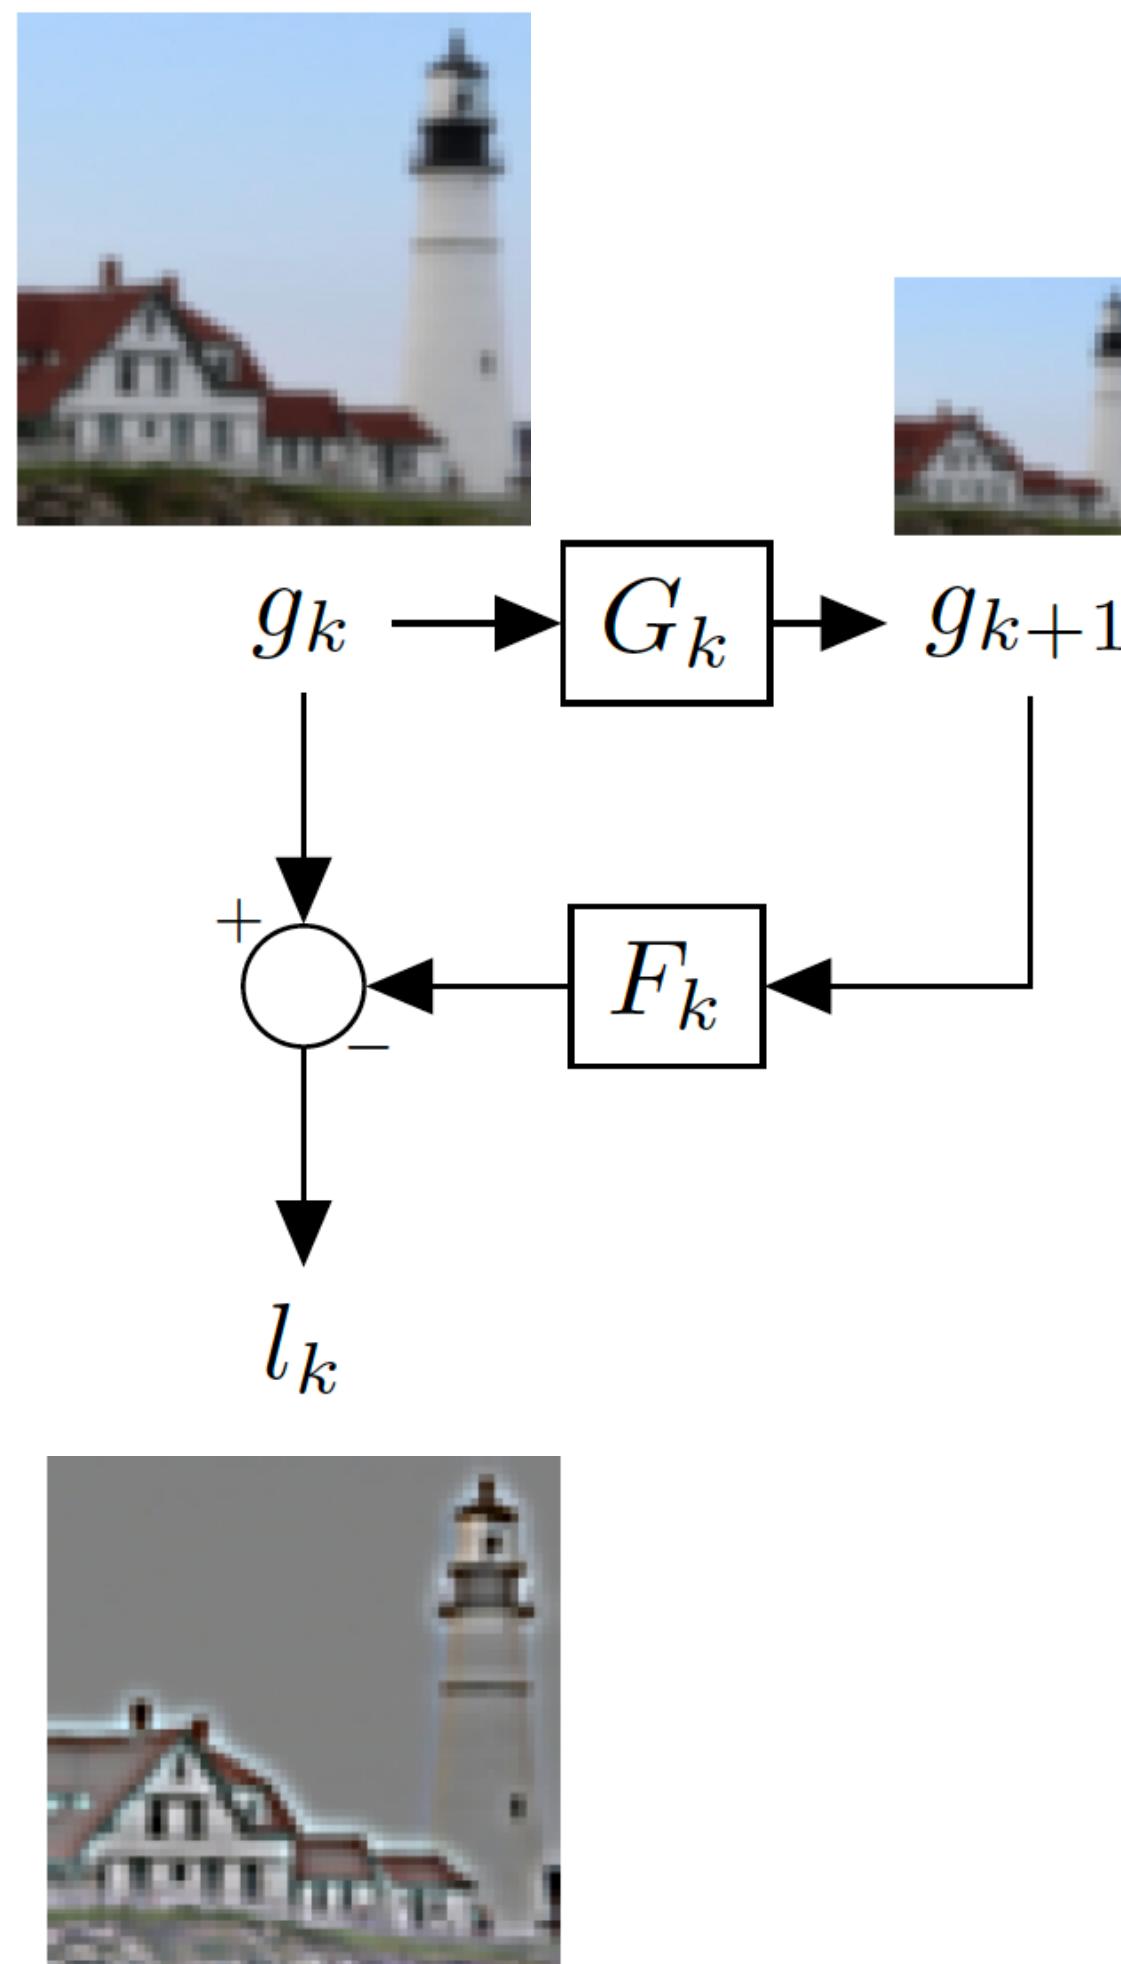
# Upsampling

1	0	1	0	1
0	0	0	0	0
1	0	1	0	1
0	0	0	0	0
1	0	1	0	1

$$\odot \begin{bmatrix} 0.25 & 0.5 & 0.25 \\ 0.5 & 1 & 0.5 \\ 0.25 & 0.5 & 0.25 \end{bmatrix} =$$

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

# The Laplacian Pyramid



# Blurring and downsampling

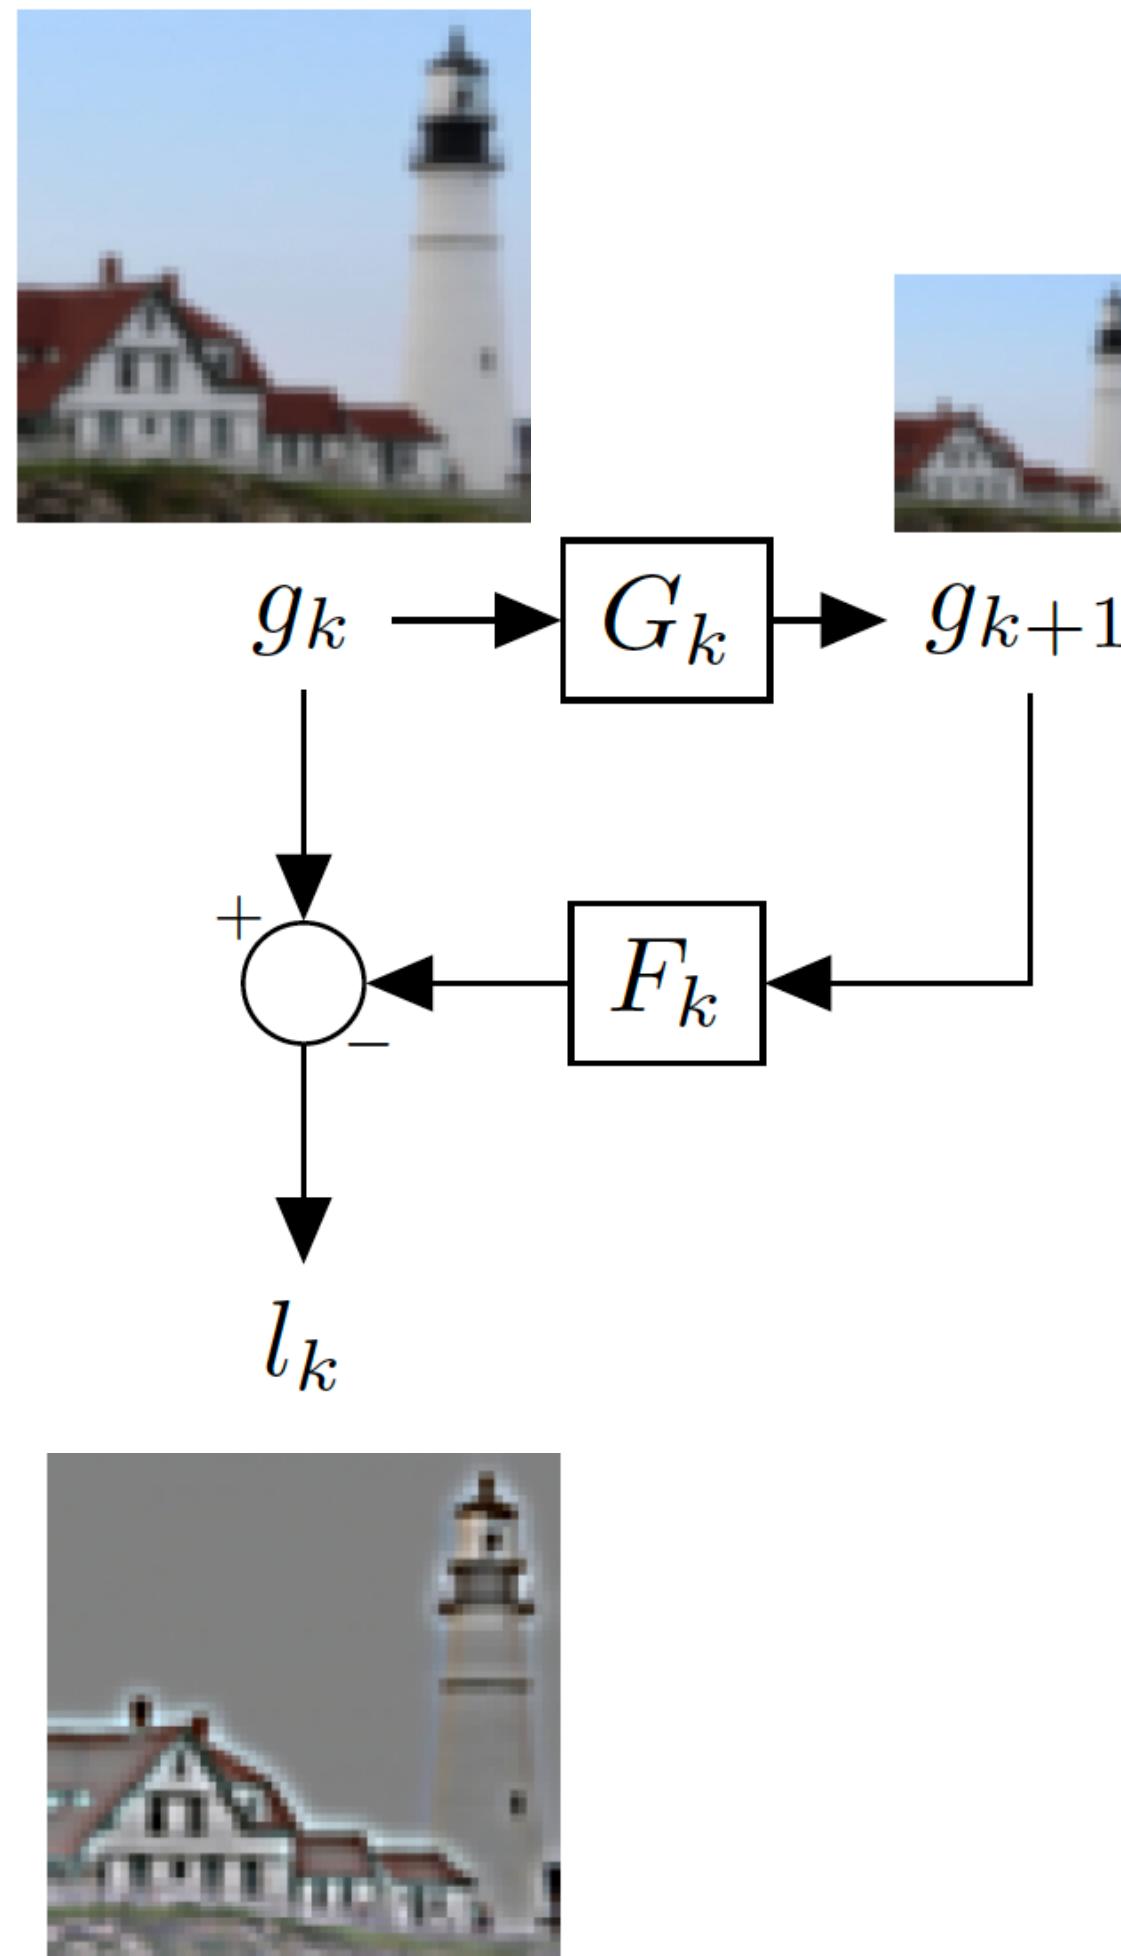
$$G_0 = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

(Downsampling by 2)

## Upsampling and blurring

$$l_0 = (I_0 - F_0 G_0) g_0$$

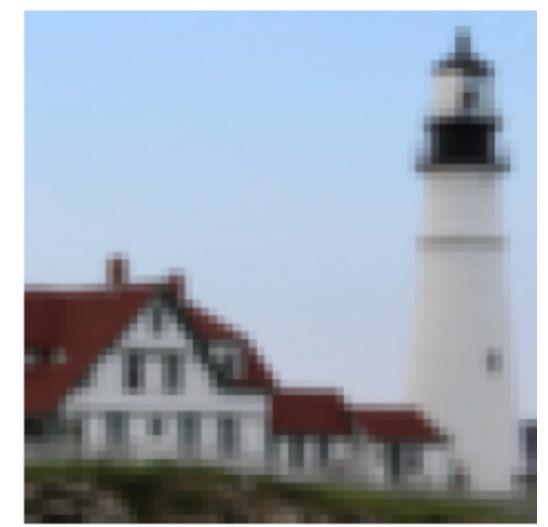
# The Laplacian Pyramid



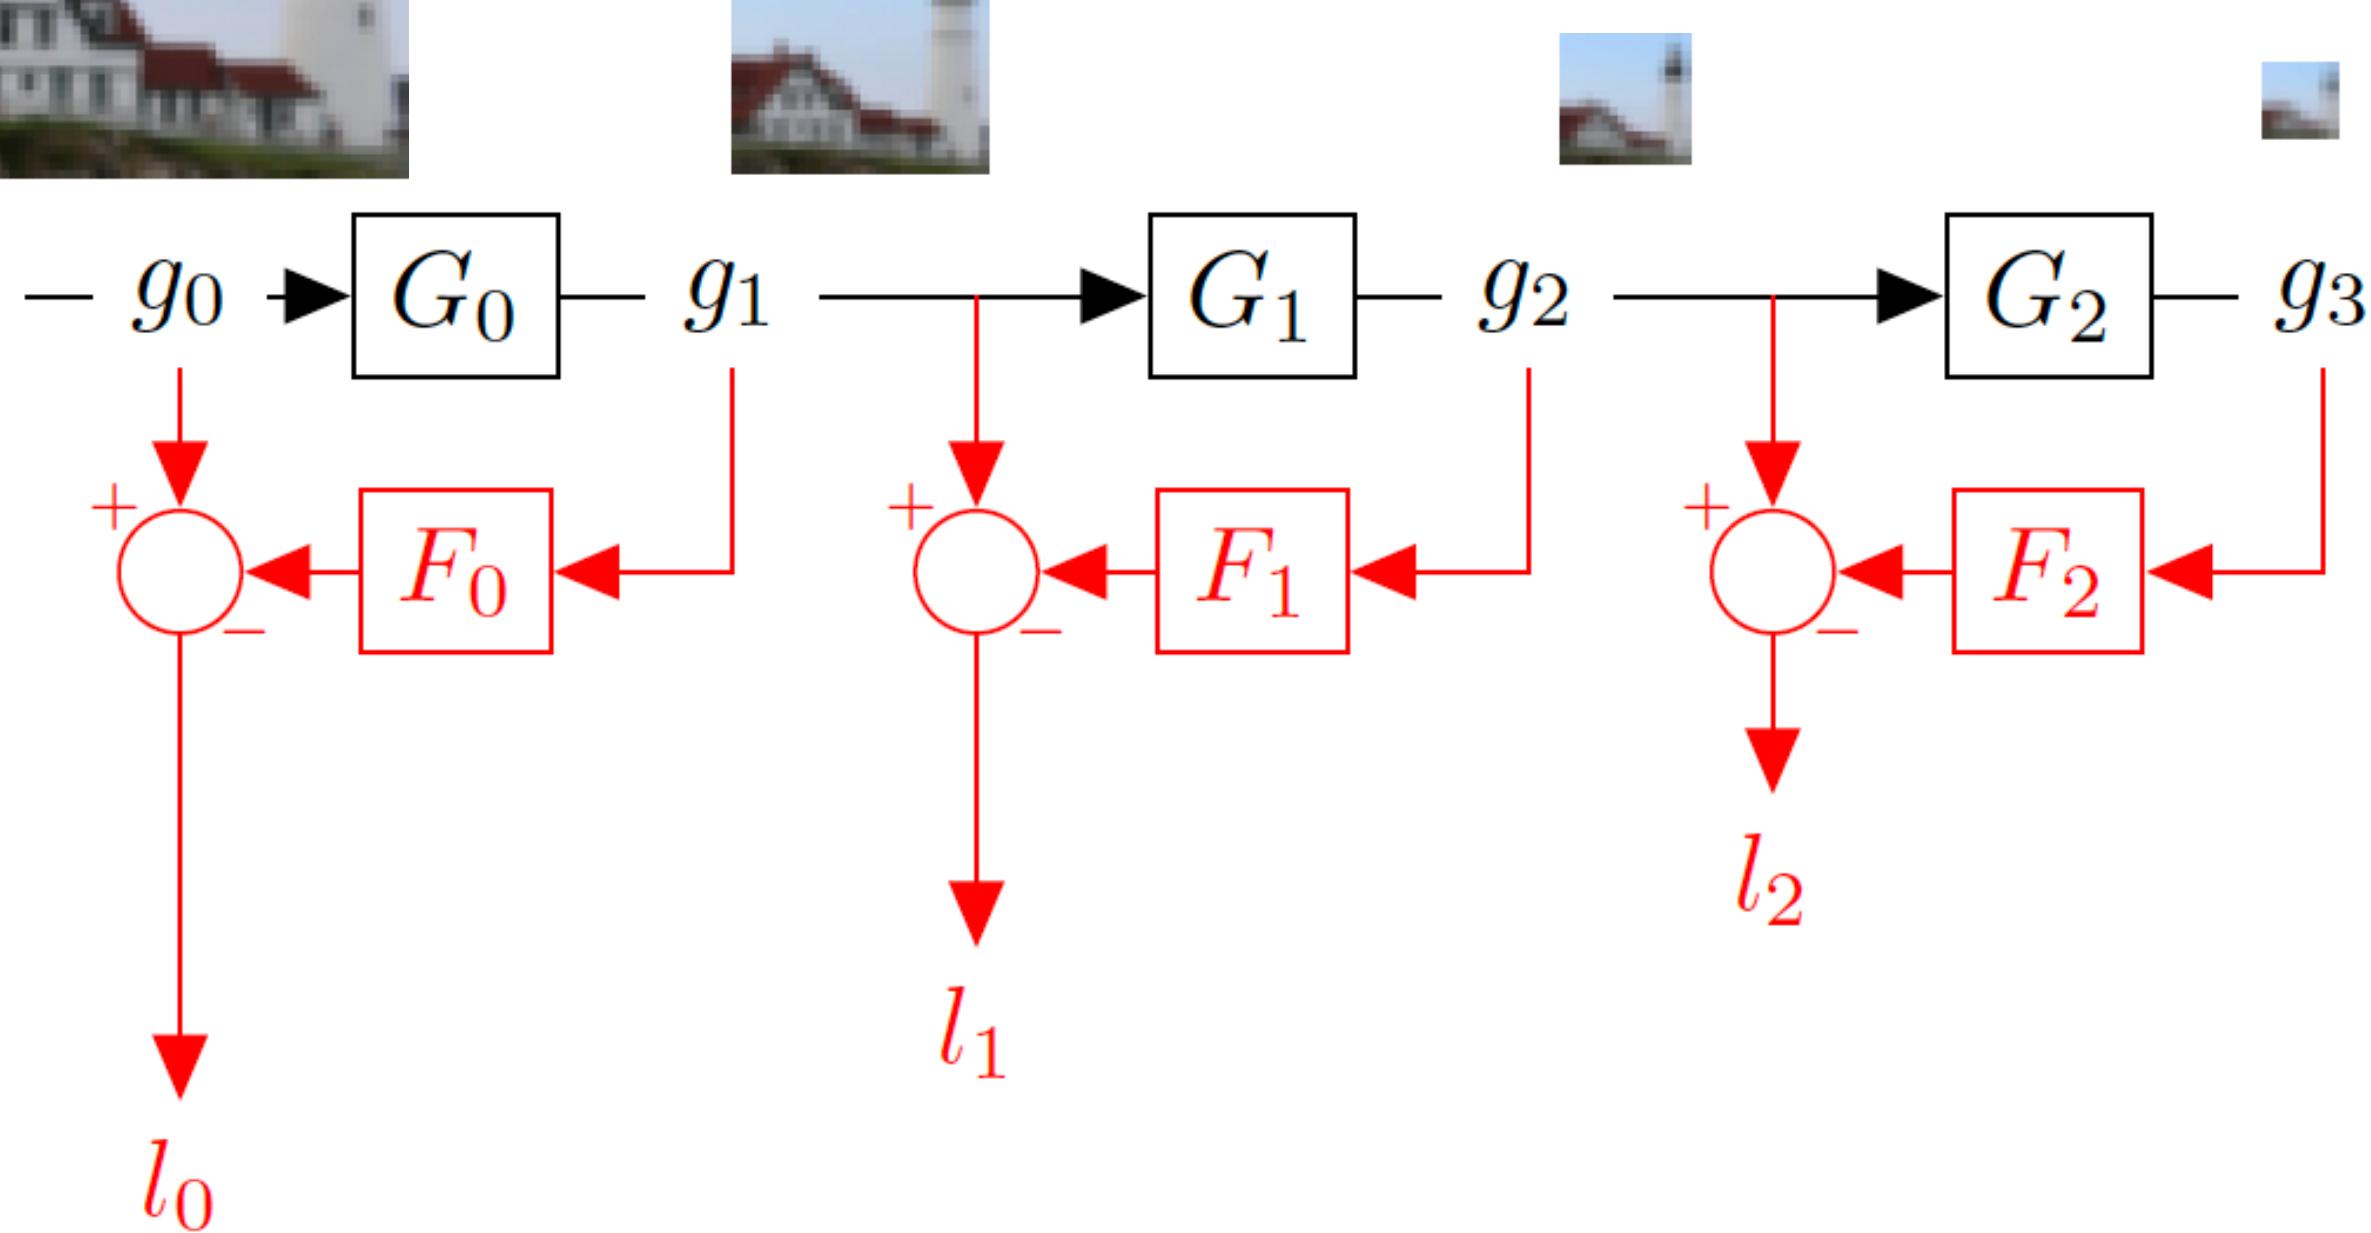
$$l_0 = (I_0 - F_0 G_0) g_0$$

$$= \frac{1}{256} \begin{bmatrix} 182 & -56 & -24 & -8 & -2 & 0 & 0 & 0 \\ -56 & 192 & -56 & -32 & -8 & 0 & 0 & 0 \\ -24 & -56 & 180 & -56 & -24 & -8 & -2 & 0 \\ -8 & -32 & -56 & 192 & -56 & -32 & -8 & 0 \\ -2 & -8 & -24 & -56 & 180 & -56 & -24 & -8 \\ 0 & 0 & -8 & -32 & -56 & 192 & -56 & -32 \\ 0 & 0 & -2 & -8 & -24 & -56 & 182 & -48 \\ 0 & 0 & 0 & 0 & -8 & -32 & -48 & 224 \end{bmatrix} x$$

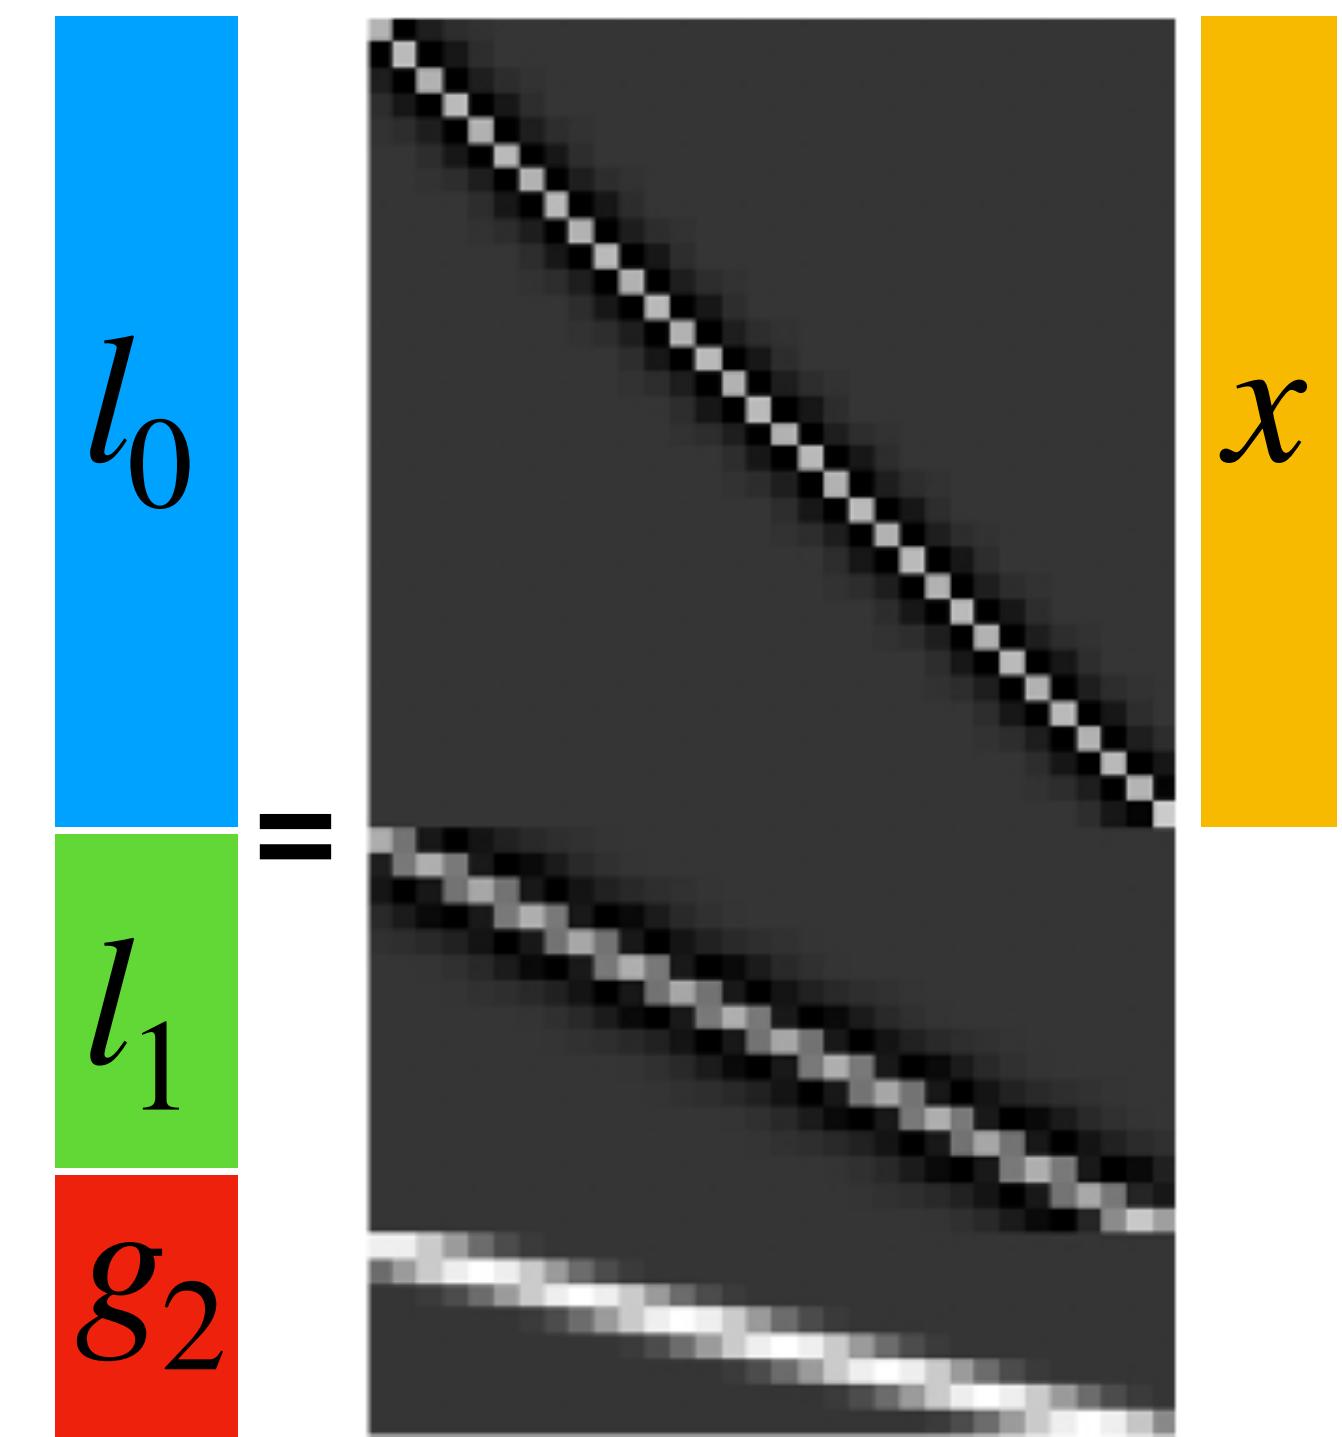
# The Laplacian Pyramid



Gaussian pyramid



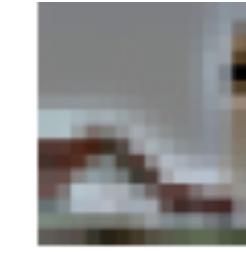
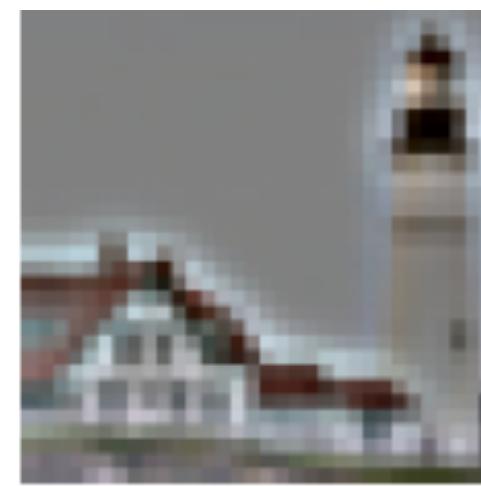
Laplacian pyramid



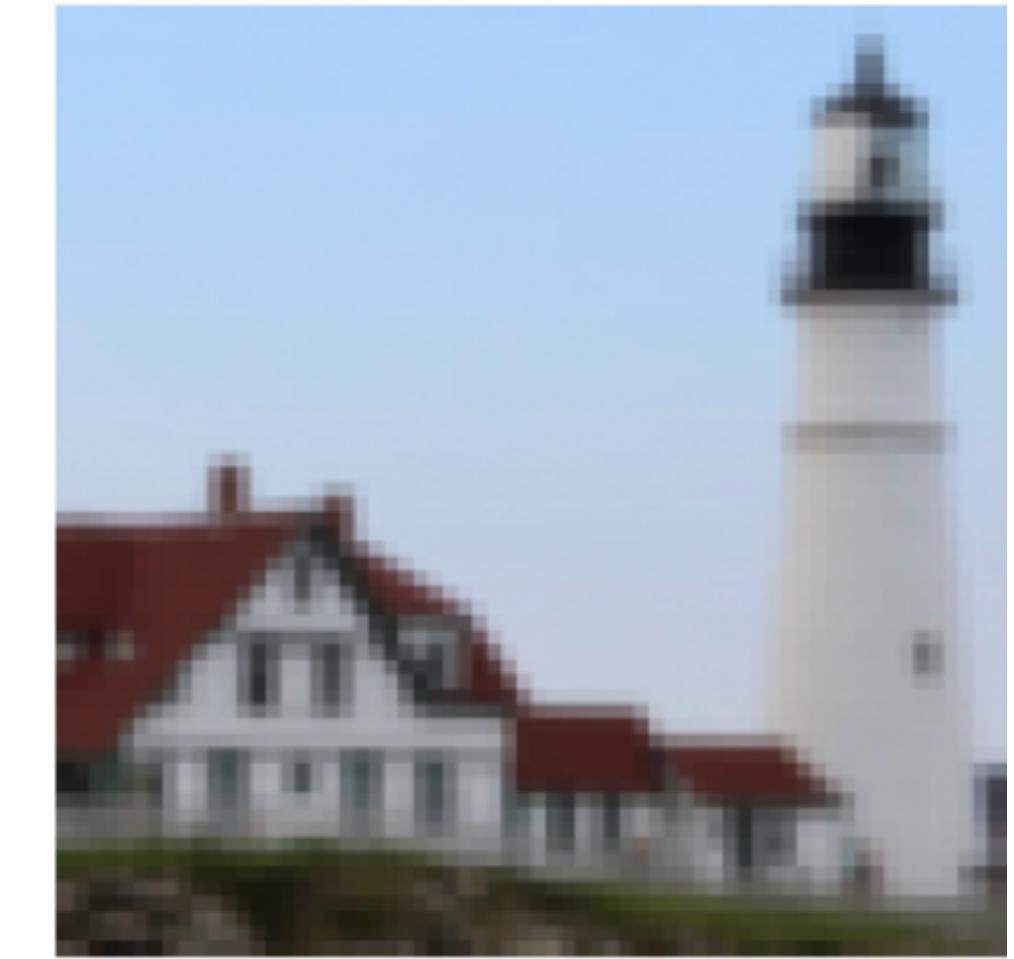
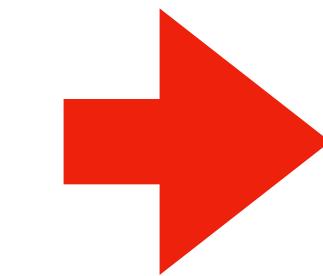
# The Laplacian Pyramid



Laplacian pyramid



Gaussian  
residual

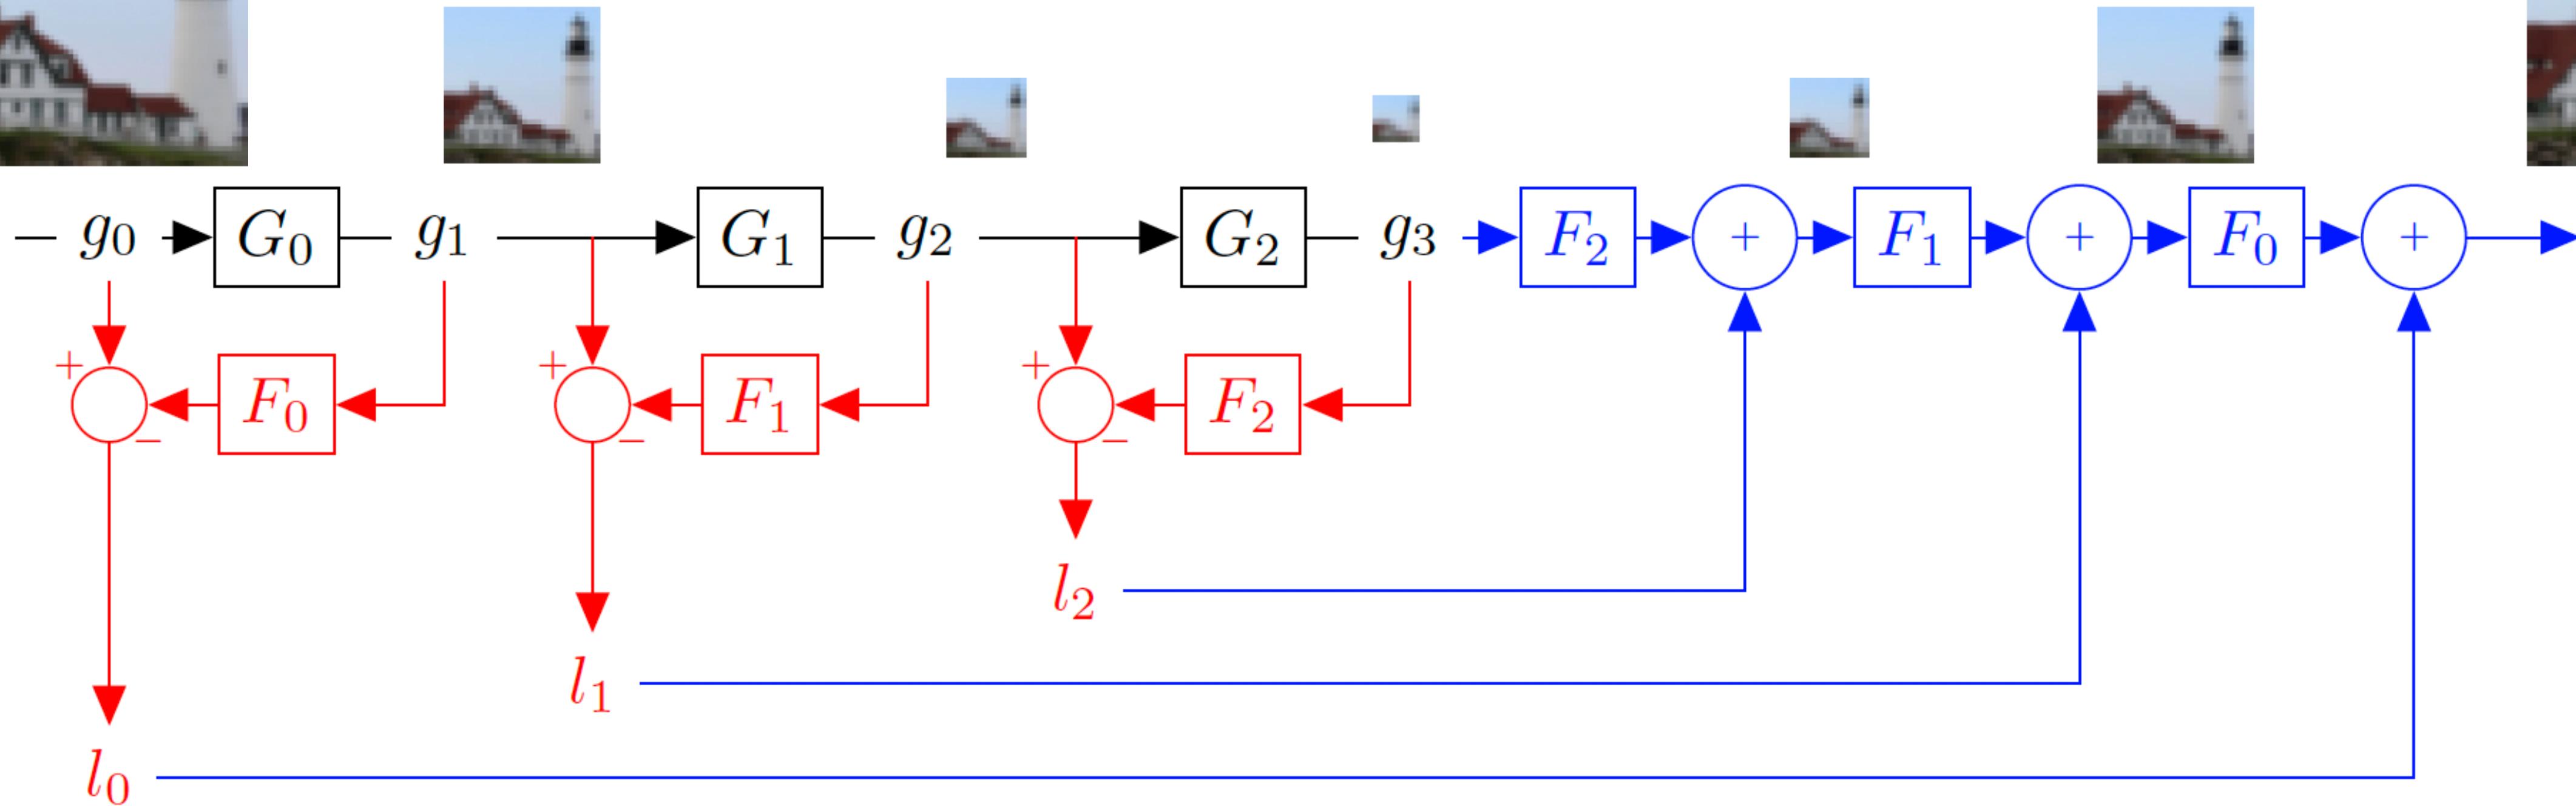


Can we invert the  
Laplacian Pyramid?

# The Laplacian Pyramid

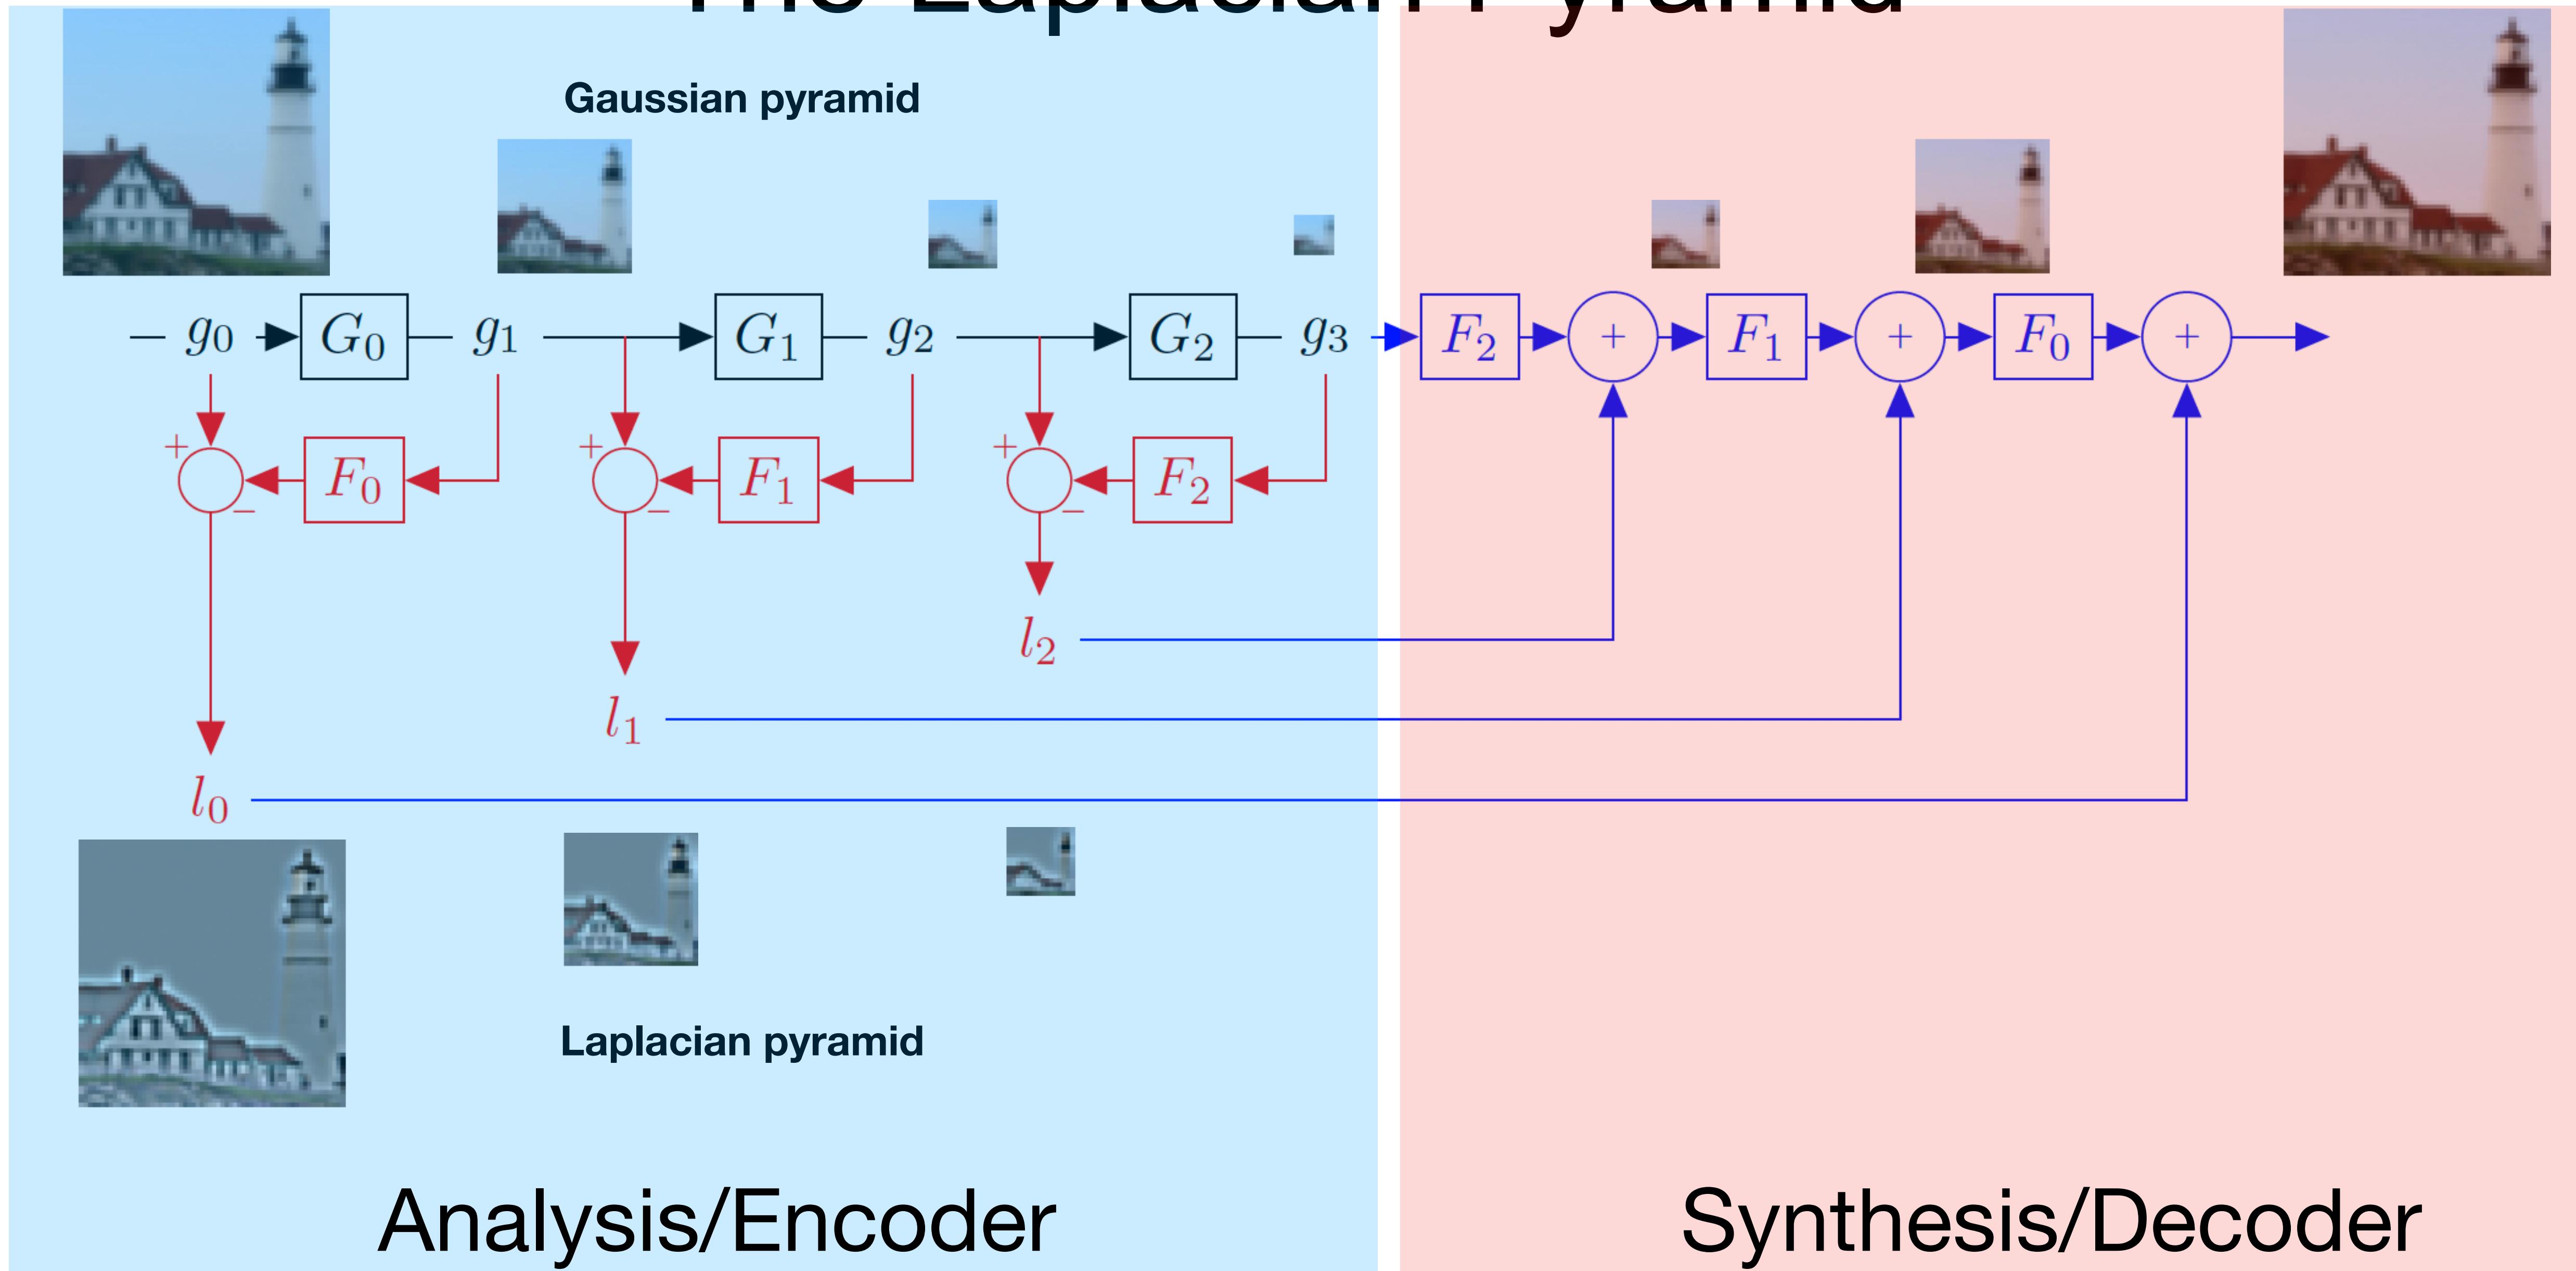


Gaussian pyramid



Laplacian pyramid

# The Laplacian Pyramid



Analysis/Encoder

Synthesis/Decoder

# The Laplacian Pyramid

Laplacian pyramid



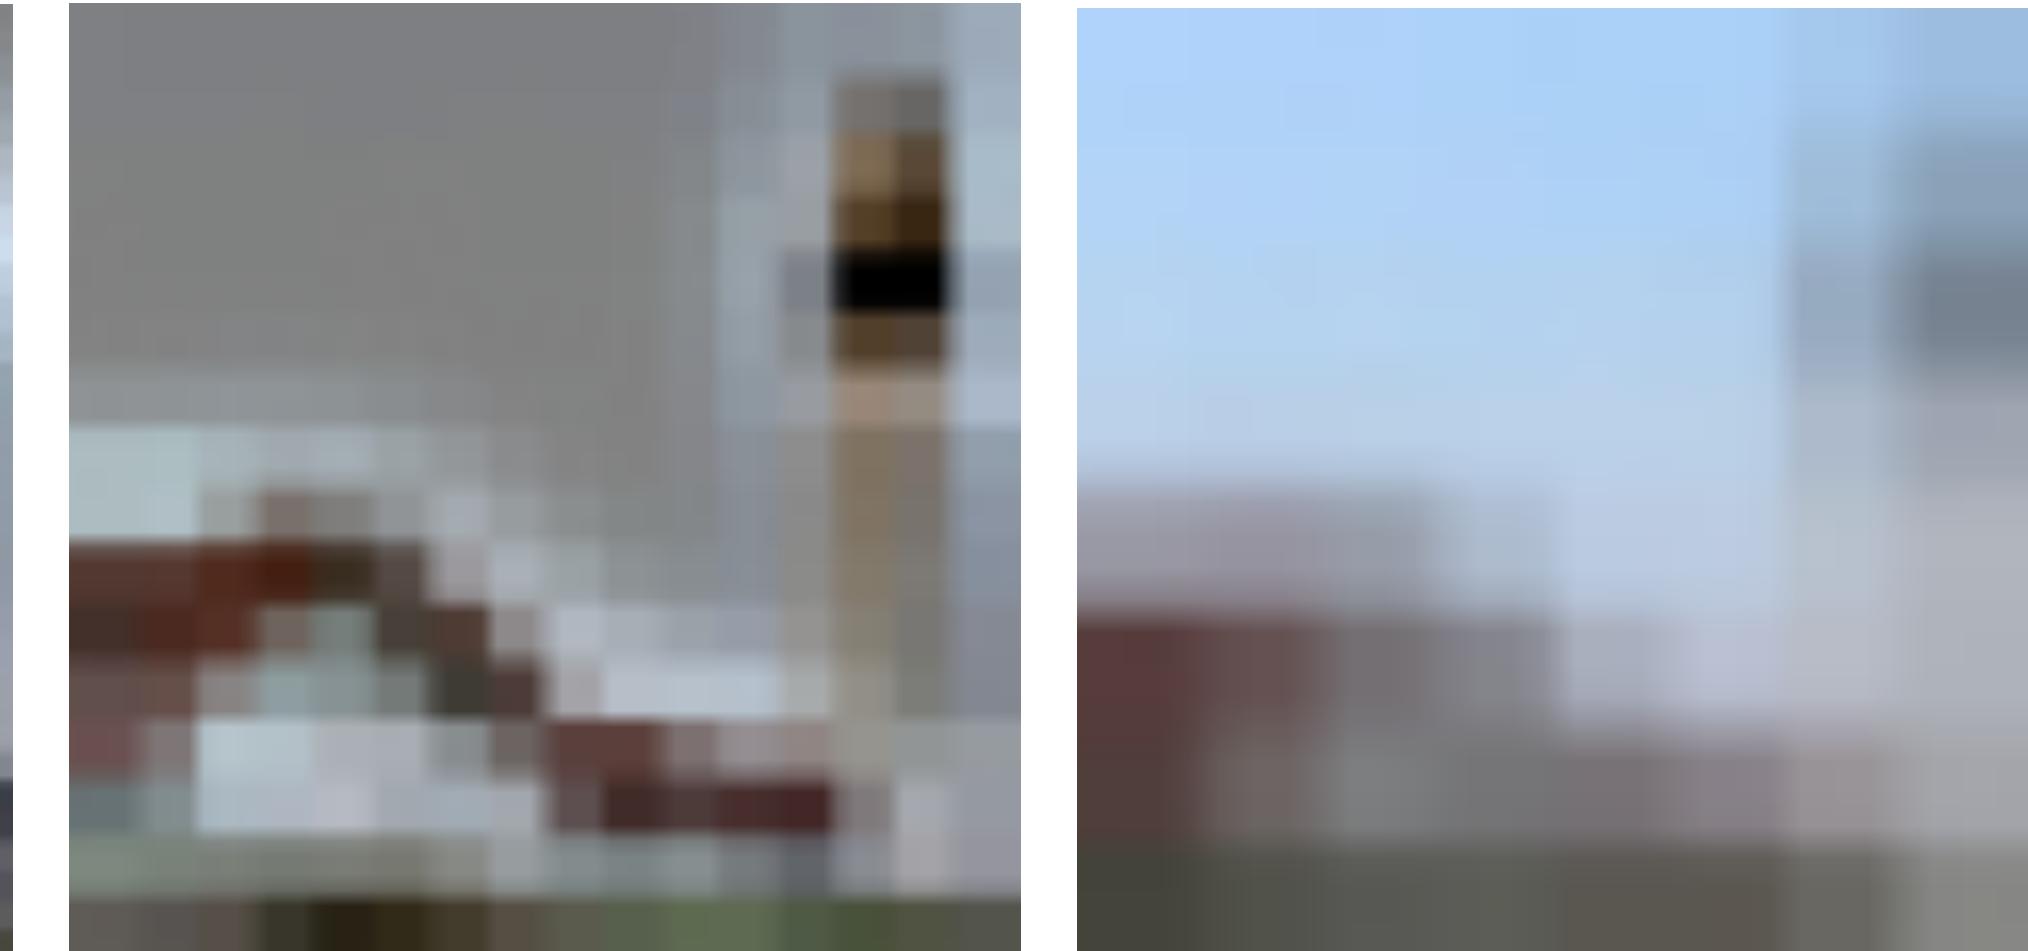
Gaussian residual



Laplacian pyramid



Gaussian residual



# Laplacian pyramid applications

- Image compression
- Noise removal
- Computing image features (e.g., SIFT)
- Connections to common neural network architectures: ResNets, U-Nets, skip connections in transformers...

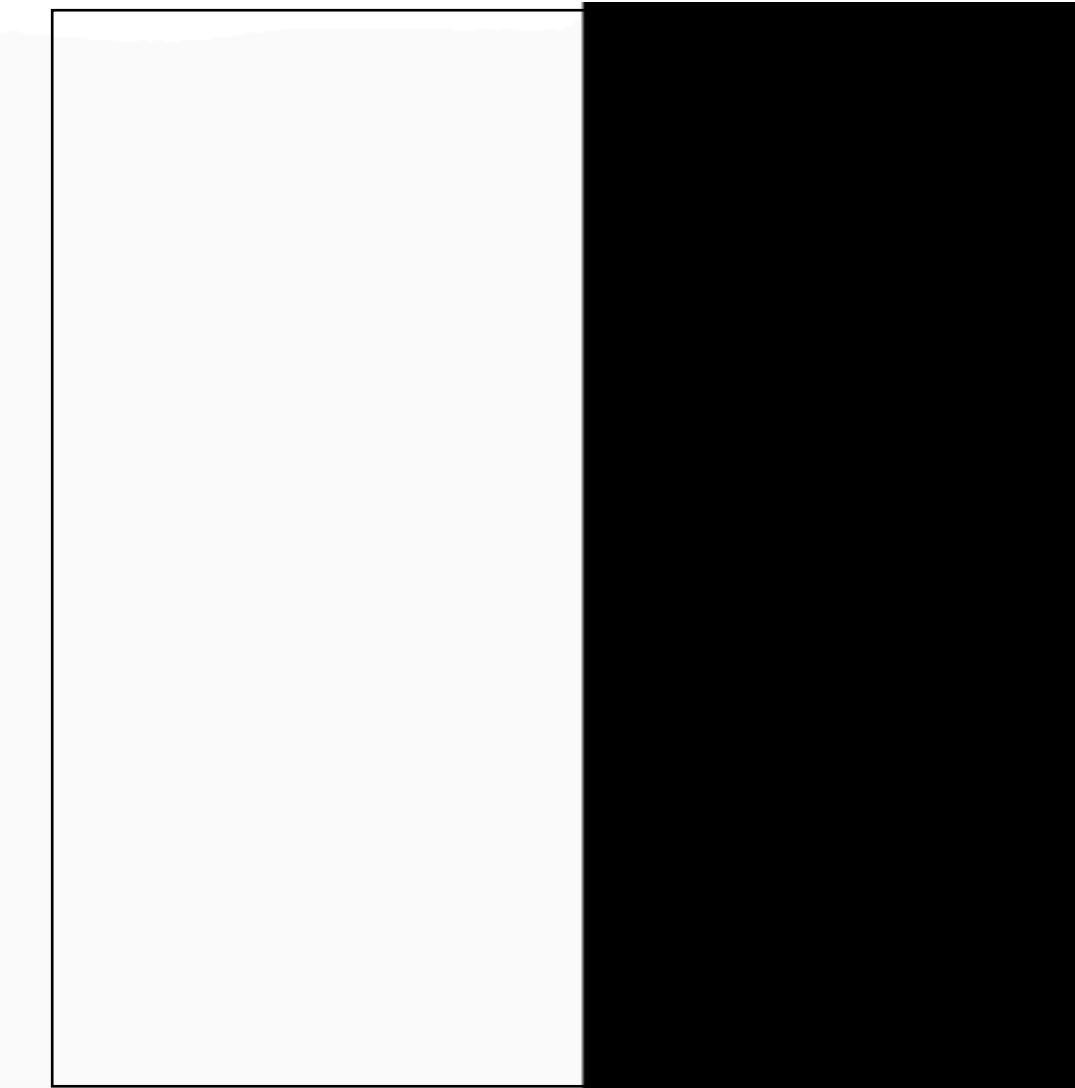
# Image Blending



# Image Blending

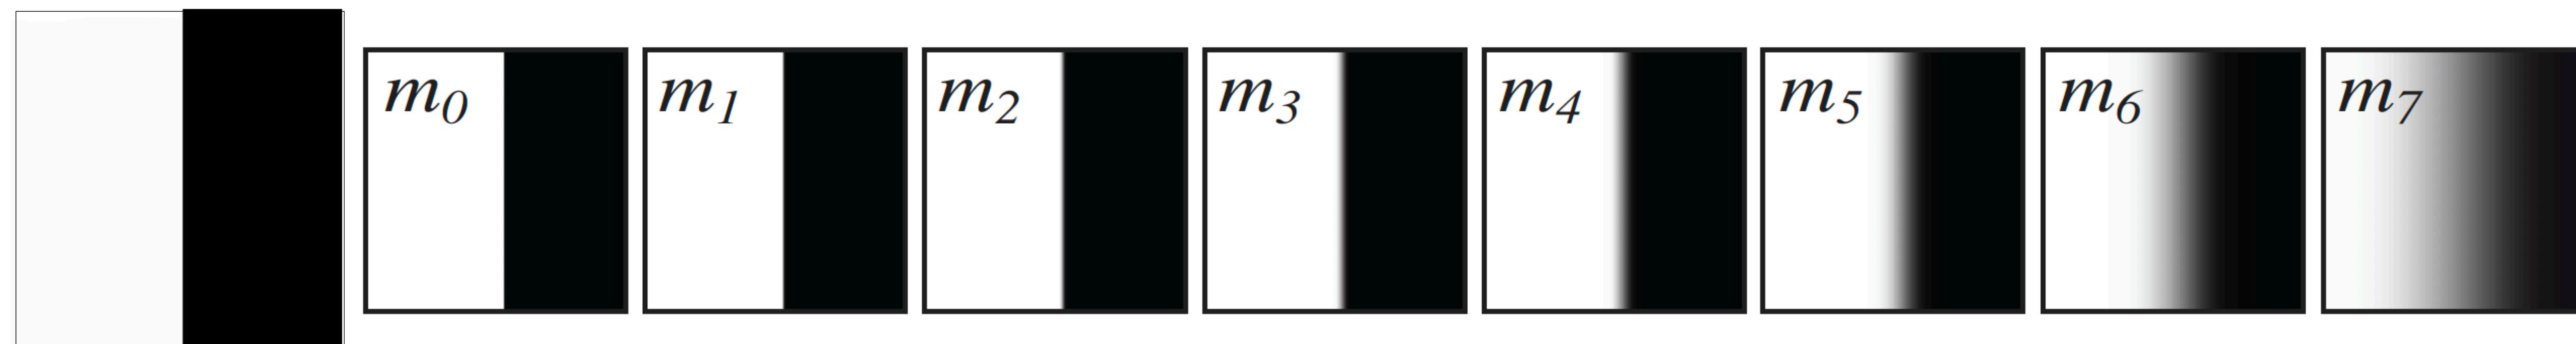
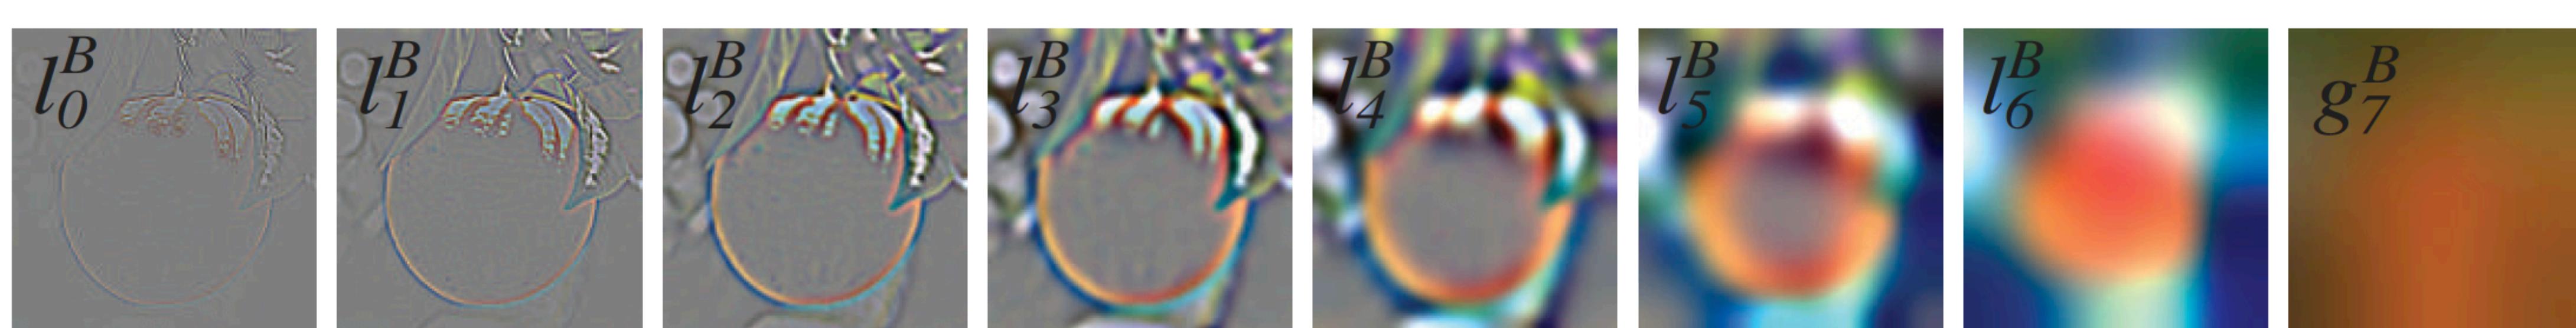
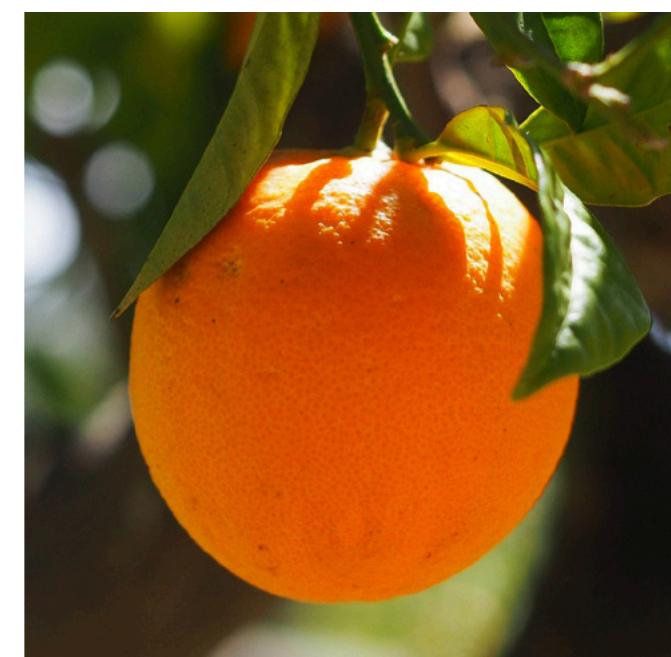
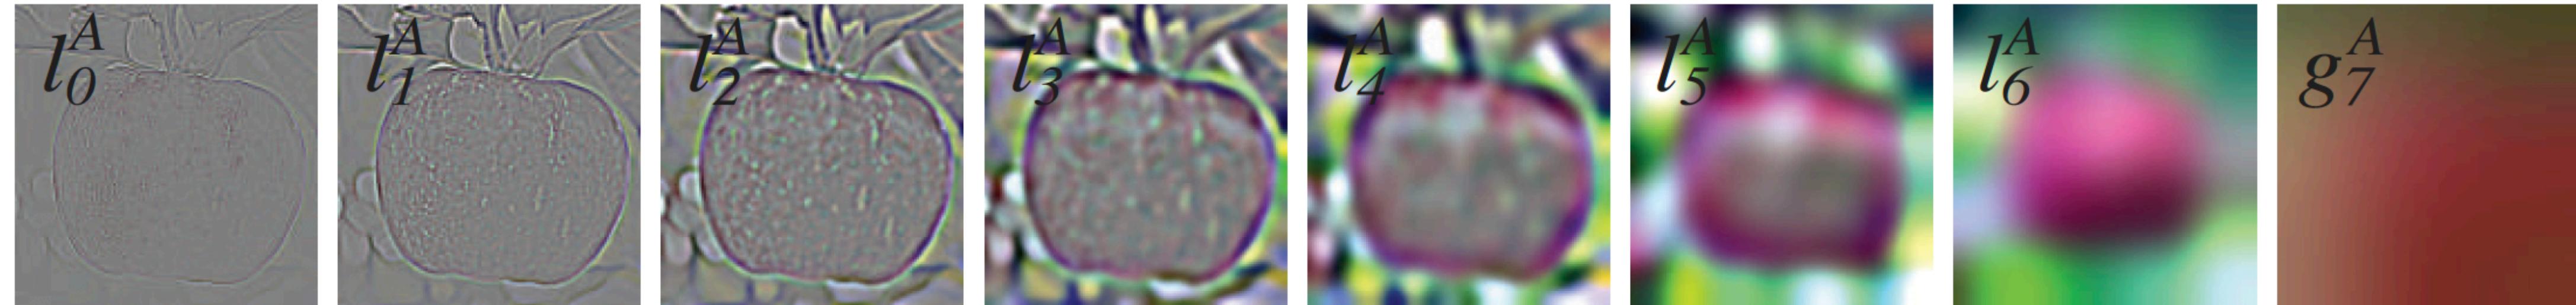


# Image Blending

 $I^A$  $I^B$  $m$  $I$ 

$$I = m * I^A + (1 - m) * I^B$$

# Image Blending with the Laplacian Pyramid



$$l_k = l_k^A * m_k + l_i^B * (1 - m_k)$$

# Image Blending with the Laplacian Pyramid



# Image Blending with the Laplacian Pyramid

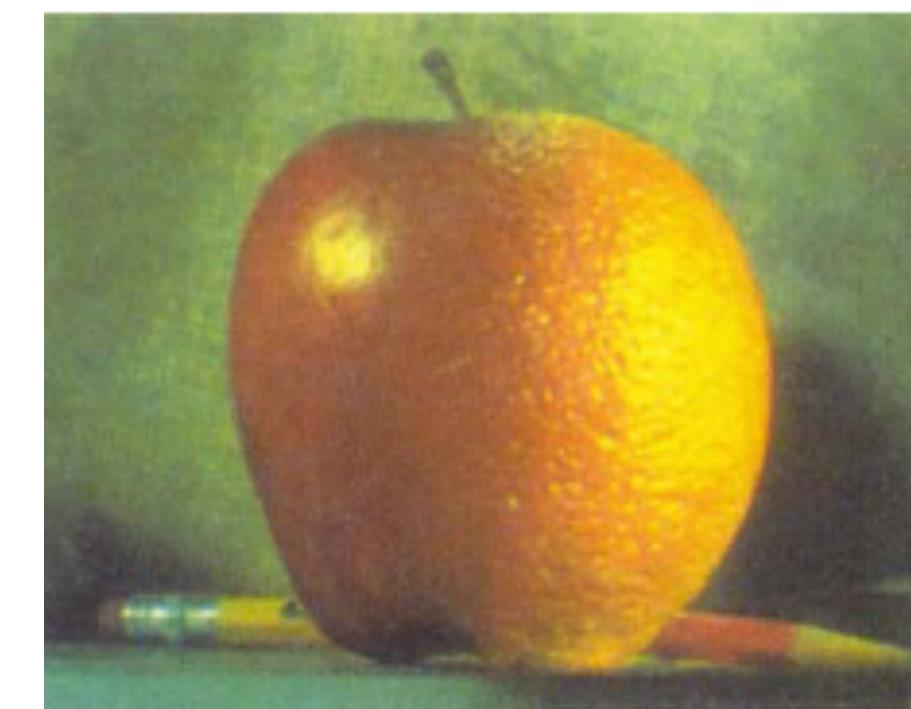
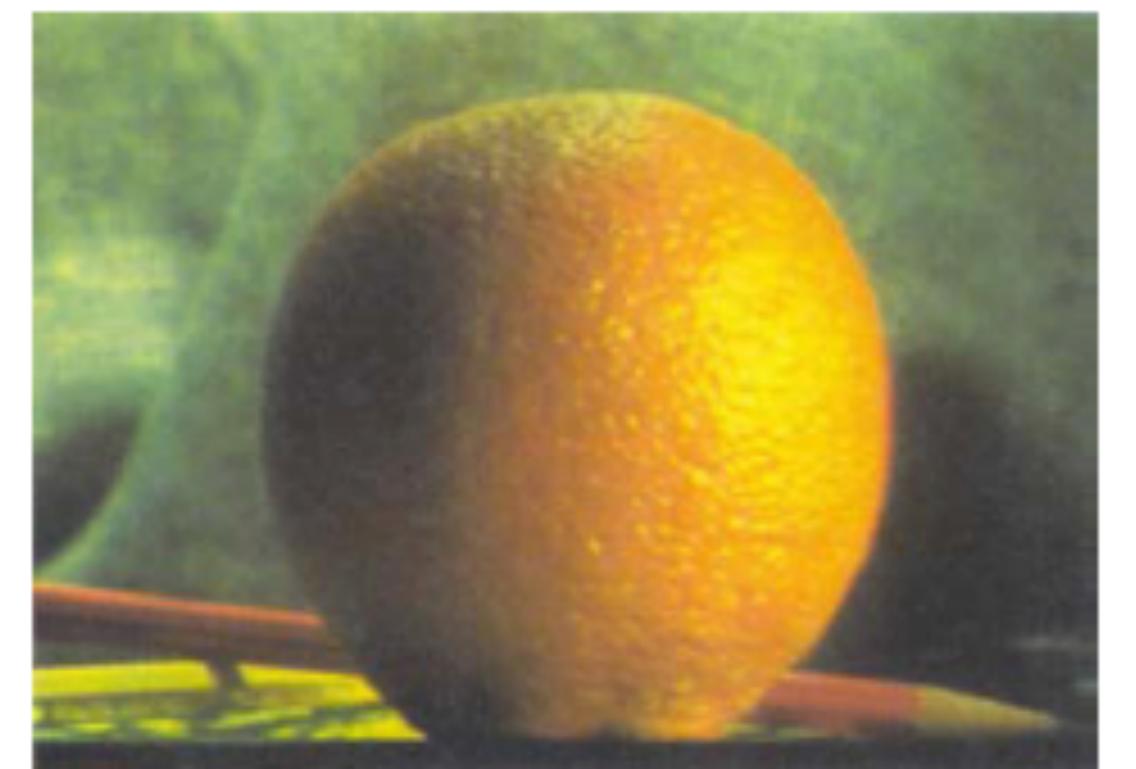
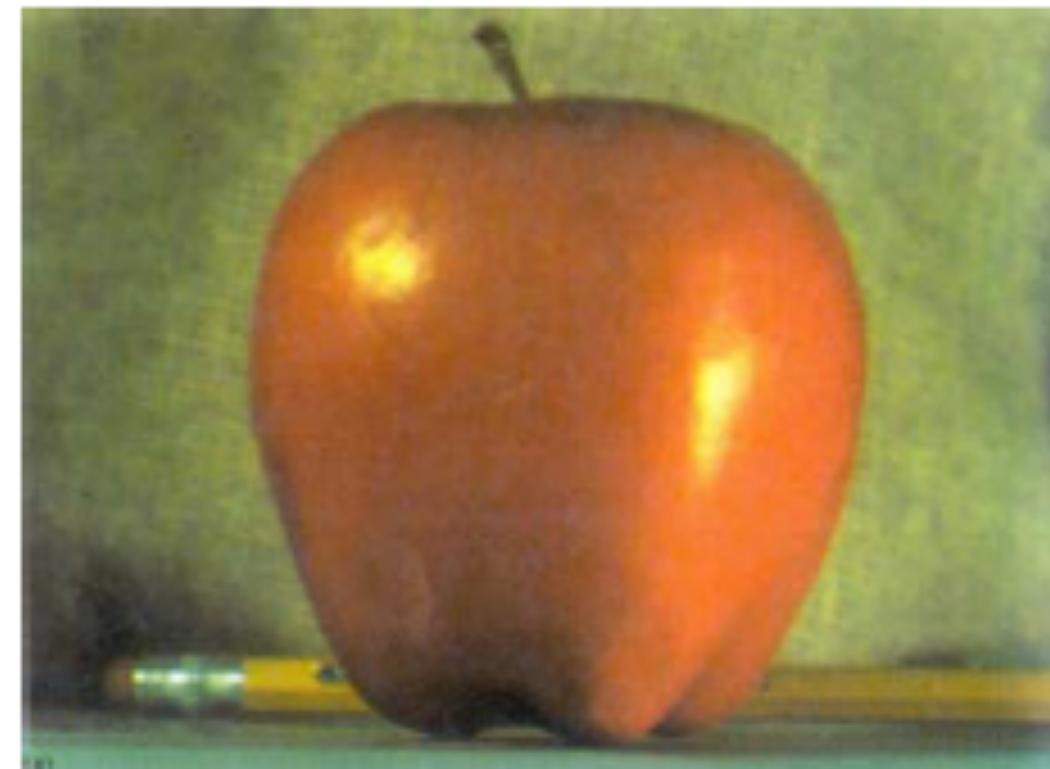
532

IEEE TRANSACTIONS ON COMMUNICATIONS, VOL. COM-31, NO. 4, APRIL 1983

## The Laplacian Pyramid as a Compact Image Code

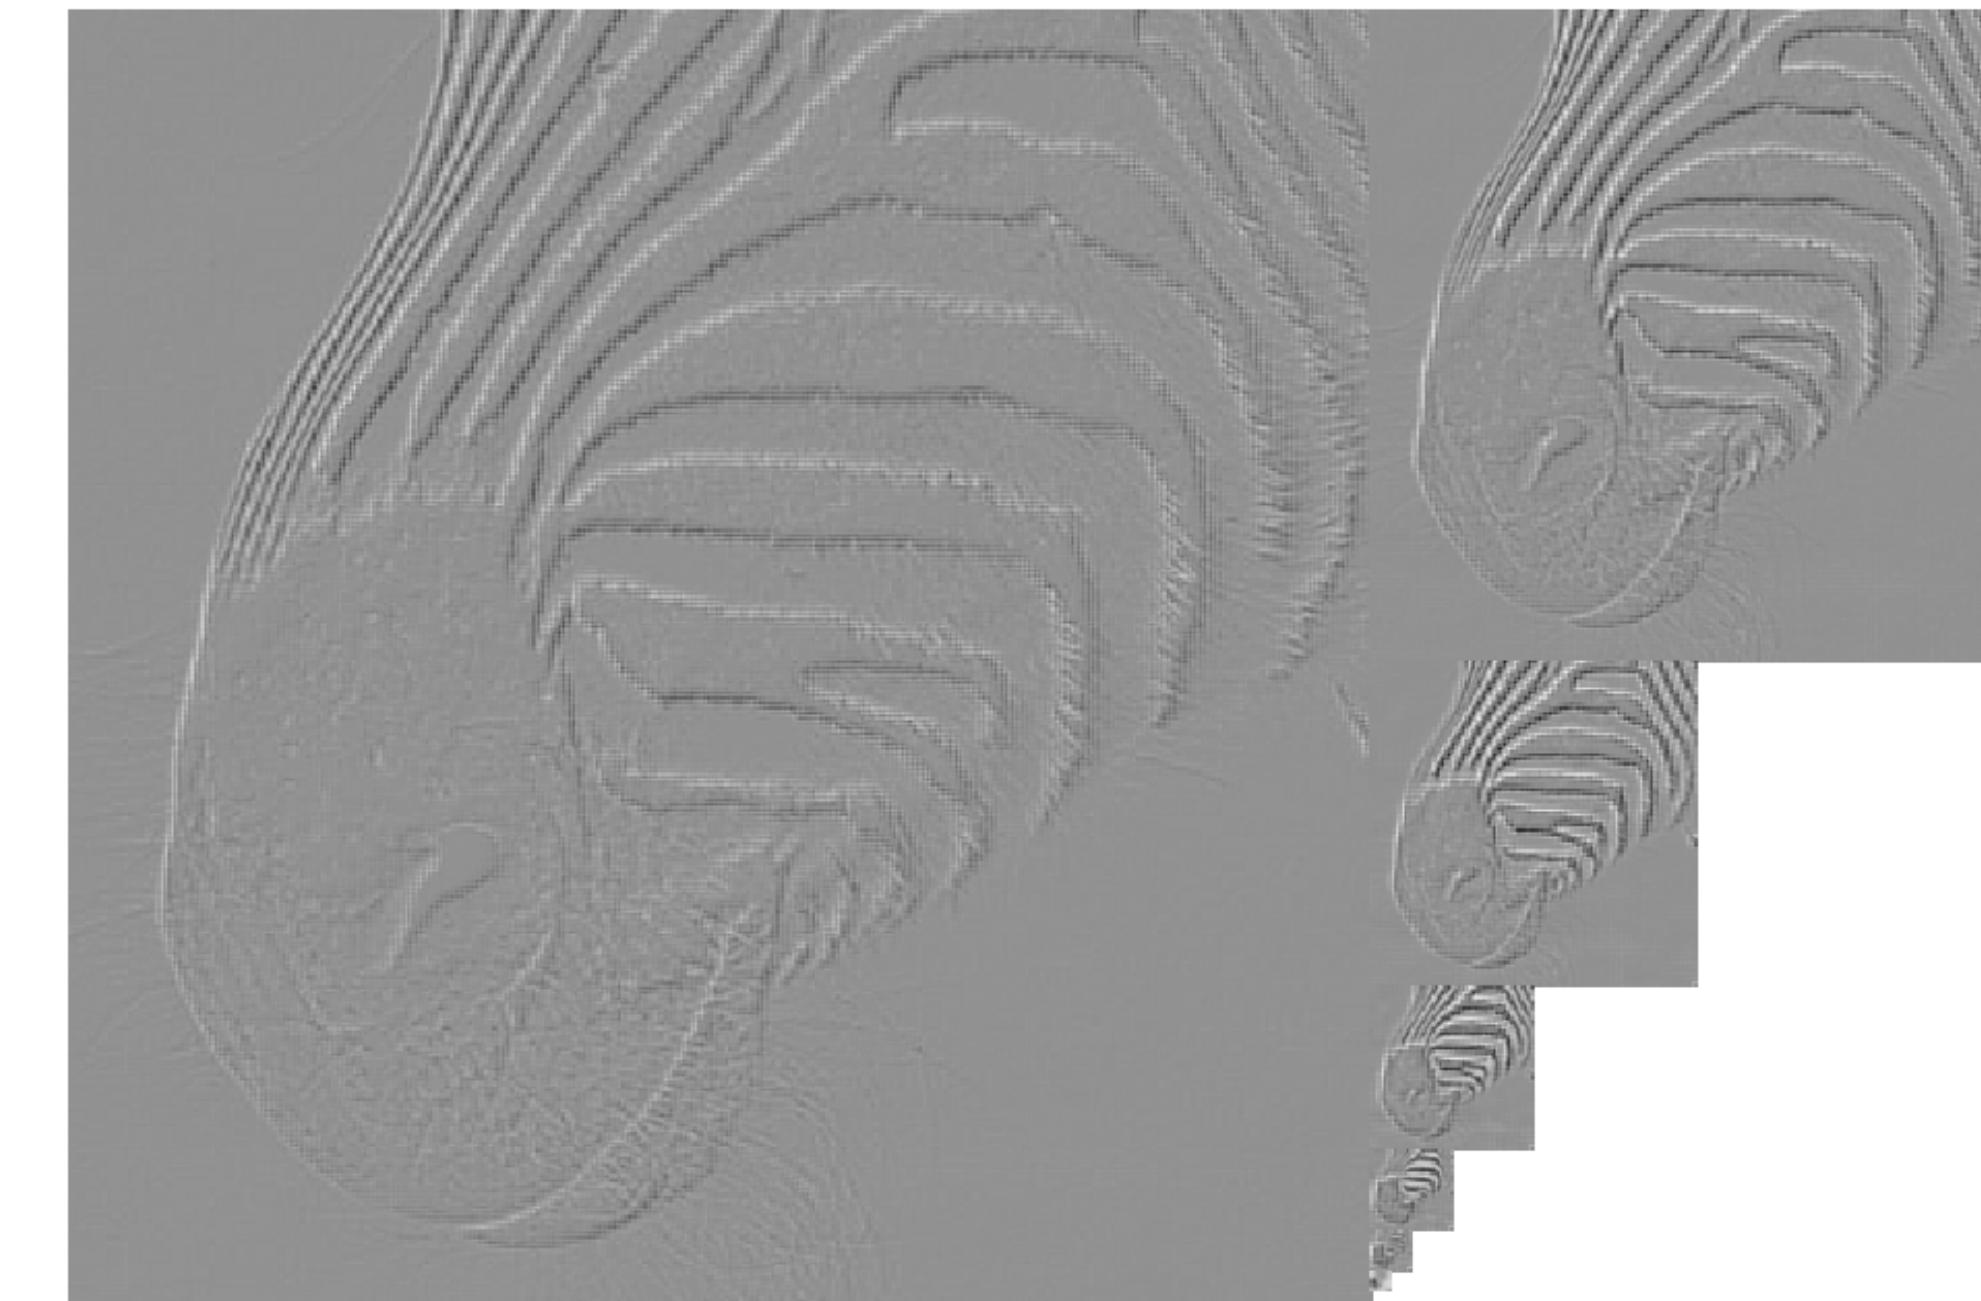
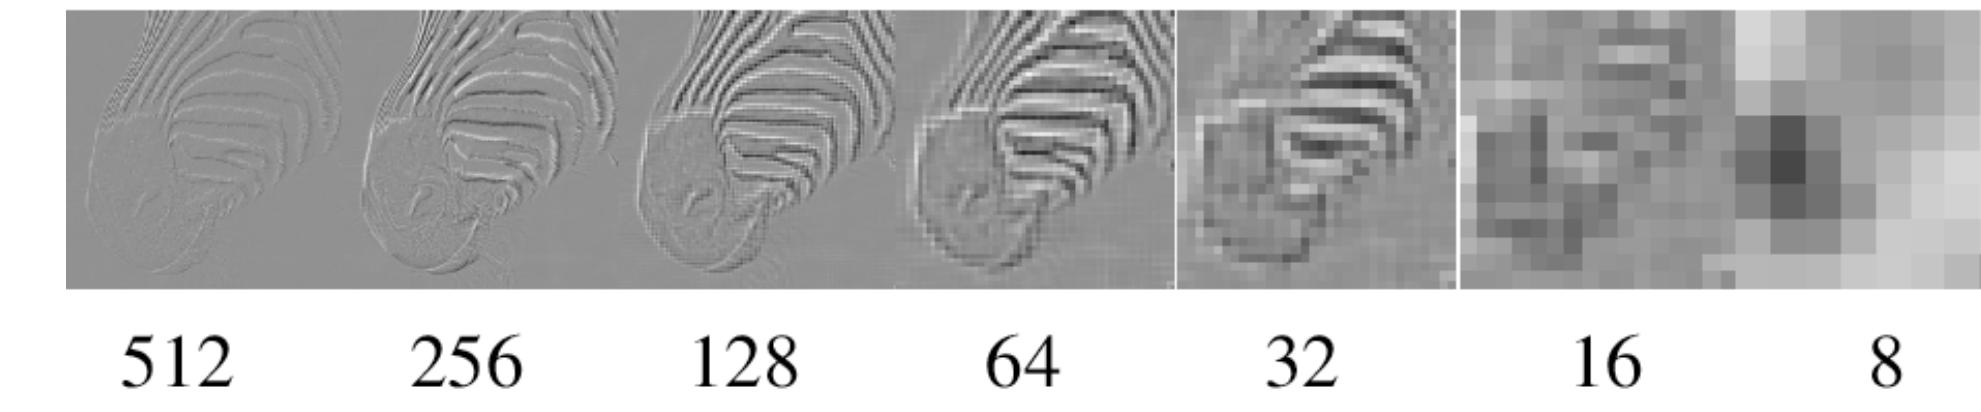
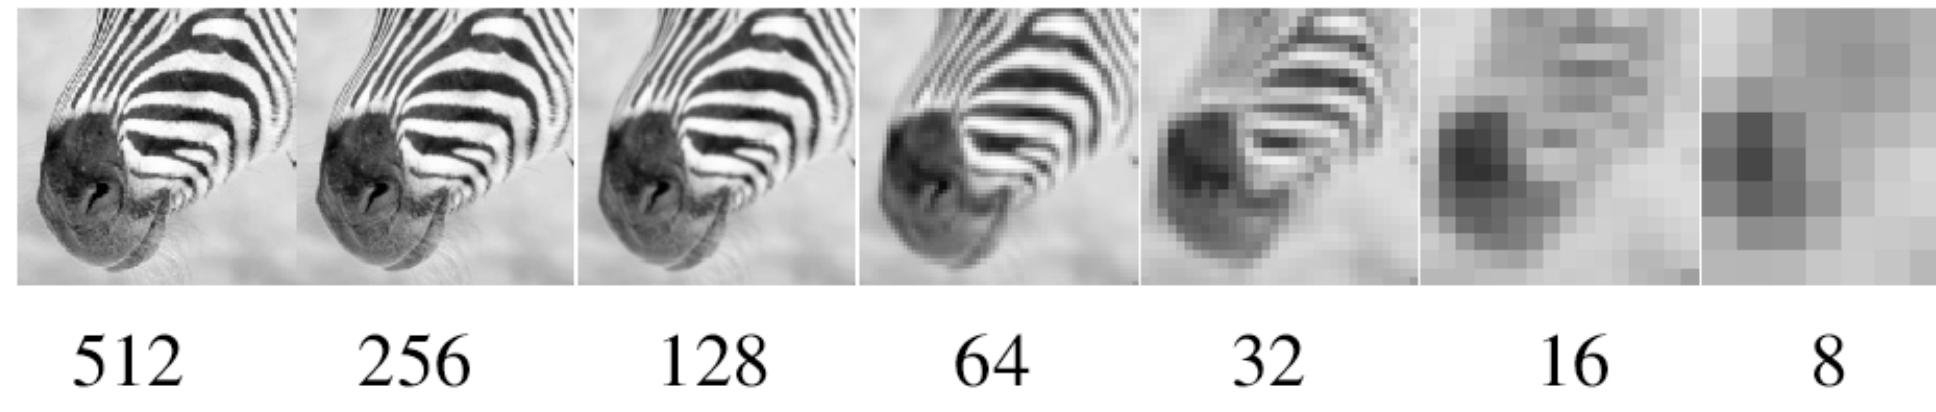
PETER J. BURT, MEMBER, IEEE, AND EDWARD H. ADELSON

- Build Laplacian pyramid for both images: LA, LB
- Build Gaussian pyramid for mask: G
- Build a combined Laplacian pyramid:
- Collapse L to obtain the blended image



[http://persci.mit.edu/pub\\_pdfs/pyramid83.pdf](http://persci.mit.edu/pub_pdfs/pyramid83.pdf)

# Image pyramids

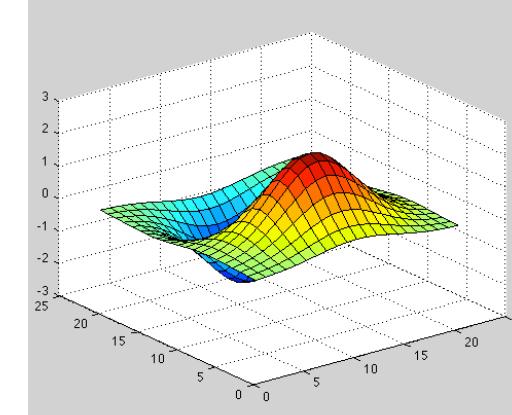
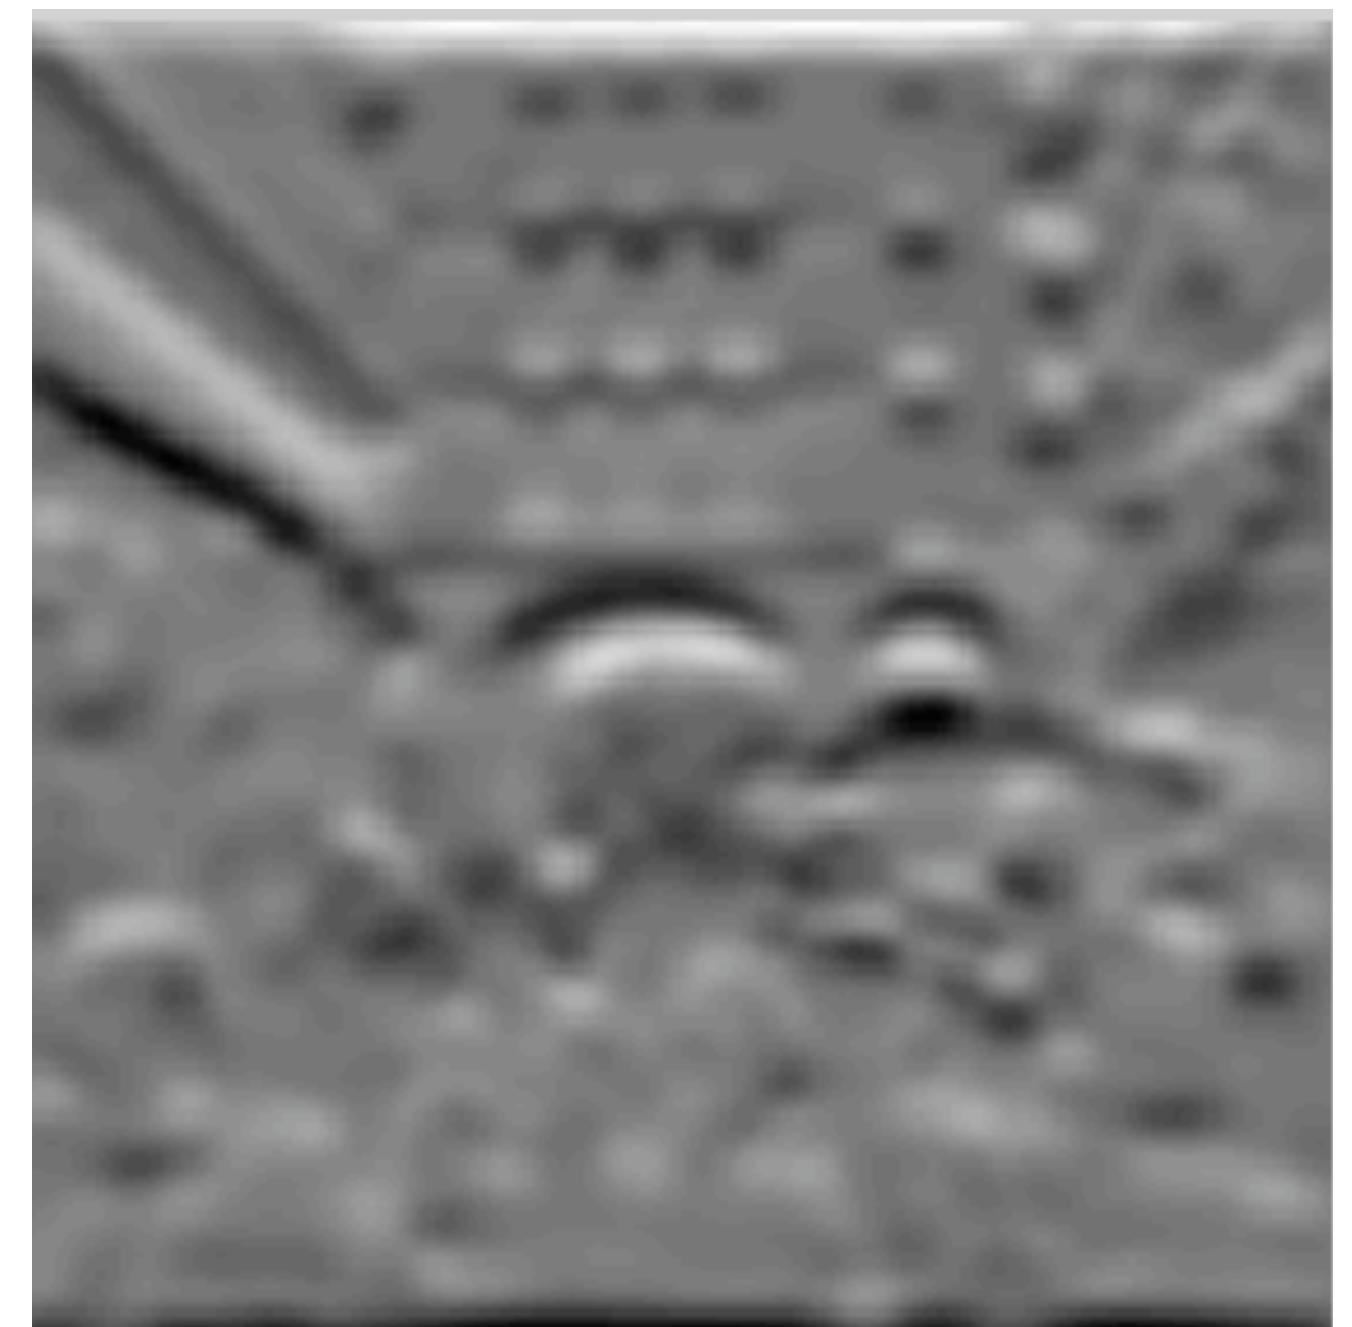
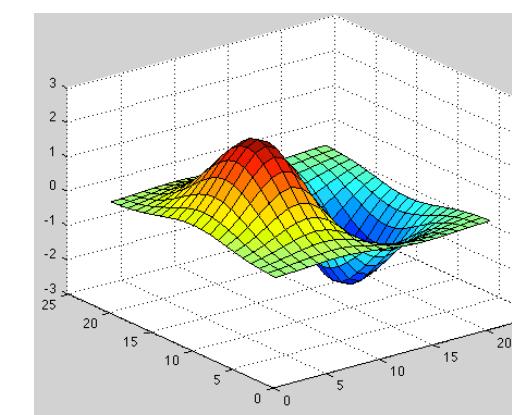
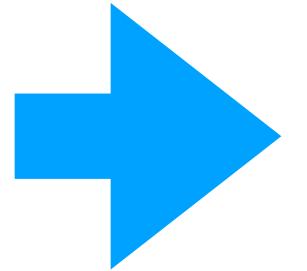


Gaussian Pyr

Laplacian Pyr

And many more: QMF, steerable, ... Convnets!

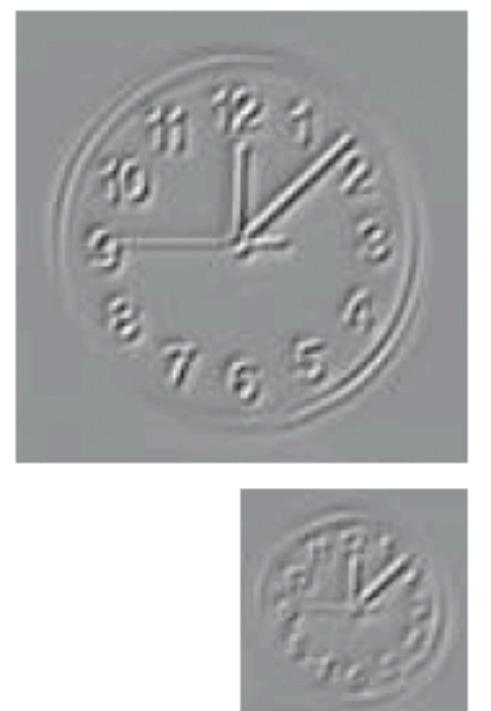
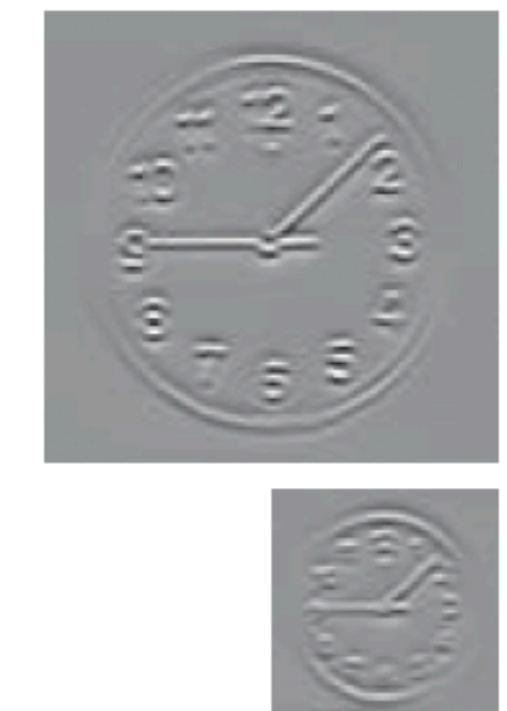
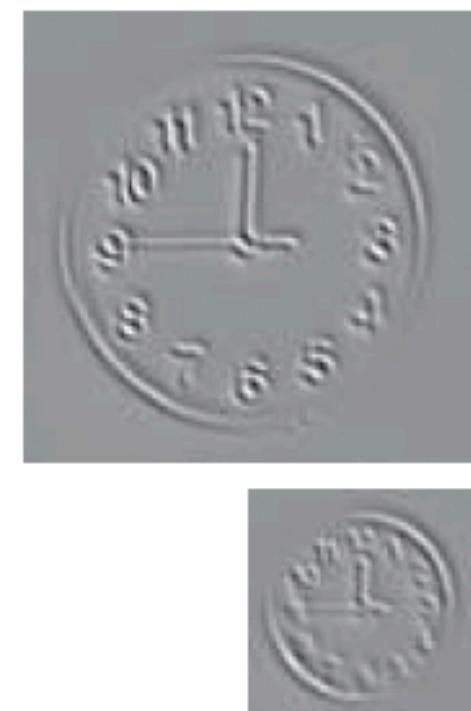
# Orientations



# Orientations

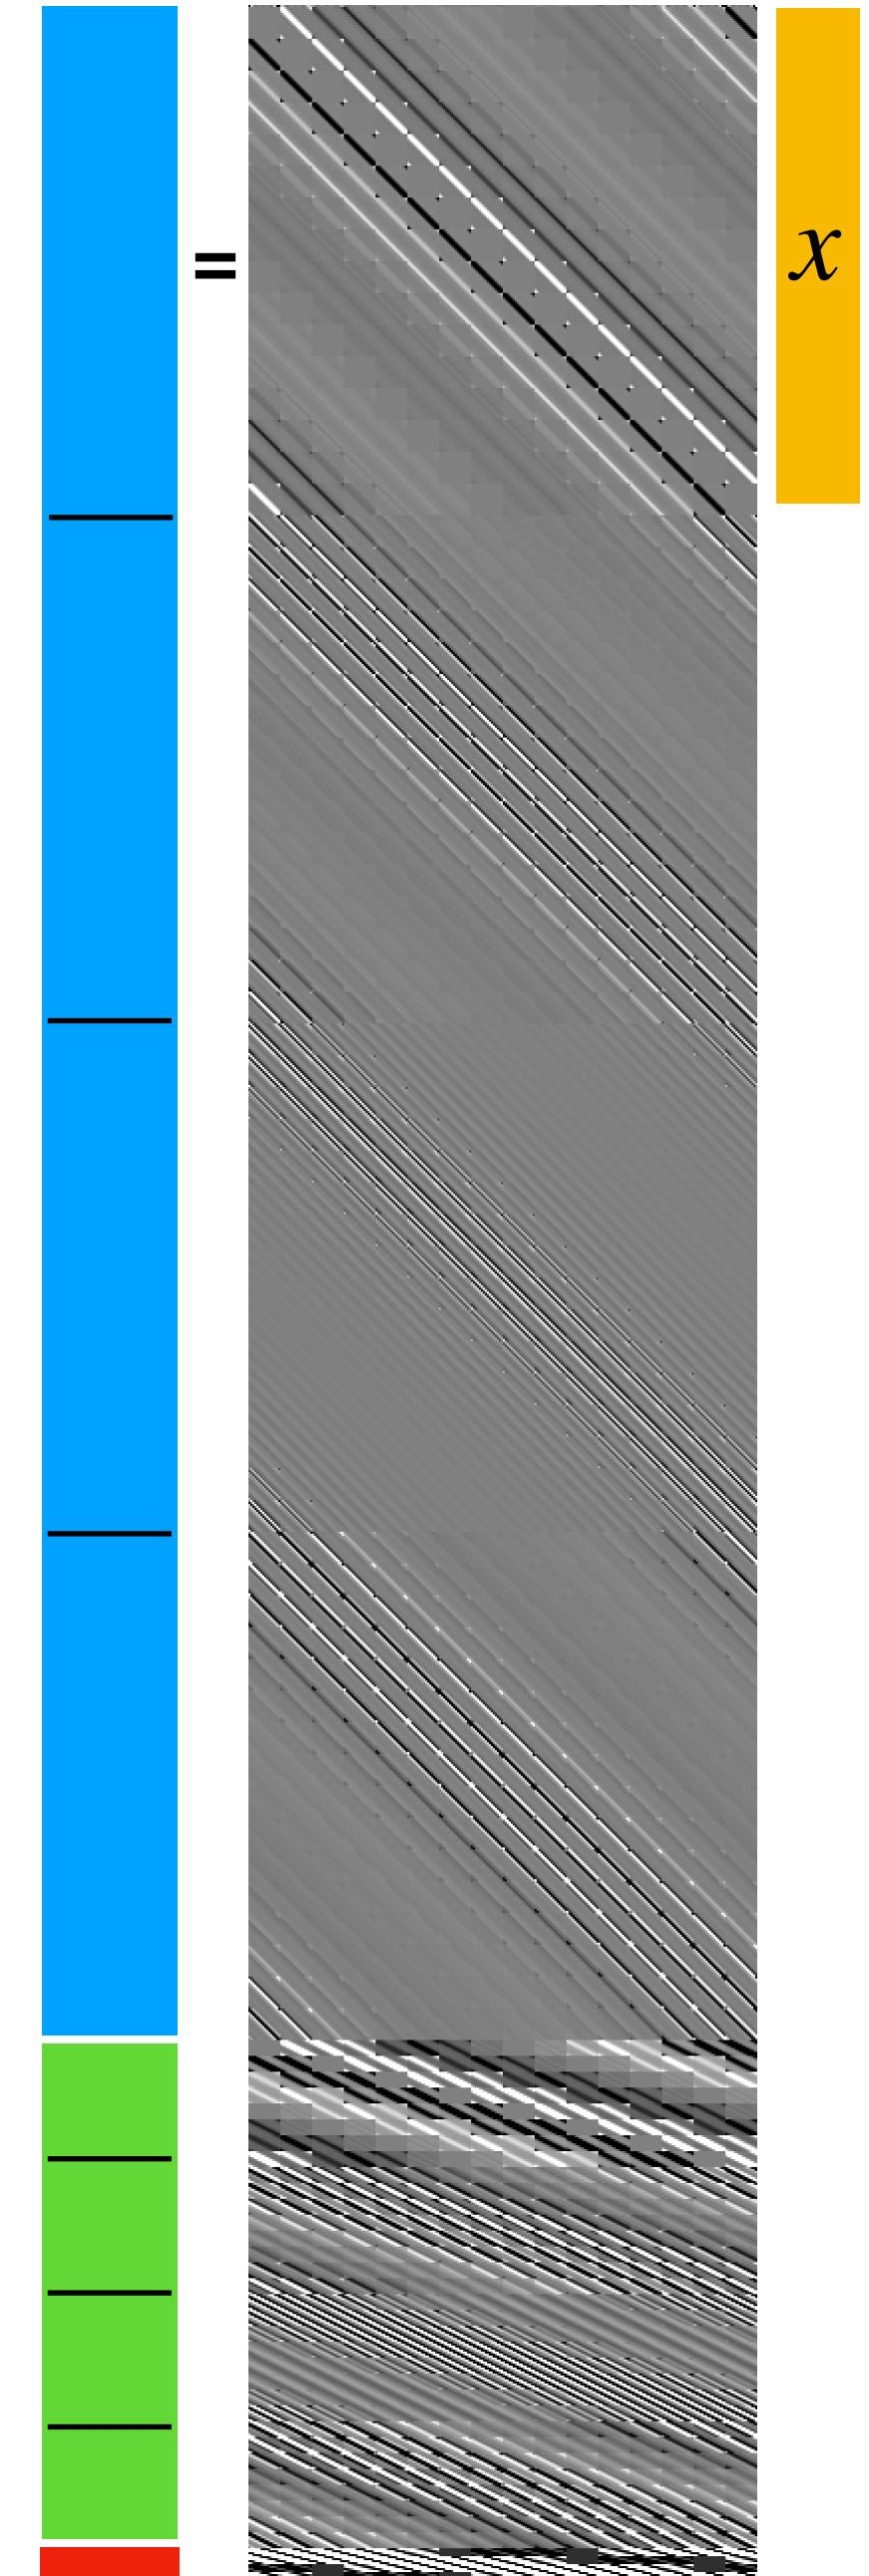
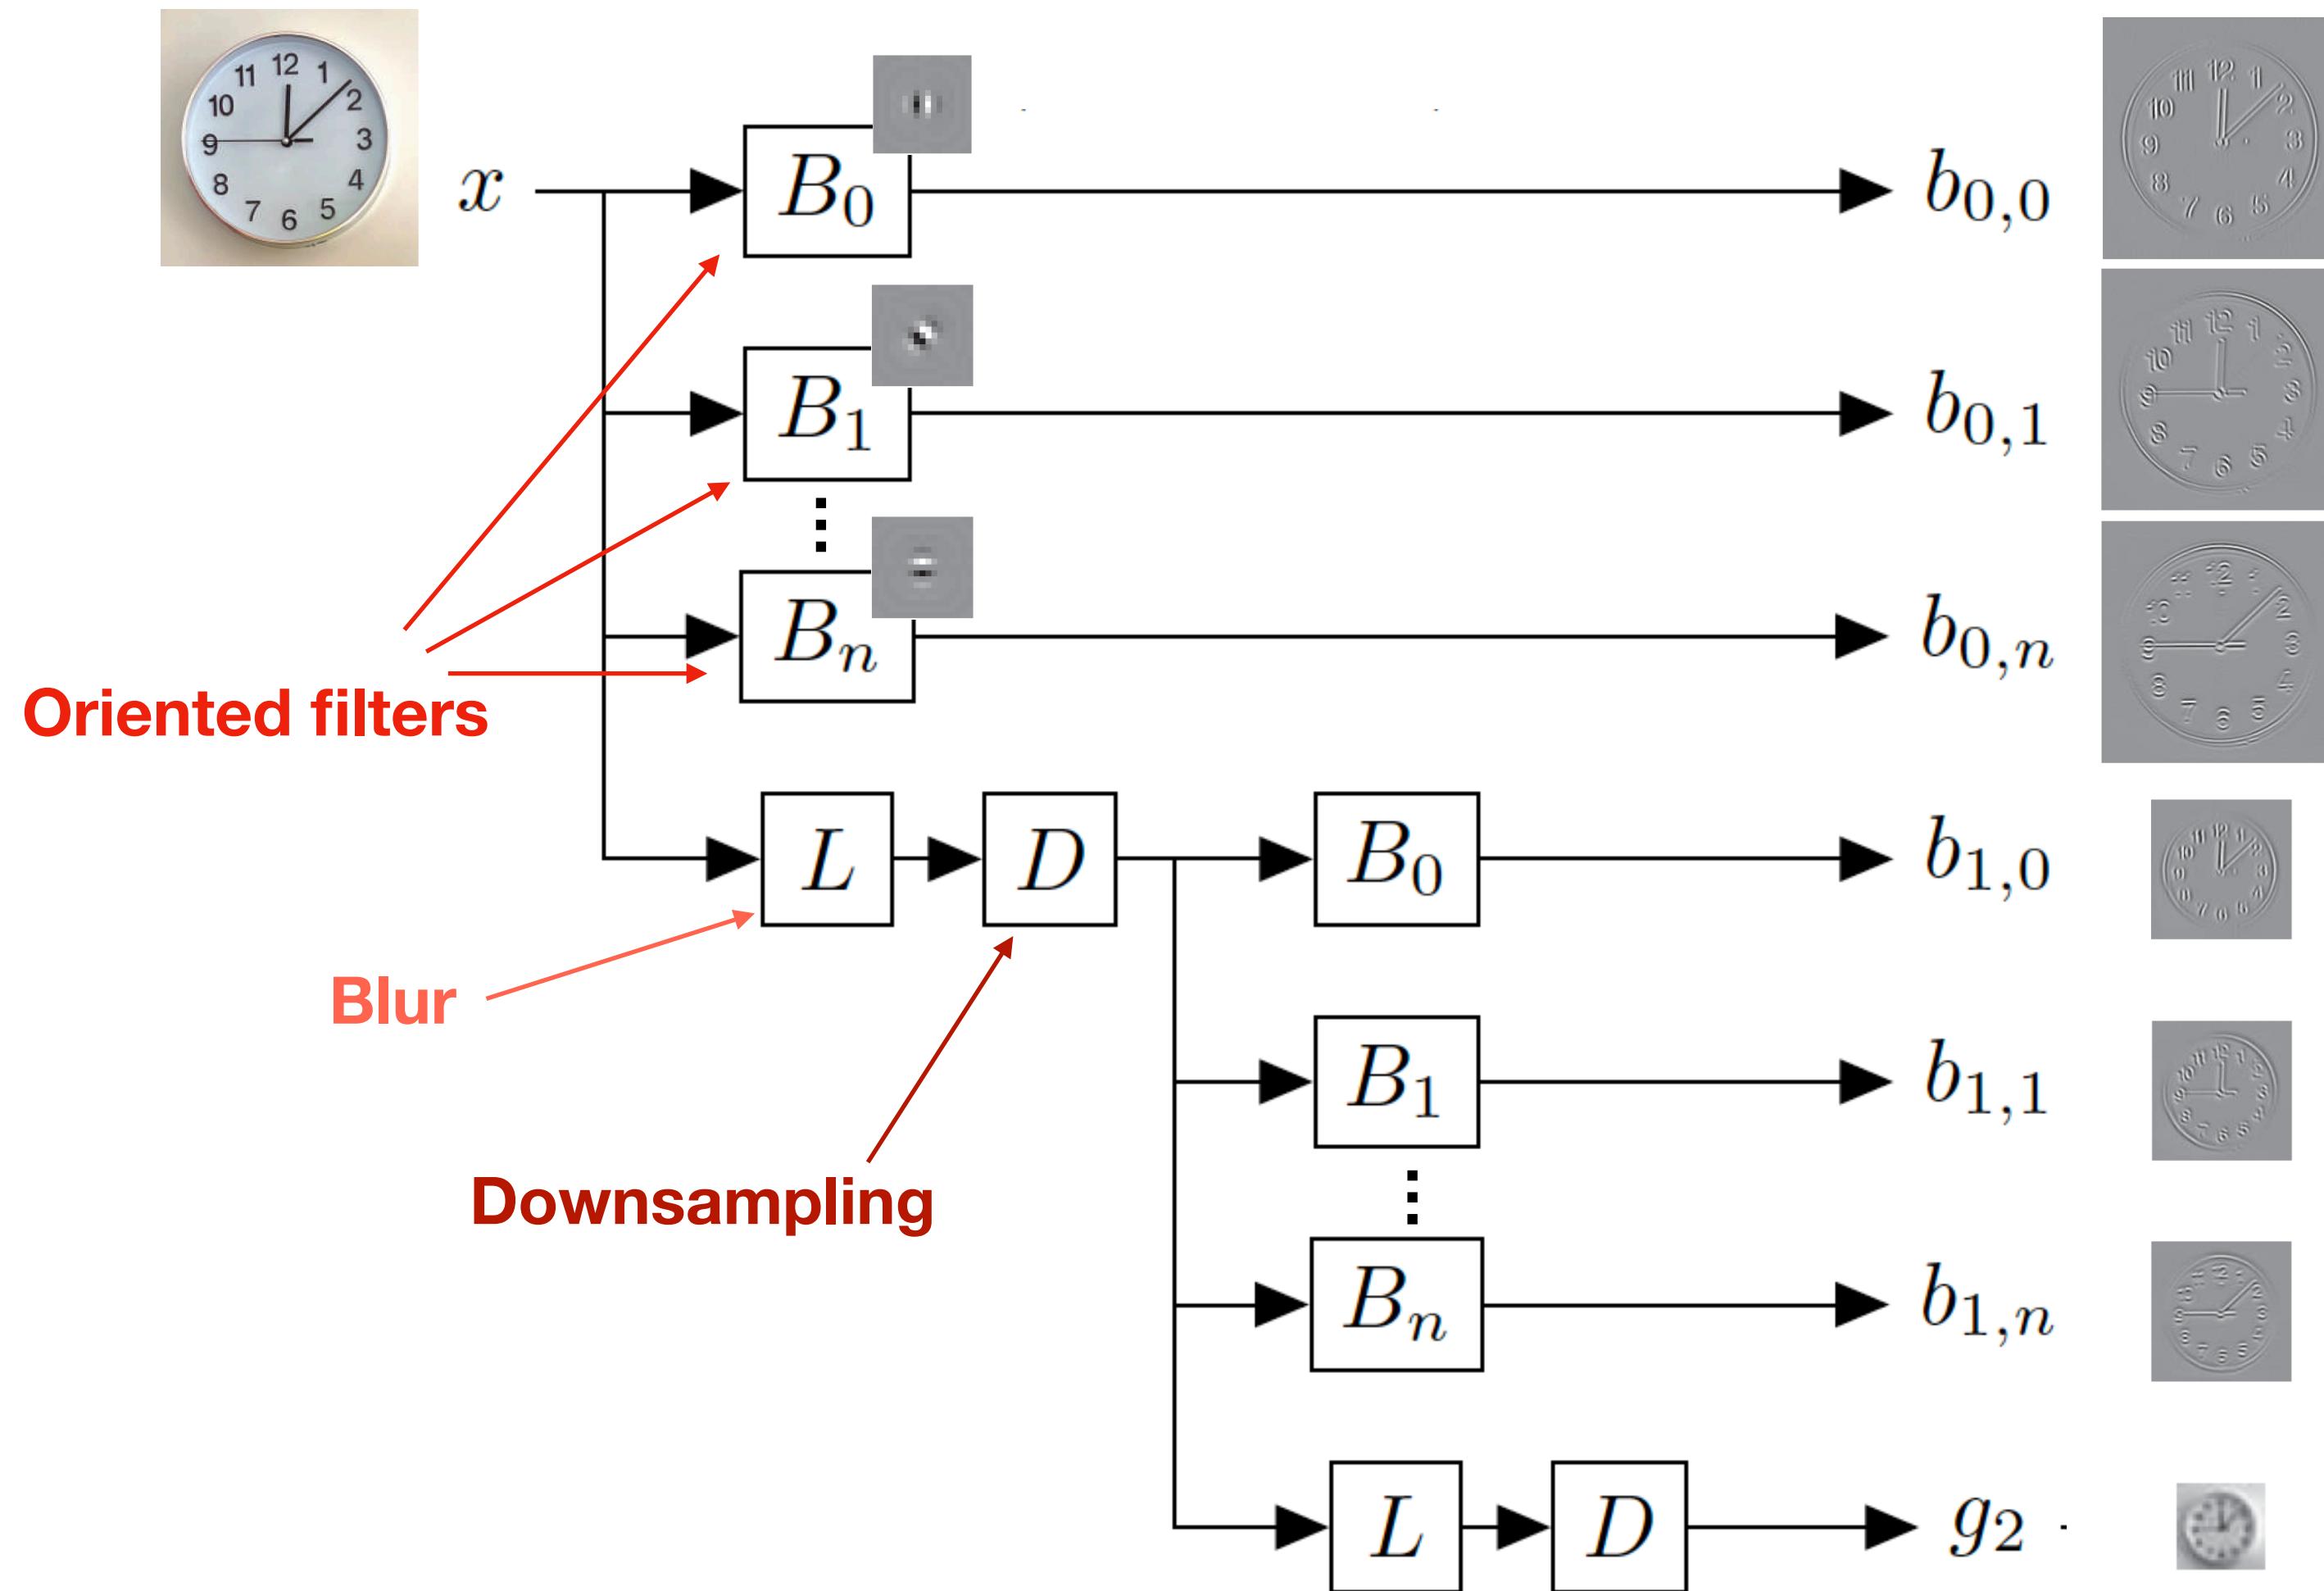


# Steerable Pyramid

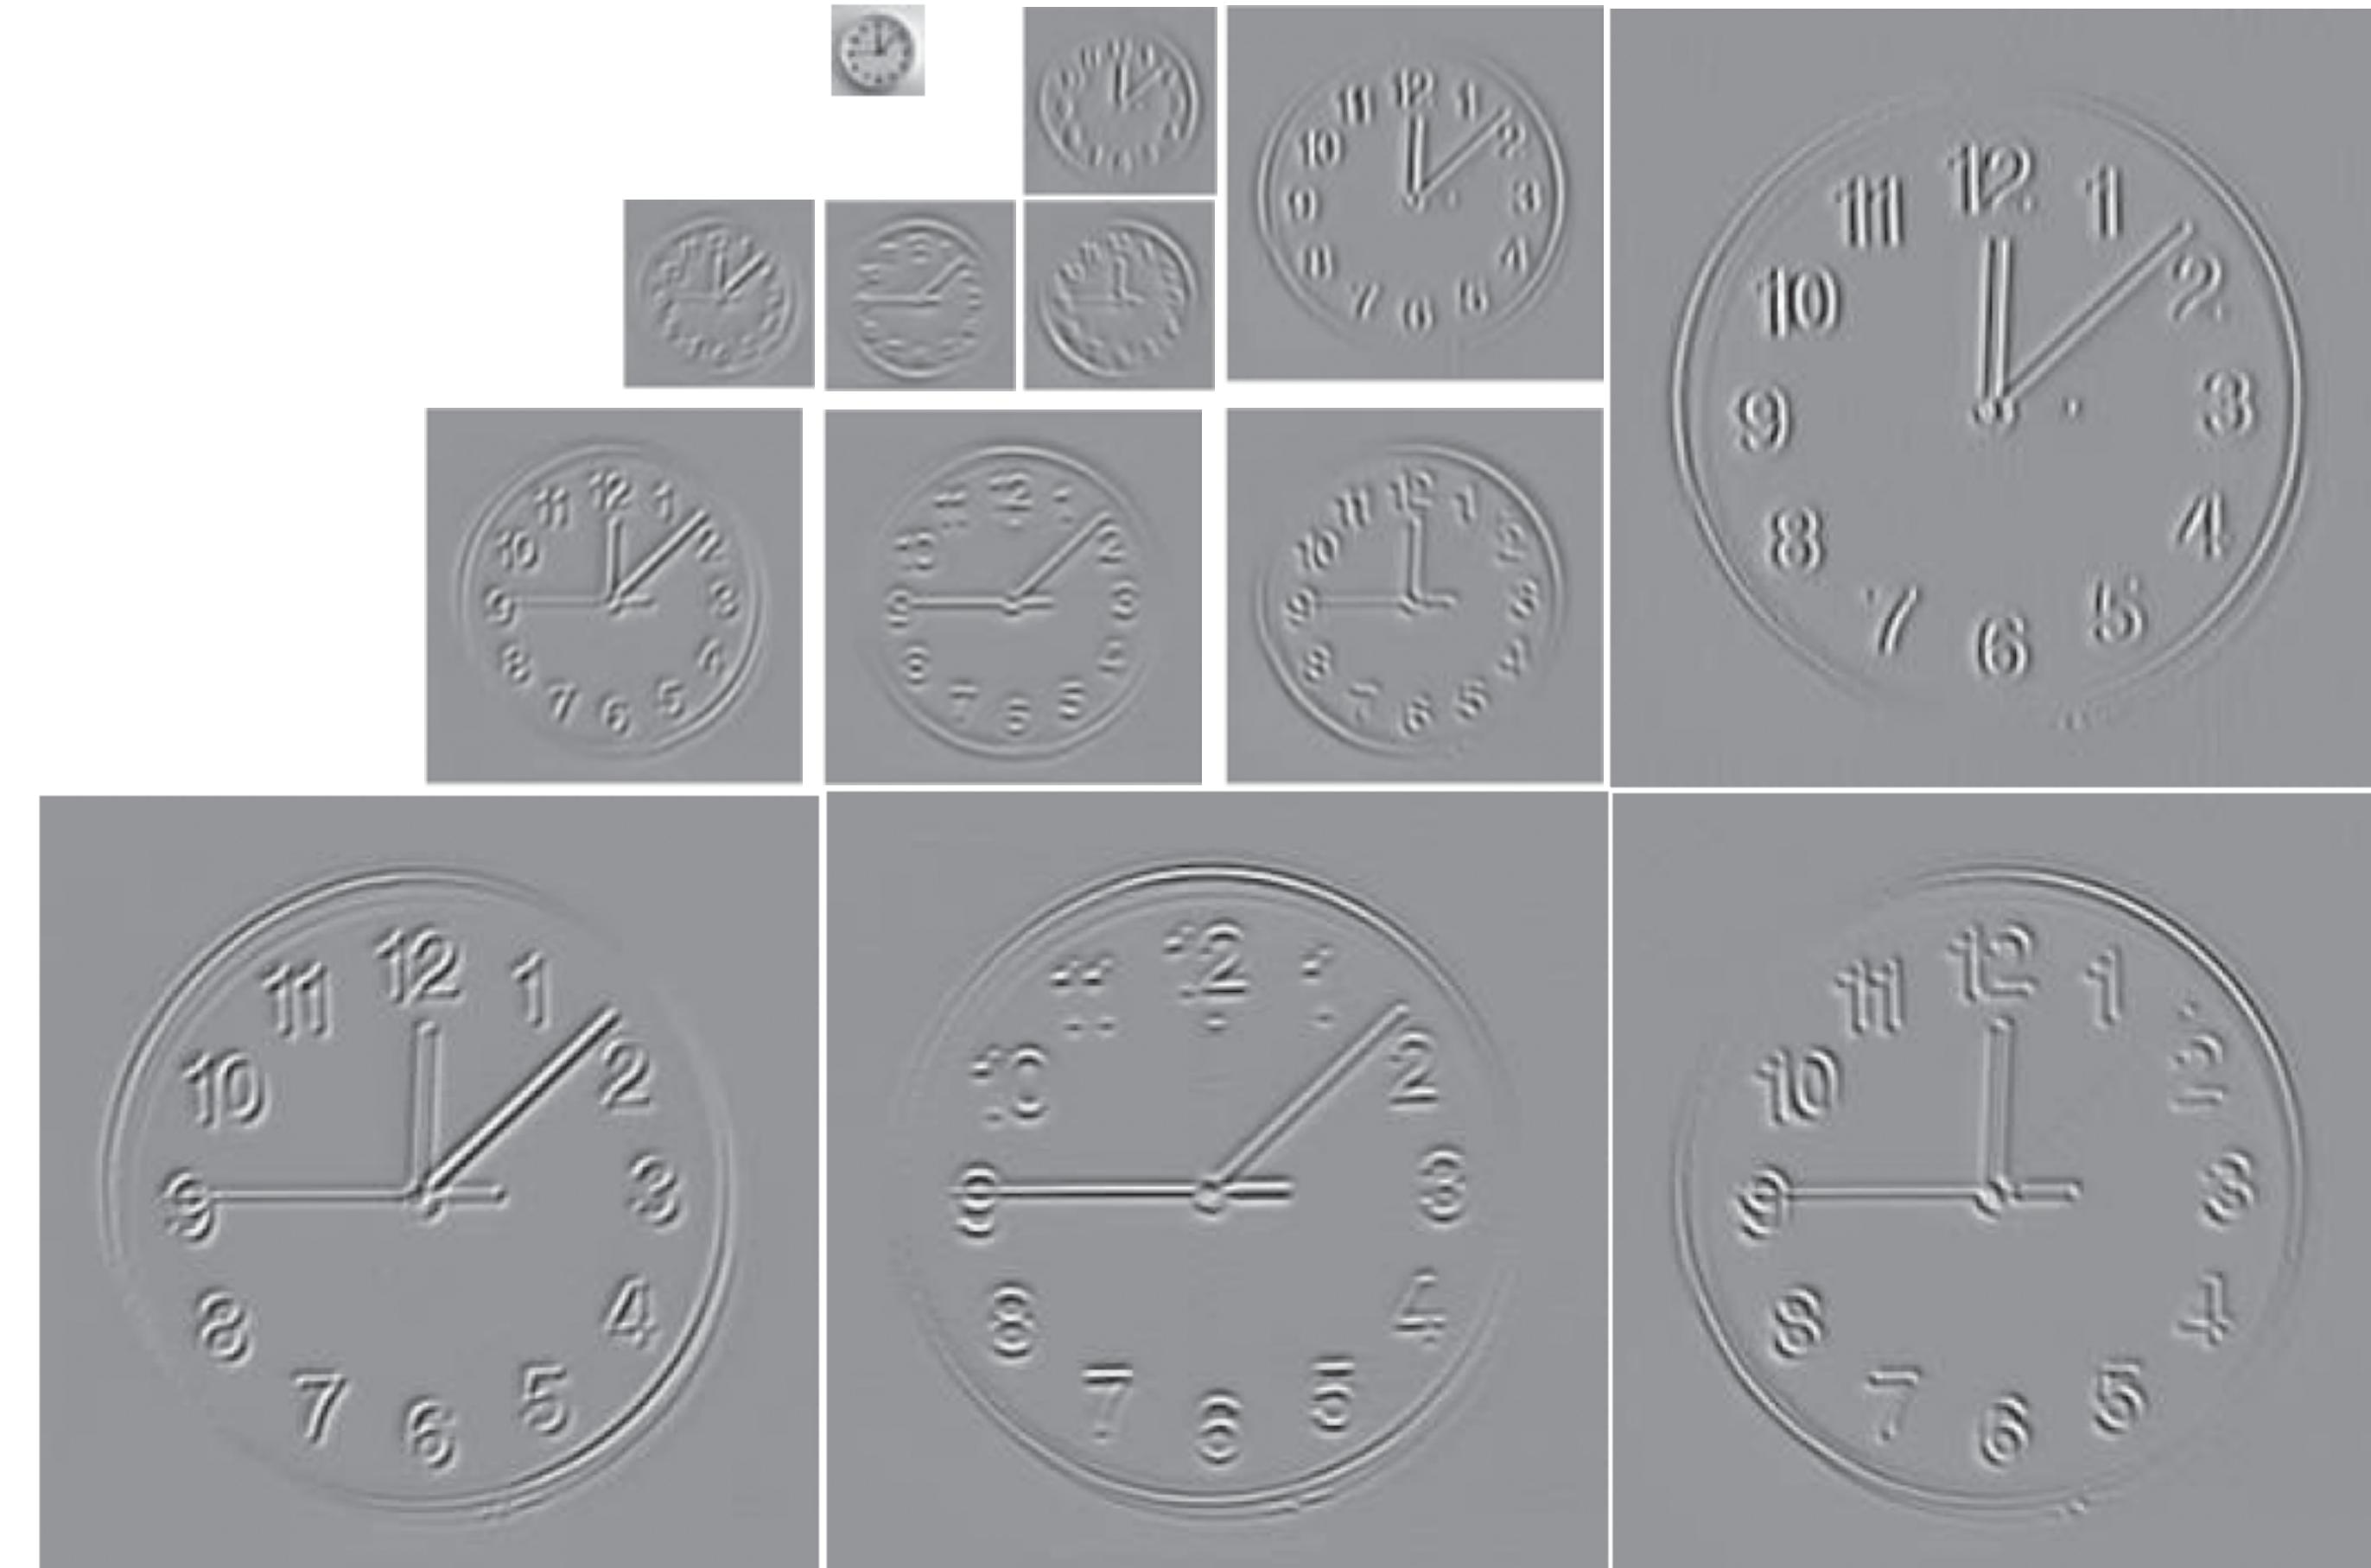
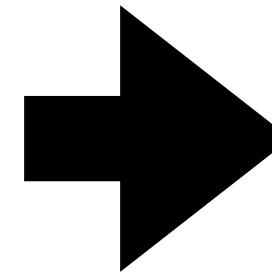


Low pass residual

# Steerable Pyramid

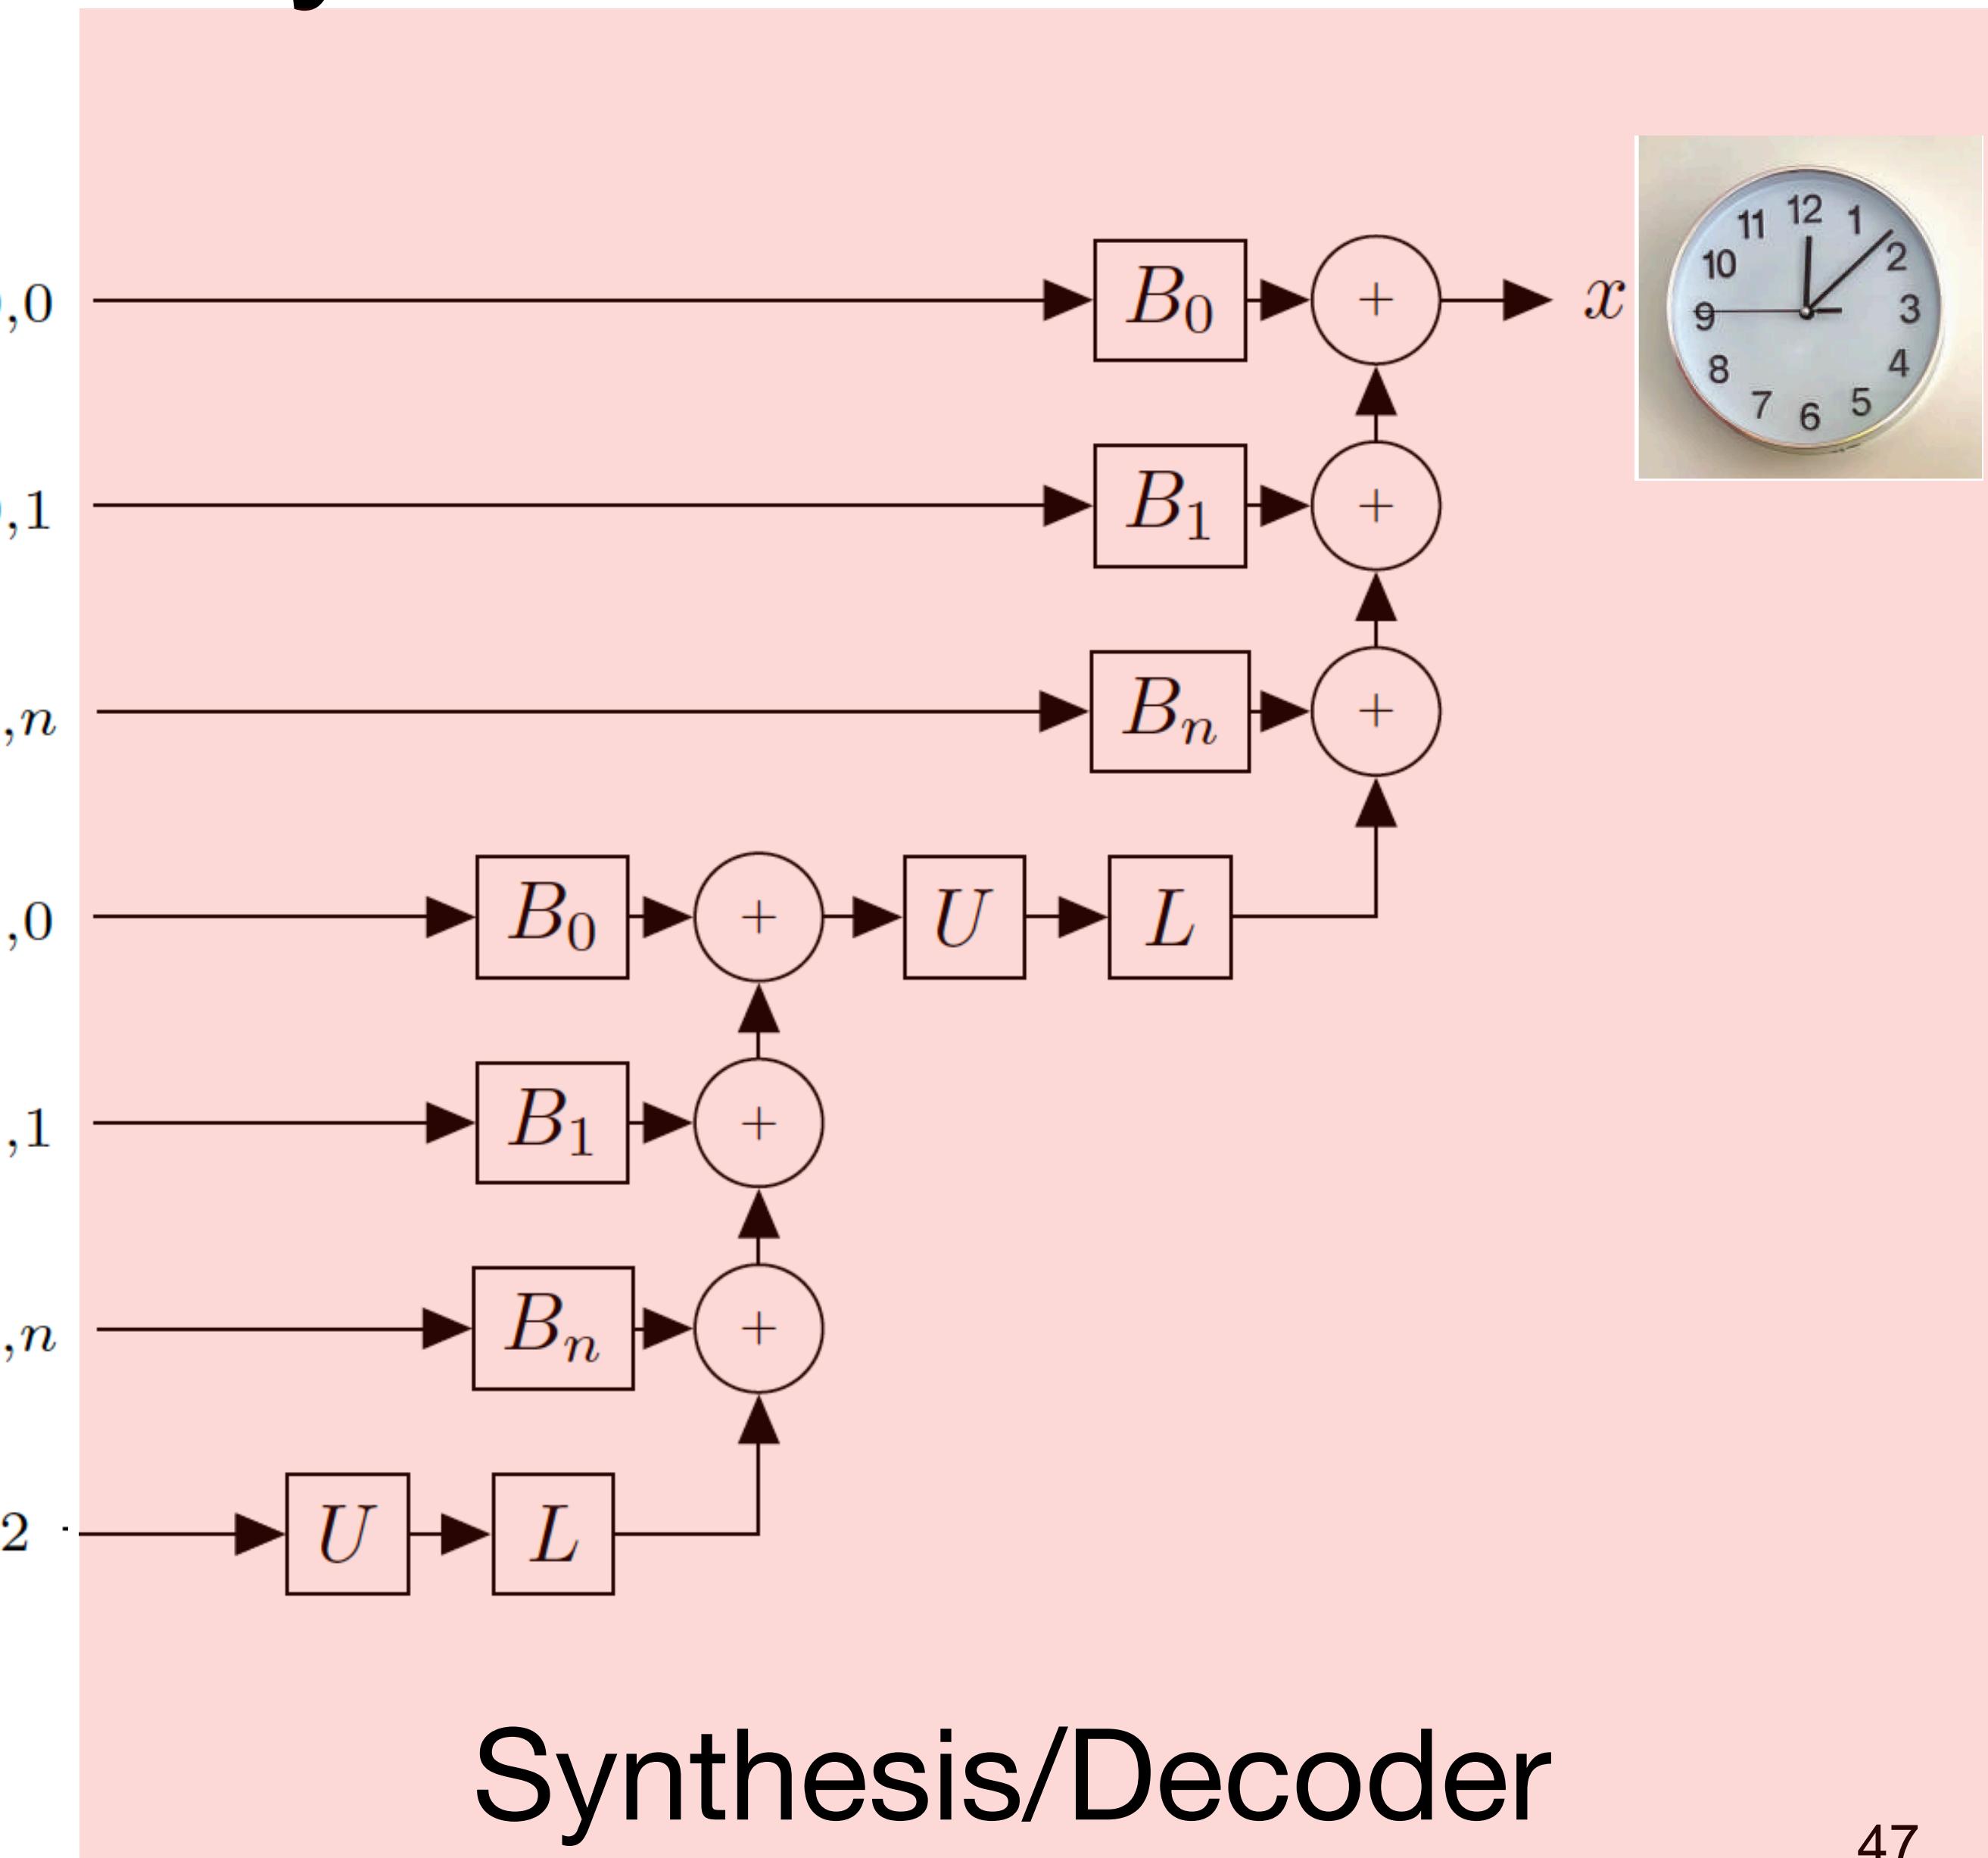
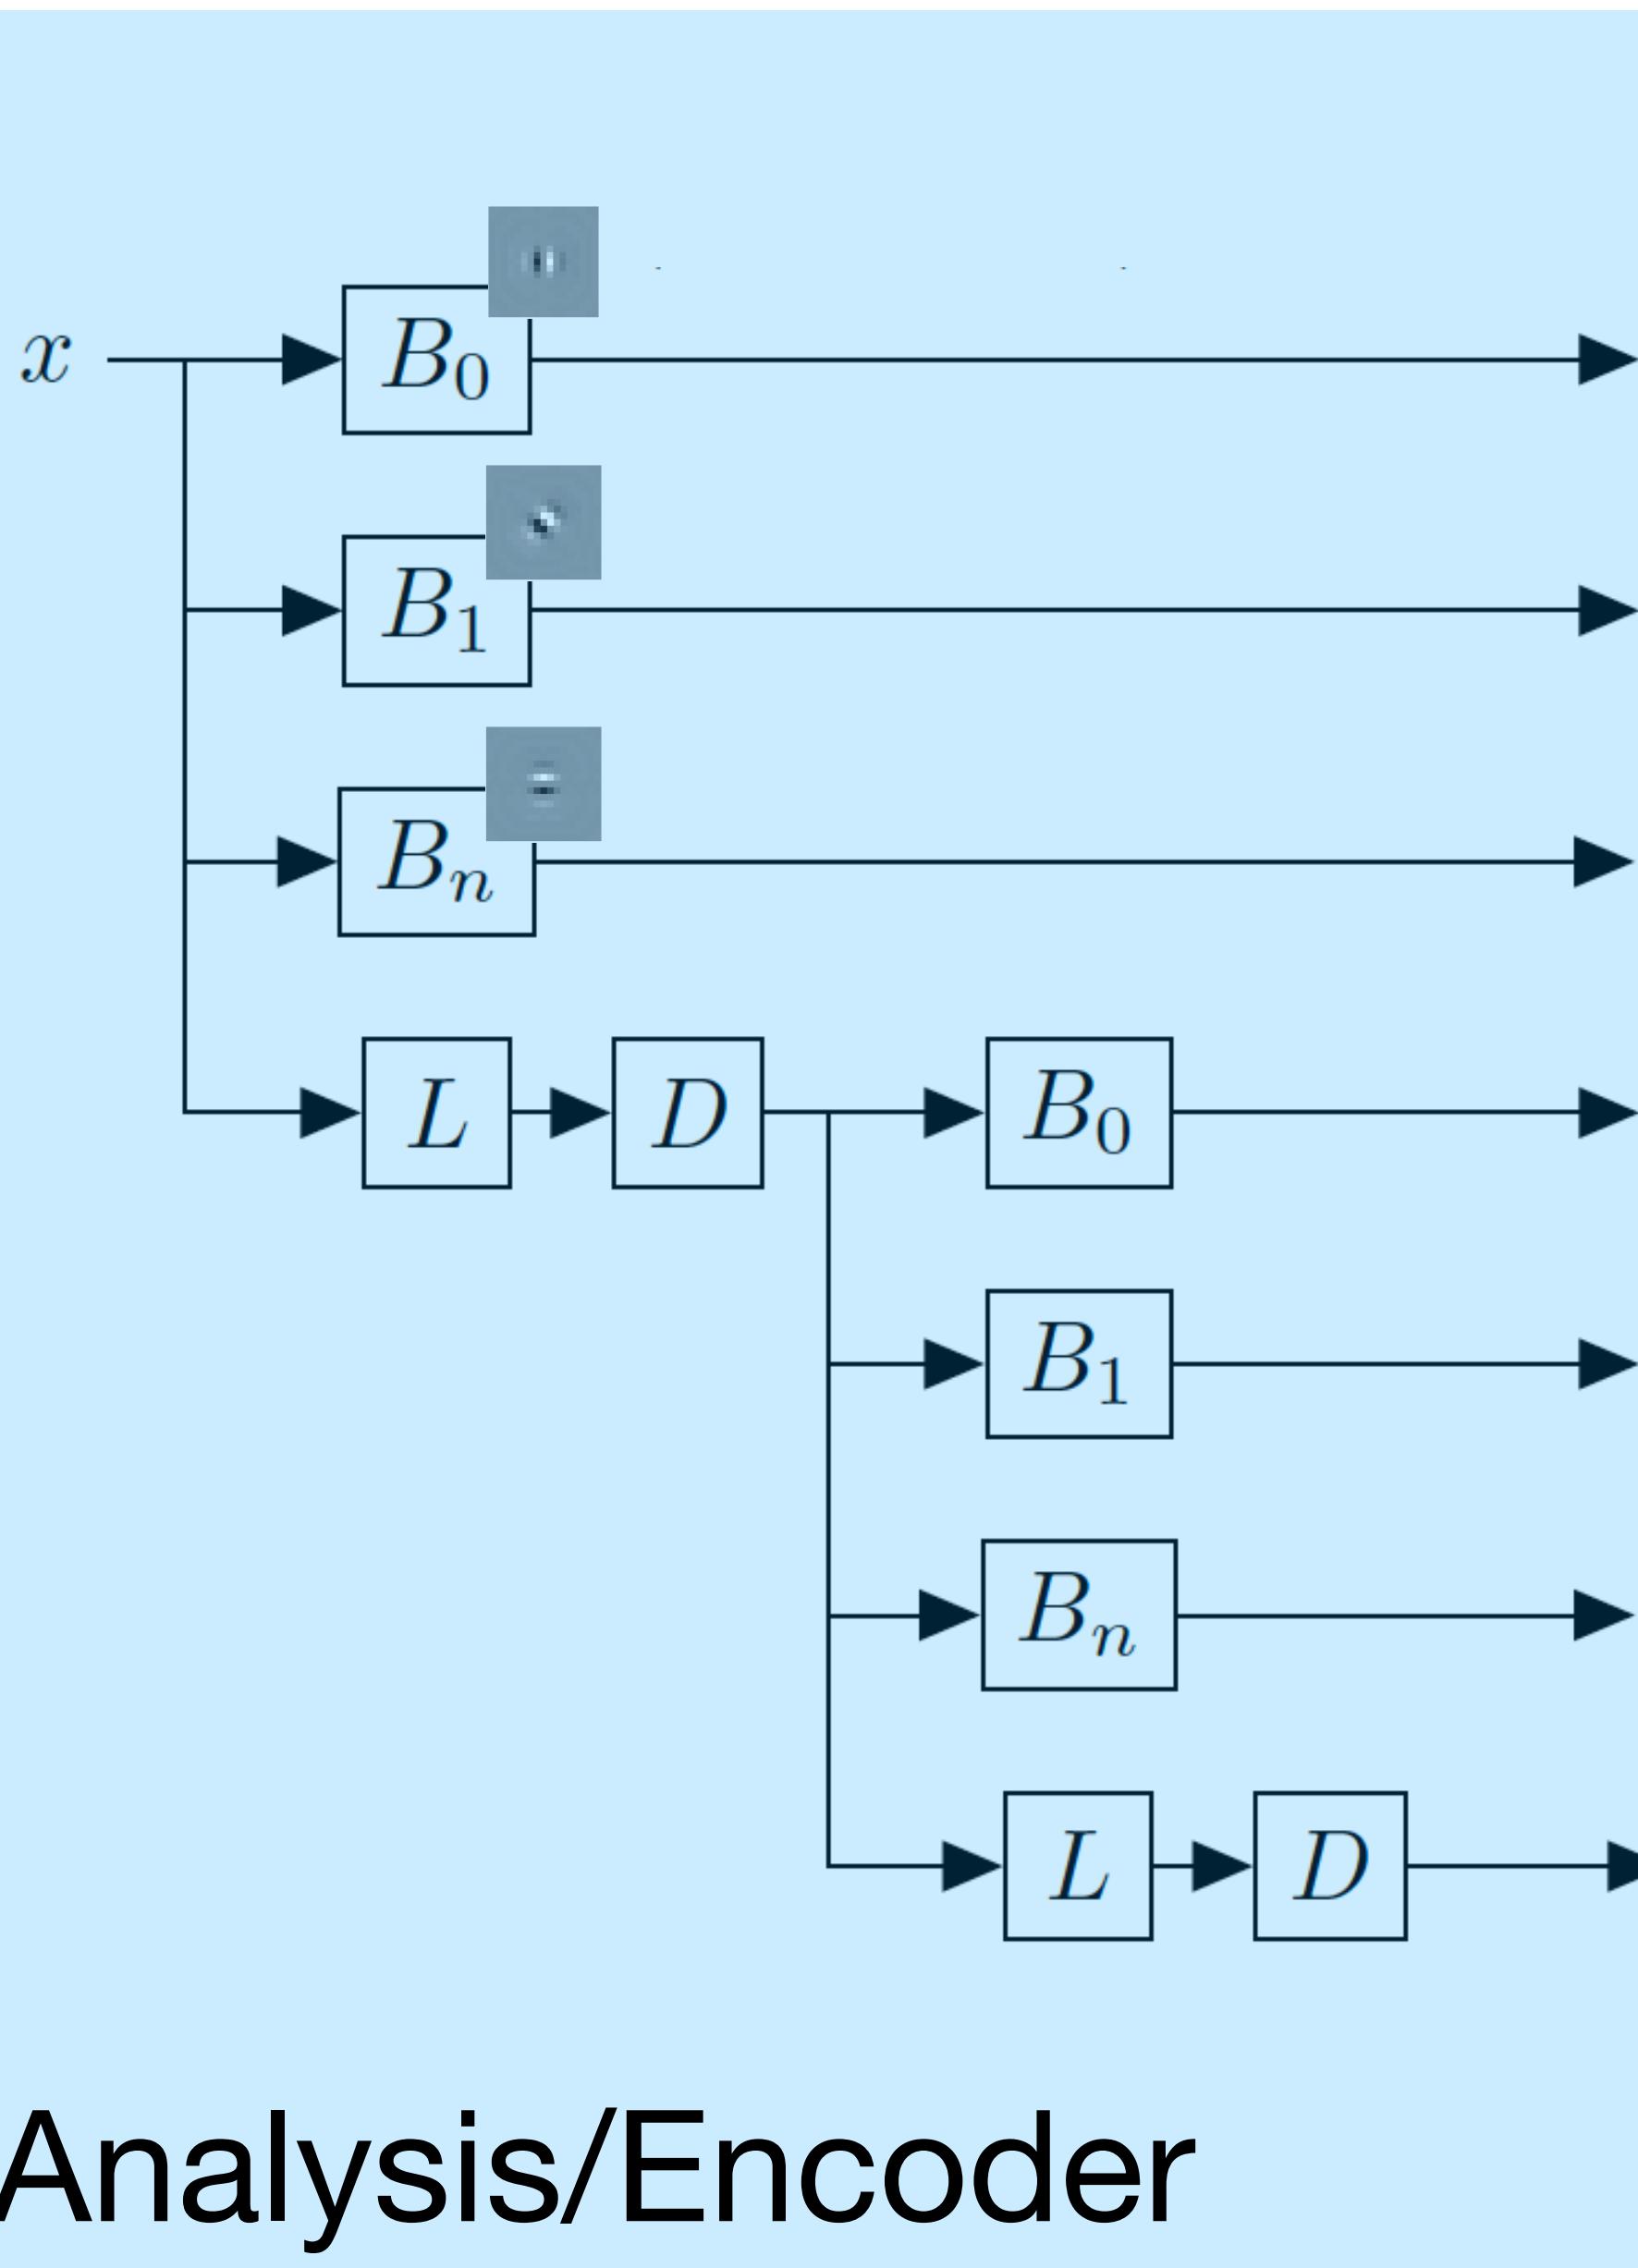


# Steerable Pyramid



3 scales, 4 orientations

# Steerable Pyramid



# Visual summary: Linear Image Transforms

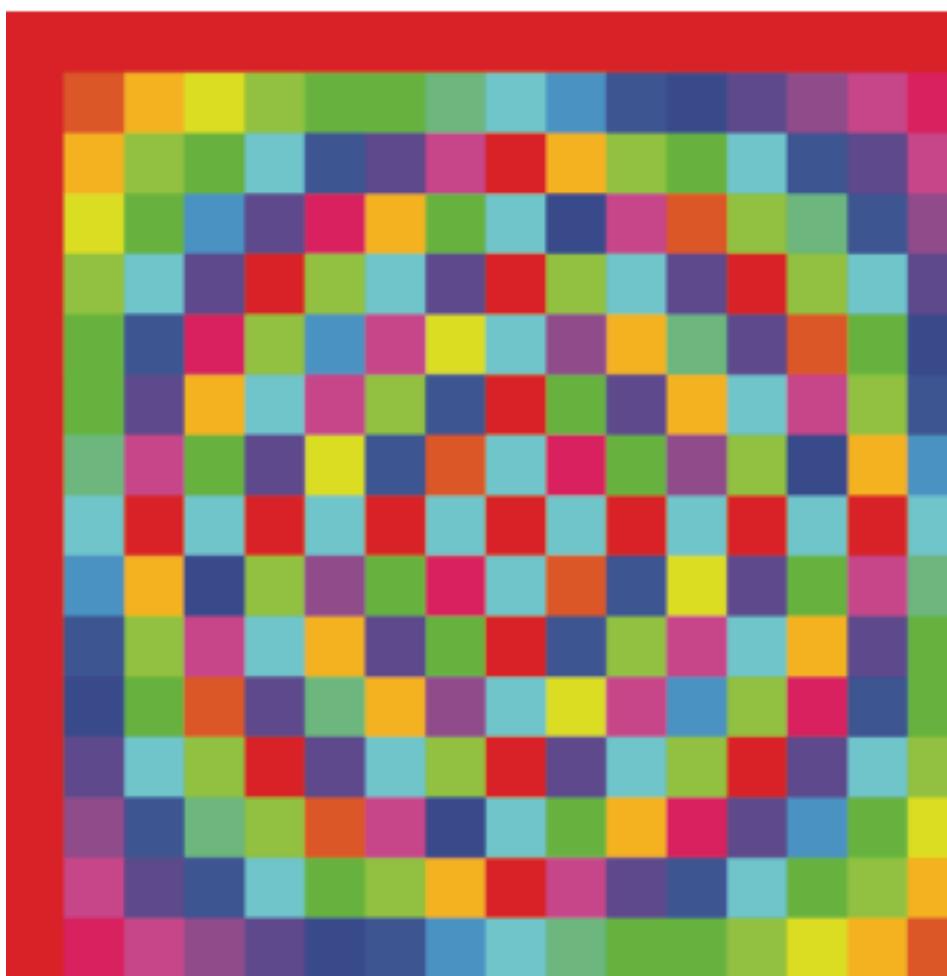
1297x256



Steerable Pyramid

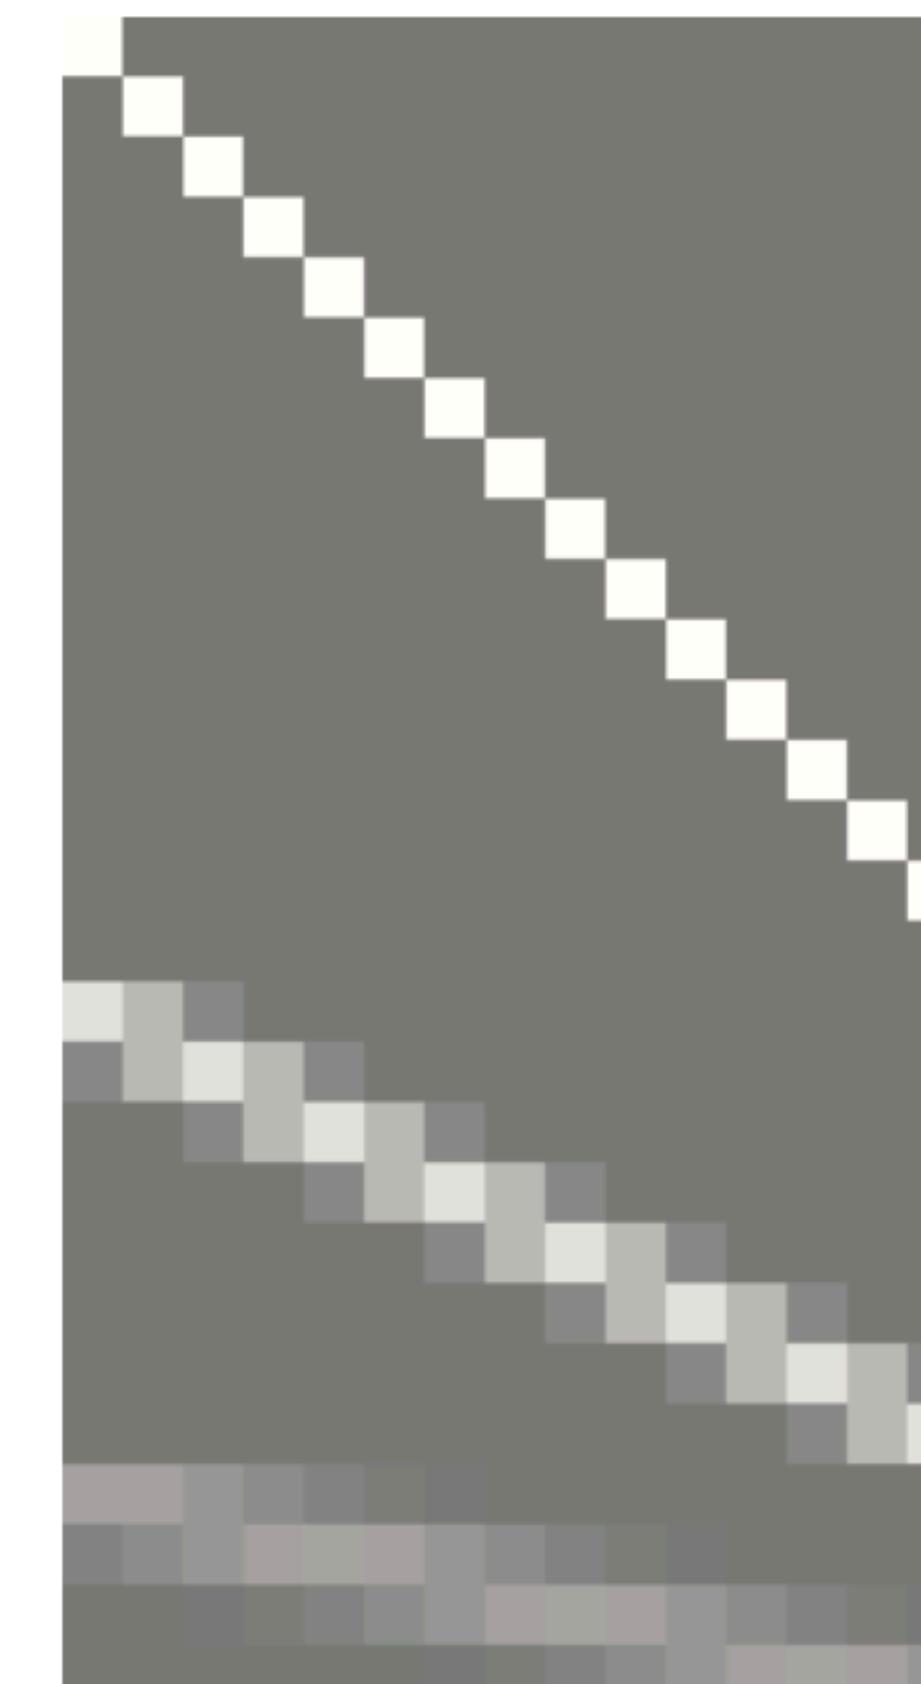
Fourier Transform

16x16



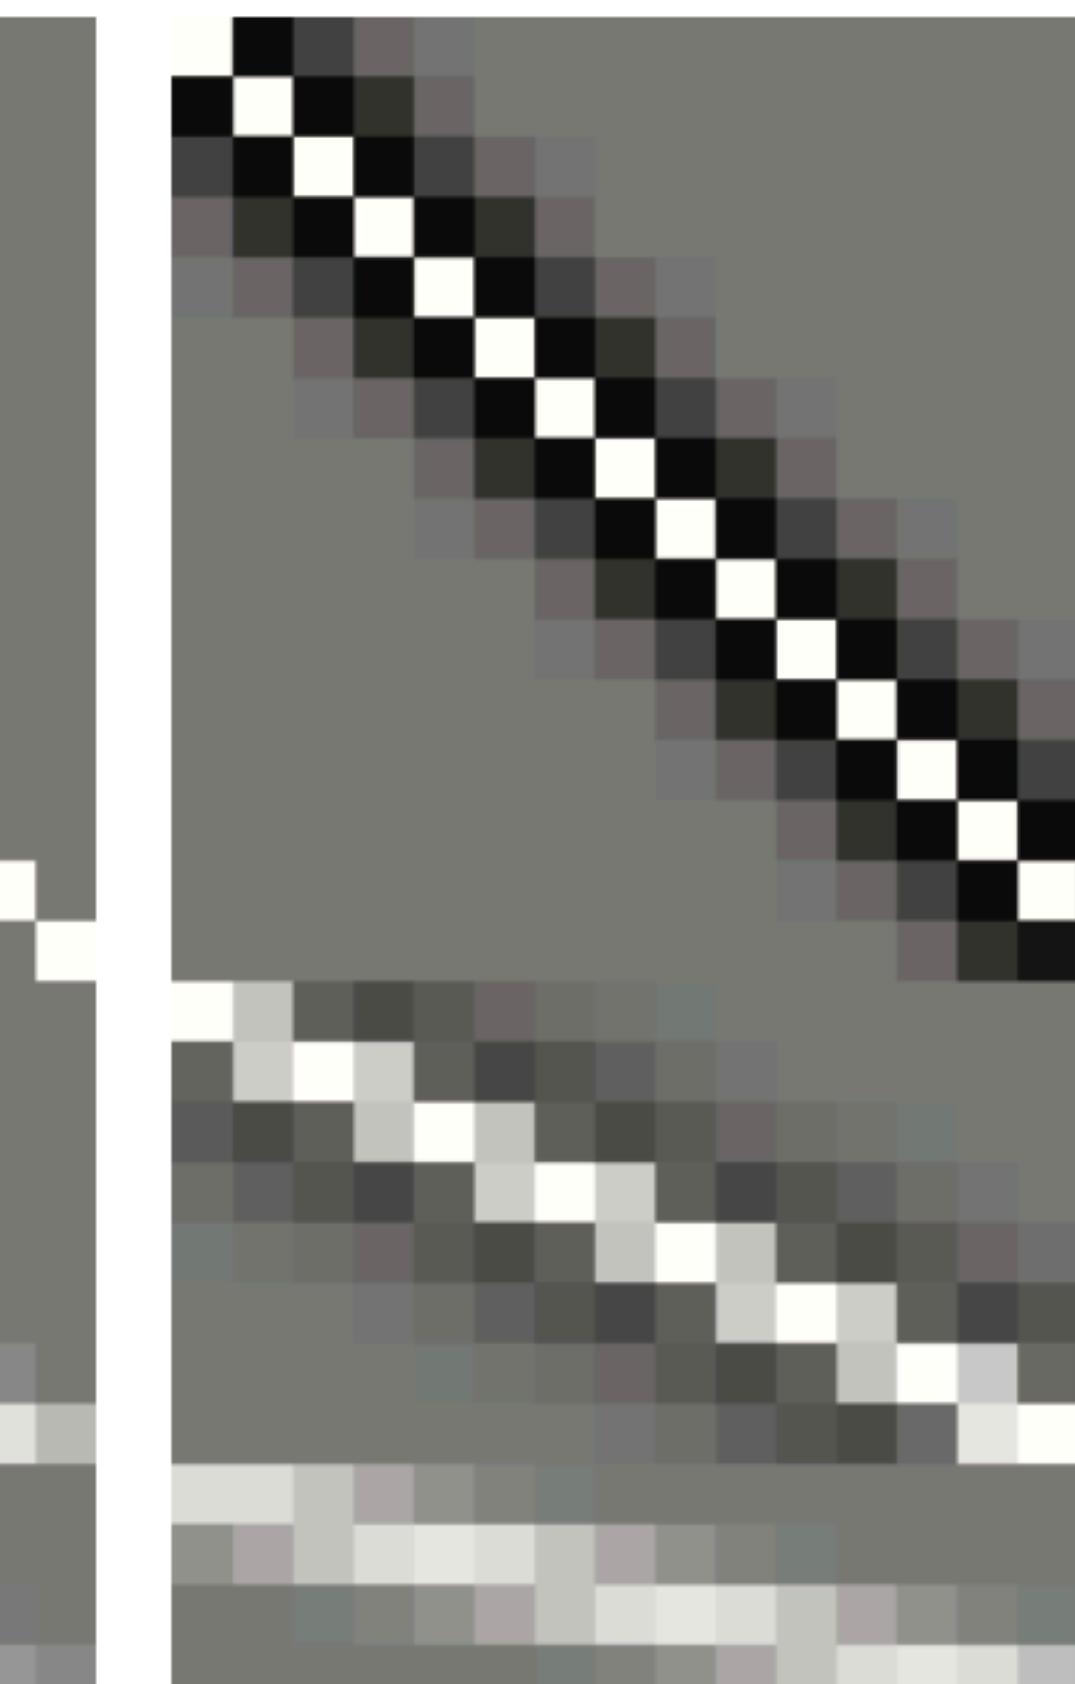
Gaussian Pyramid

28x16



Laplacian Pyramid

28x16



1D

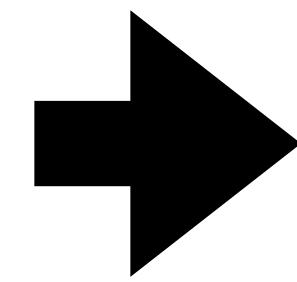
2D

# Steerable pyramid applications

- Texture recognition
- Image compression
- Noise removal
- Computing image features (e.g., SIFT)
- Object recognition

## Dall-E 2

An astronaut riding a horse  
in a photorealistic style

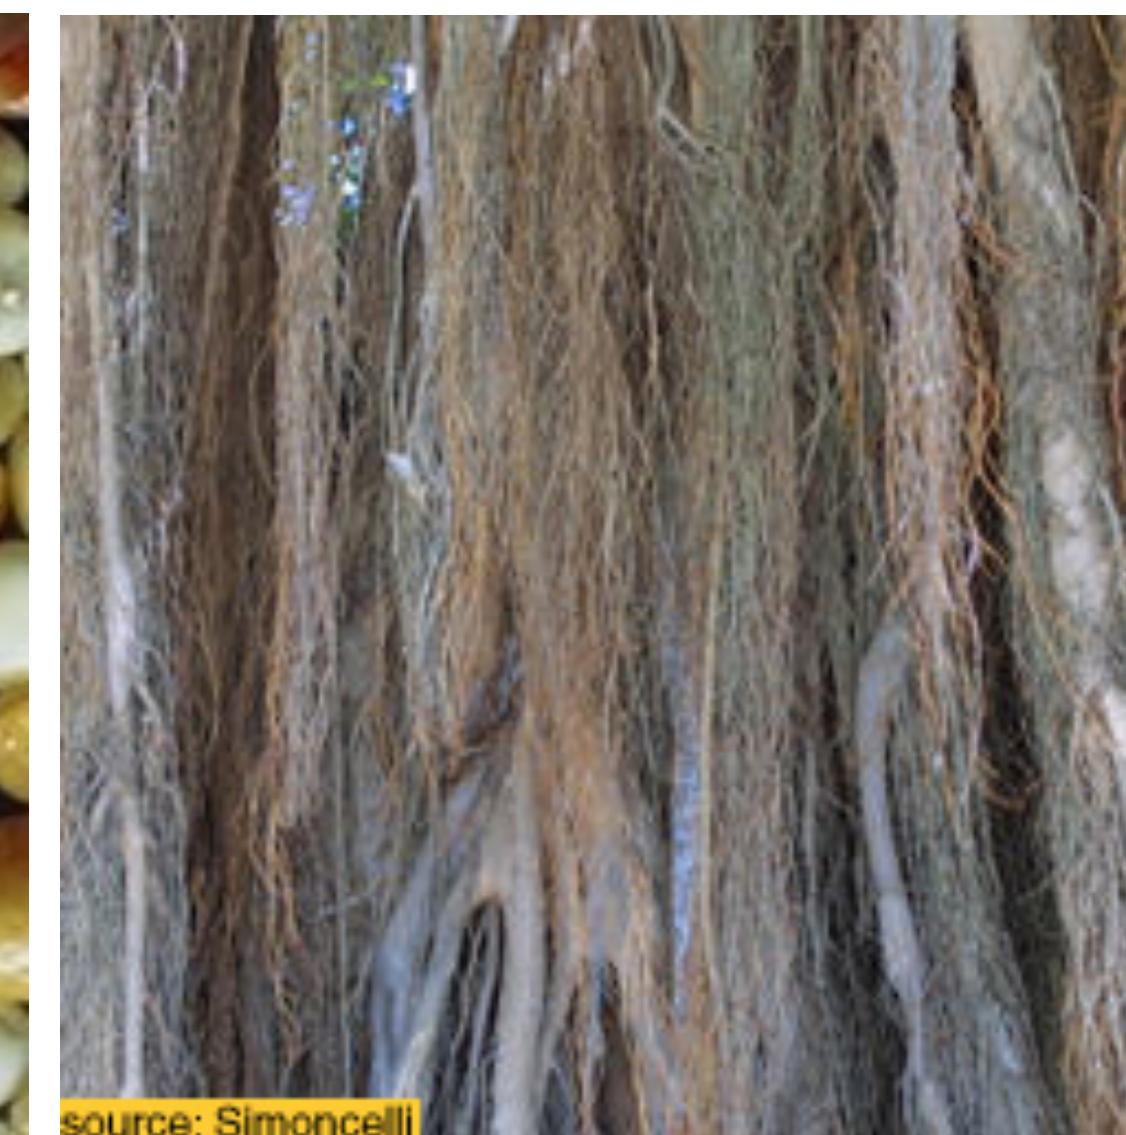
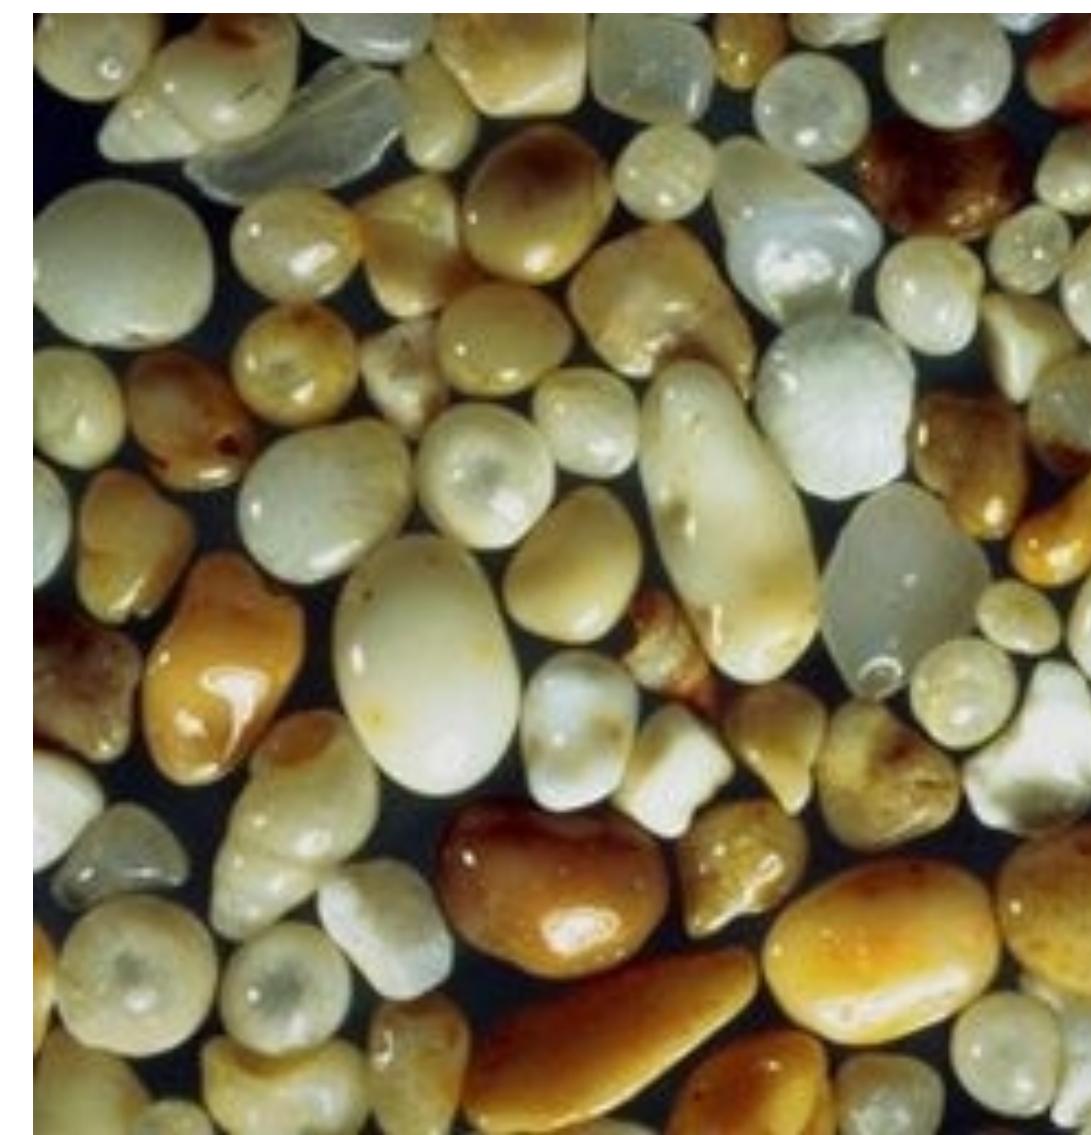


# Making textures



# Textures

Stationary  
Stochastic





## REVIEW ARTICLES

# Textons, the elements of texture perception, and their interactions

Bela Julesz

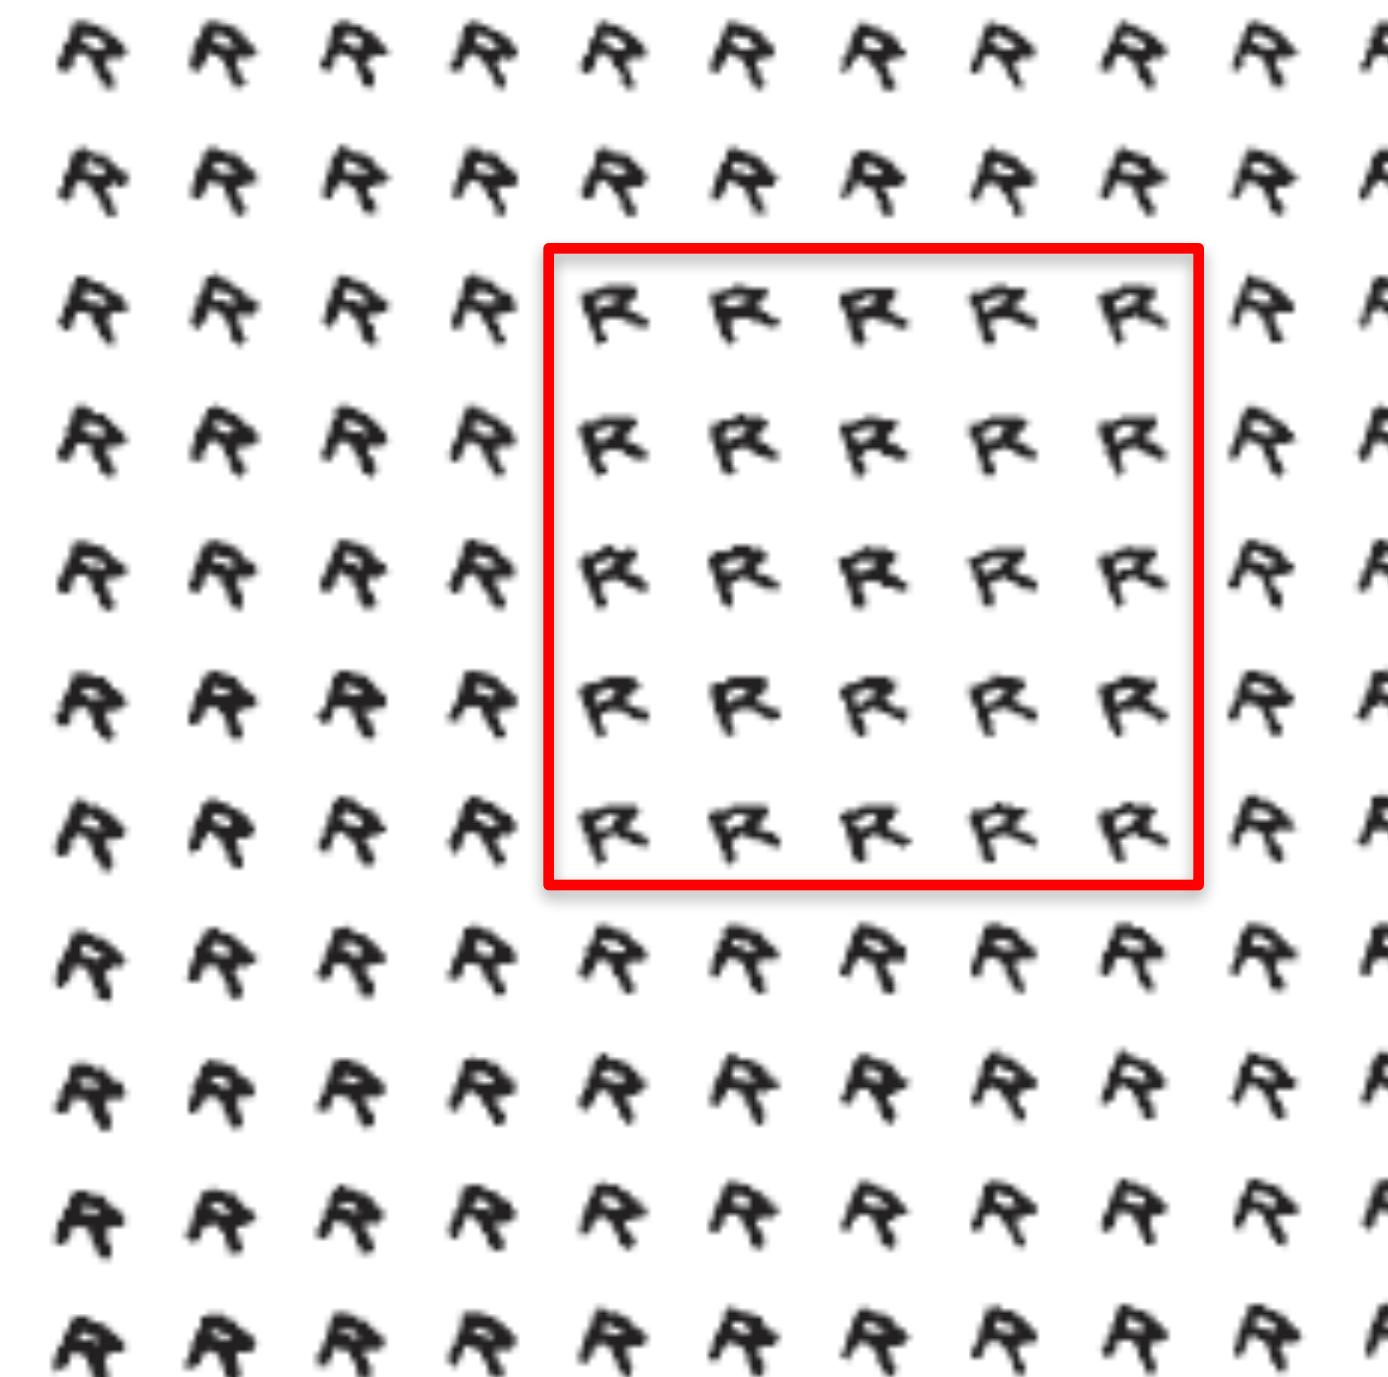
Bell Laboratories, Murray Hill, New Jersey 07974, USA

*Research with texture pairs having identical second-order statistics has revealed that the pre-attentive texture discrimination system cannot globally process third- and higher-order statistics, and that discrimination is the result of a few local conspicuous features, called textons. It seems that only the first-order statistics of these textons have perceptual significance, and the relative phase between textons cannot be perceived without detailed scrutiny by focal attention.*



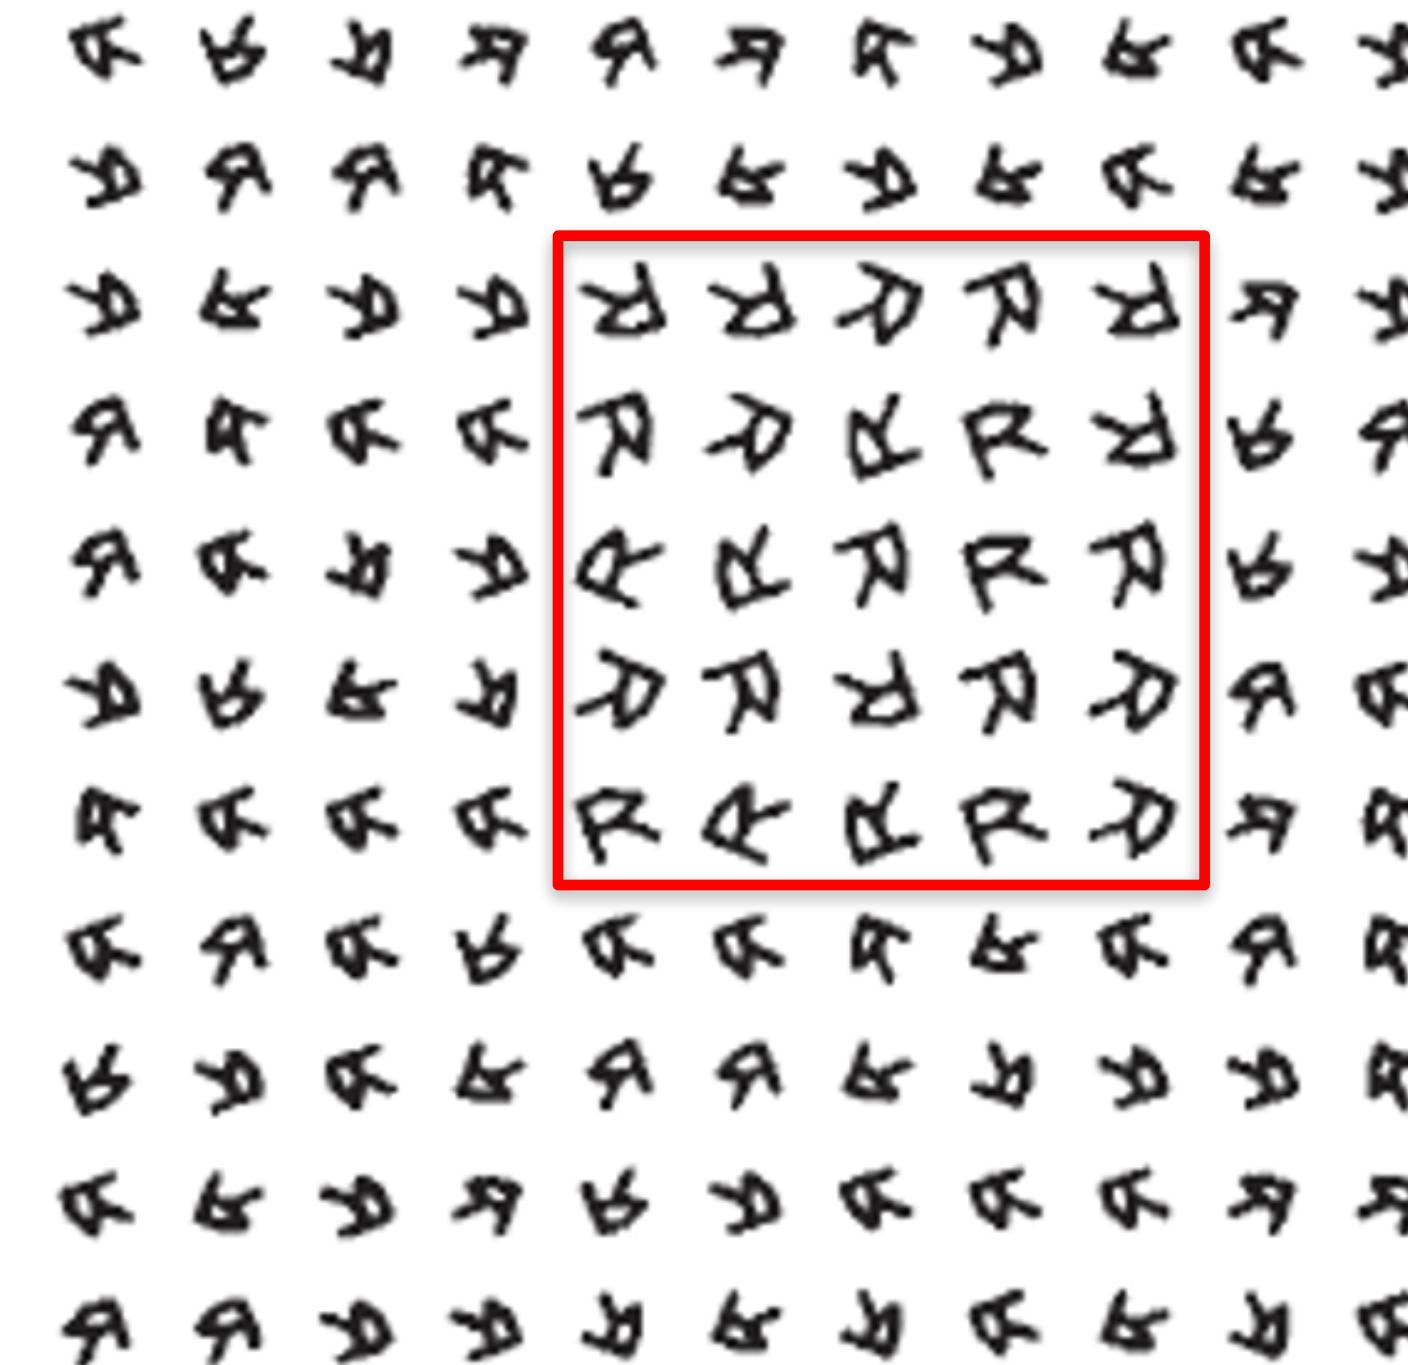
Bela Julesz, "Textons, the Elements of Texture Perception, and their Interactions". Nature 290: 91-97. March, 1981.

# Pre-attentive texture discrimination



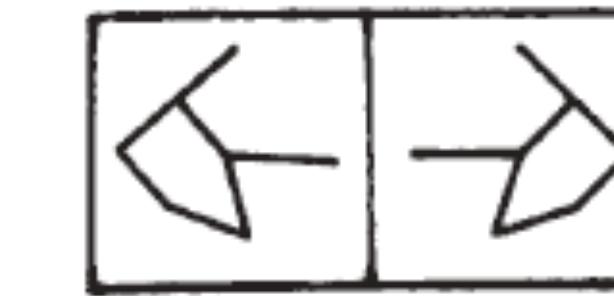
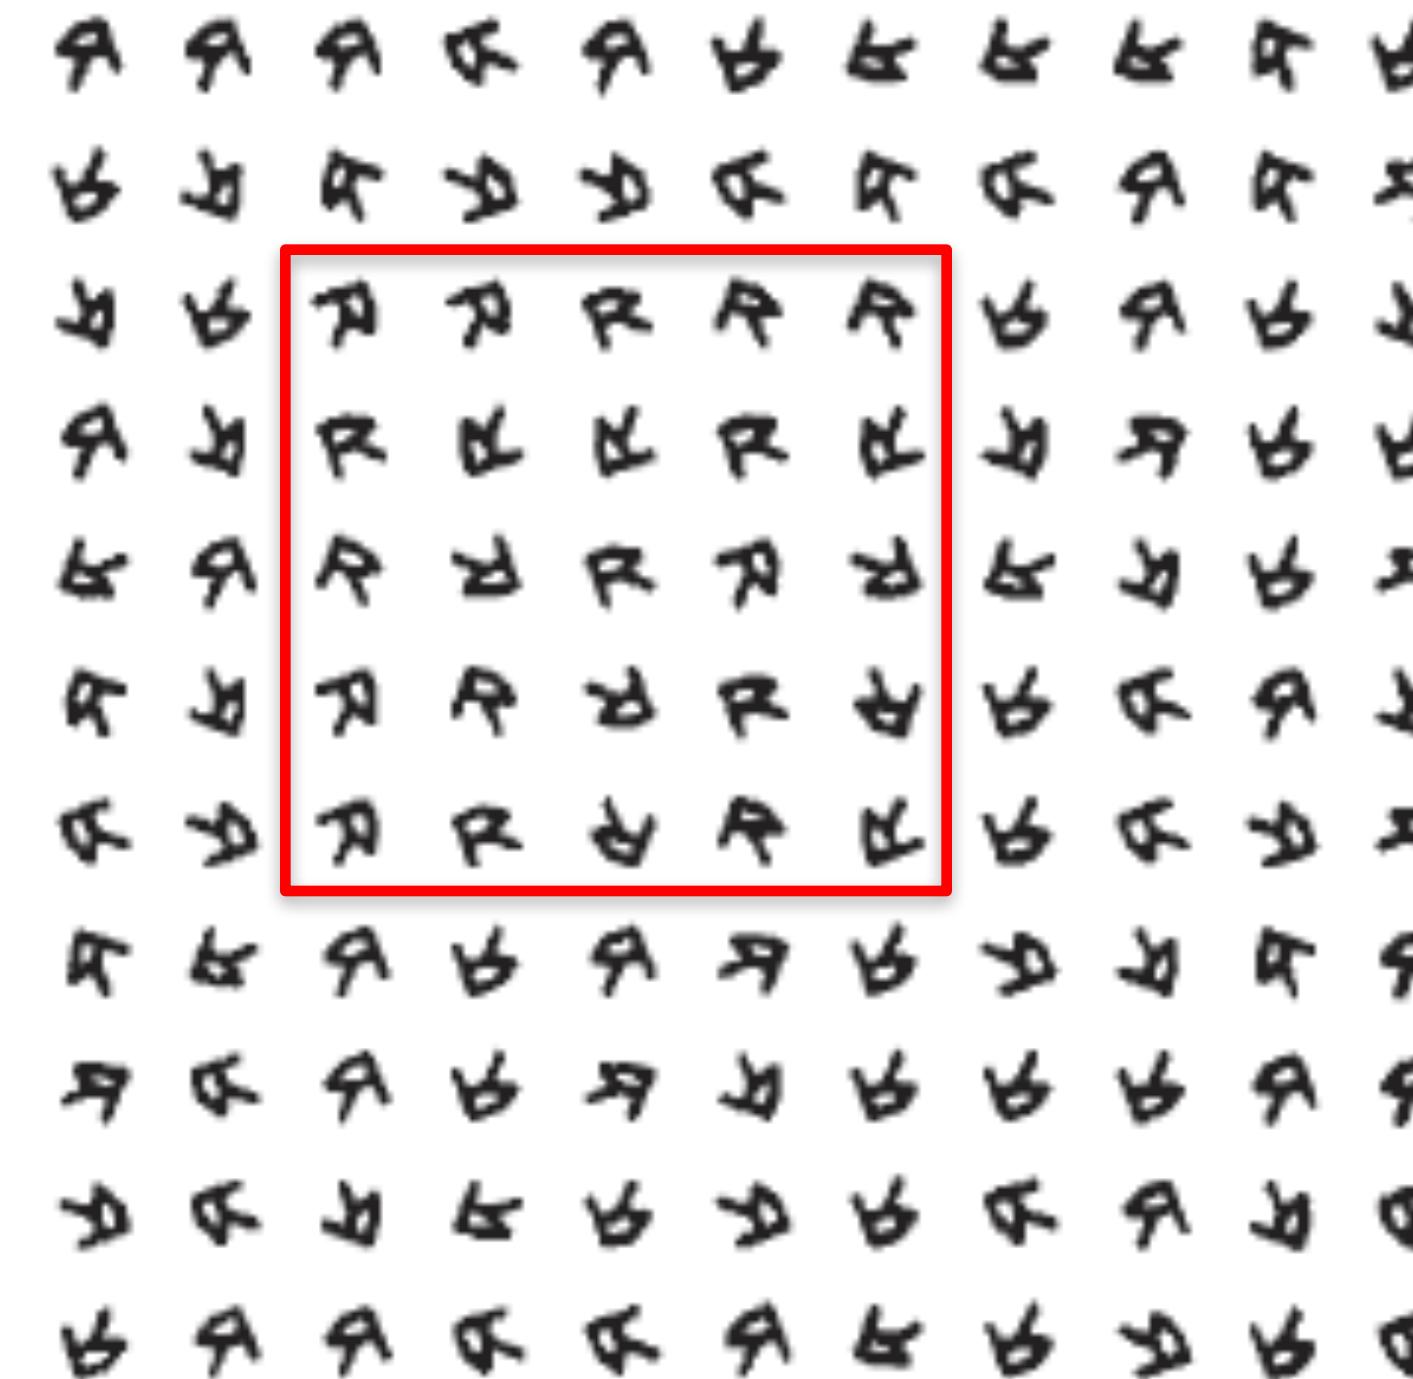
Bela Julesz, "Textons, the Elements of Texture Perception, and their Interactions". Nature 290: 91-97. March, 1981.

# Pre-attentive texture discrimination



Bela Julesz, "Textons, the Elements of Texture Perception, and their Interactions". Nature 290: 91-97. March, 1981.

# Pre-attentive texture discrimination

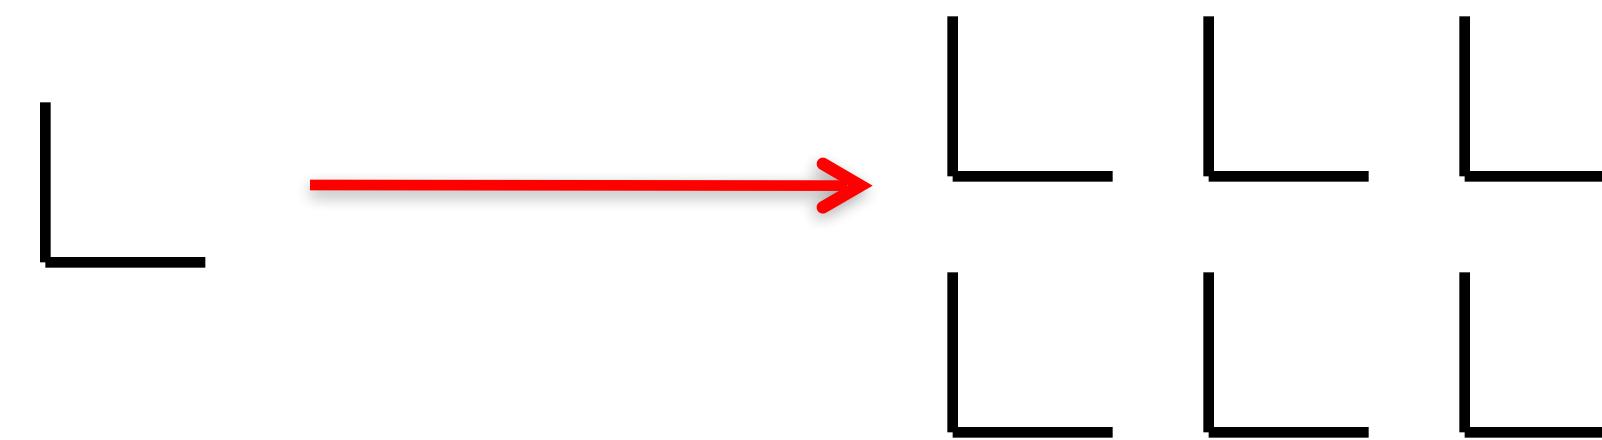


This texture pair is pre-attentively indistinguishable. Why?

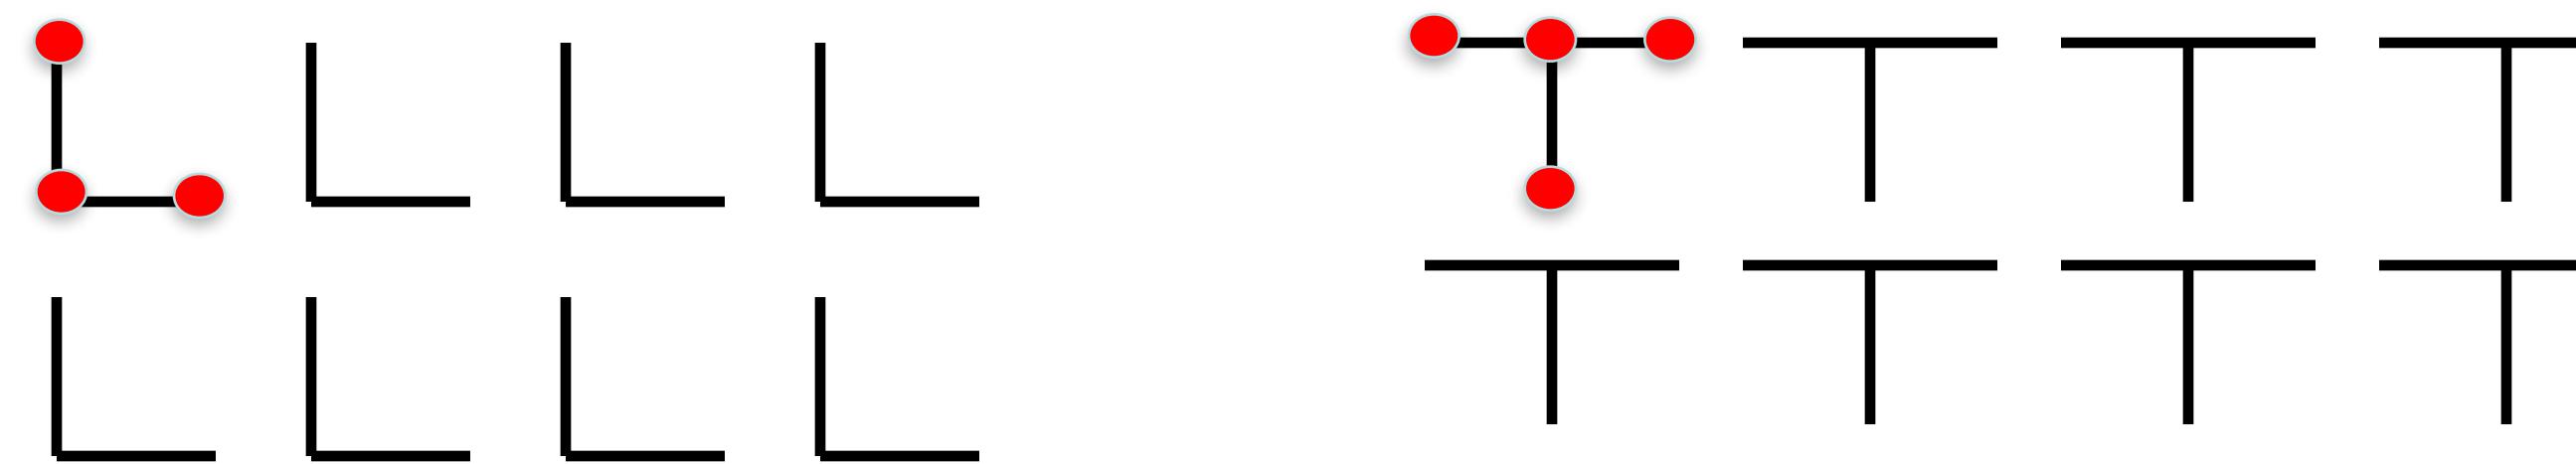
Bela Julesz, "Textons, the Elements of Texture Perception, and their Interactions". Nature 290: 91-97. March, 1981.

# Julesz - Textons

Textons: fundamental texture elements.



Textons might be represented by features such as terminators, corners, and intersections within the patterns...



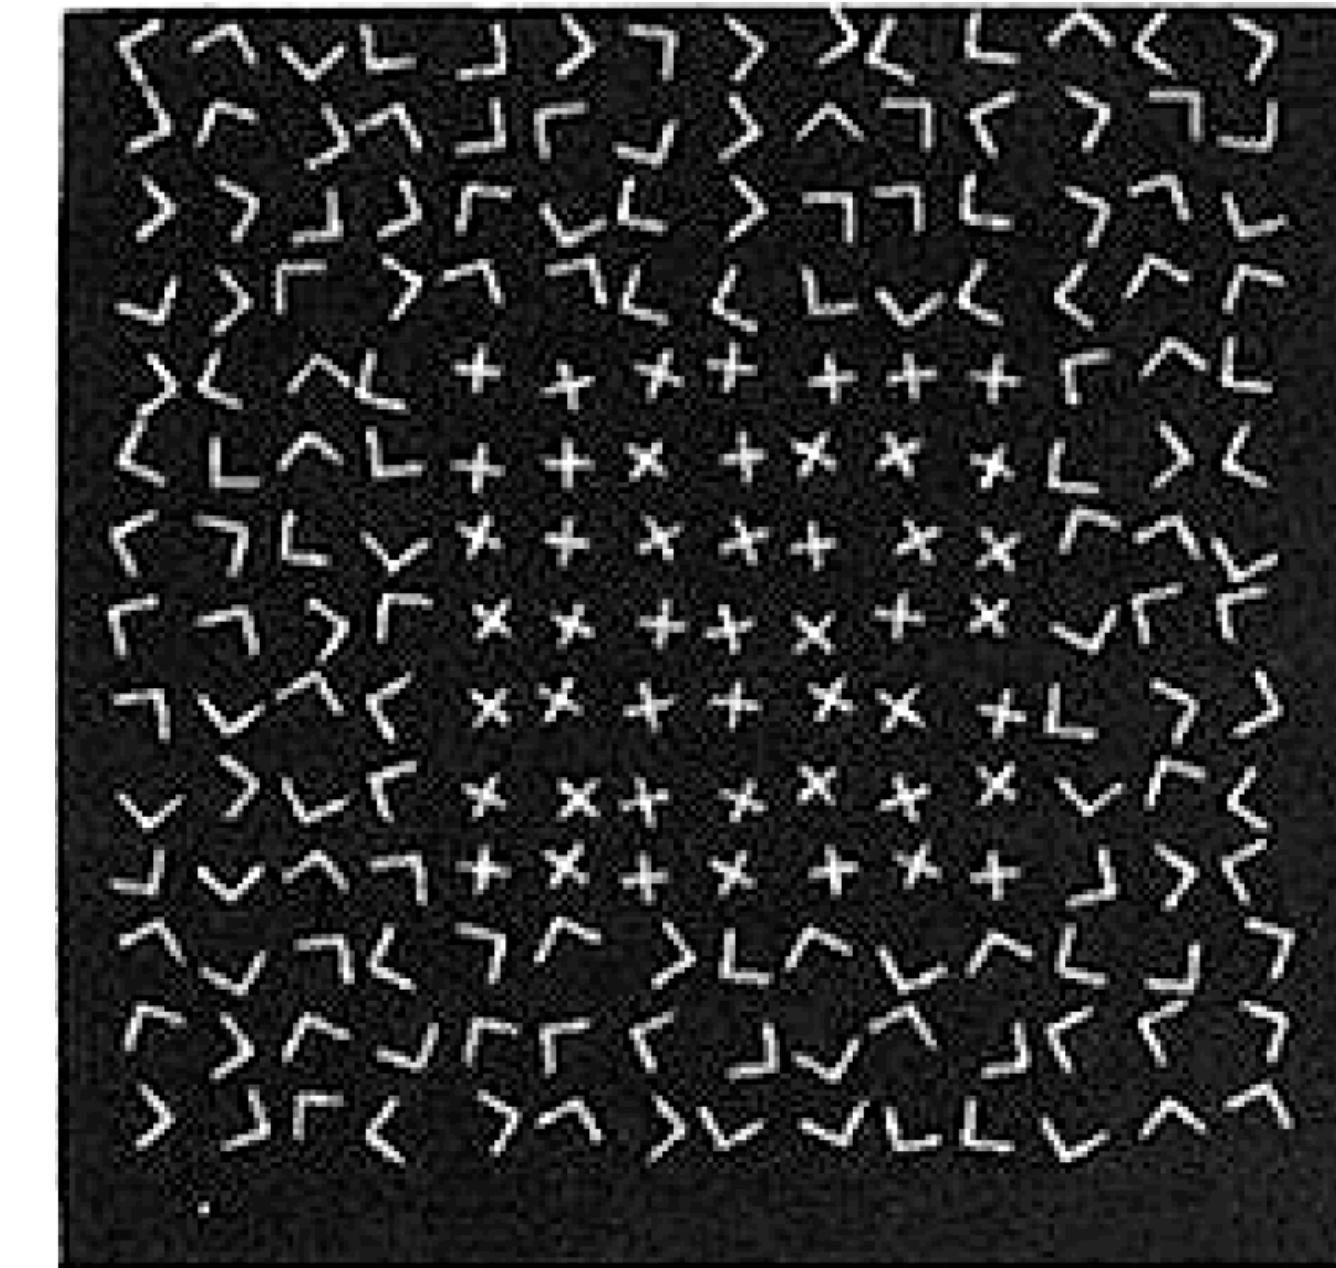
Nature, Vol. 333. No. 6171. pp. 363-364, 26 May 1988

# Early vision and texture perception

James R. Bergen\* & Edward H. Adelson\*\*

\* SRI David Sarnoff Research Center, Princeton,  
New Jersey 08540, USA

\*\* Media Lab and Department of Brain and Cognitive Science,



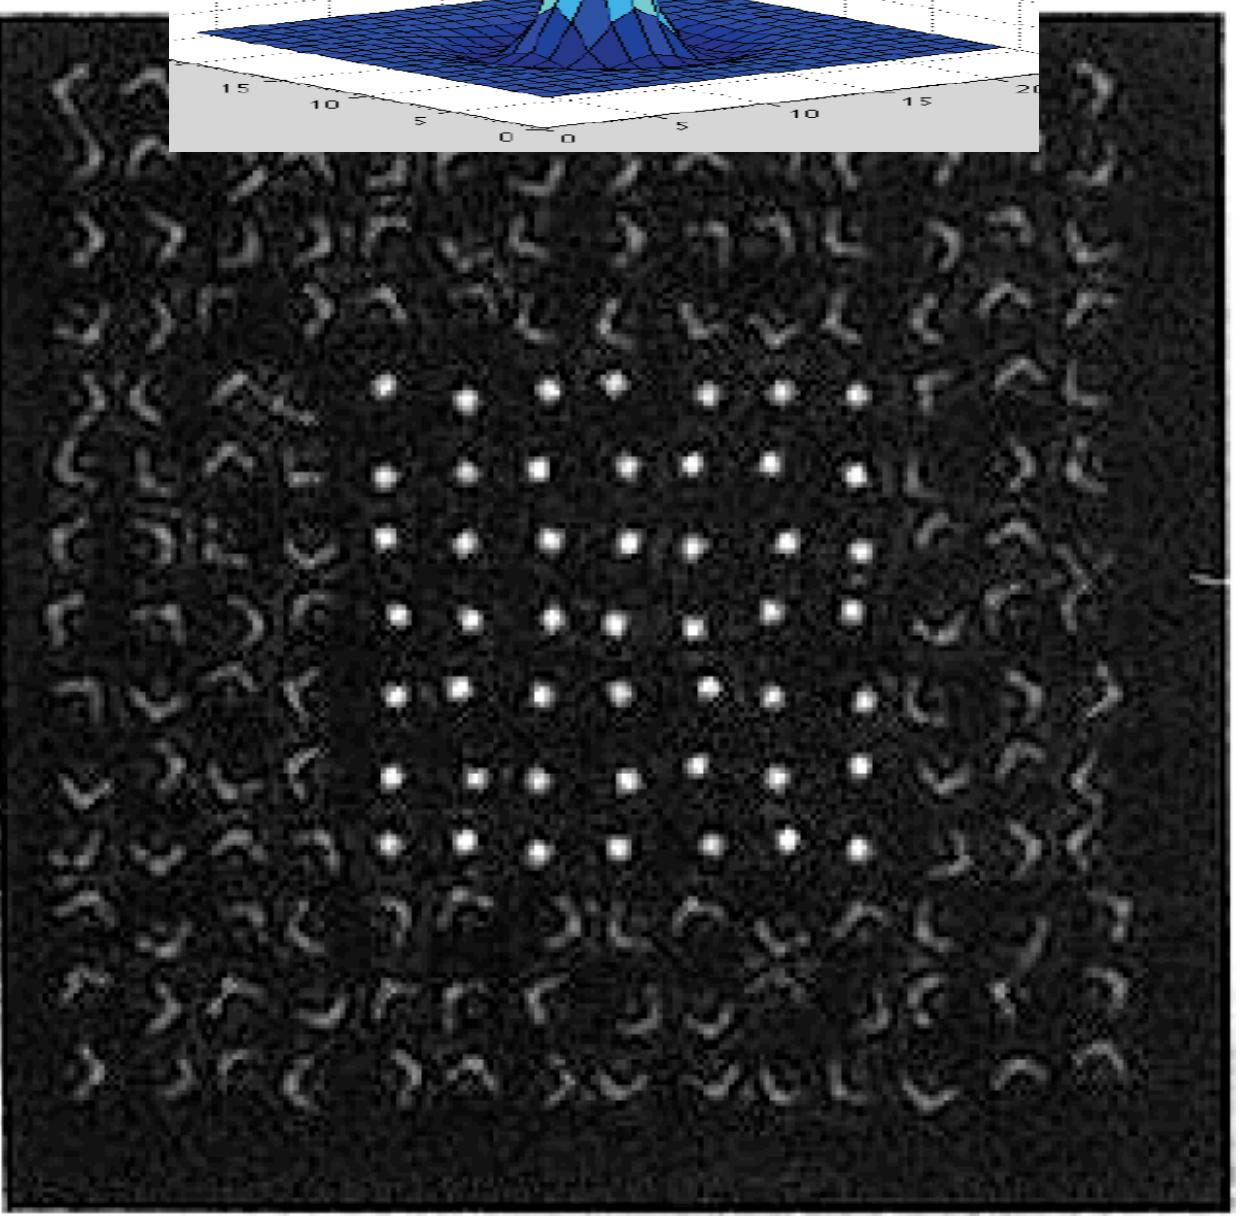
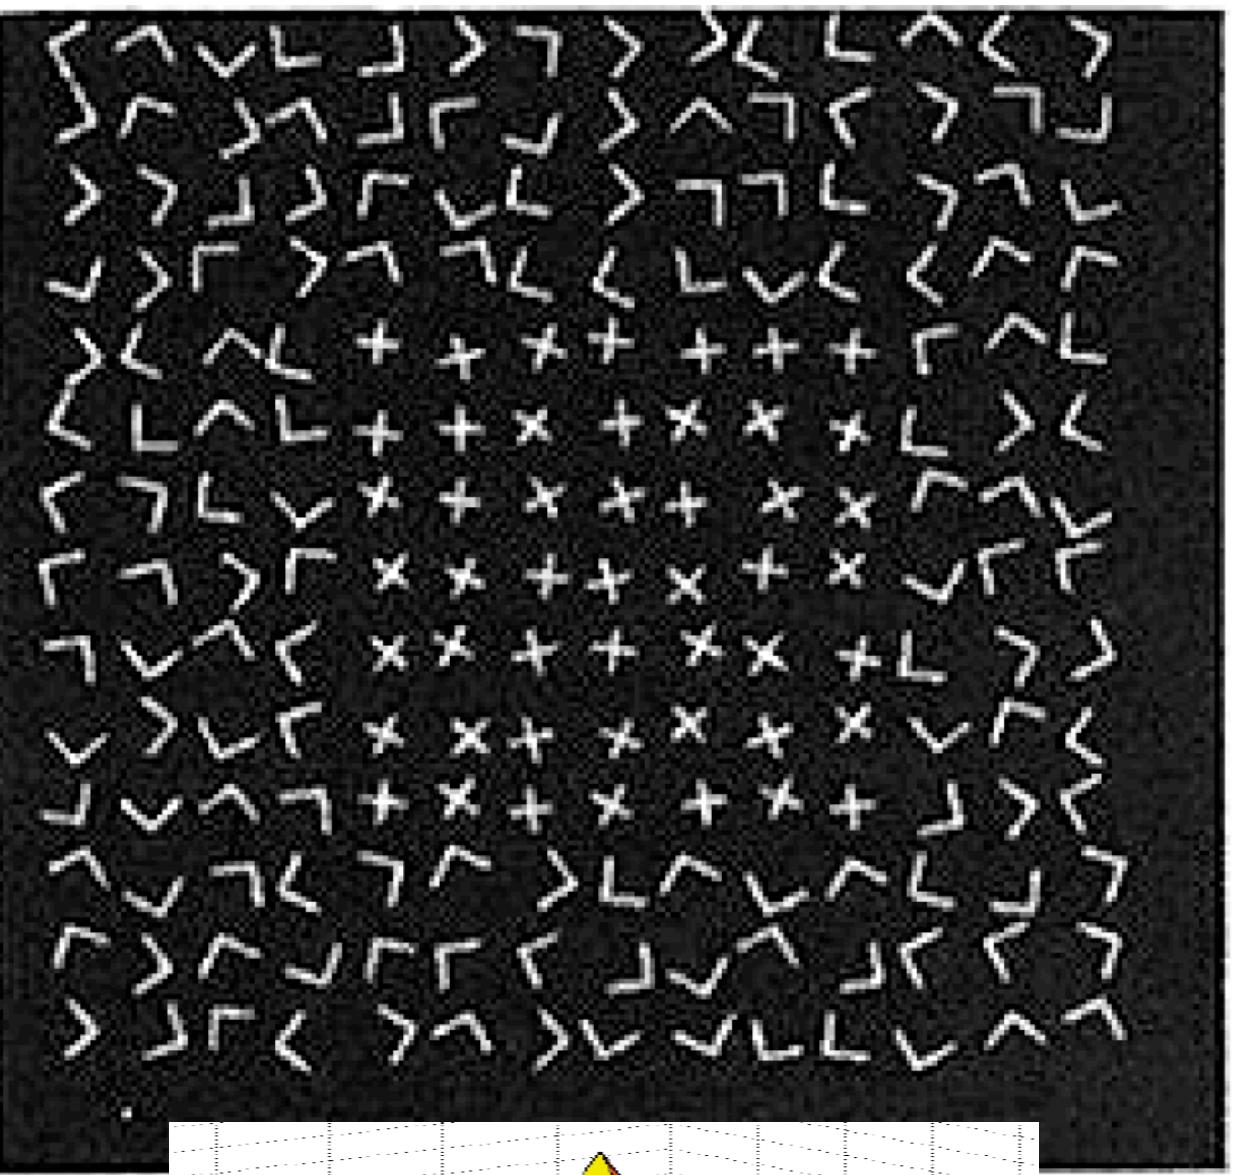
Observation: the Xs  
look smaller than the Ls.

**“We note here that simpler, lower-level mechanisms tuned for size may be sufficient to explain this discrimination.”**

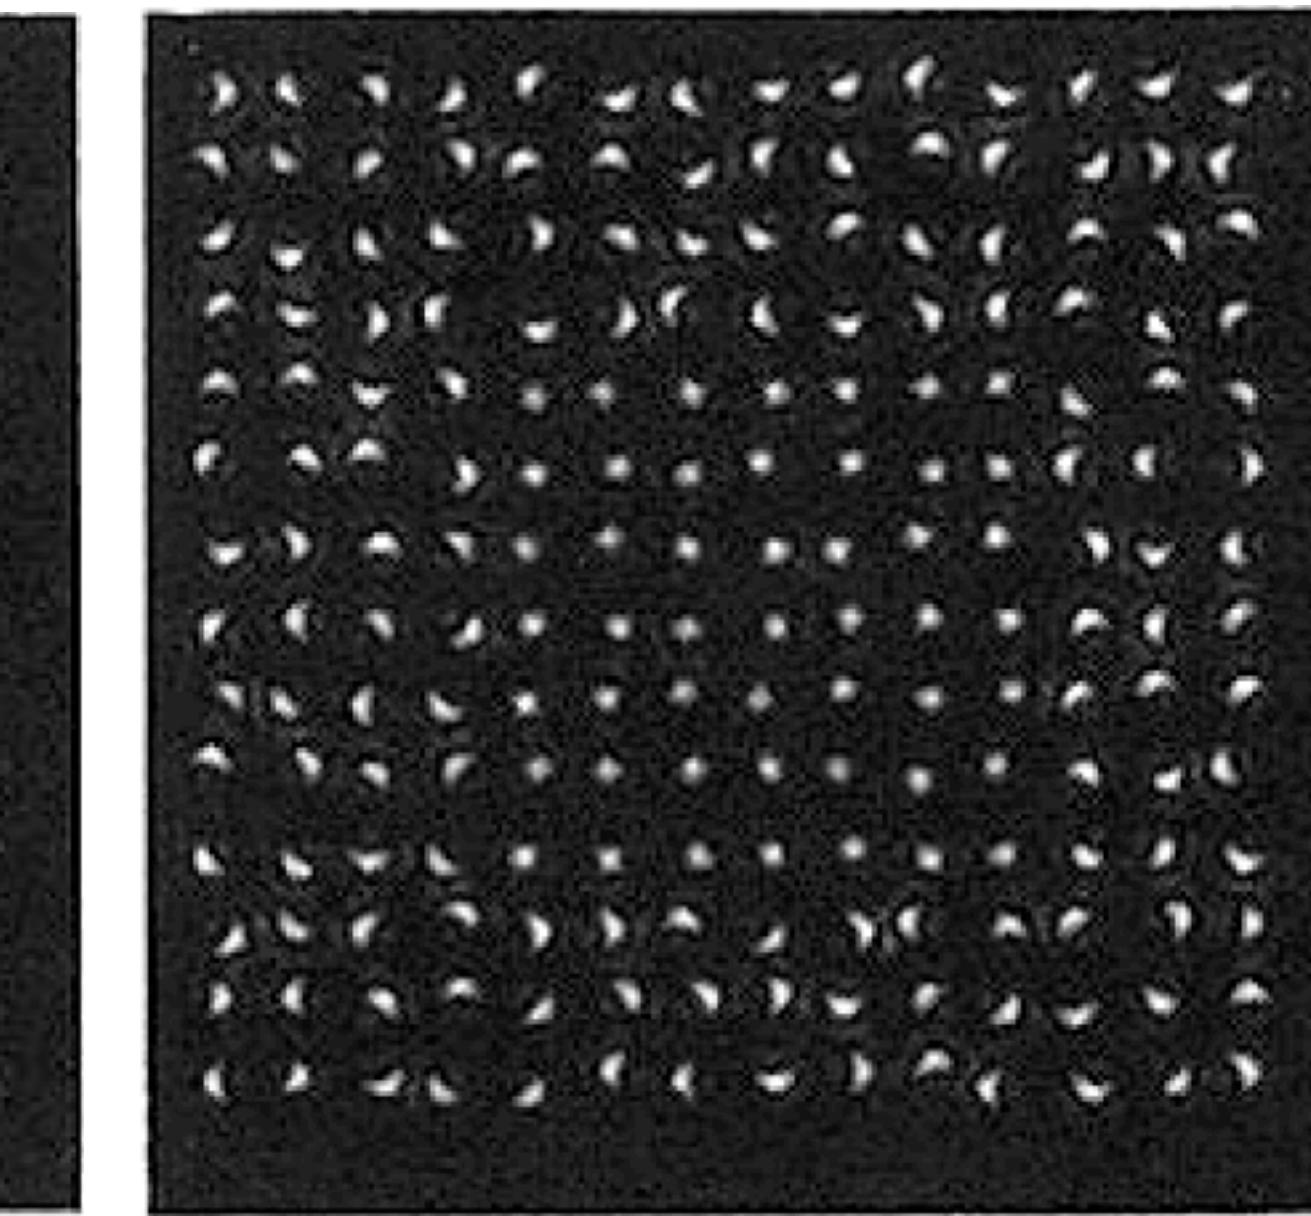
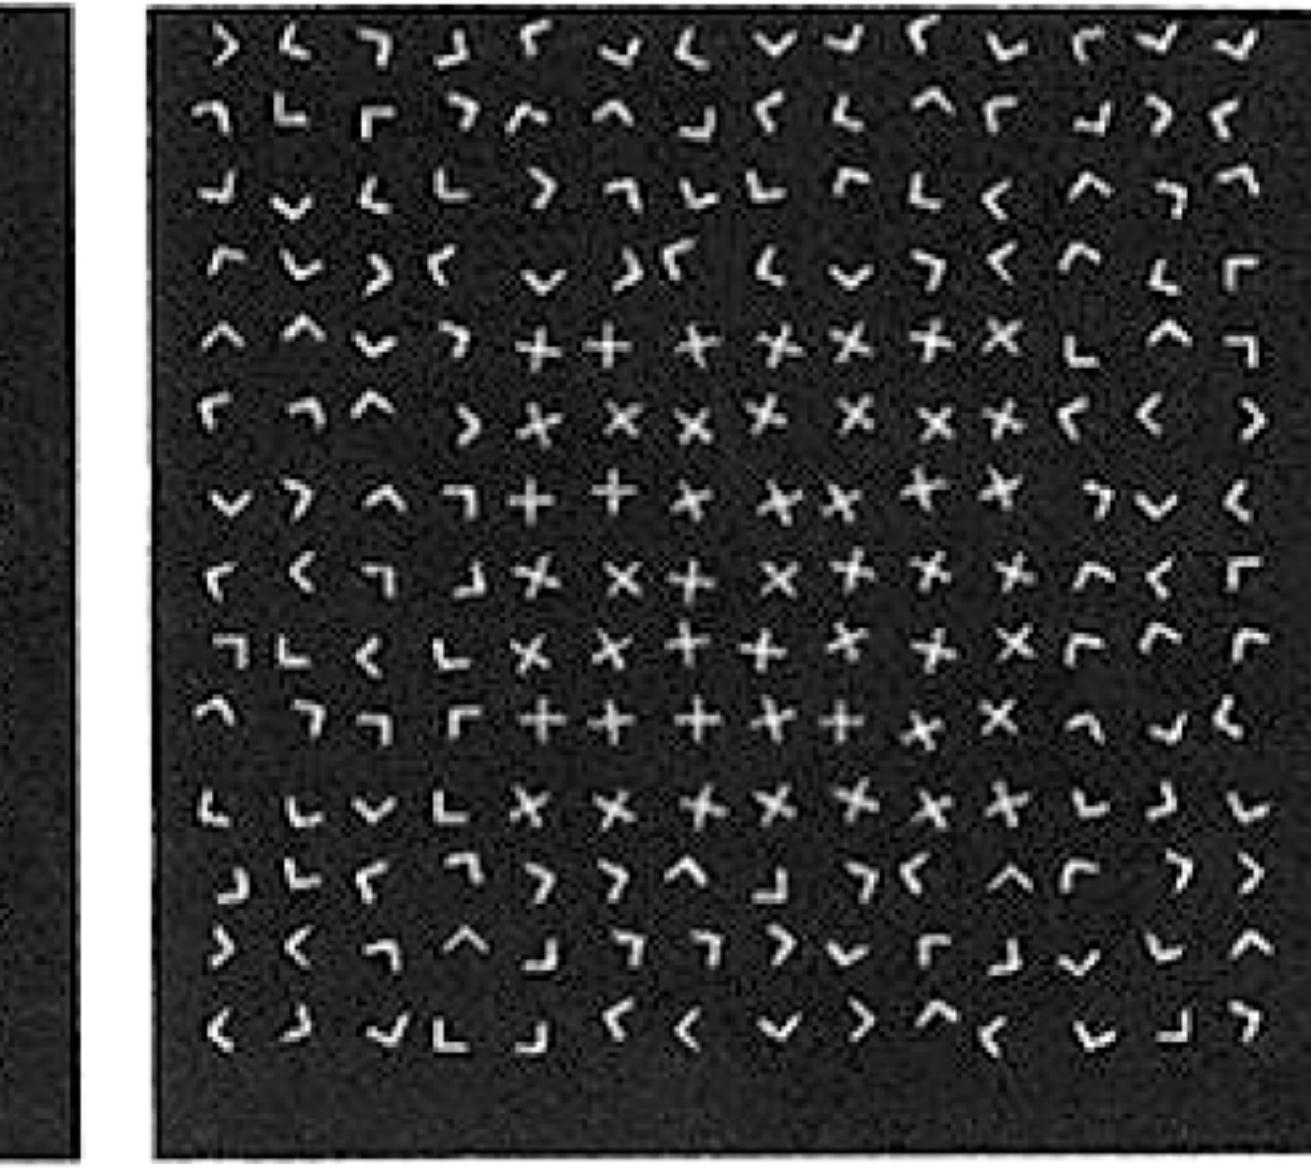
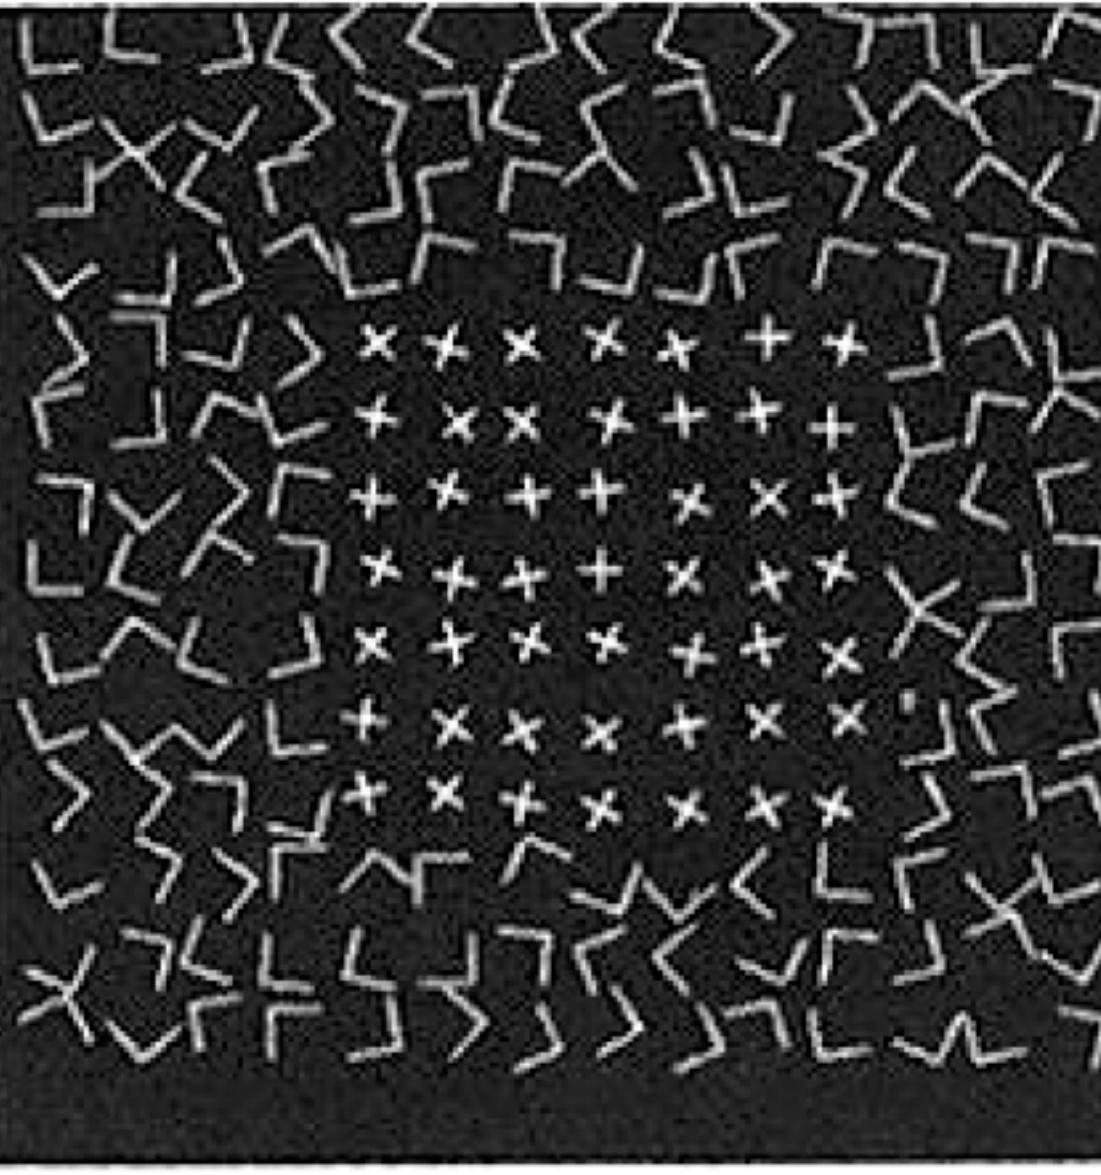
# Early vision and texture perception

James R. Bergen\* & Edward H. Adelson\*\*

Ls 25% larger  
contrast adjusted to keep mean constant



Ls 25% shorter



# Pyramid-Based Texture Analysis/Synthesis

David J. Heeger\*  
Stanford University

James R. Bergen†  
SRI David Sarnoff Research Center

## Abstract

This paper describes a method for synthesizing images that match the texture appearance of a given digitized sample. This synthesis is completely automatic and requires only the “target” texture as input. It allows generation of as much texture as desired so that any object can be covered. It can be used to produce solid textures for creating textured 3-d objects without the distortions inherent in texture mapping. It can also be used to synthesize texture mixtures, images that look a bit like each of several digitized samples. The approach is based on a model of human texture perception, and has potential to be a practically useful tool for graphics applications.

## 1 Introduction

Computer renderings of objects with surface texture are more interesting and realistic than those without texture. Texture mapping [15] is a technique for adding the appearance of surface detail by wrapping or projecting a digitized texture image onto a surface. Digitized textures can be obtained from a variety of sources, e.g., cropped from a photoCD image, but the resulting texture chip may not have the desired size or shape. To cover a large object you may need to repeat the texture; this can lead to unacceptable artifacts either in the form of visible seams, visible repetition, or both.

Texture mapping suffers from an additional fundamental problem: often there is no natural map from the (planar) texture image to the geometry/topology of the surface, so the texture may be distorted unnaturally when mapped. There are some partial solutions to this distortion problem [15] but there is no universal solution for mapping an image onto an arbitrarily shaped surface.

An alternative to texture mapping is to create (paint) textures by hand directly onto the 3-d surface model [14], but this process is both very labor intensive and requires considerable artistic skill.

Another alternative is to use computer-synthesized textures so that as much texture can be generated as needed. Furthermore, some of the synthesis techniques produce textures that tile seamlessly.

Using synthetic textures, the distortion problem has been solved in two different ways. First, some techniques work by synthesizing texture directly on the object surface (e.g., [31]). The second solution is to use *solid textures* [19, 23, 24]. A solid texture is a 3-d array of color values. A point on the surface of an object is colored by the value of the solid texture at the corresponding 3-d point. Solid texturing can be a very natural solution to the distortion problem:

there is no distortion because there is no mapping. However, existing techniques for synthesizing solid textures can be quite cumbersome. One must learn how to tweak the parameters or procedures of the texture synthesizer to get a desired effect.

This paper presents a technique for synthesizing an image (or solid texture) that matches the appearance of a given texture sample. The key advantage of this technique is that it works entirely from the example texture, requiring no additional information or adjustment. The technique starts with a digitized image and analyzes it to compute a number of texture parameter values. Those parameter values are then used to synthesize a new image (of any size) that looks (in its color and texture properties) like the original. The analysis phase is inherently two-dimensional since the input digitized images are 2-d. The synthesis phase, however, may be either two- or three-dimensional. For the 3-d case, the output is a solid texture such that planar slices through the solid look like the original scanned image. In either case, the (2-d or 3-d) texture is synthesized so that it tiles seamlessly.

## 2 Texture Models

Textures have often been classified into two categories, deterministic textures and stochastic textures. A deterministic texture is characterized by a set of primitives and a placement rule (e.g., a tile floor). A stochastic texture, on the other hand, does not have easily identifiable primitives (e.g., granite, bark, sand). Many real-world textures have some mixture of these two characteristics (e.g. woven fabric, woodgrain, plowed fields).

Much of the previous work on texture analysis and synthesis can be classified according to what type of texture model was used. Some of the successful texture models include reaction-diffusion [31, 34], frequency domain [17], fractal [9, 18], and statistical/random field [1, 6, 8, 10, 12, 13, 21, 26] models. Some (e.g., [10]) have used hybrid models that include a deterministic (or periodic) component and a stochastic component. In spite of all this work, scanned images and hand-drawn textures are still the principle source of texture maps in computer graphics.

This paper focuses on the synthesis of stochastic textures. Our approach is motivated by research on human texture perception. Current theories of texture discrimination are based on the fact that two textures are often difficult to discriminate when they produce a similar distribution of responses in a bank of (orientation and spatial-frequency selective) linear filters [2, 3, 7, 16, 20, 32]. The method described here, therefore, synthesizes textures by matching distributions (or histograms) of filter outputs. This approach depends on the principle (not entirely correct as we shall see) that all of the spatial information characterizing a texture image can be captured in the first order statistics of an appropriately chosen set of linear filter outputs. Nevertheless, this model (though incomplete) captures an interesting set of texture properties.

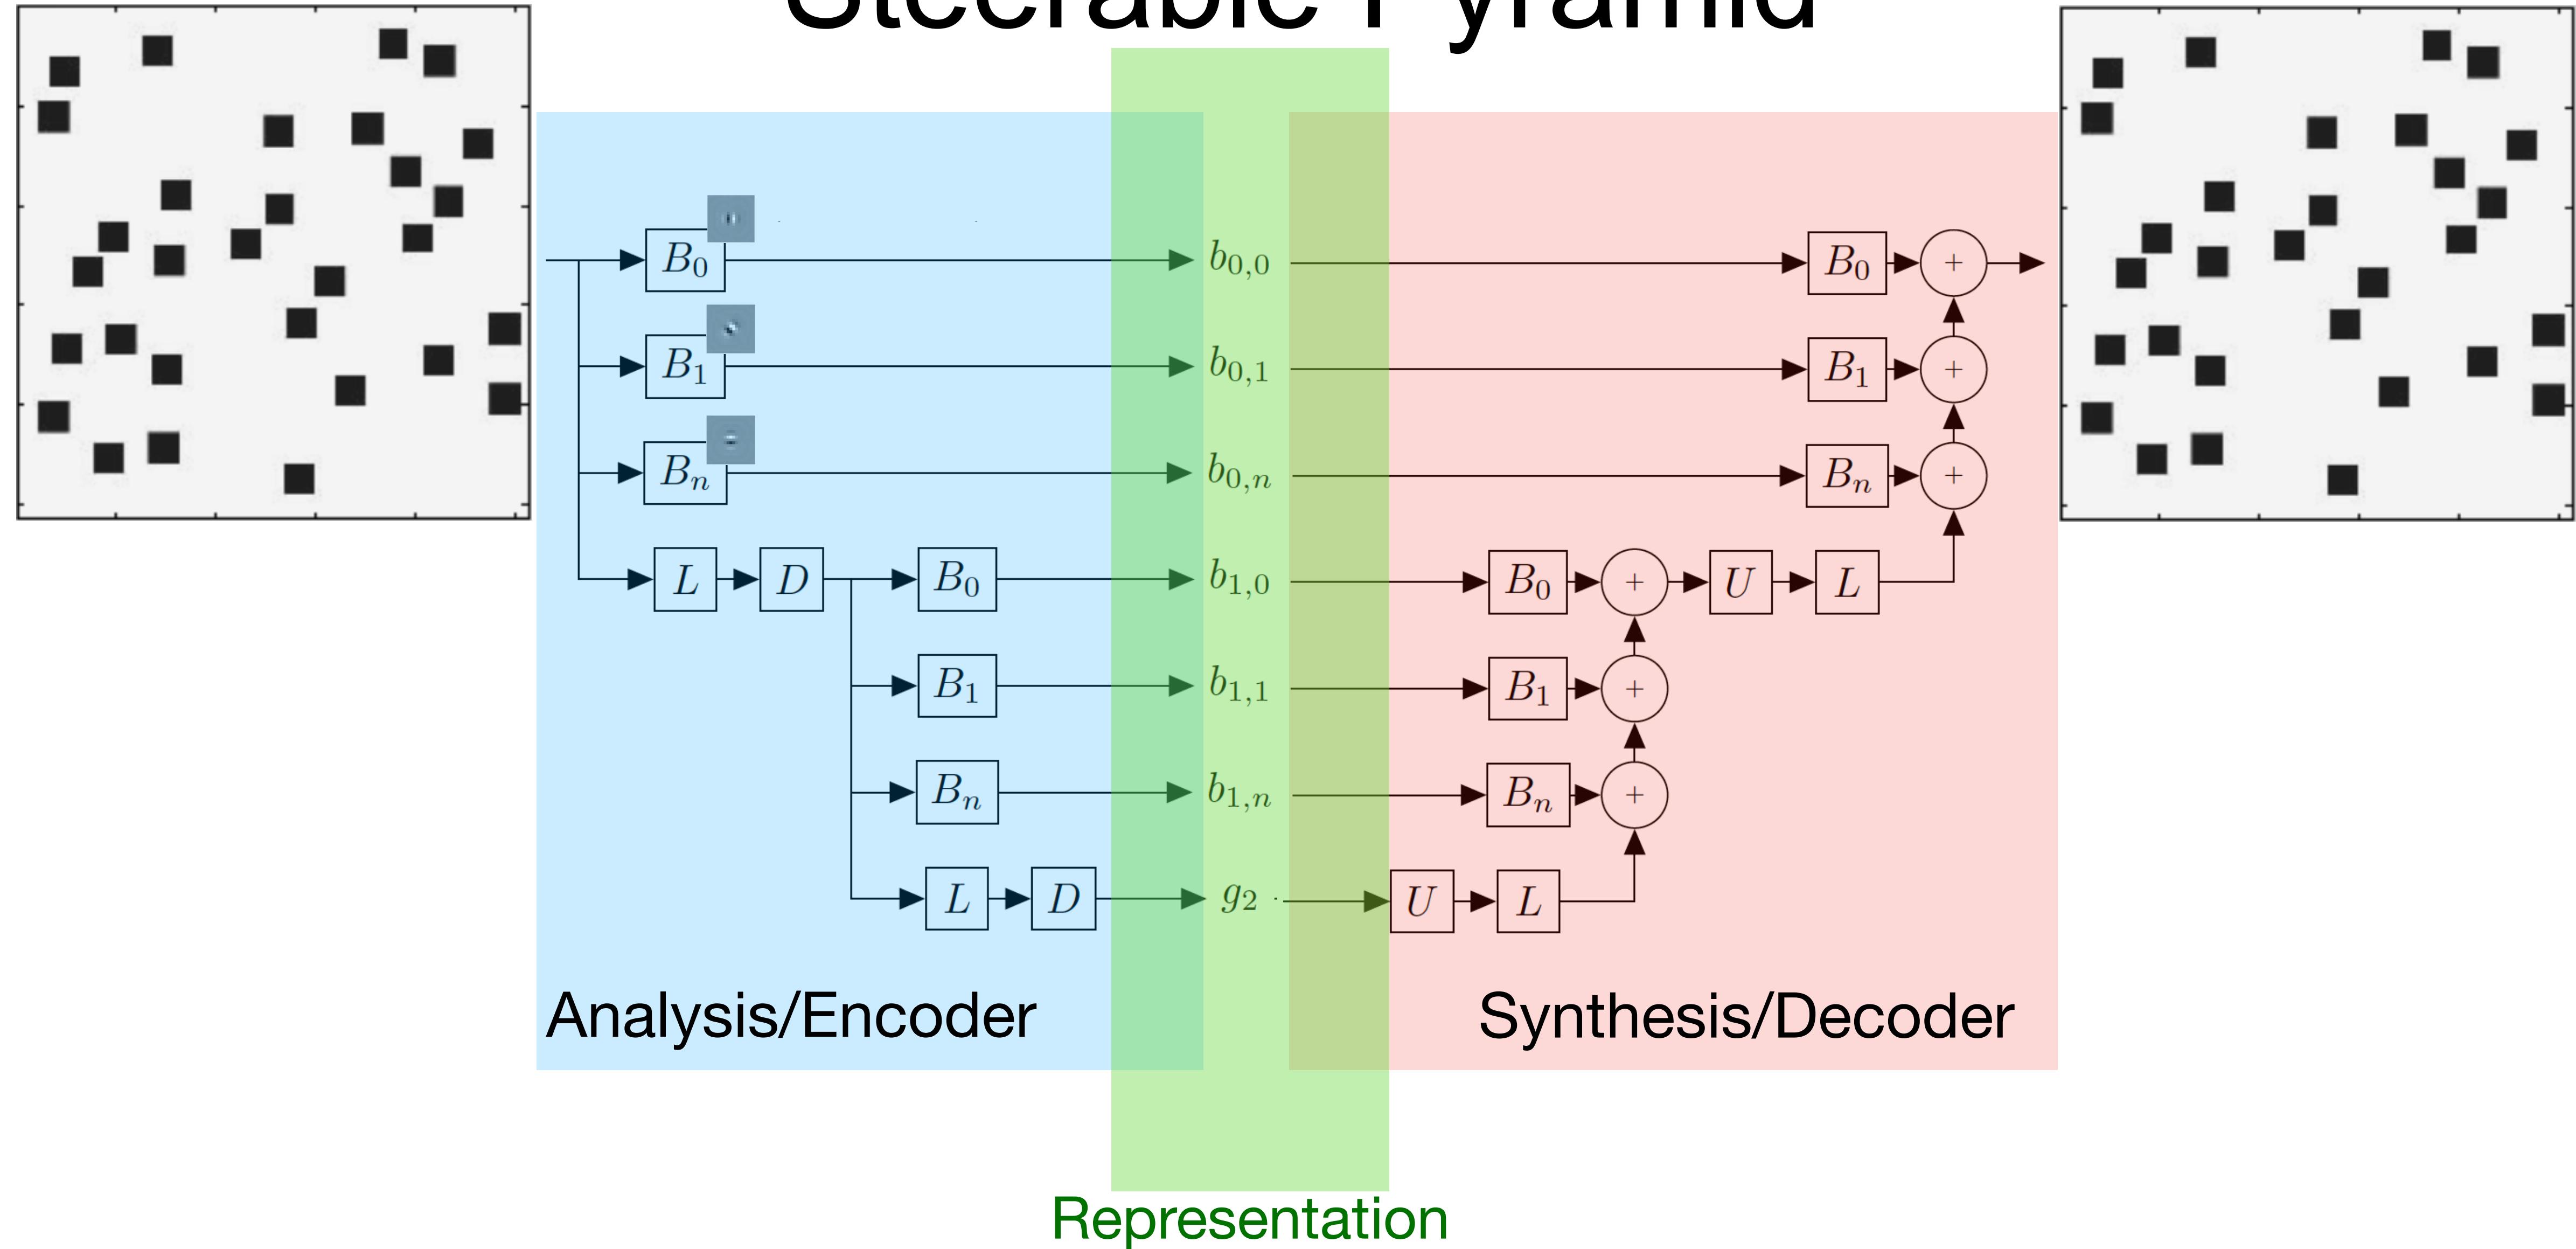


Figure 5: (Top Row) Original digitized sample textures: red granite, berry bush, figured maple, yellow coral. (Bottom Rows) Synthetic solid textured teapots.

\* Department of Psychology, Stanford University, Stanford, CA 94305.  
heeger@white.stanford.edu http://white.stanford.edu

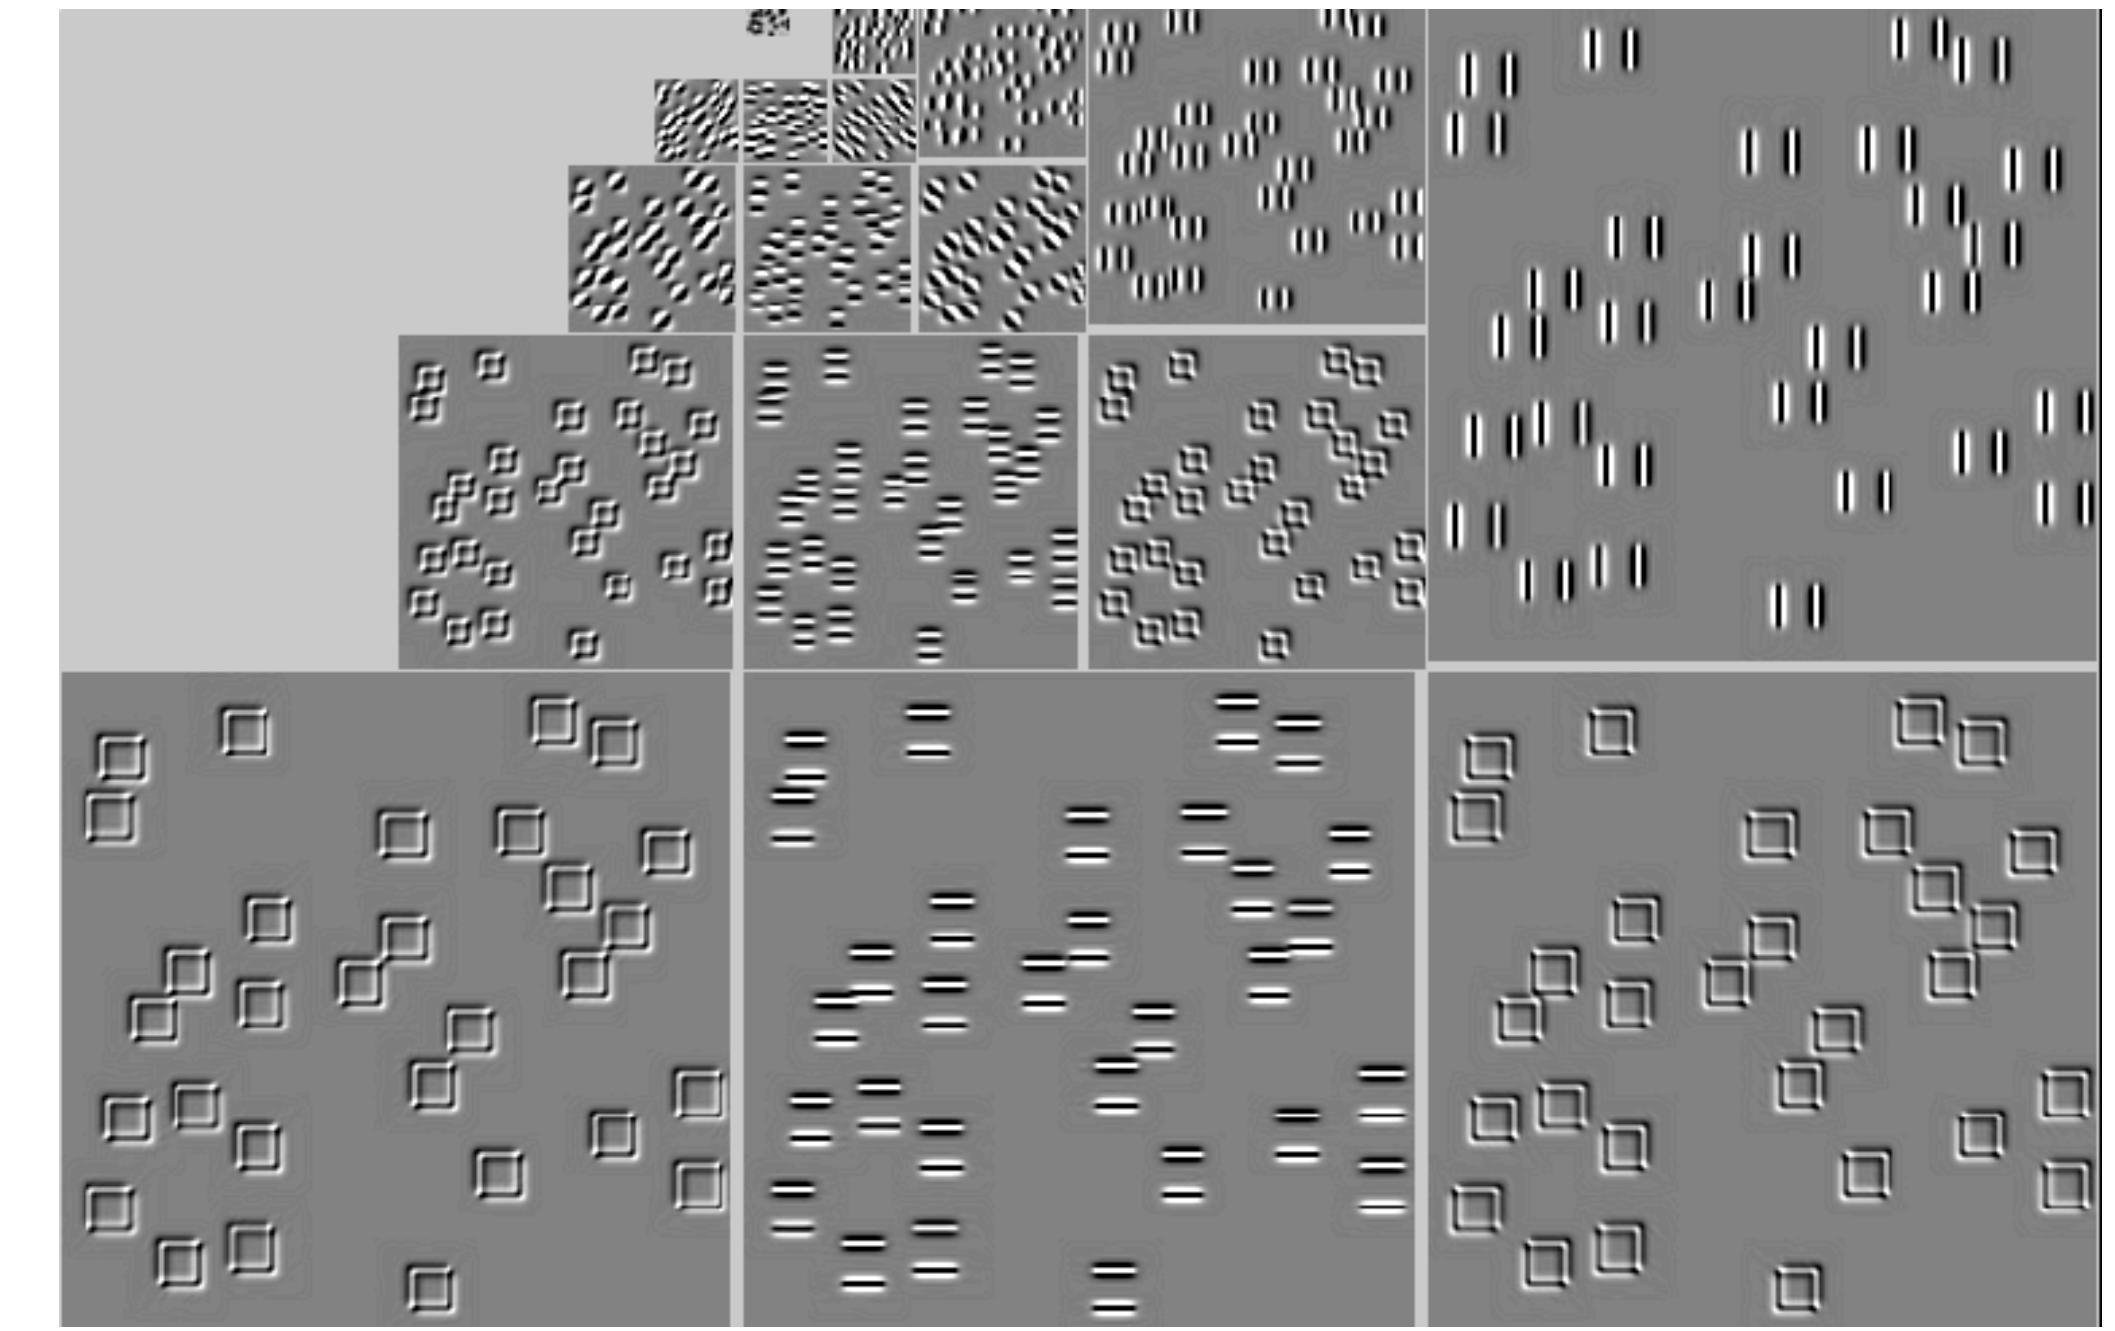
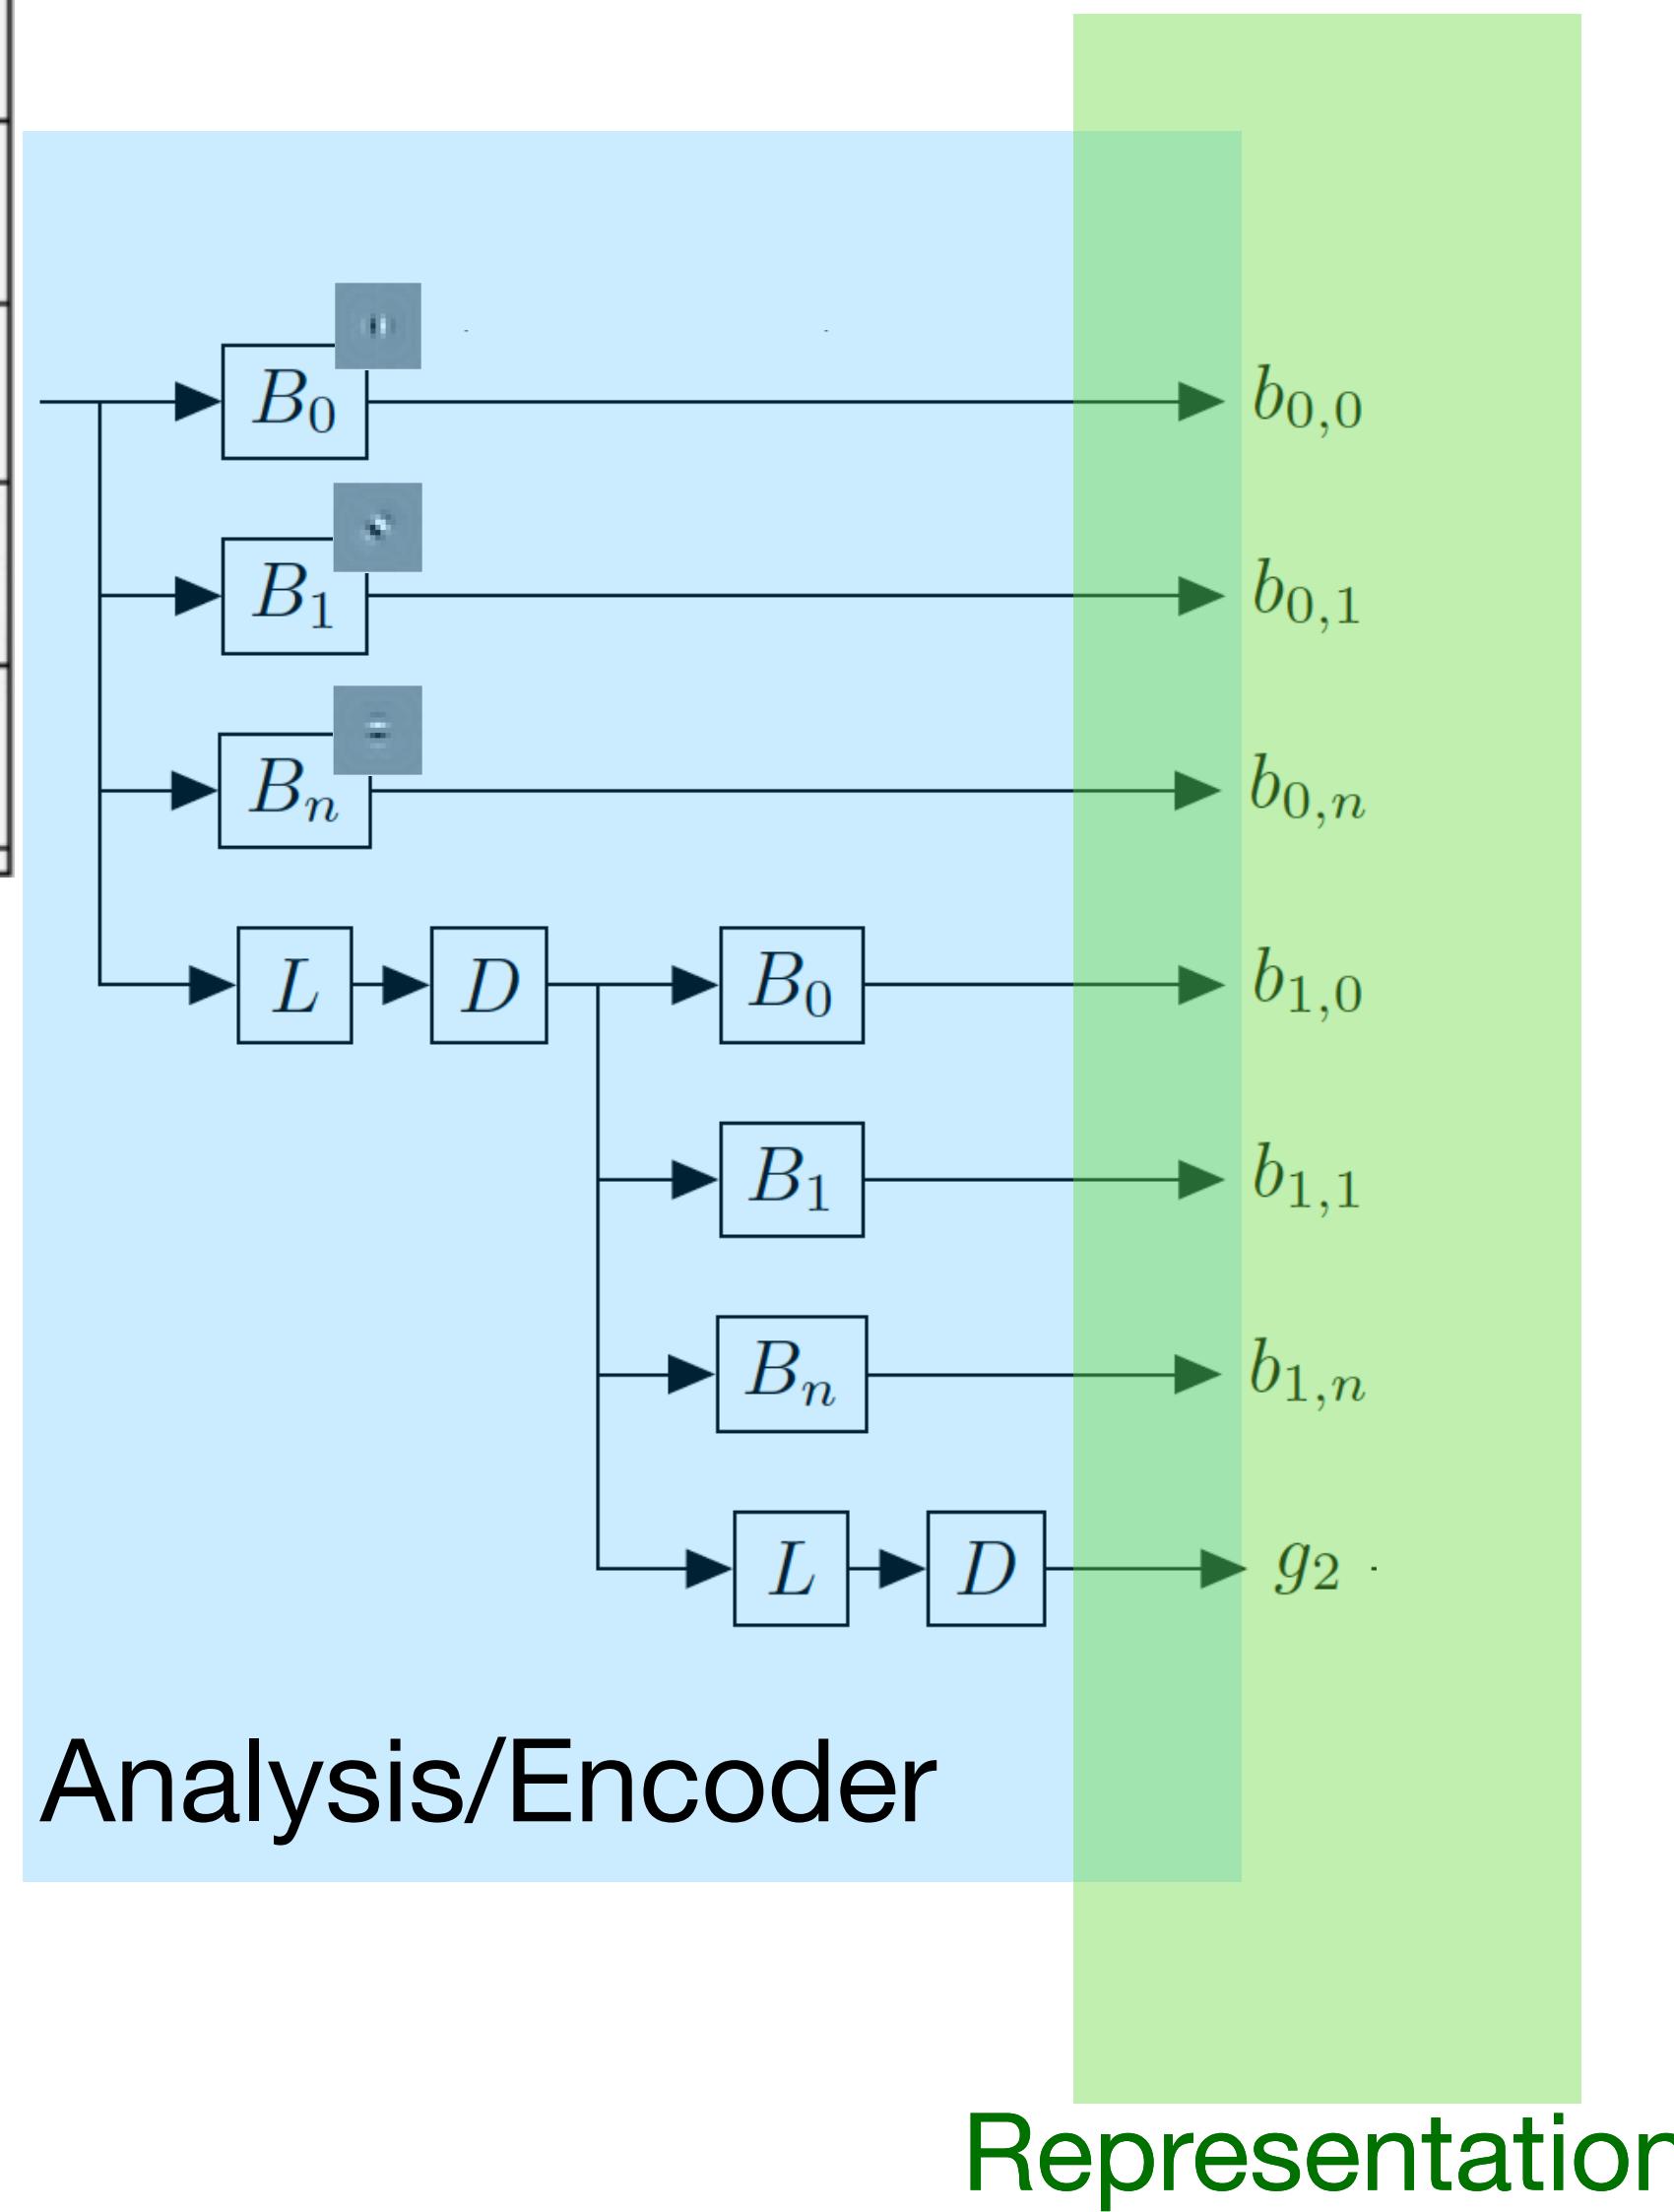
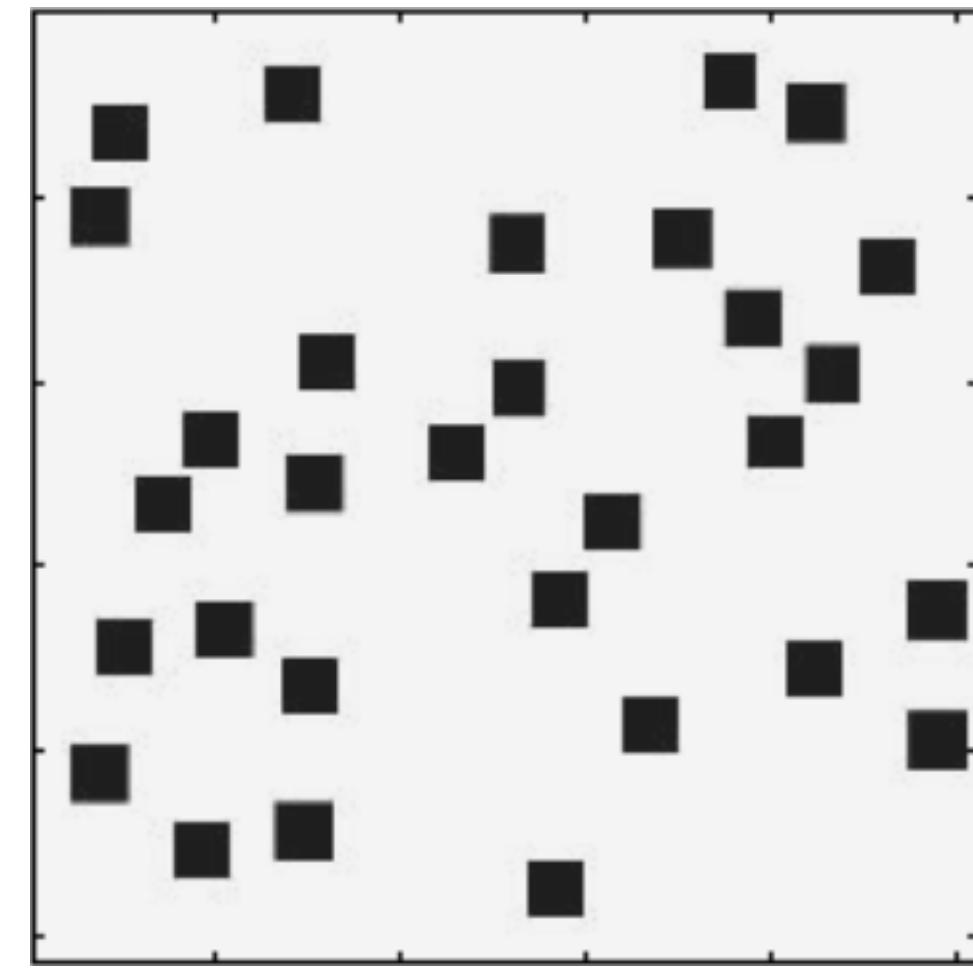
† SRI David Sarnoff Research Center, Princeton, NJ 08544.  
jrb@sarnoff.com

# Steerable Pyramid



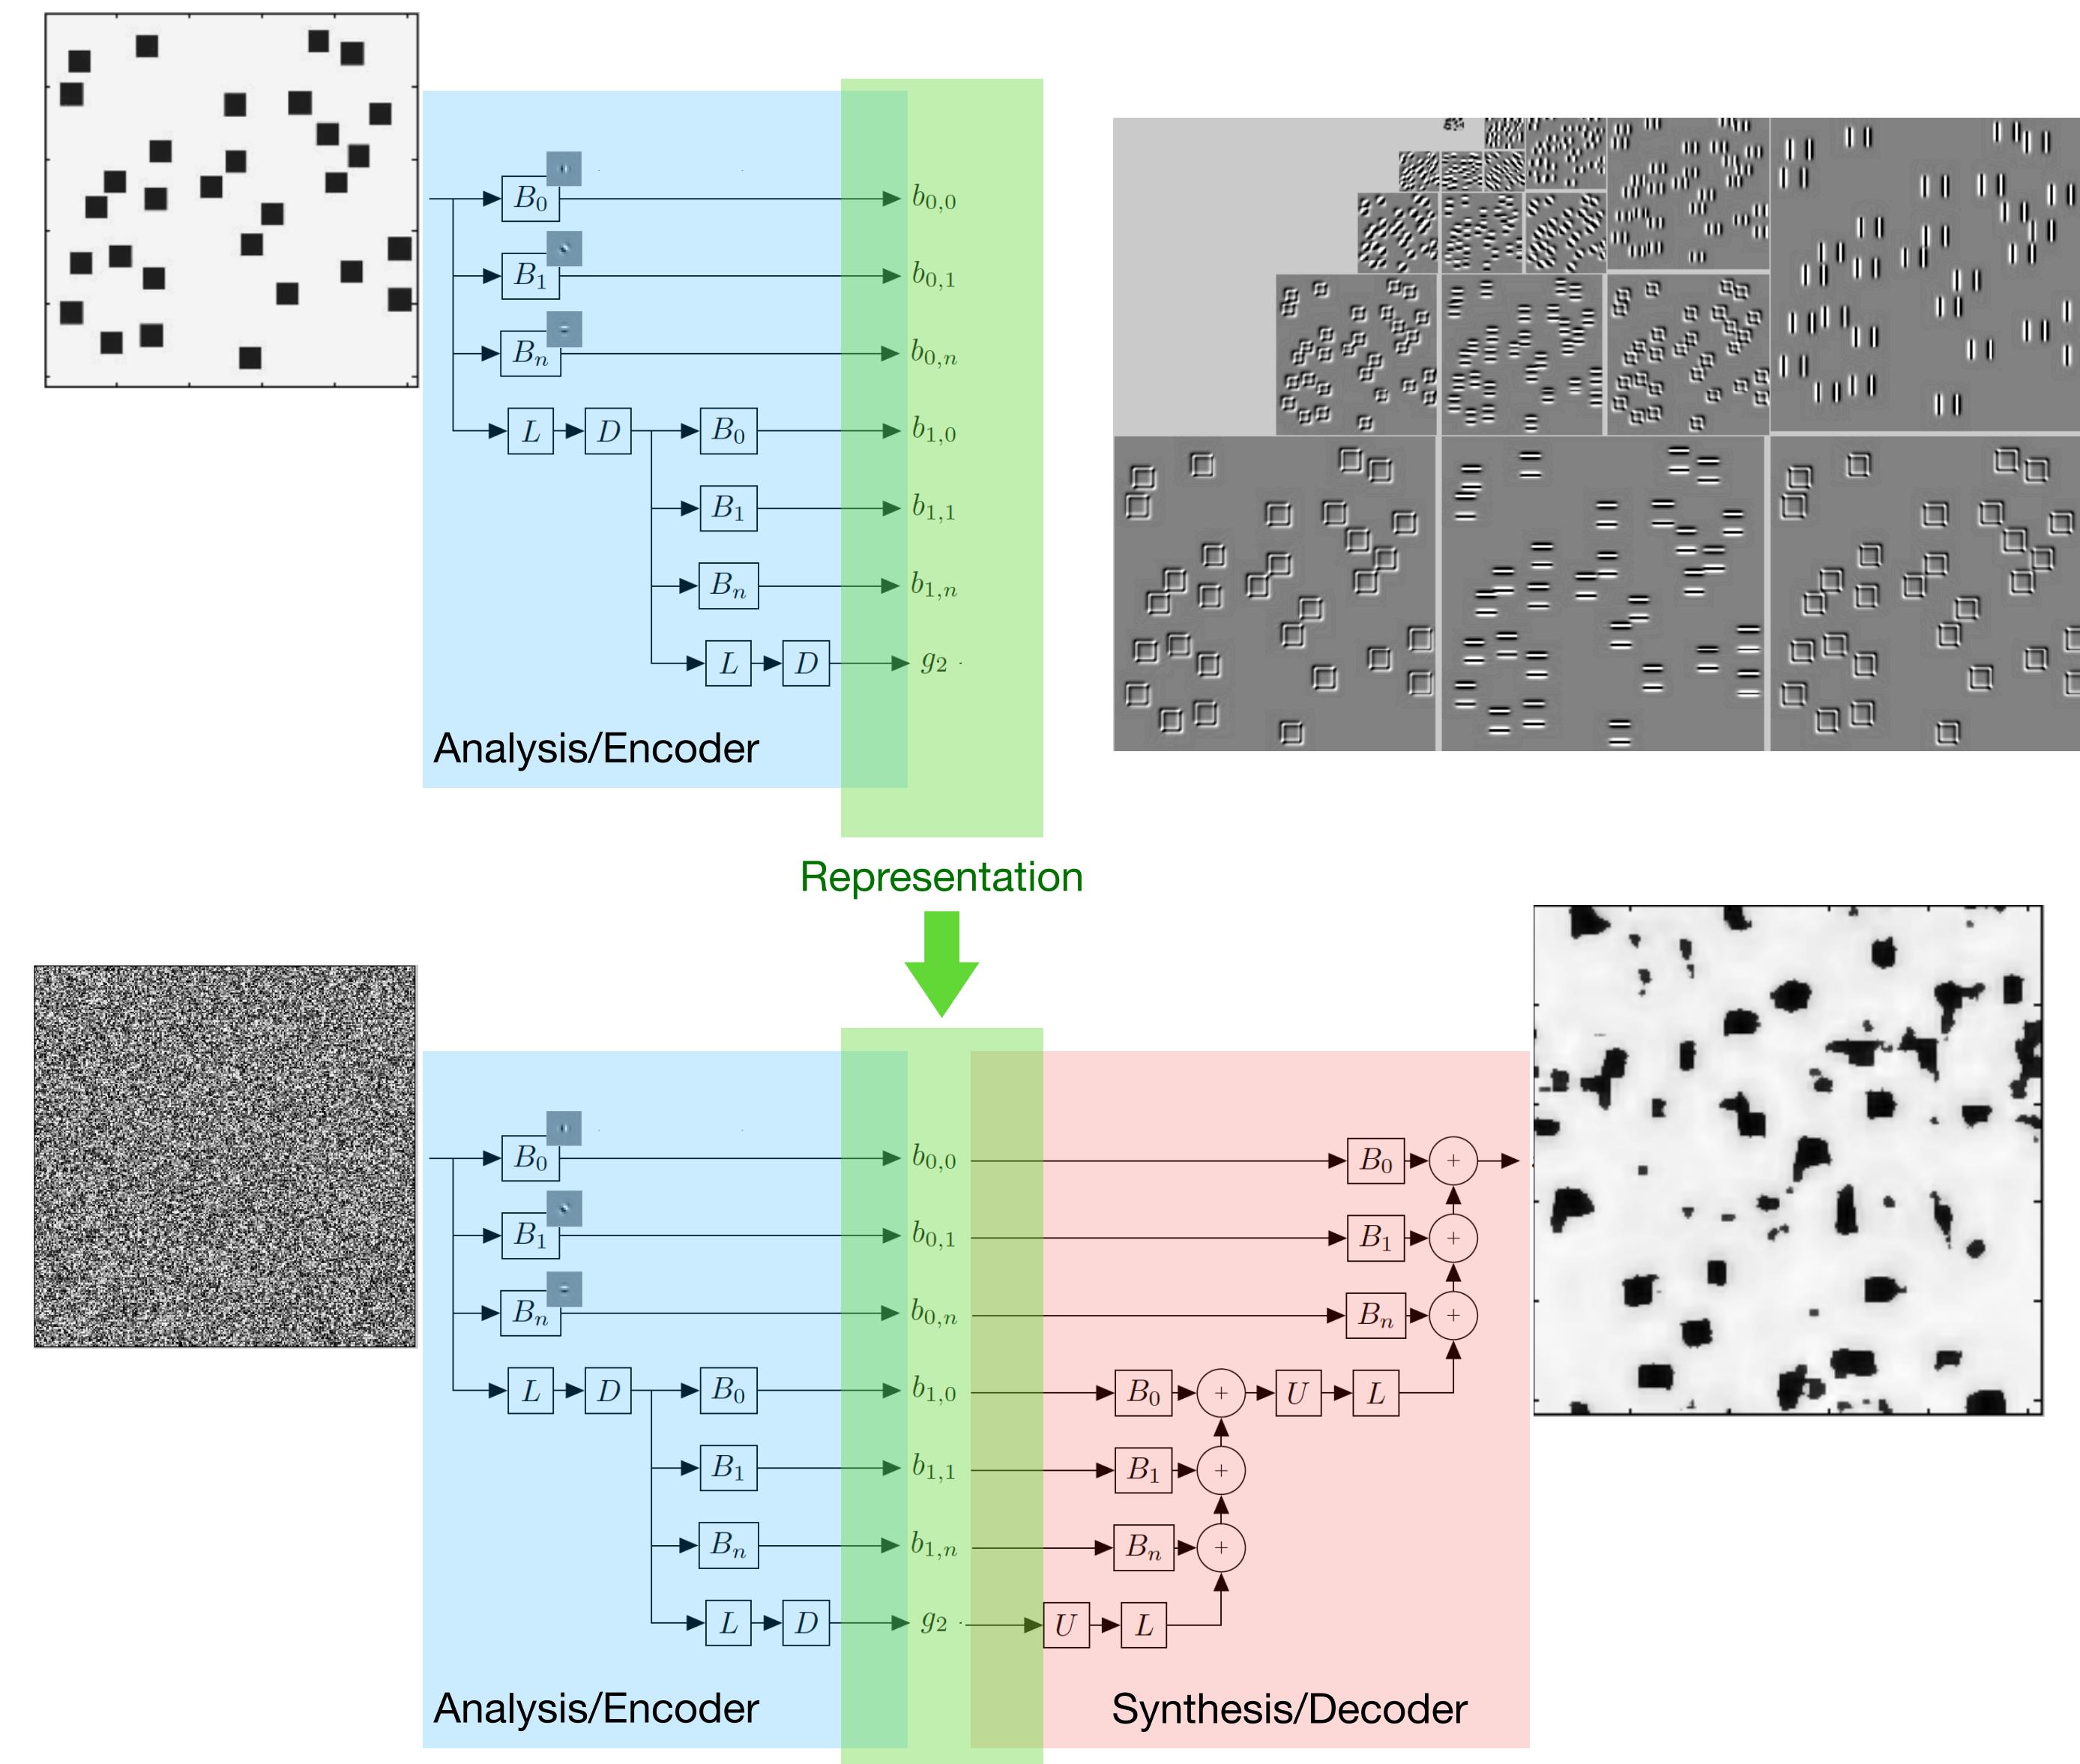
Can we use this representation to generate more samples from the same texture?

# Steerable Pyramid



Representation =  
Histograms of pixel values  
for each subband

# Transferring Image Statistics via Representation Matching



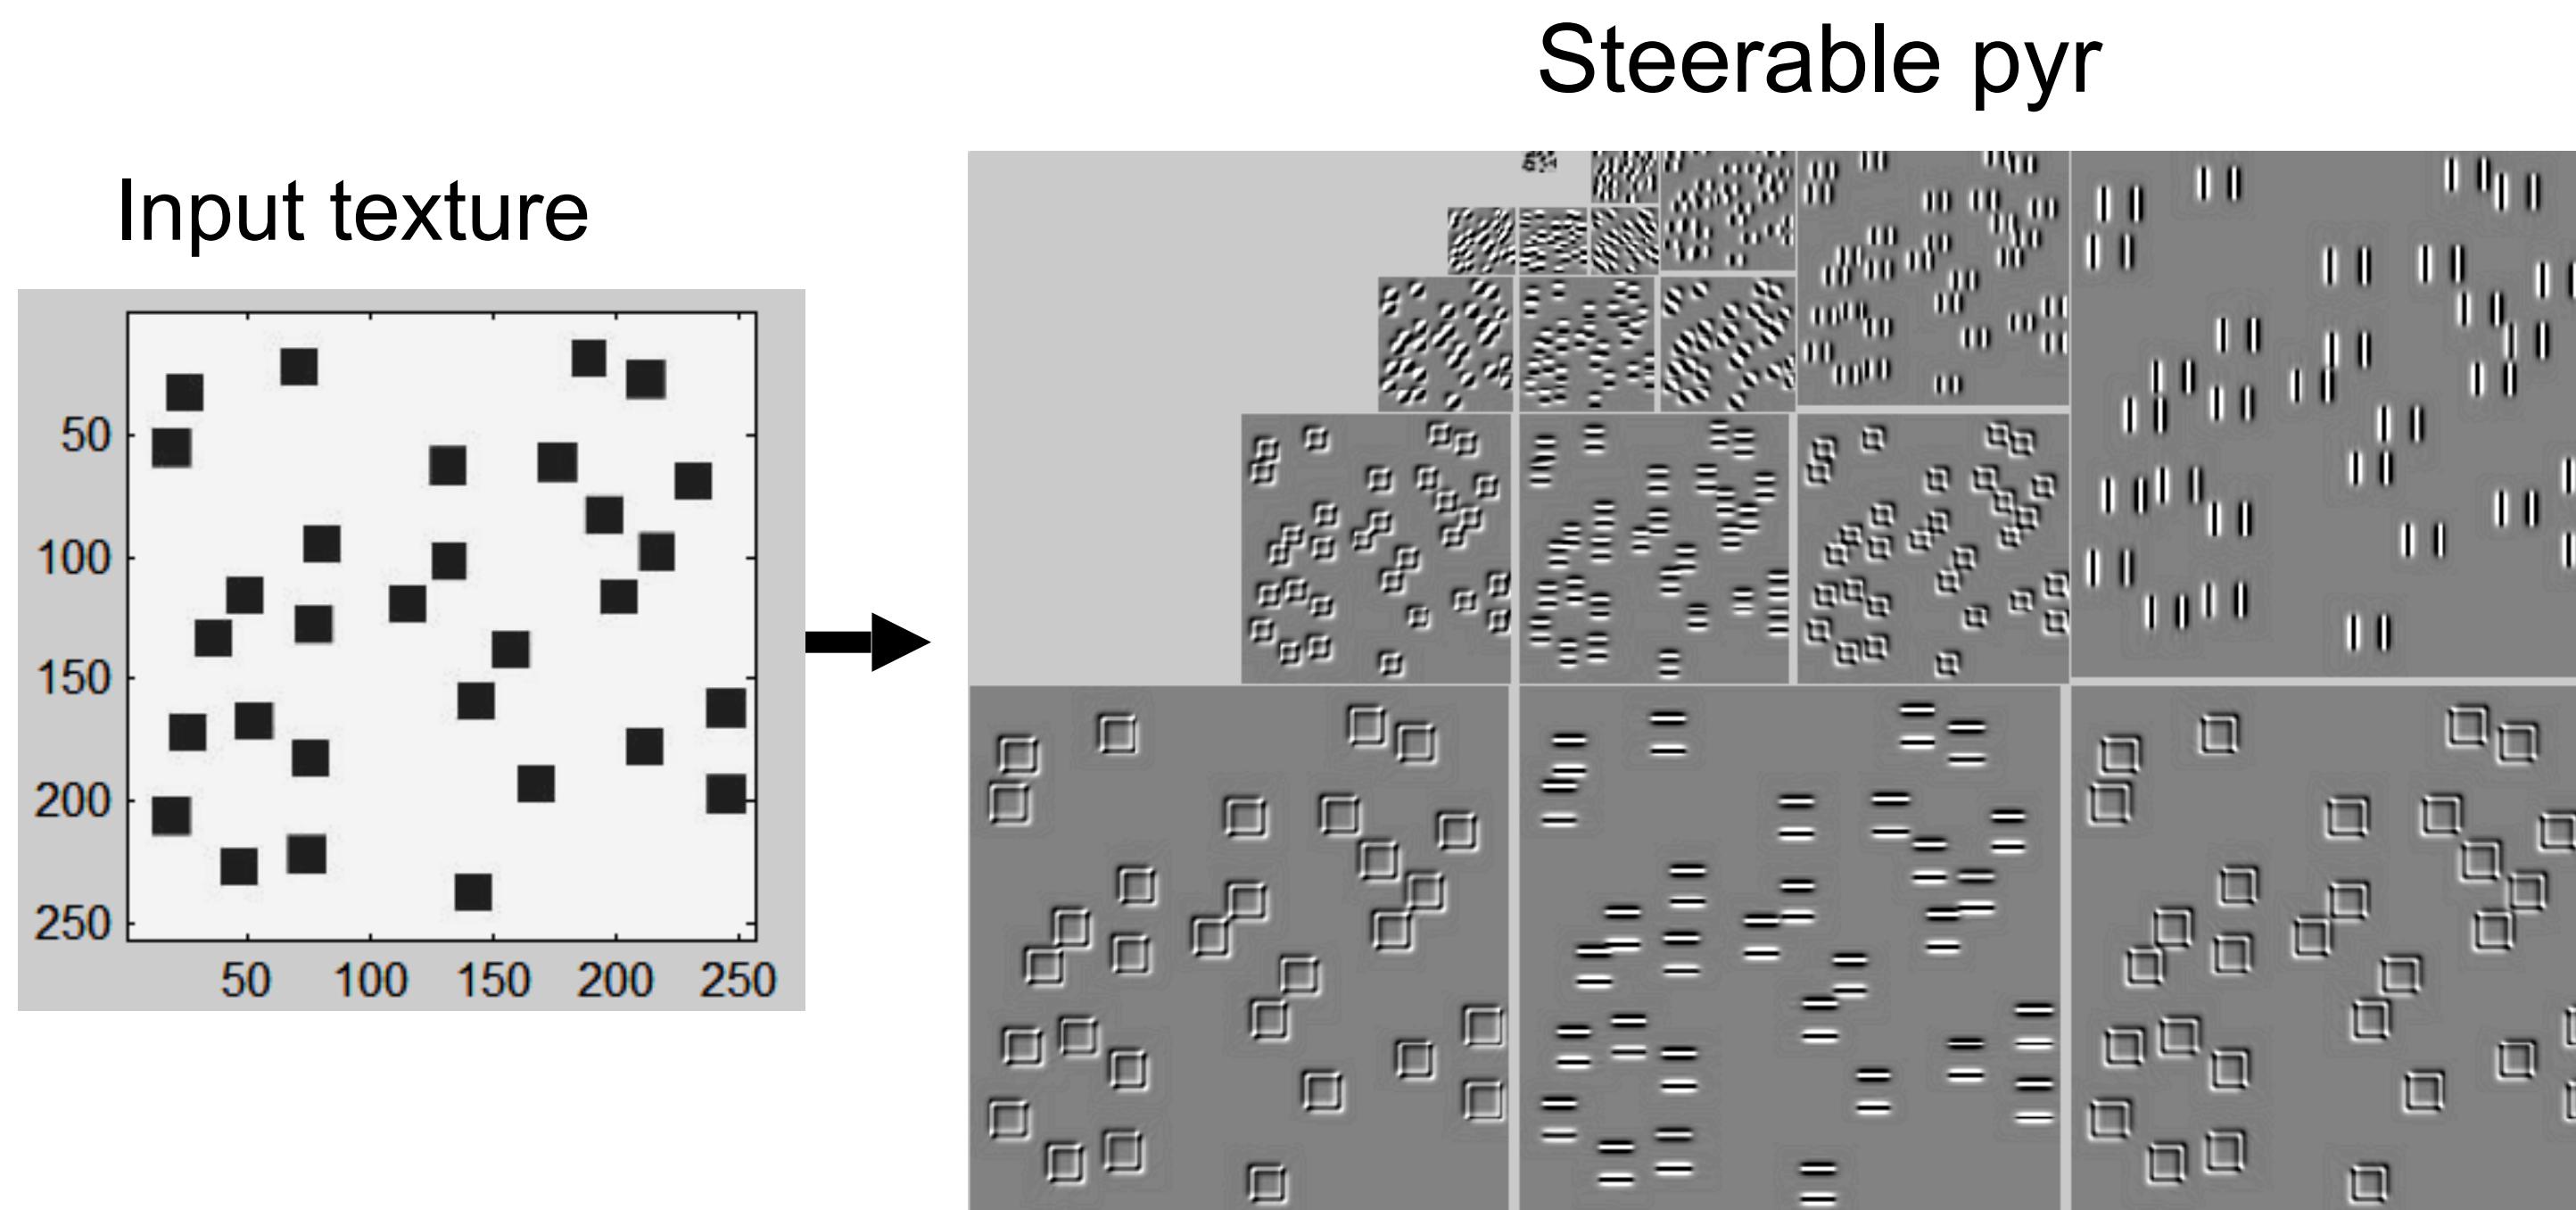
# Transferring Image Statistics via Representation Matching

```
Match-texture (noise,texture)
    Match-Histogram (noise,texture)
        analysis-pyr = Make-Pyramid (texture)
        Loop for several iterations do
            synthesis-pyr = Make-Pyramid (noise)
            Loop for a-band in subbands of analysis-pyr
                for s-band in subbands of synthesis-pyr
                    do
                        Match-Histogram (s-band, a-band)
                noise = Collapse-Pyramid (synthesis-pyr)
            Match-Histogram (noise,texture)
```

Two main tools:

- 1- steerable pyramid
- 2- matching histograms

# 1-The steerable pyramid



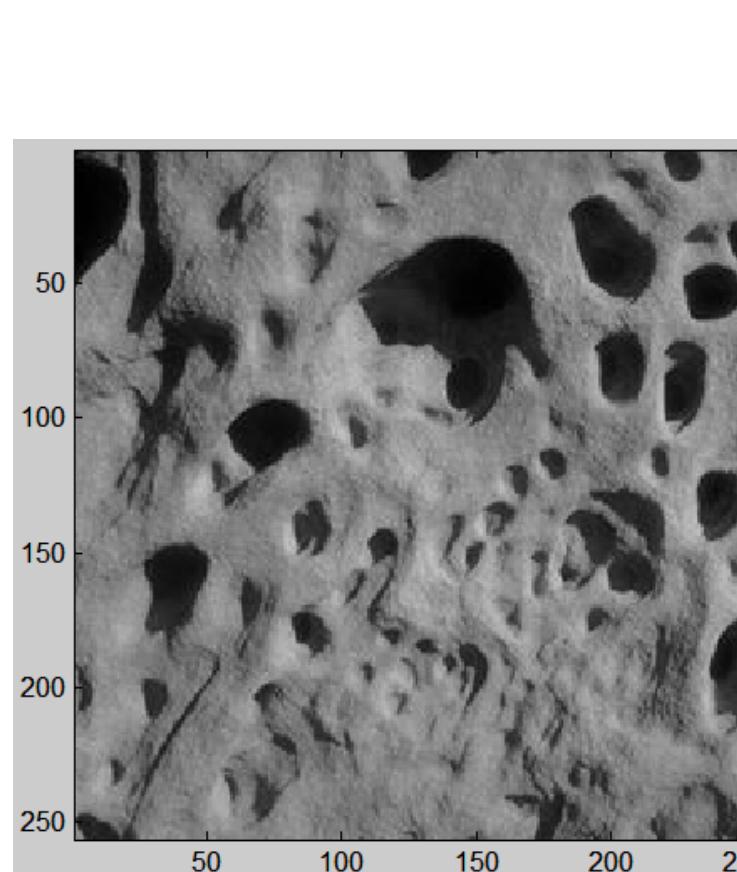
# Overview of the algorithm

```
Match-texture(noise,texture)
    Match-Histogram (noise,texture)
    analysis-pyr = Make-Pyramid (texture)
    Loop for several iterations do
        synthesis-pyr = Make-Pyramid (noise)
        Loop for a-band in subbands of analysis-pyr
            for s-band in subbands of synthesis-pyr
                do
                    Match-Histogram (s-band,a-band)
            noise = Collapse-Pyramid (synthesis-pyr)
    Match-Histogram (noise,texture)
```

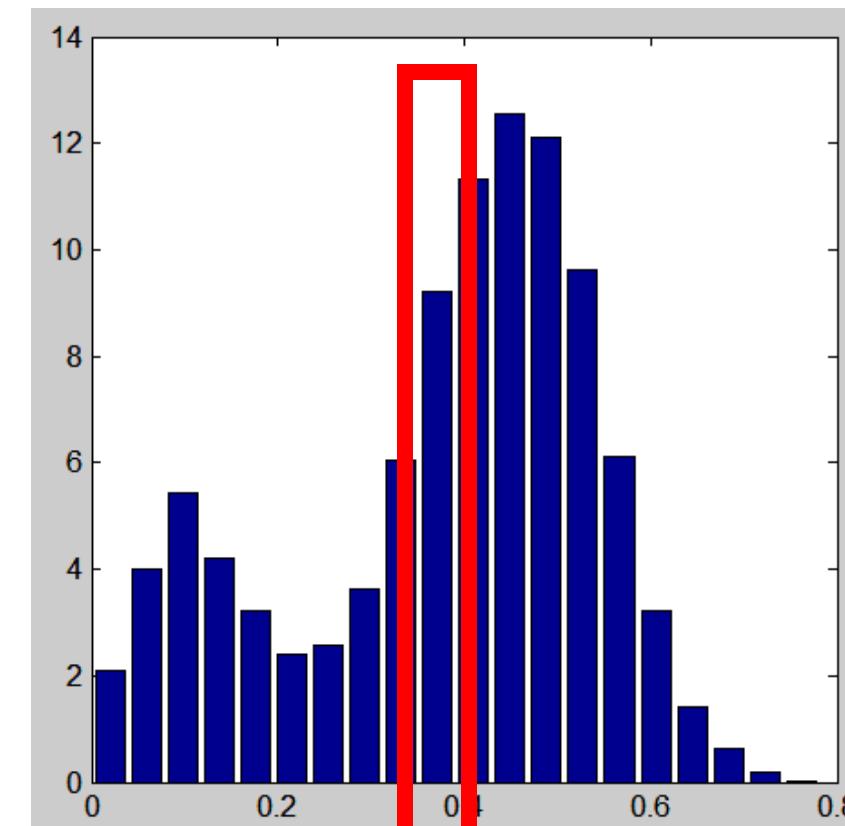
Two main tools:

- 1- steerable pyramid
- 2- matching histograms**

# 2-Matching histograms

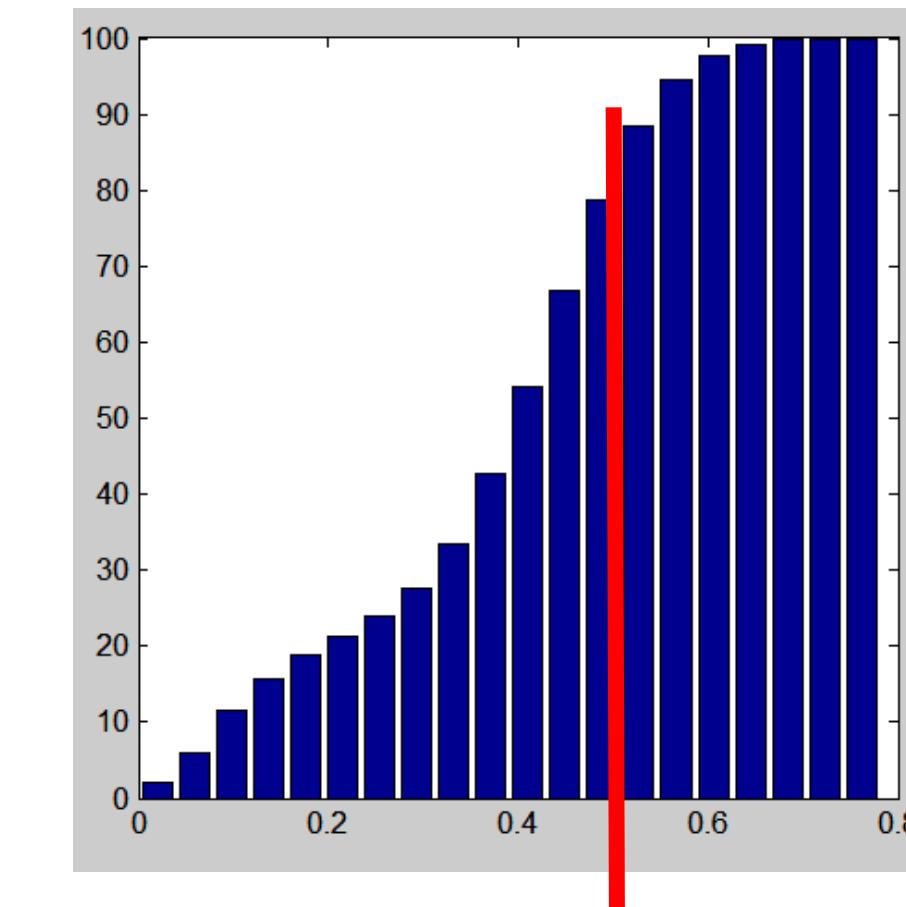


Histogram

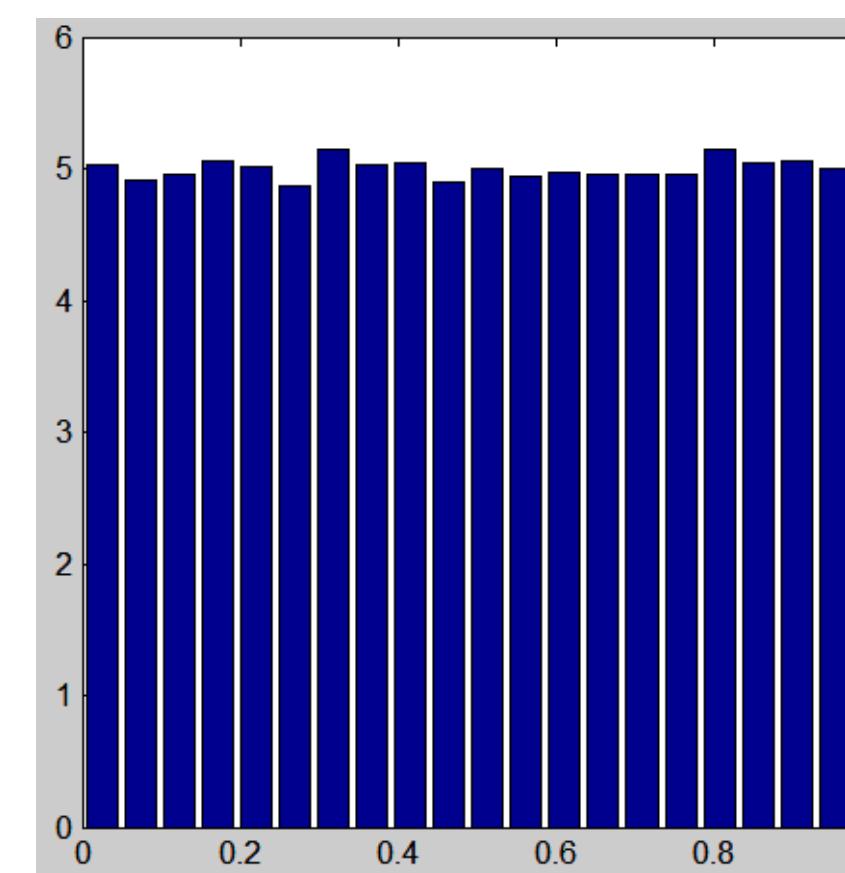
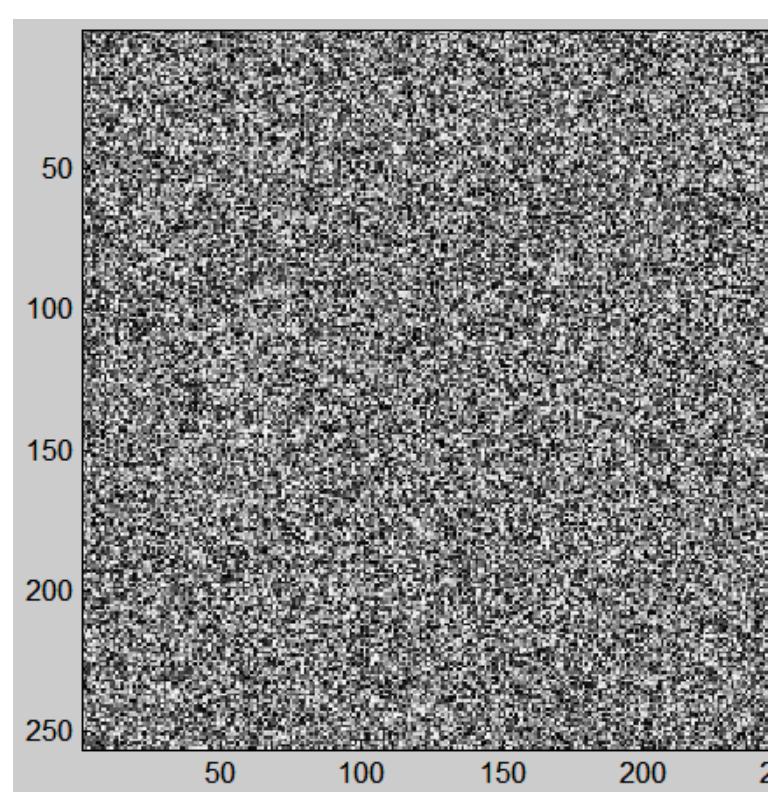


9% of pixels have an intensity value  
within the range[0.37, 0.41]

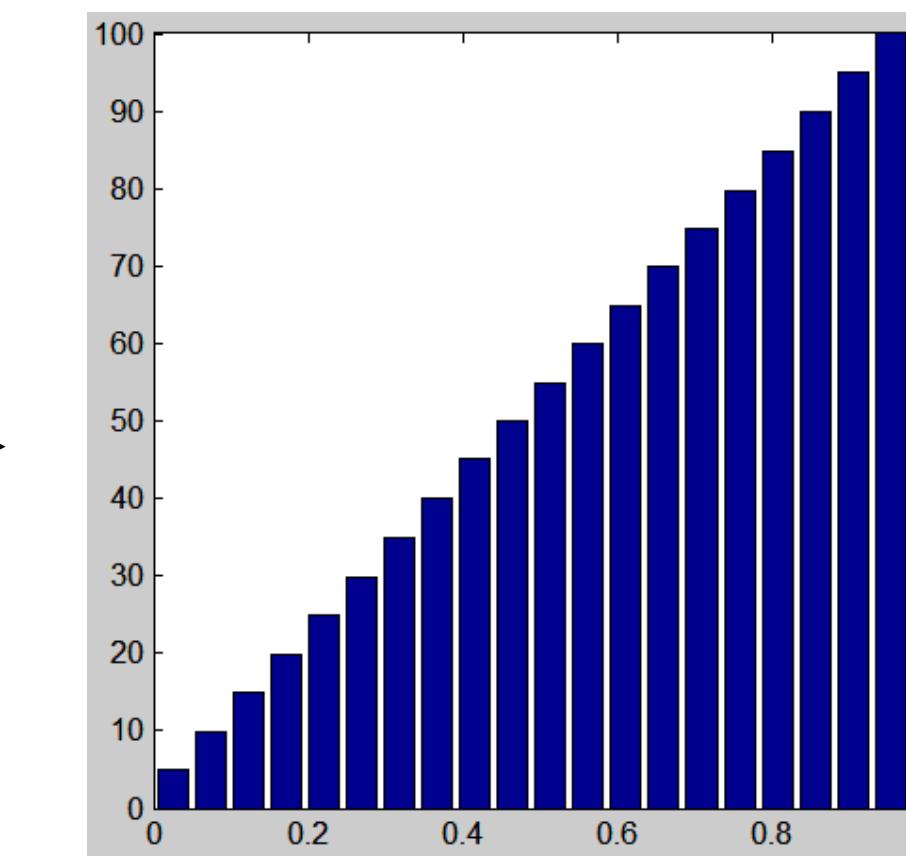
Cumulative distribution



75% of pixels have an intensity value  
smaller than 0.5

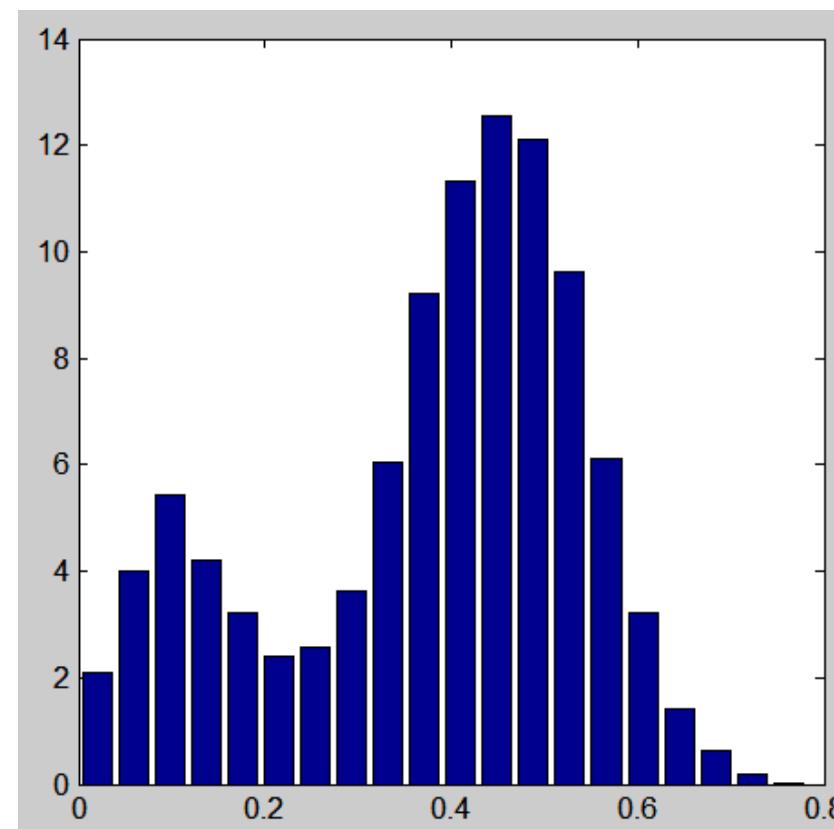
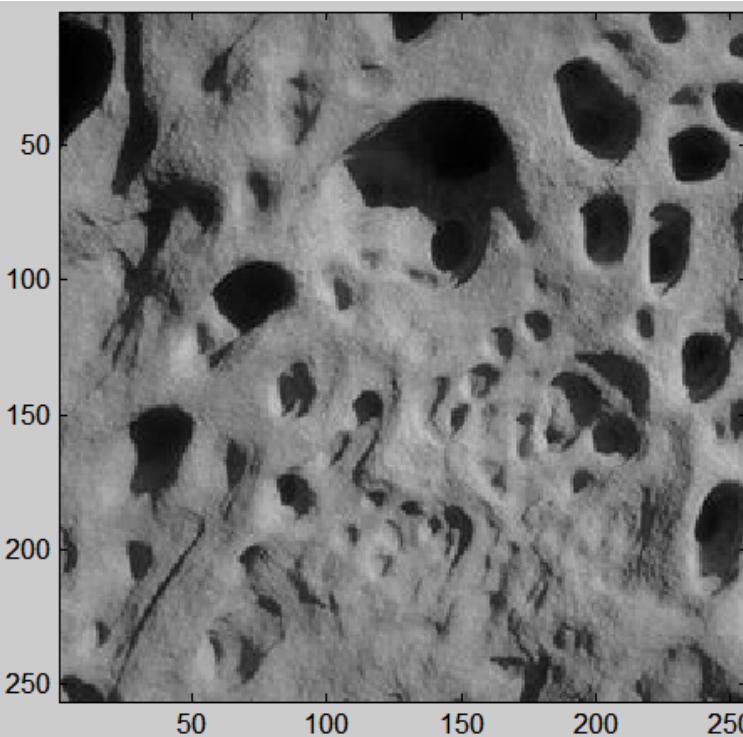


5% of pixels have an intensity value  
within the range[0.37, 0.41]

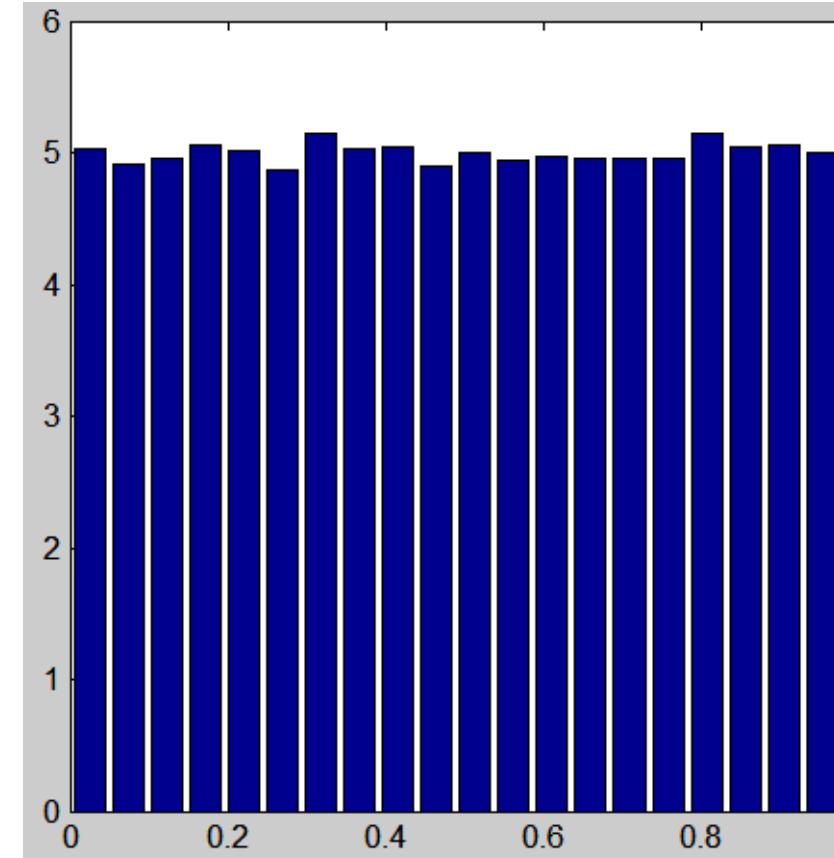
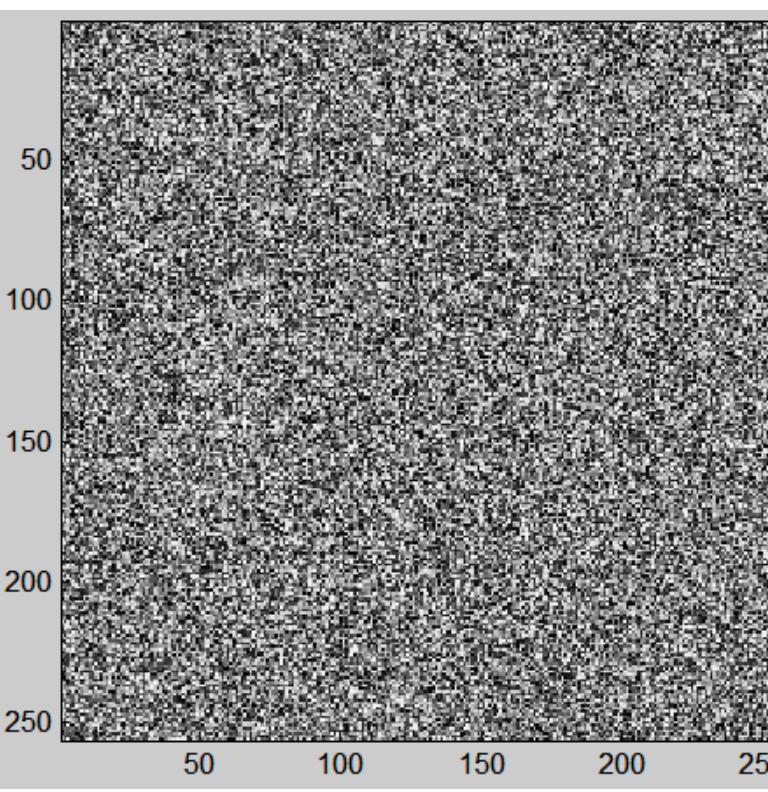


# 2-Matching histograms

Z[n,m]



Y[n,m]



↑ ?

We look for a transformation  
of the image Y

$$Y' = f(Y)$$

Such that  
 $\text{Hist}(Y) = \text{Hist}(f(Z))$

There are infinitely many functions  
that can do this transformation.

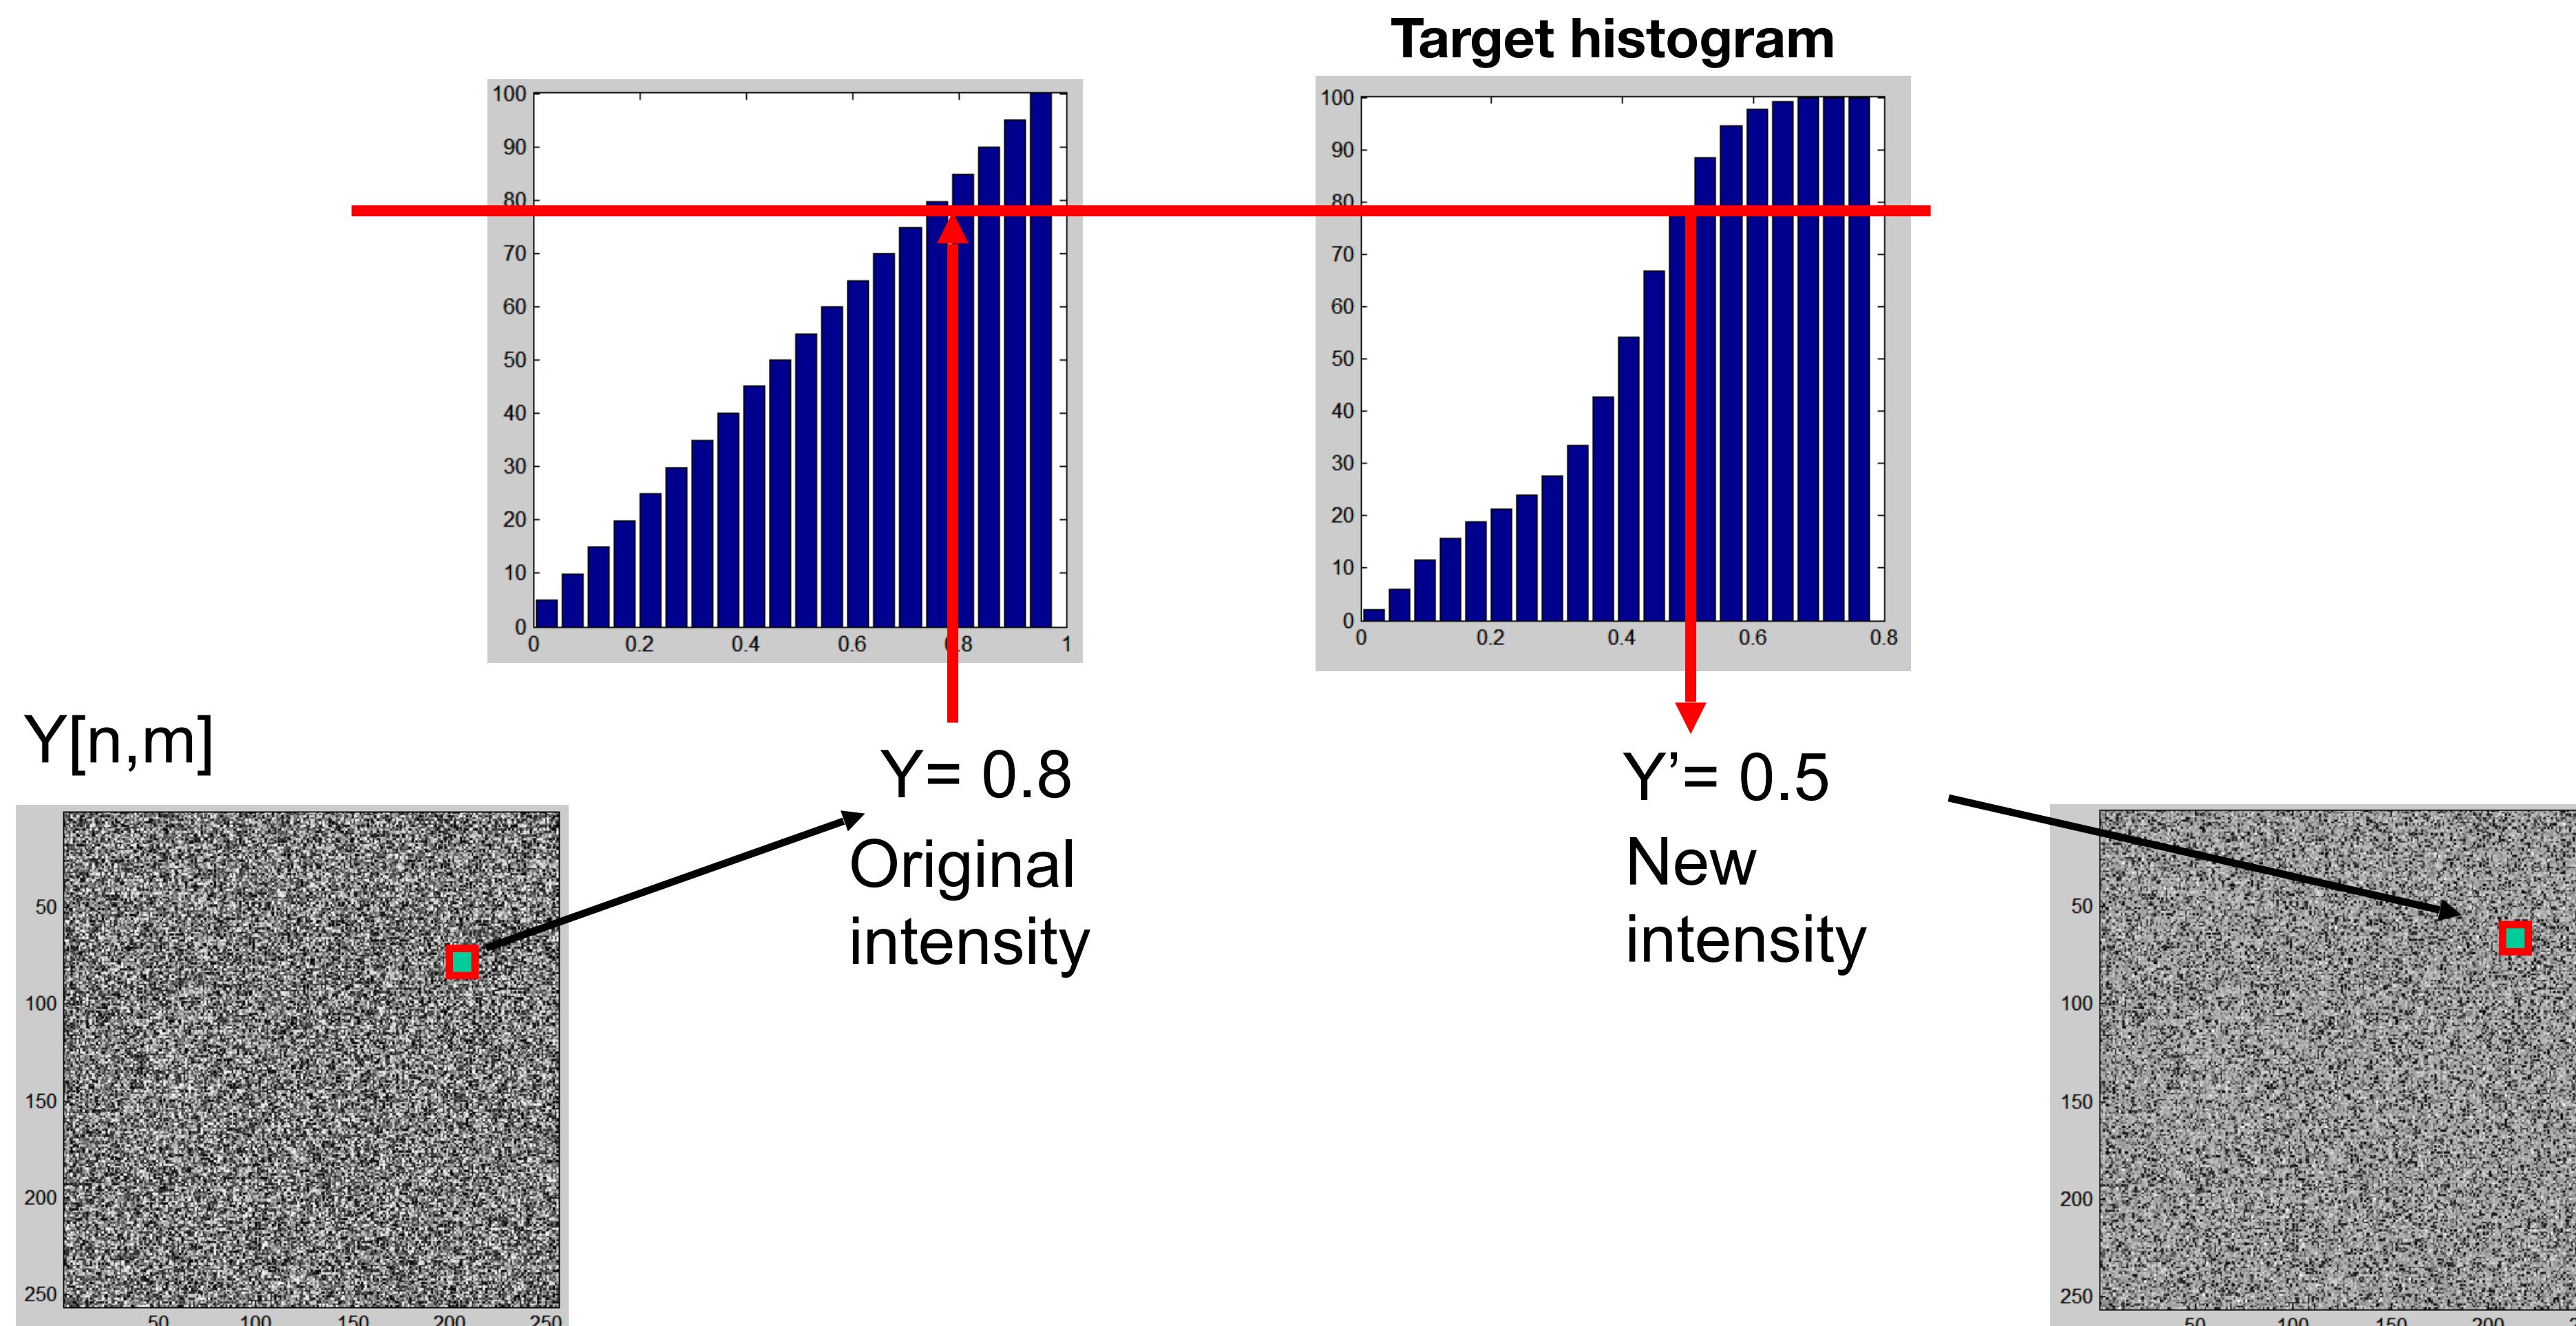
A natural choice is to use **f** being:

- point-wise non linearity
- stationary
- monotonic (most of the time invertible)

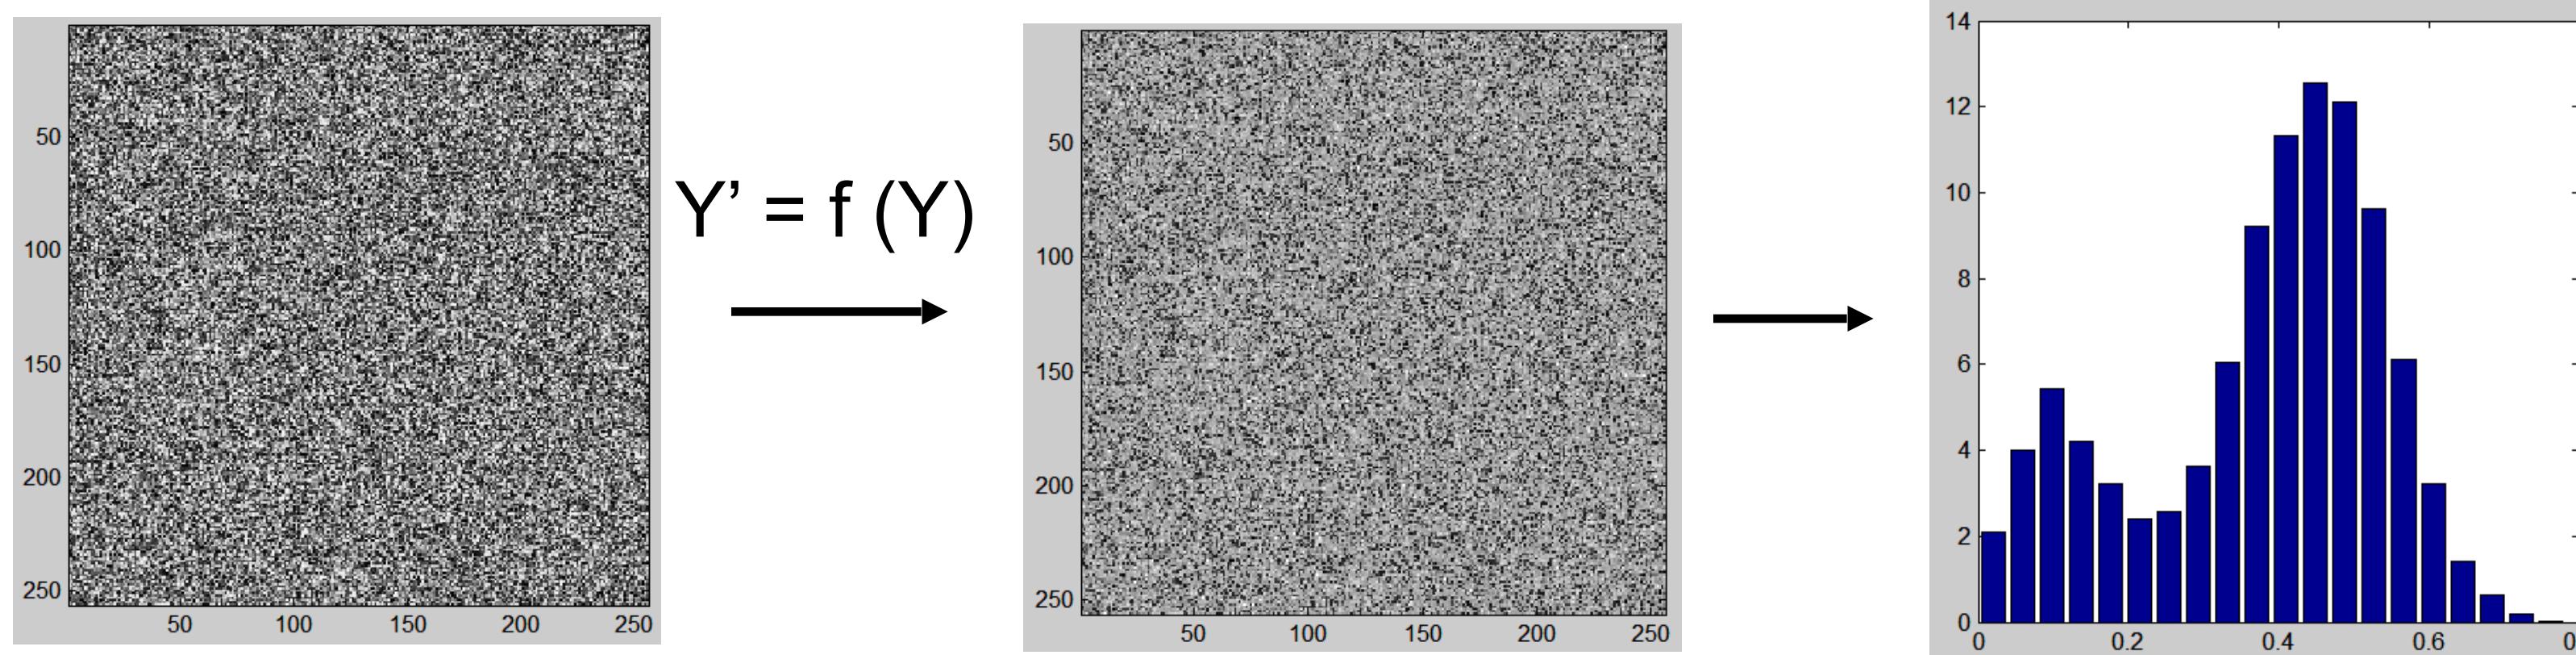
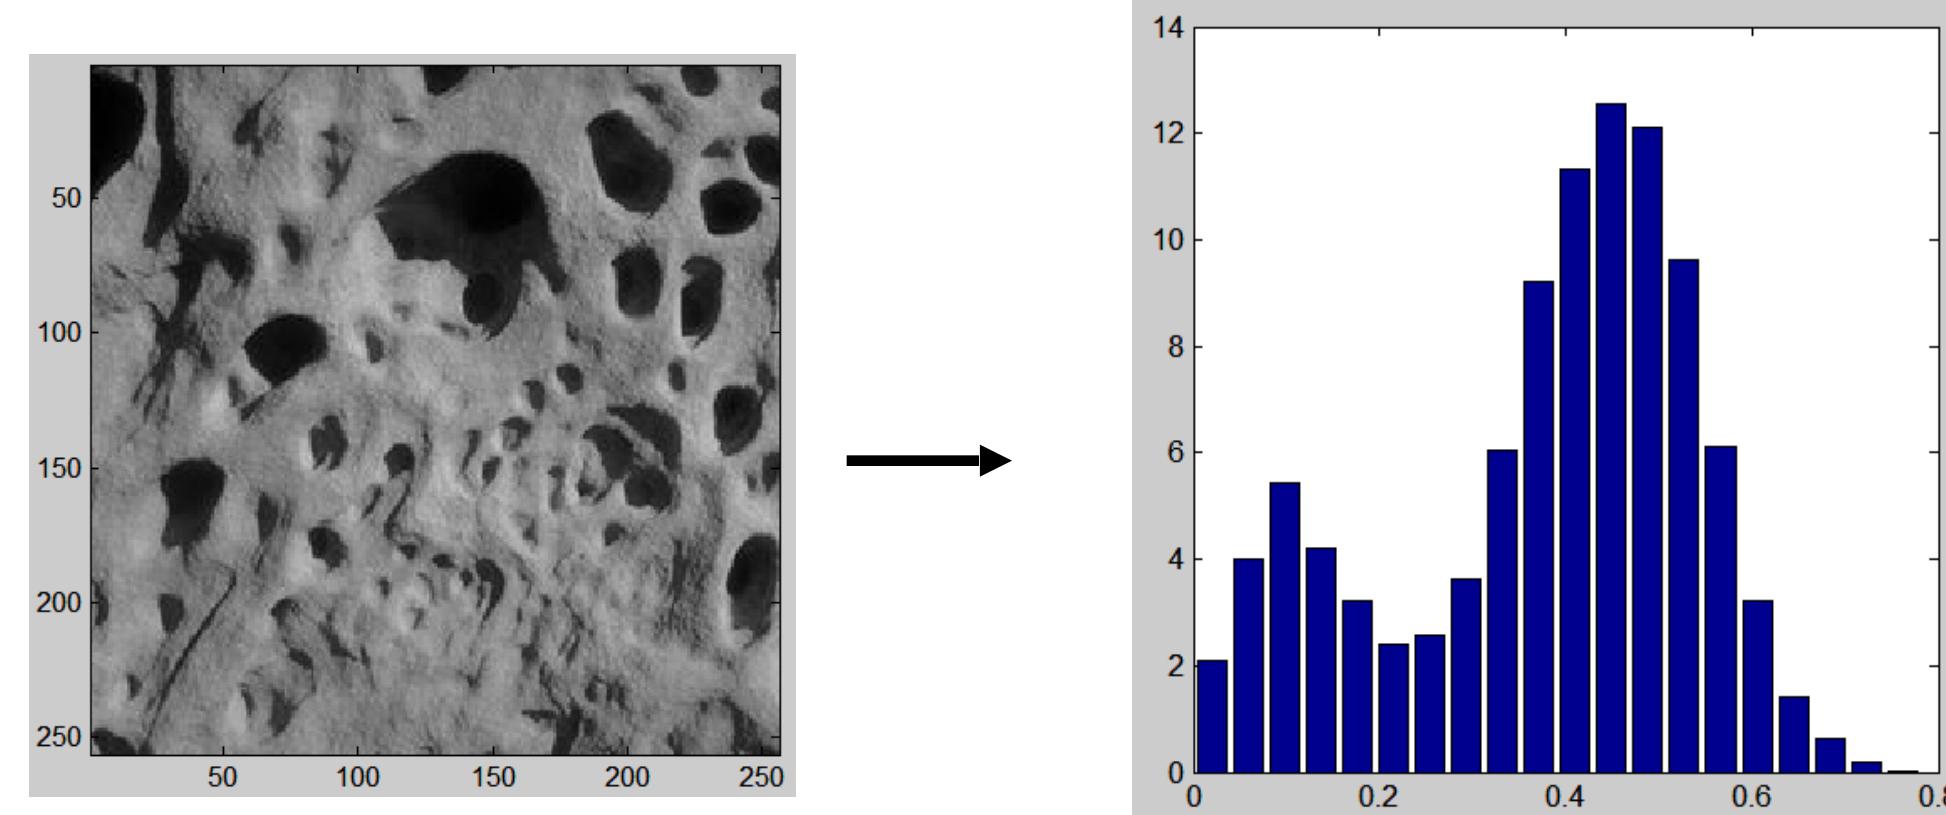
# 2-Matching histograms

The function  $f$  is just a look up table: it says, change all the pixels of value  $Y$  into a value  $f(Y)$ .

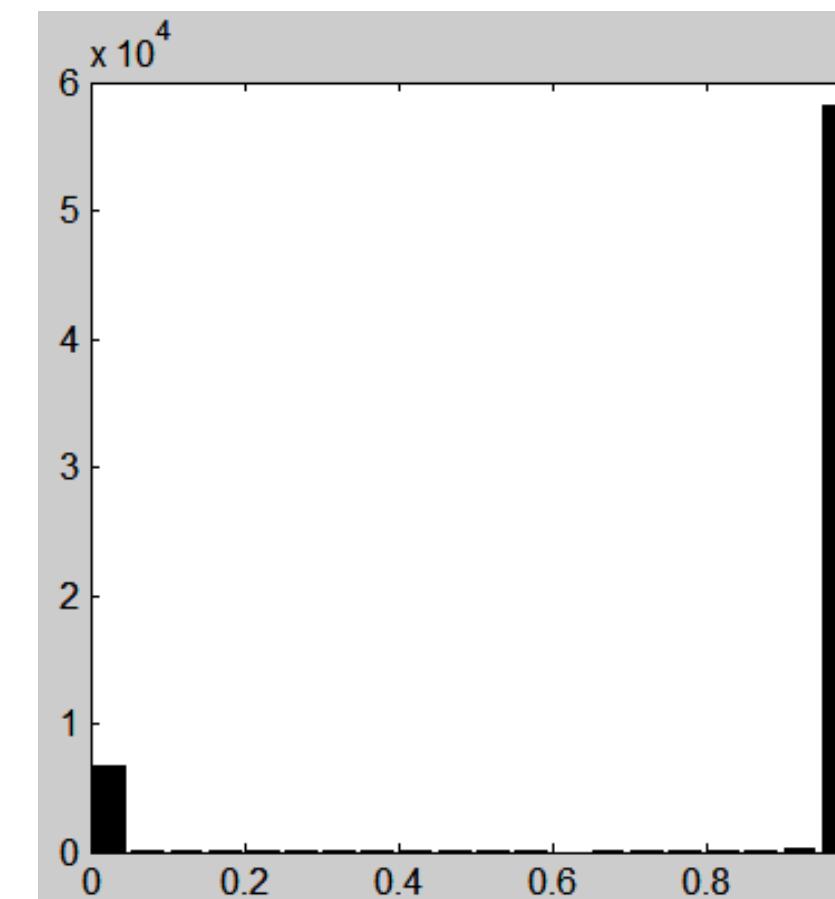
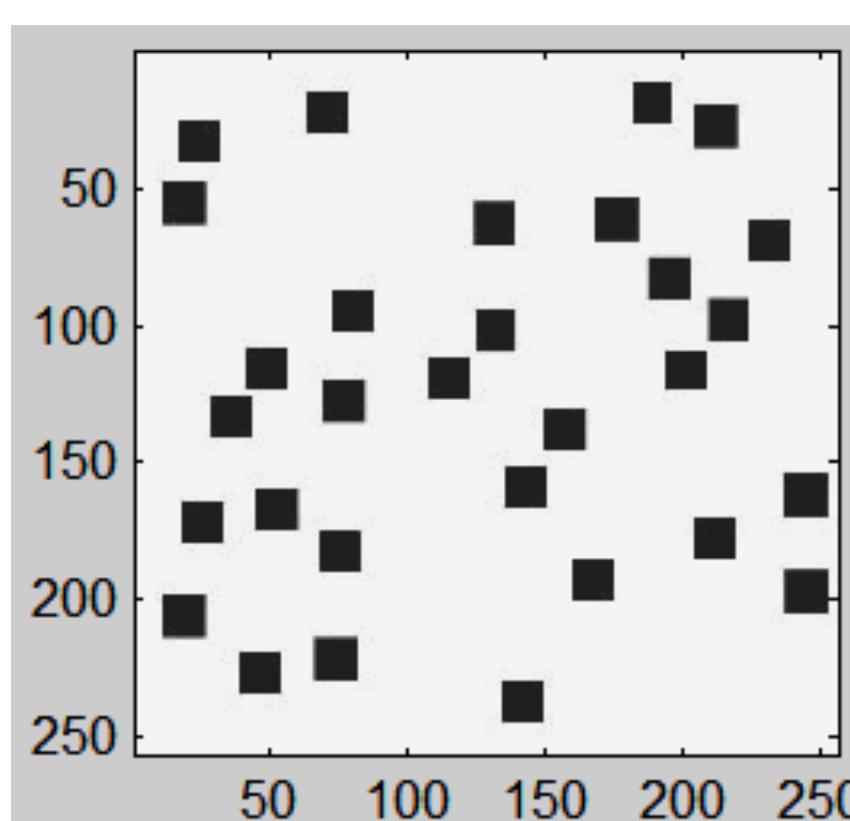
$$Y' = f(Y)$$



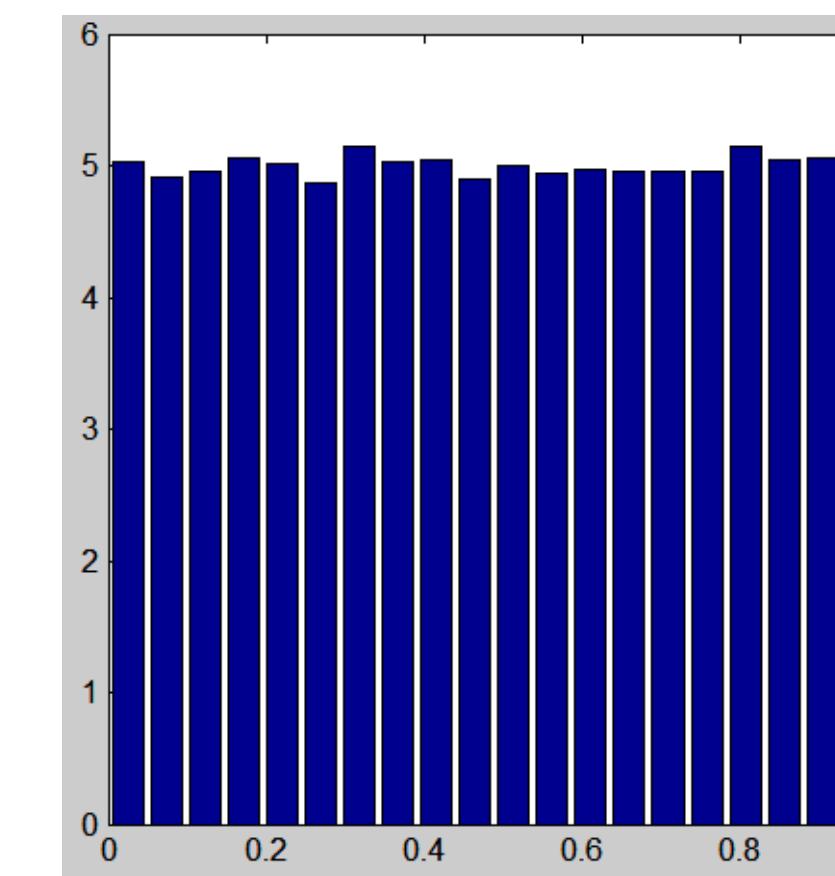
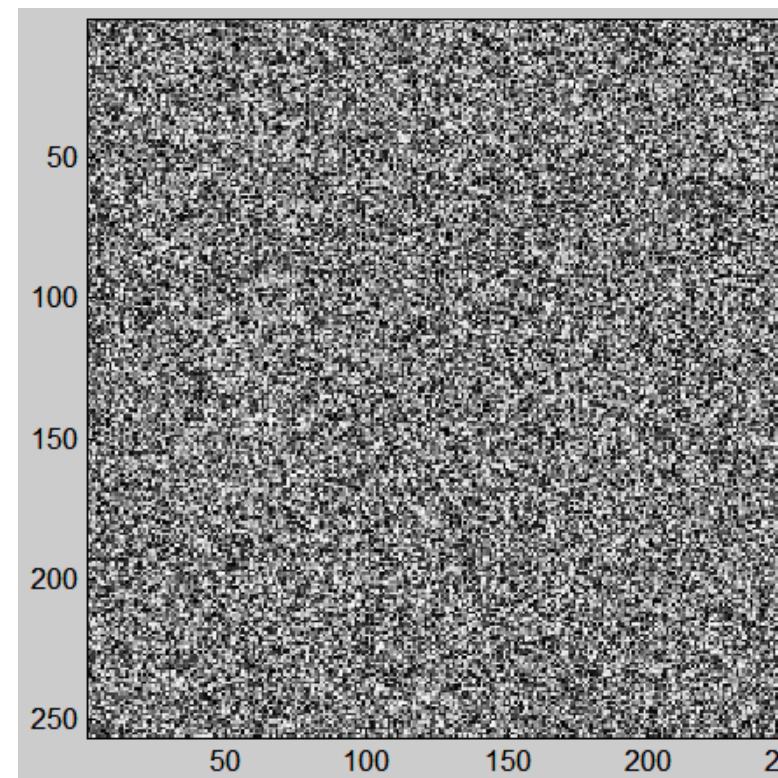
# 2-Matching histograms



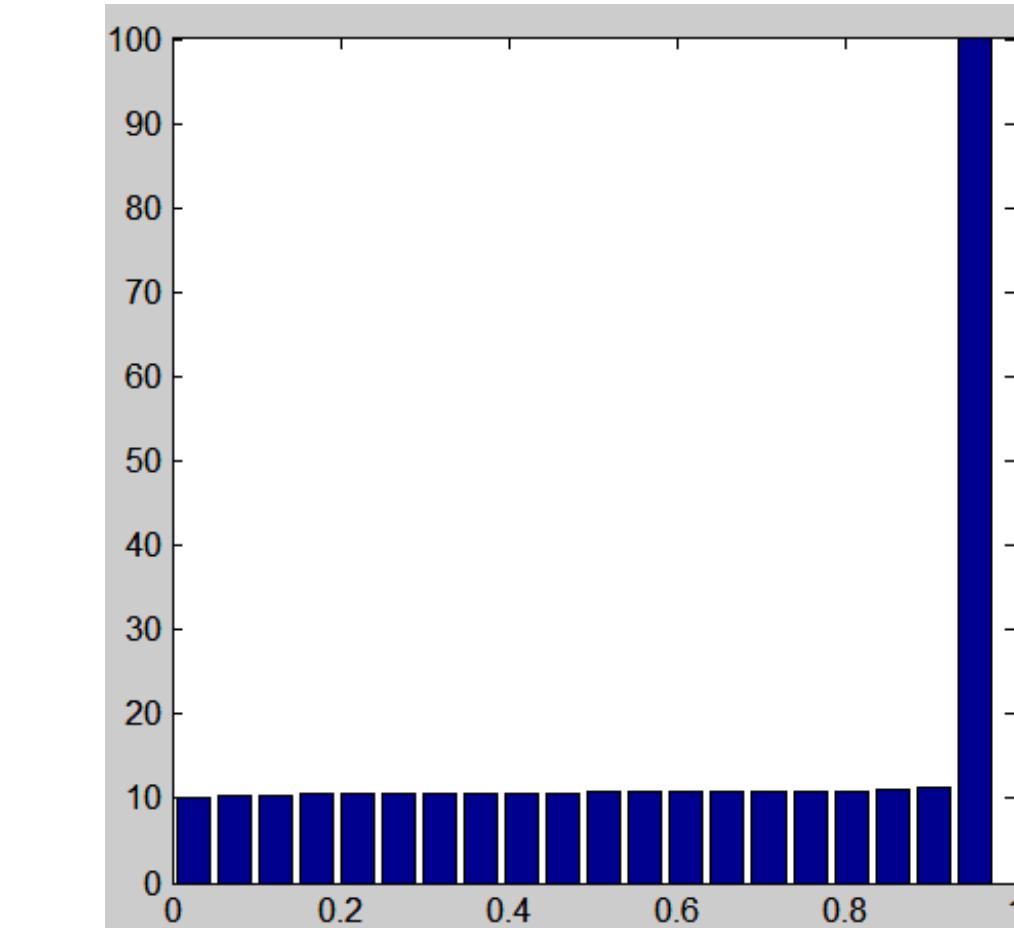
# Another example: Matching histograms



?

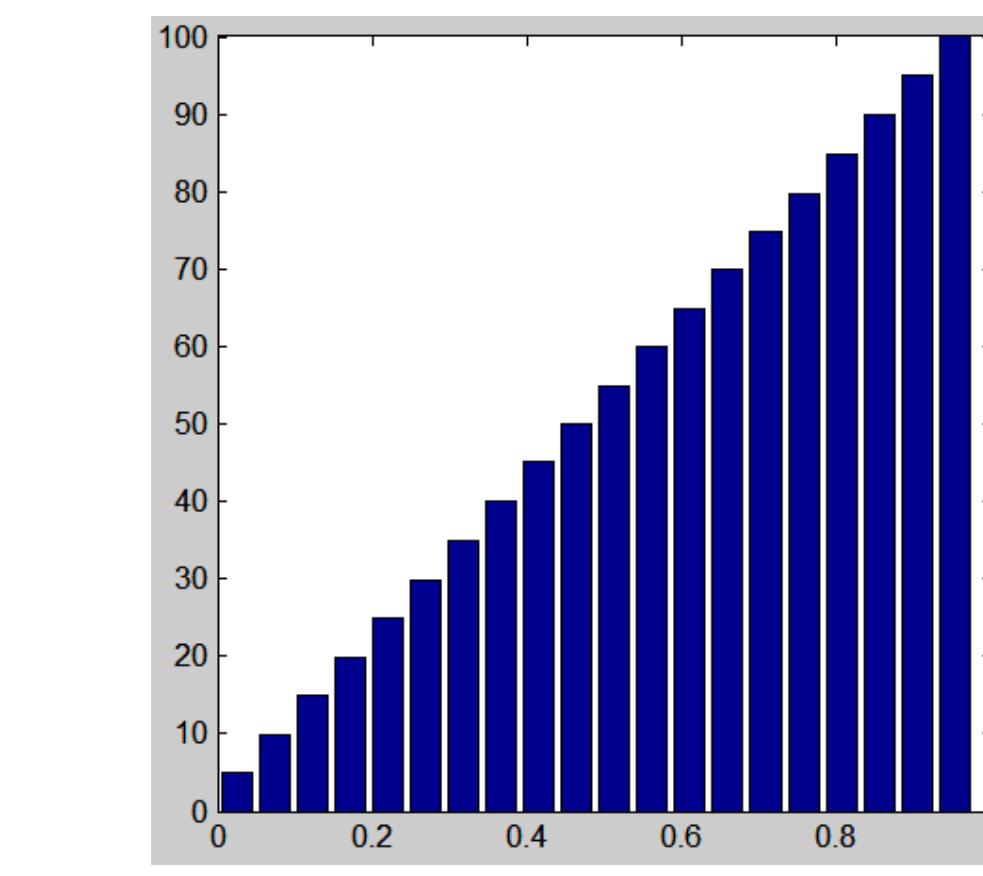


5% of pixels have an intensity value  
within the range [0.37, 0.41]



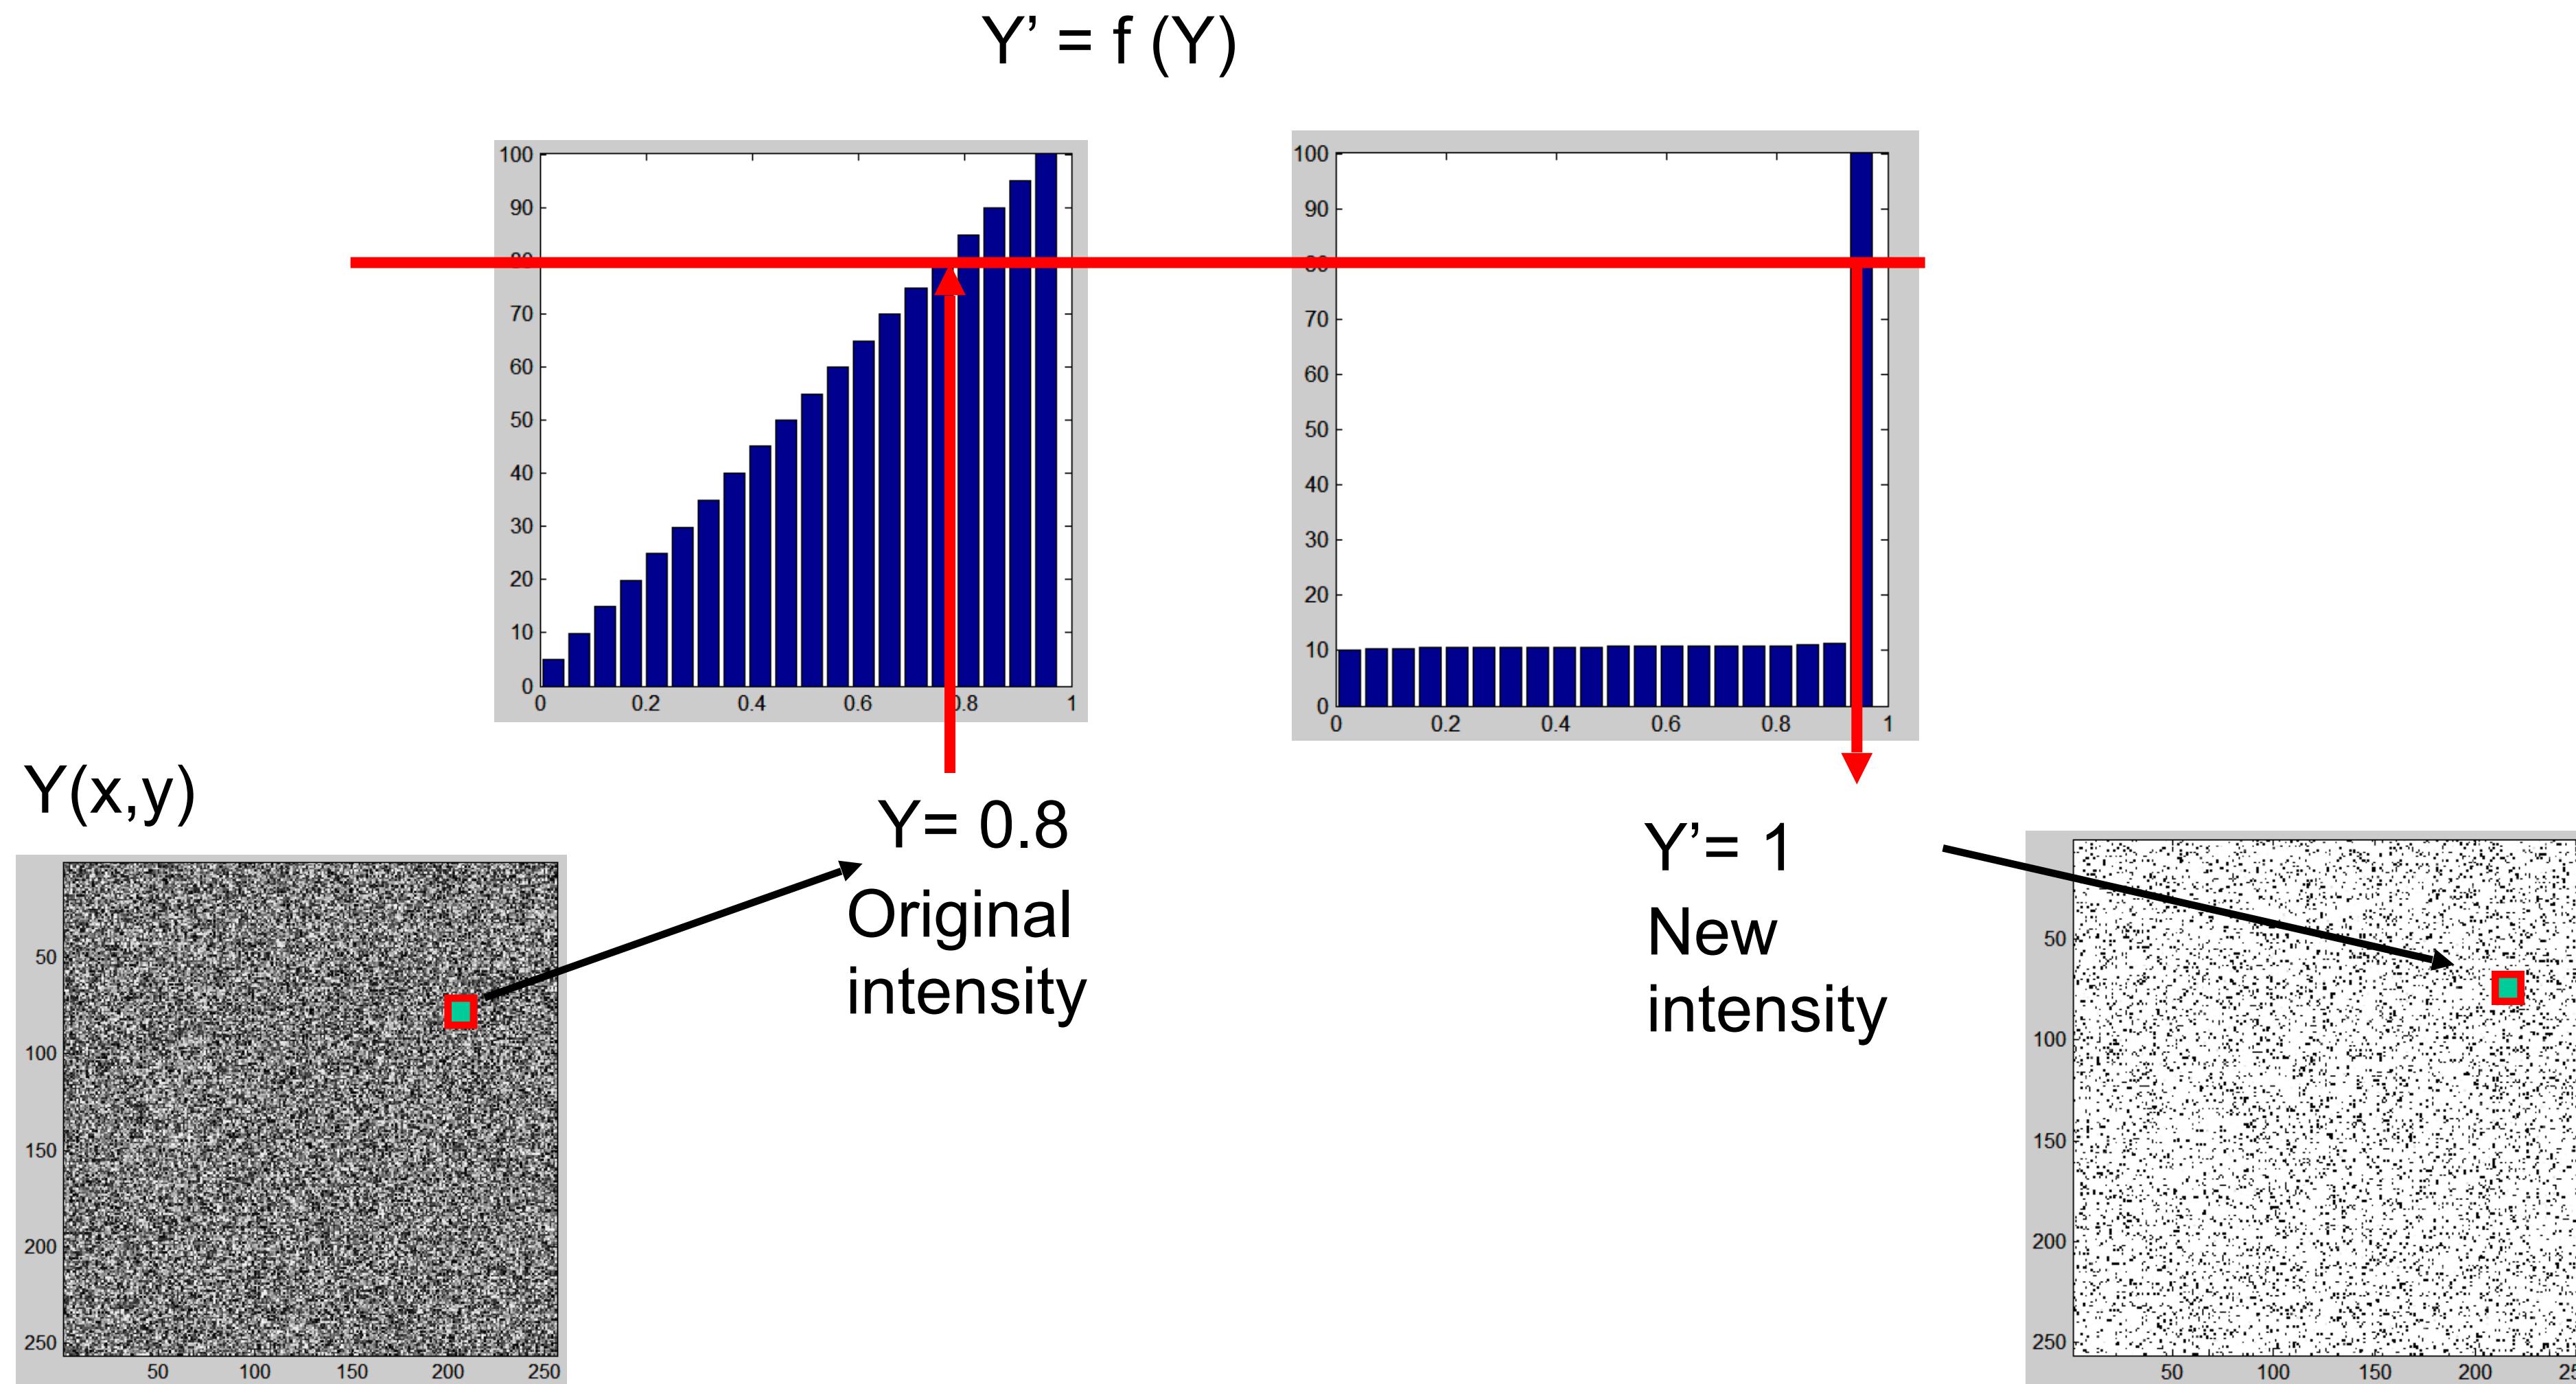
?

10% of pixels are black  
and 90% are white

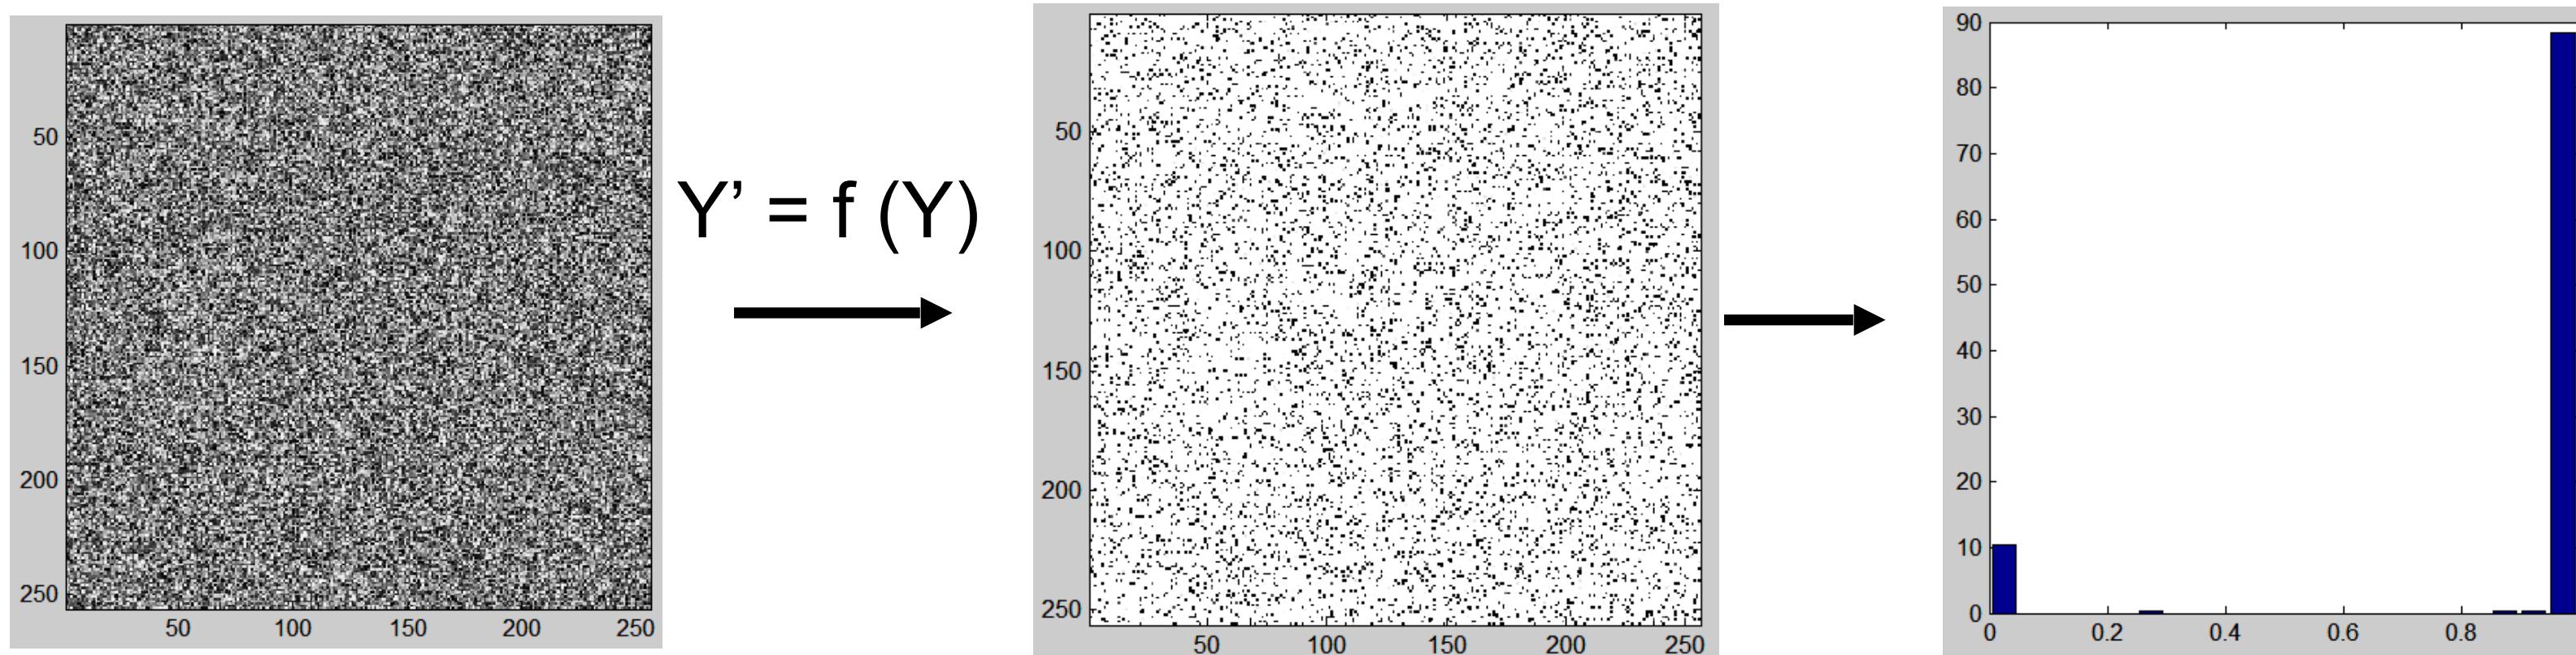
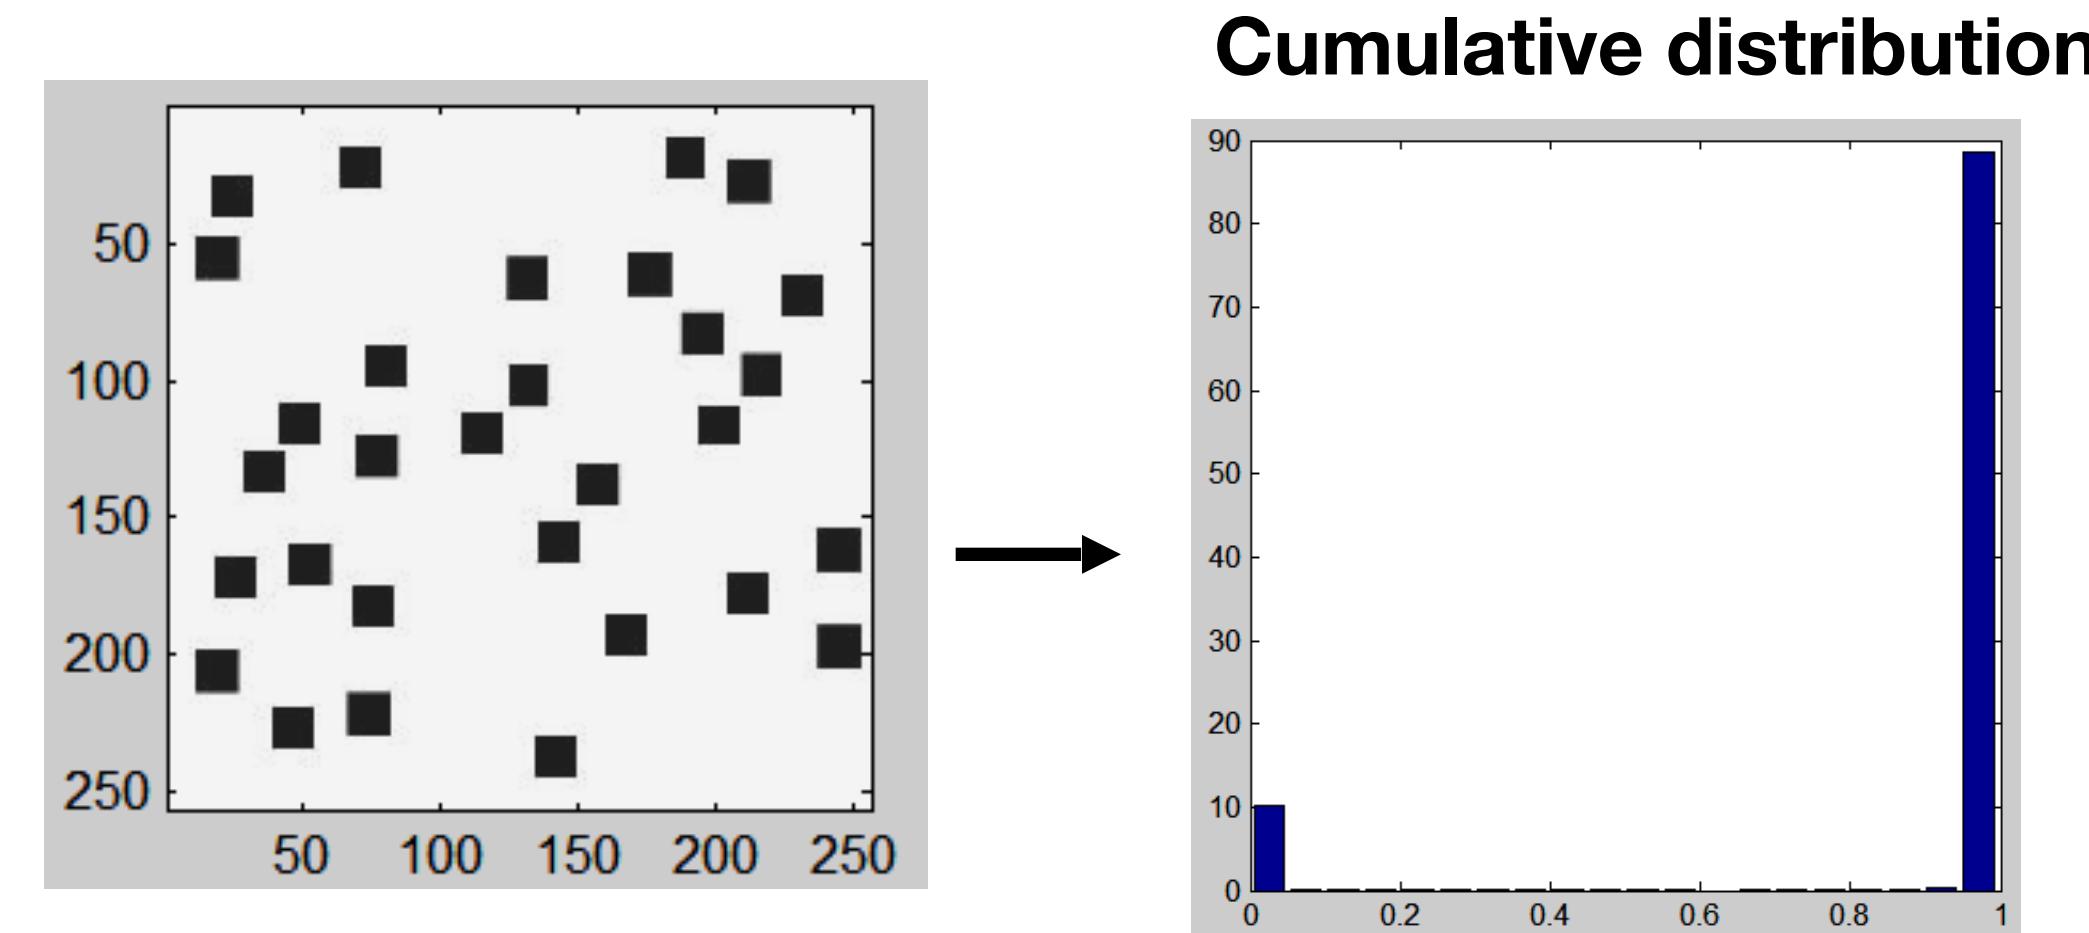


# Another example: Matching histograms

The function  $f$  is just a look up table: it says, change all the pixels of value  $Y$  into a value  $f(Y)$ .

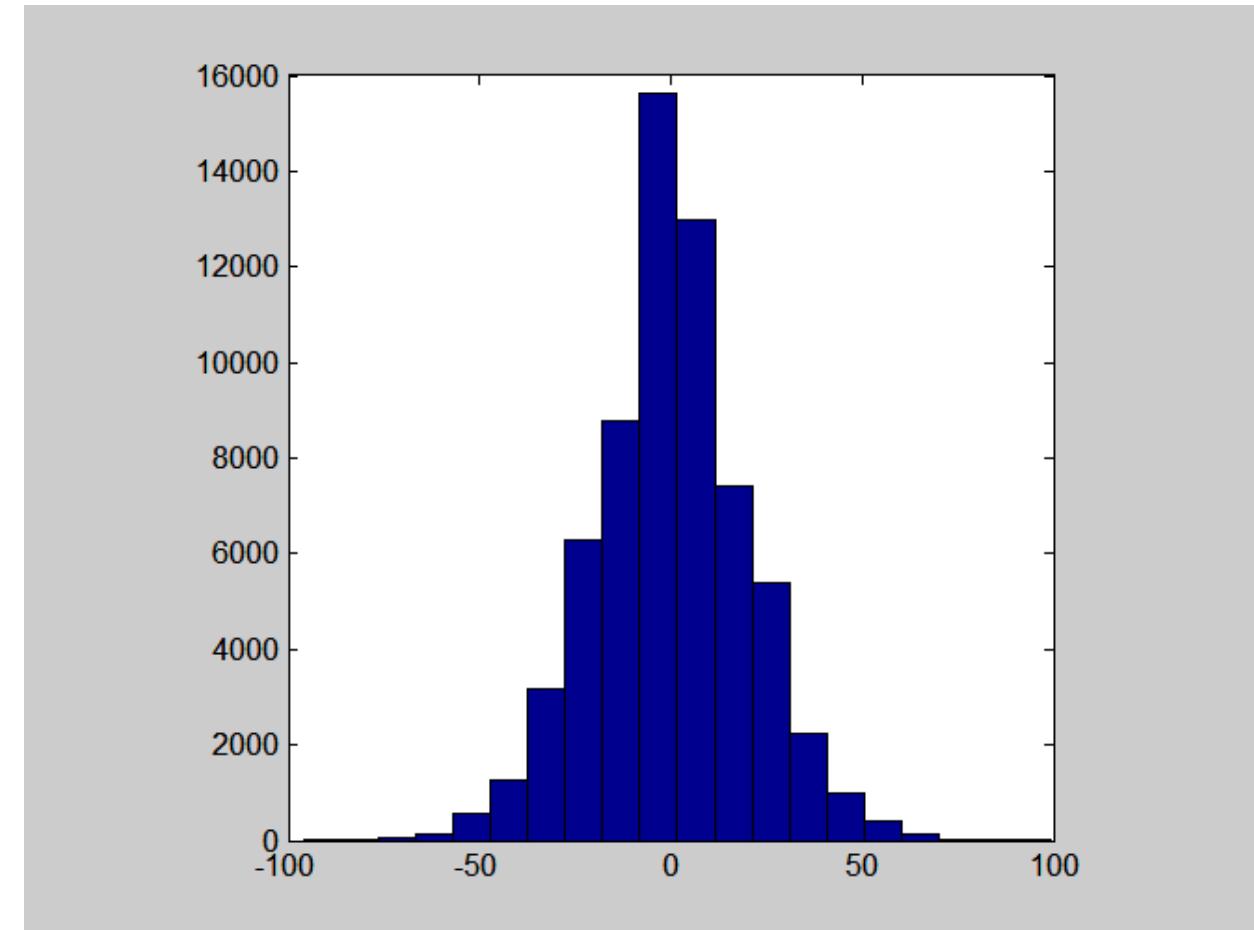
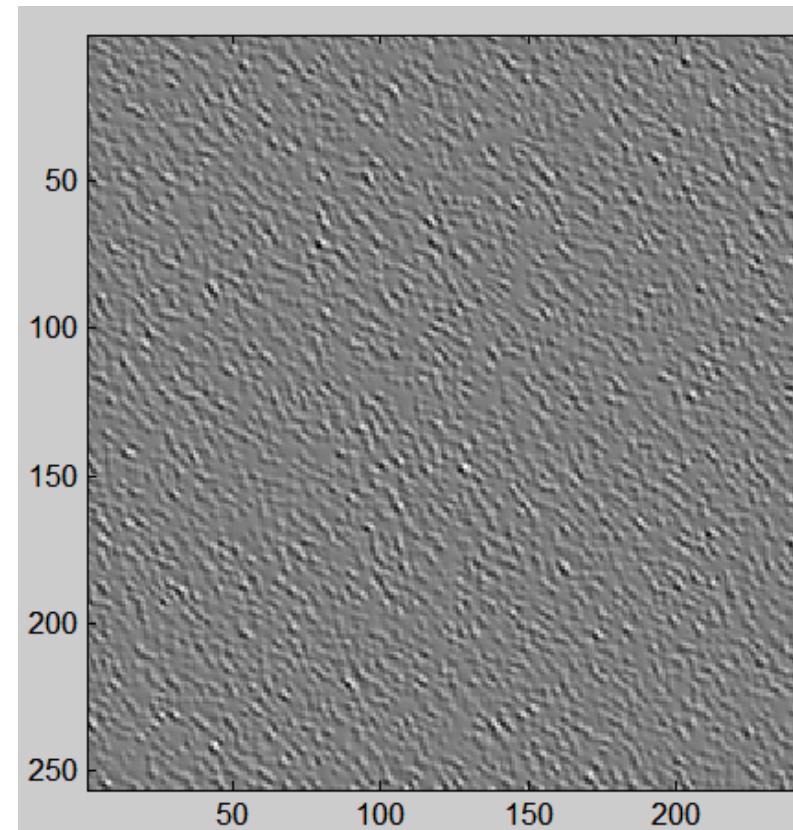
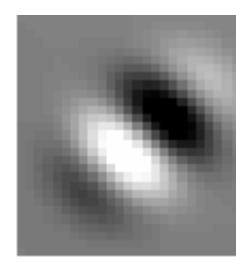
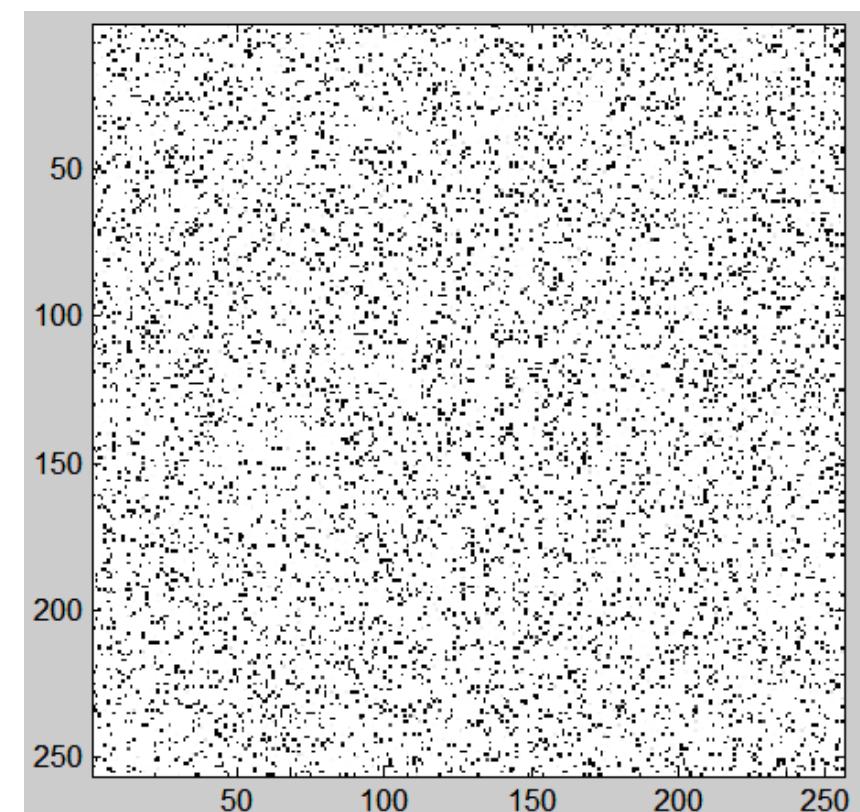
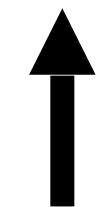
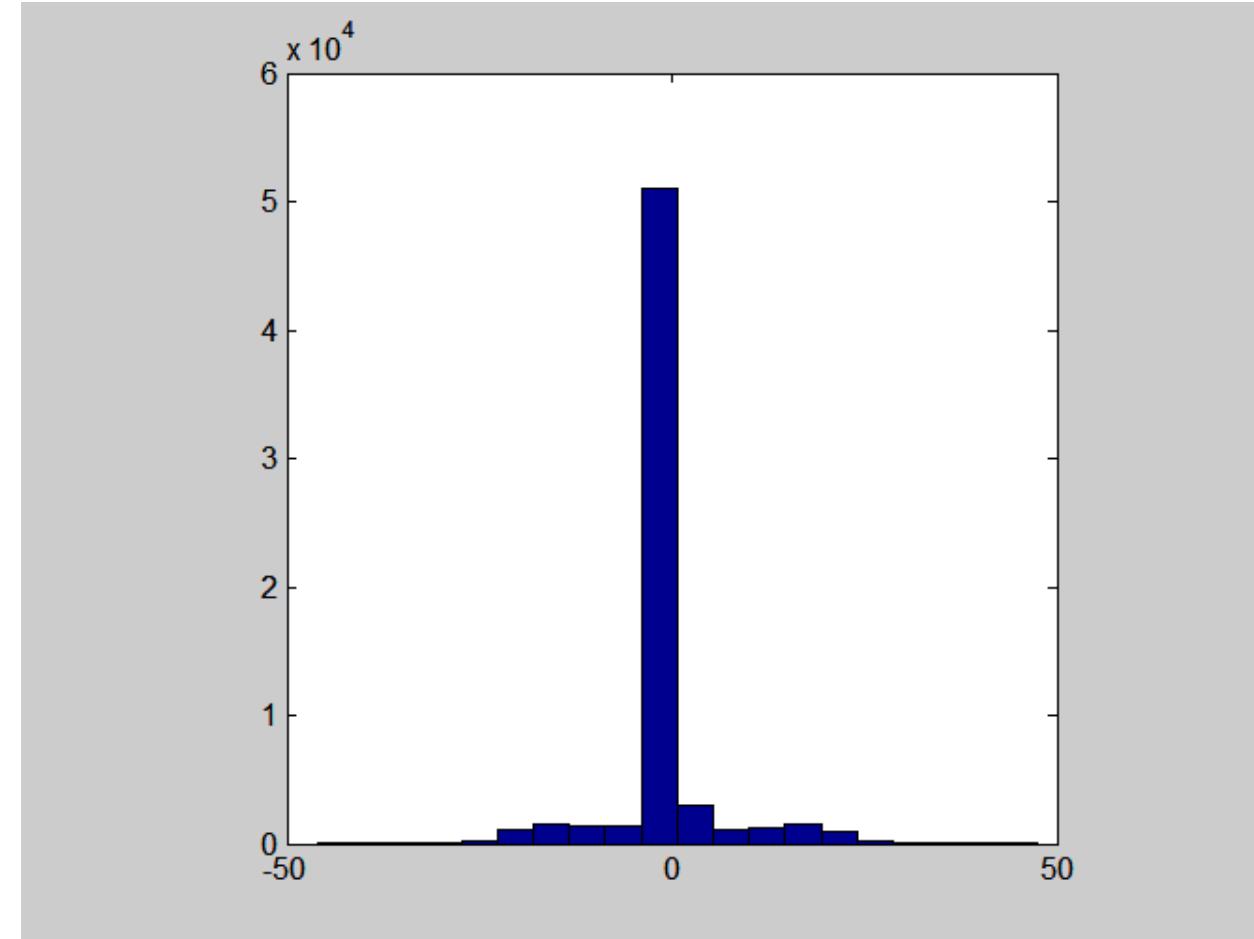
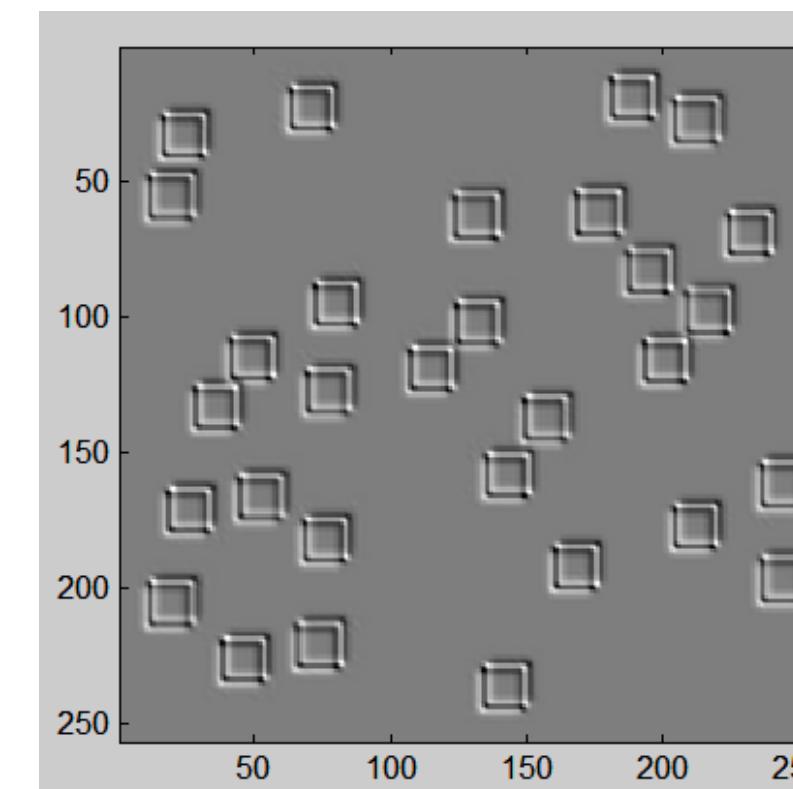
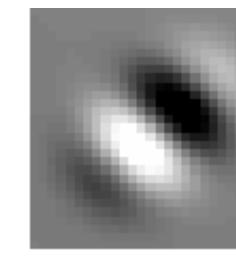
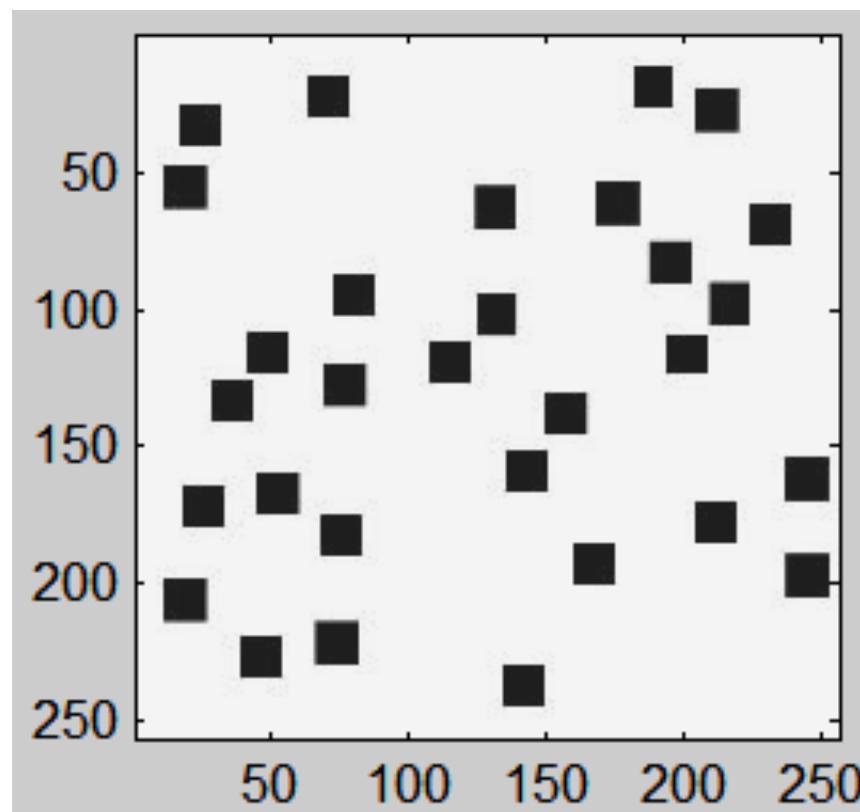


# Another example: Matching histograms

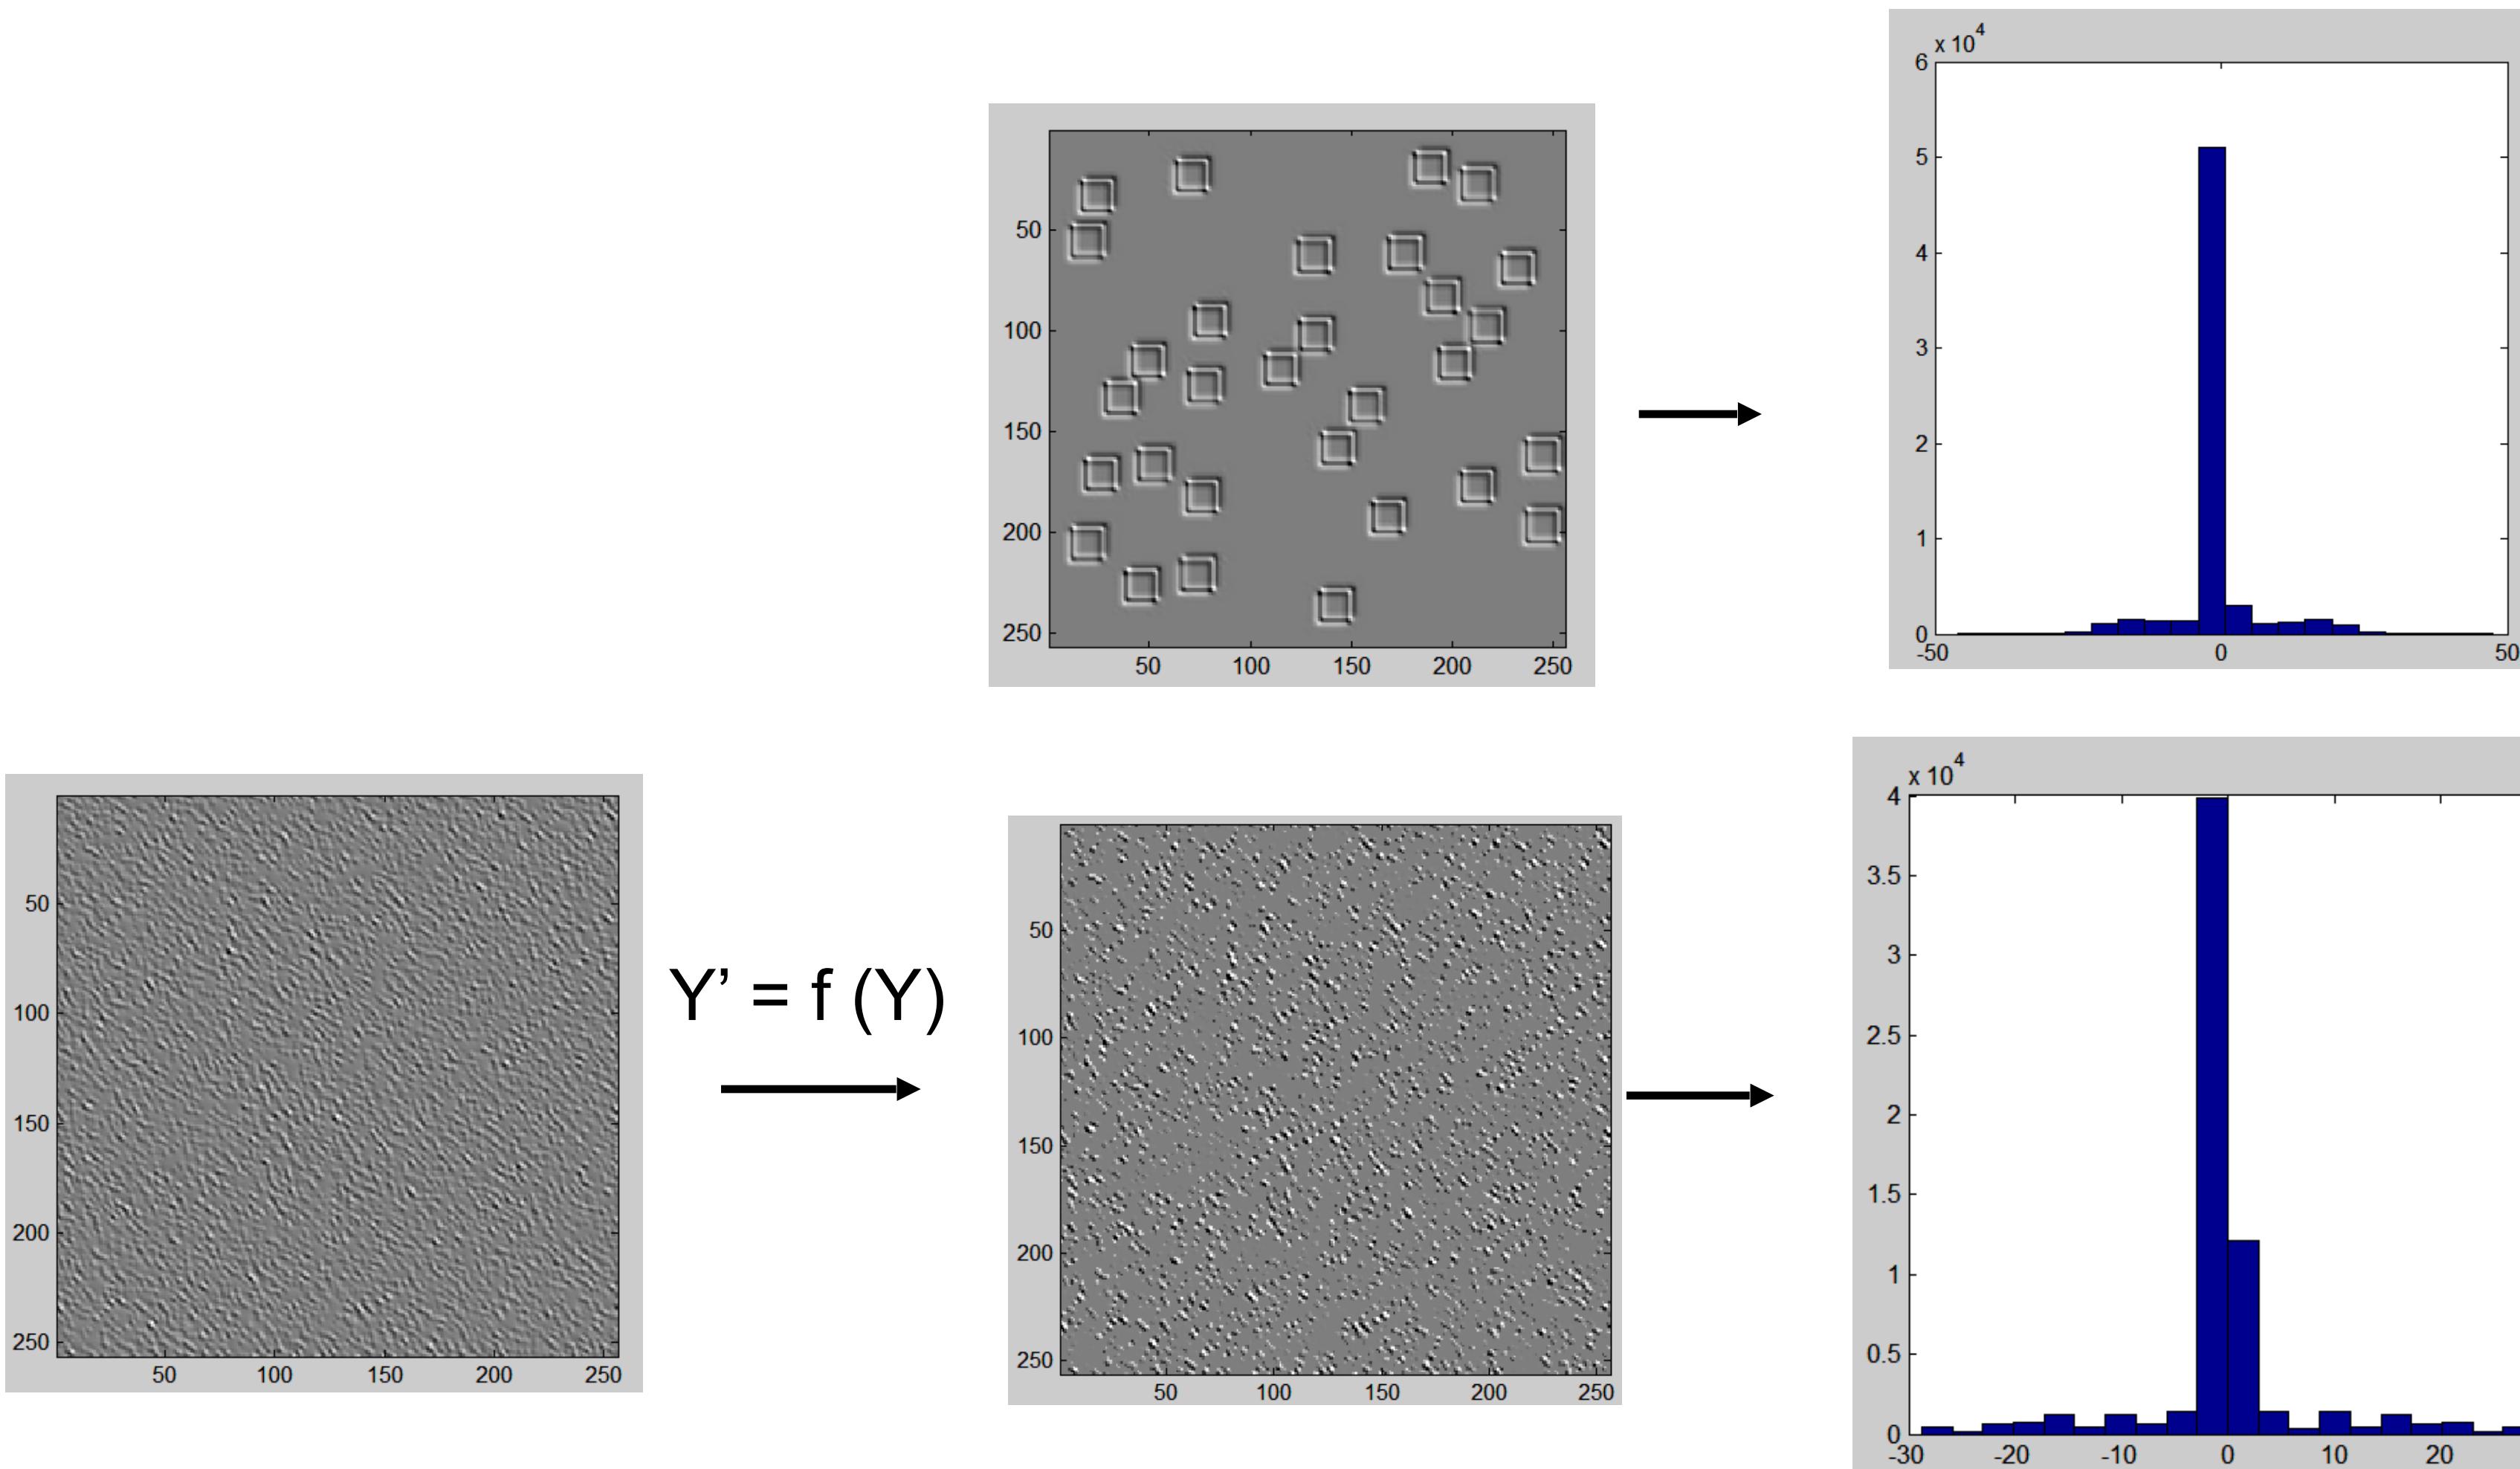


In this example,  $f$  is a step function.

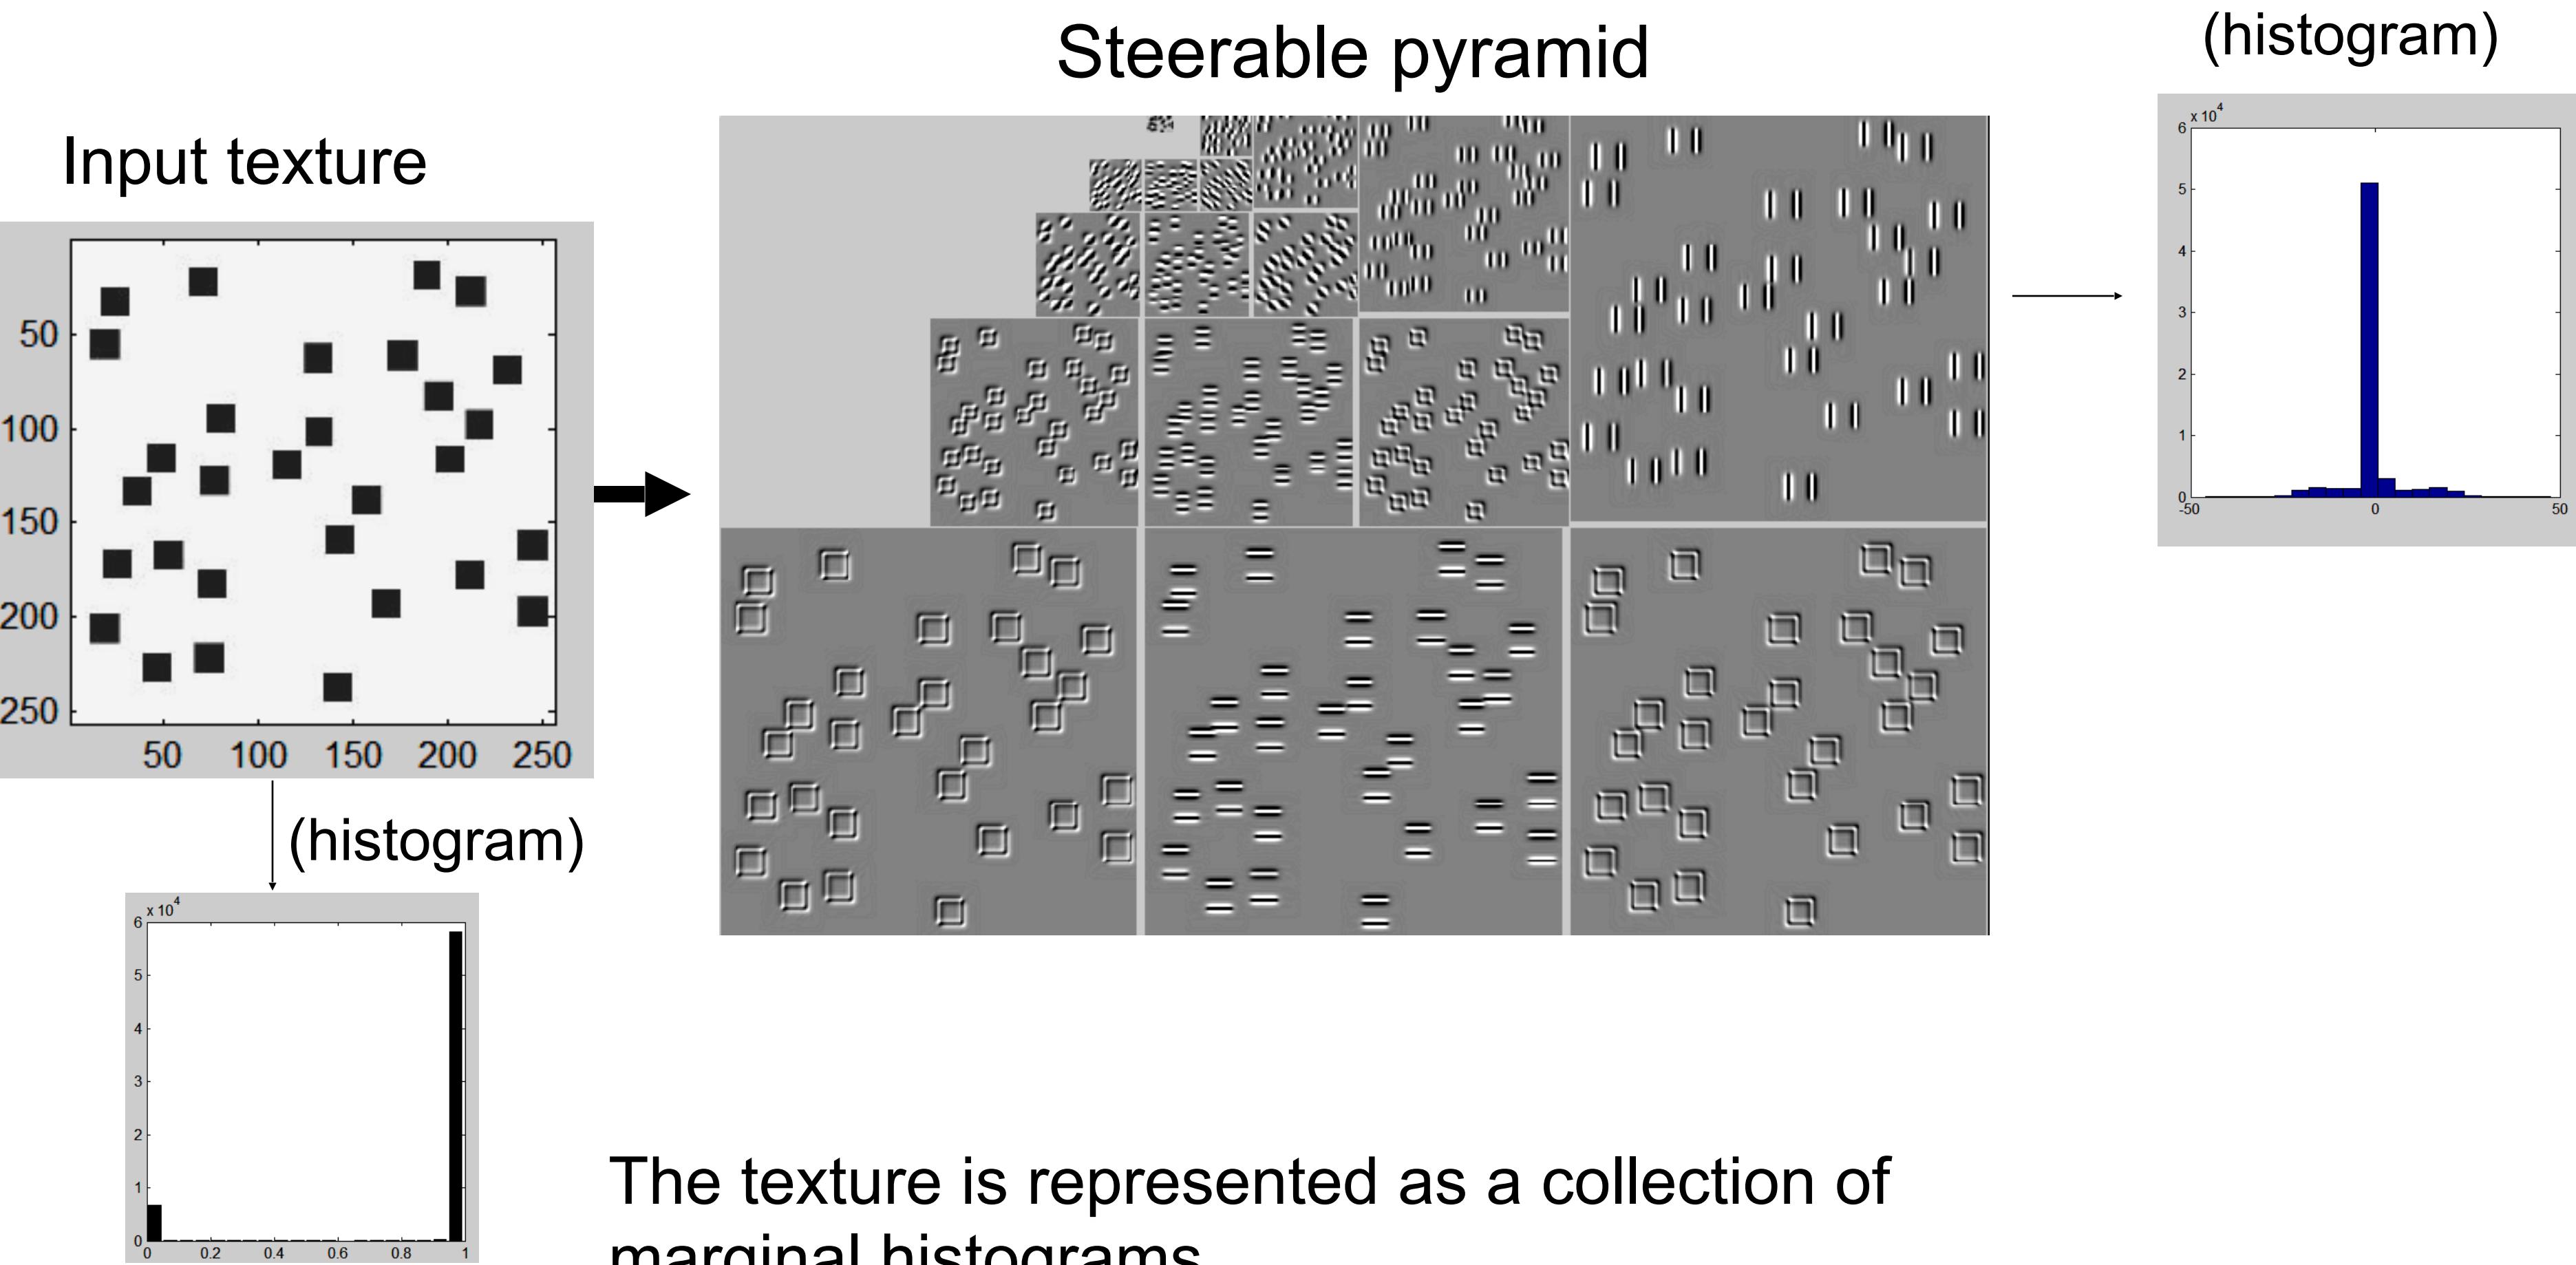
# Matching histograms of a subband



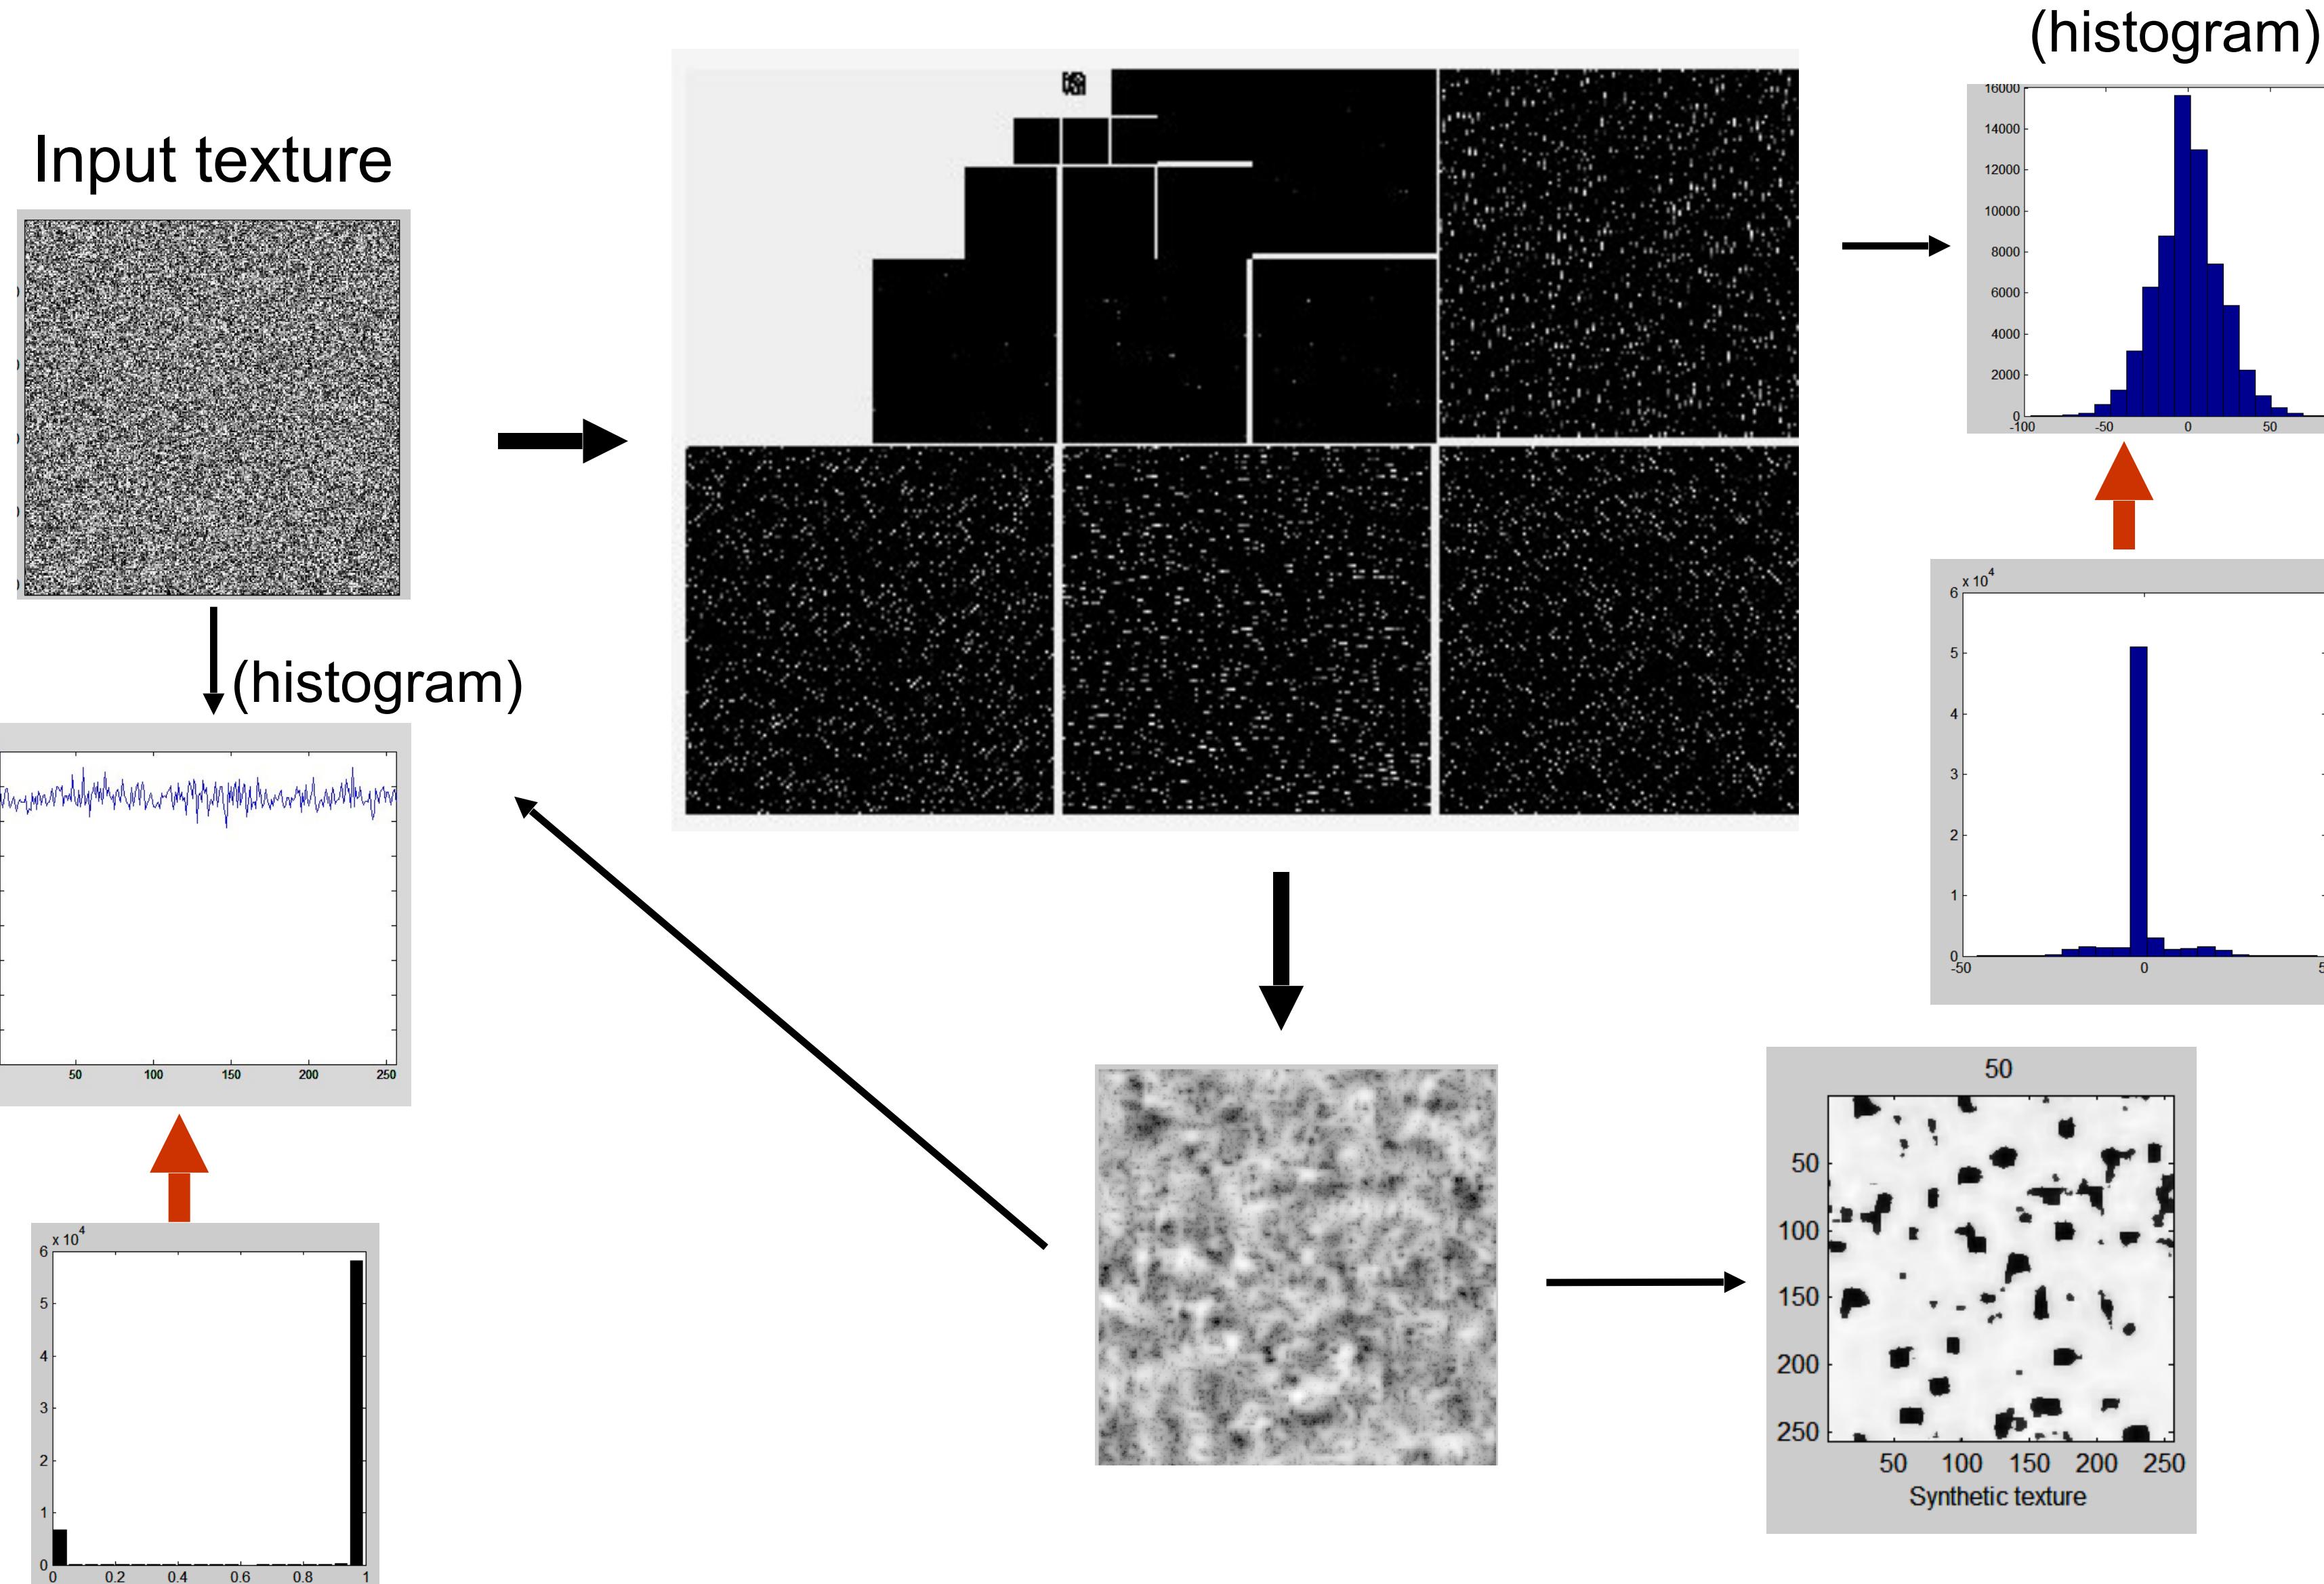
# Matching histograms of a subband



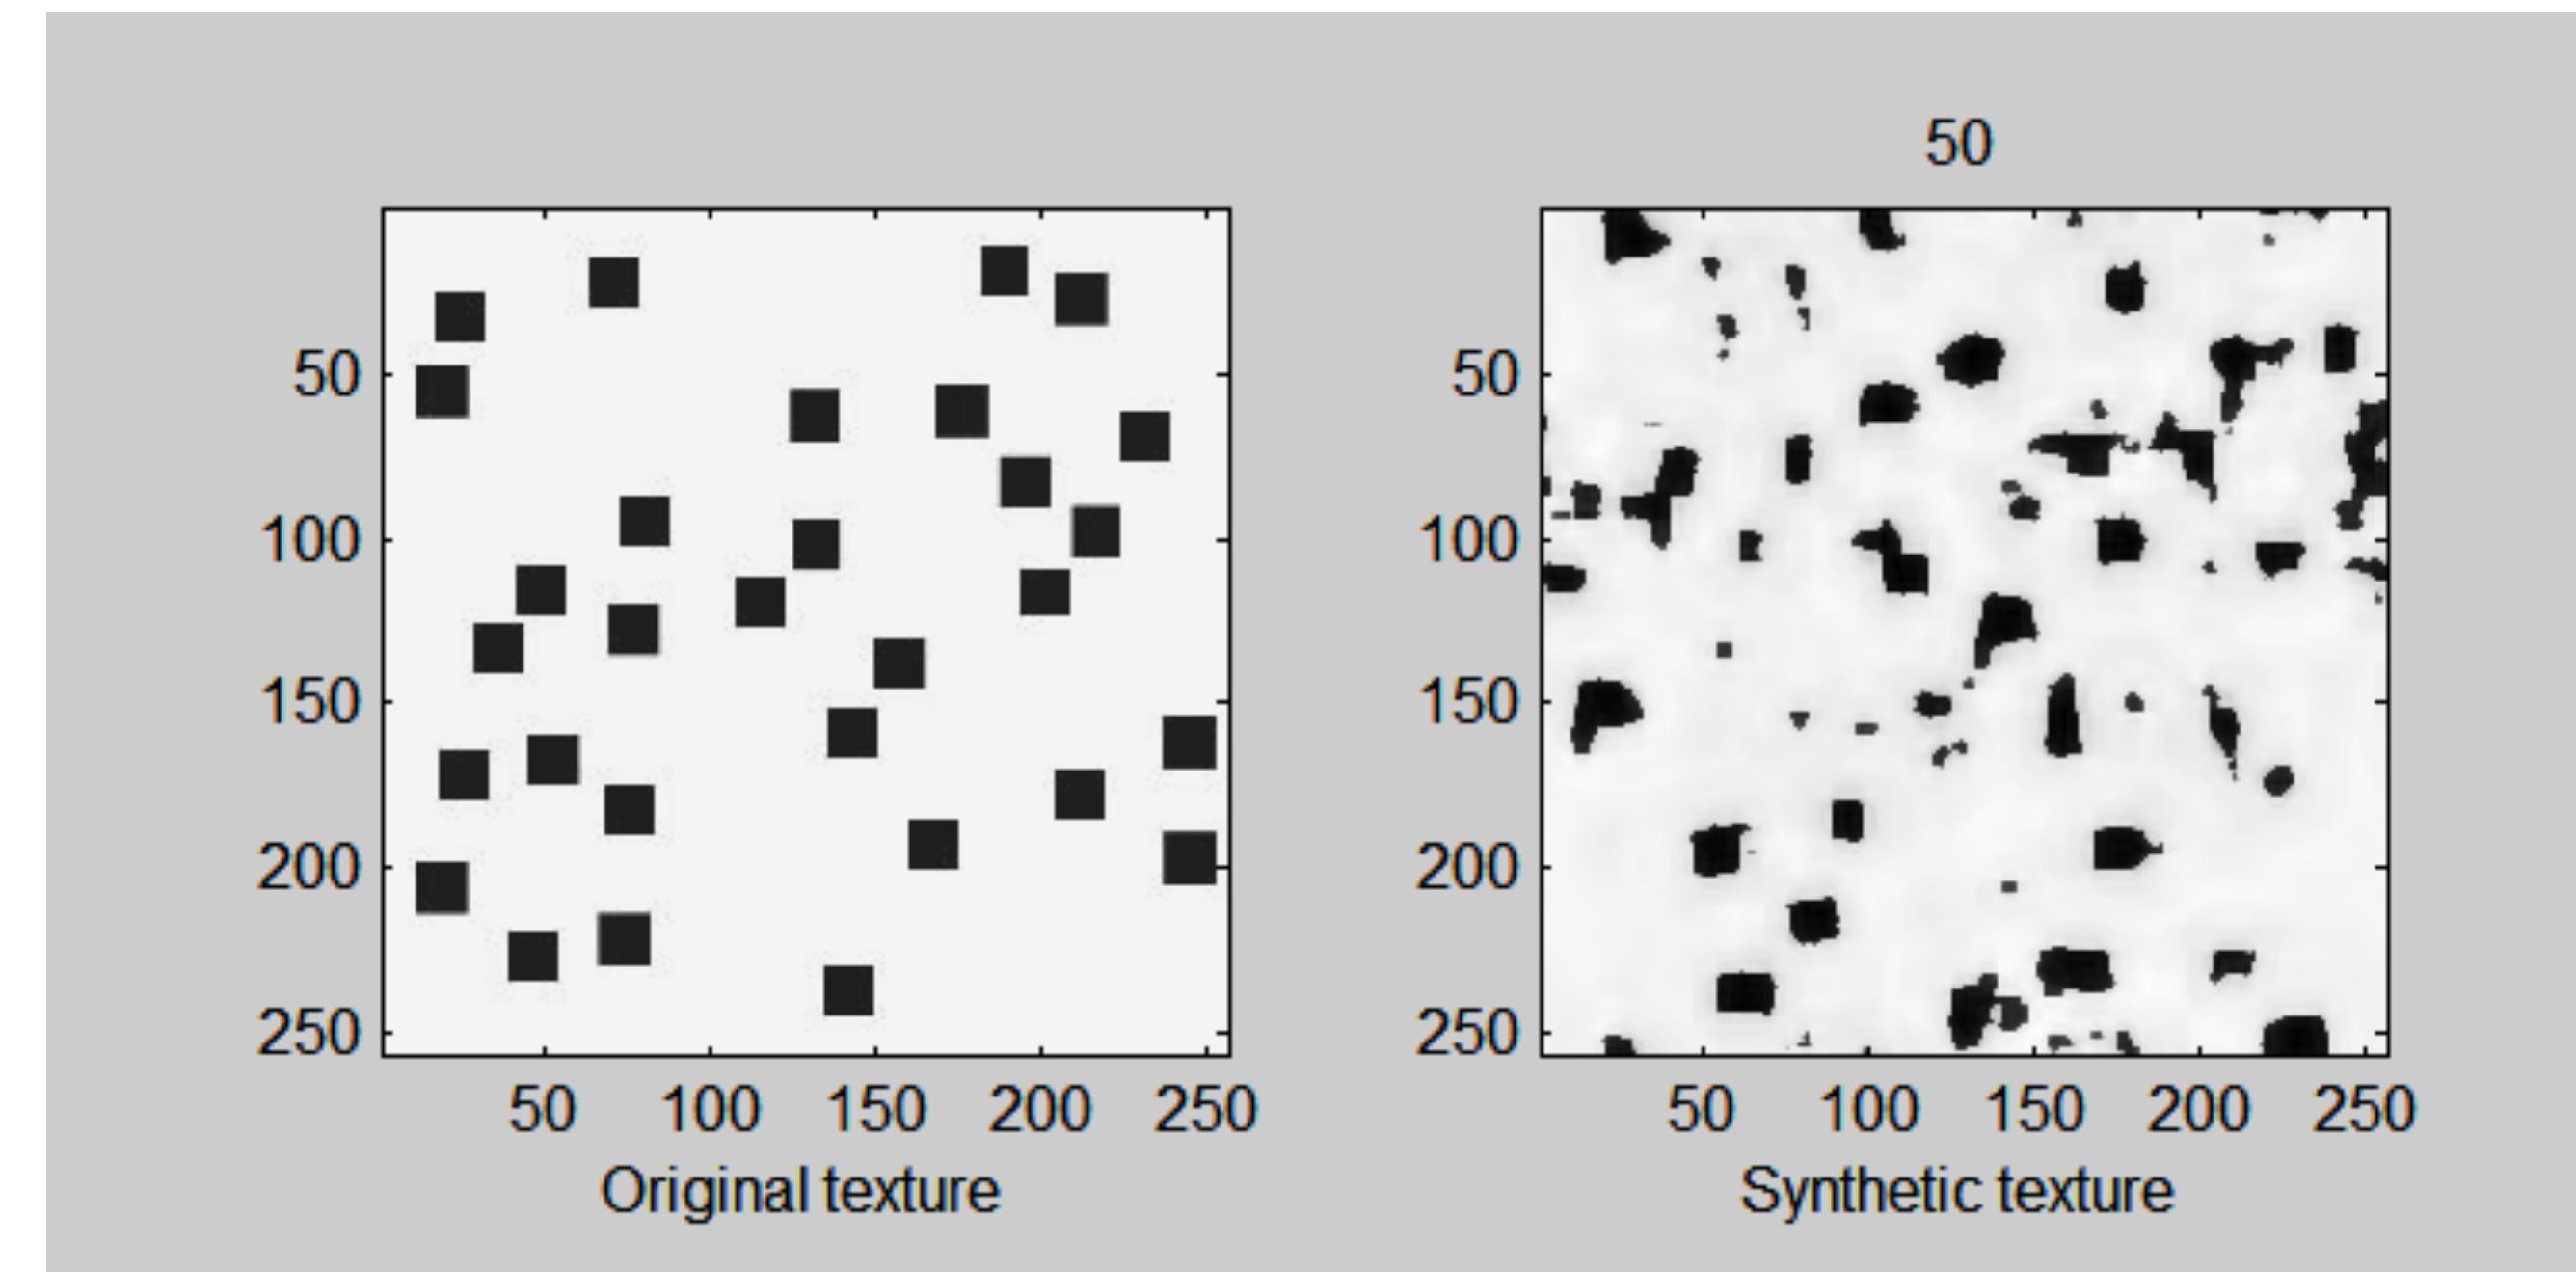
# Texture analysis

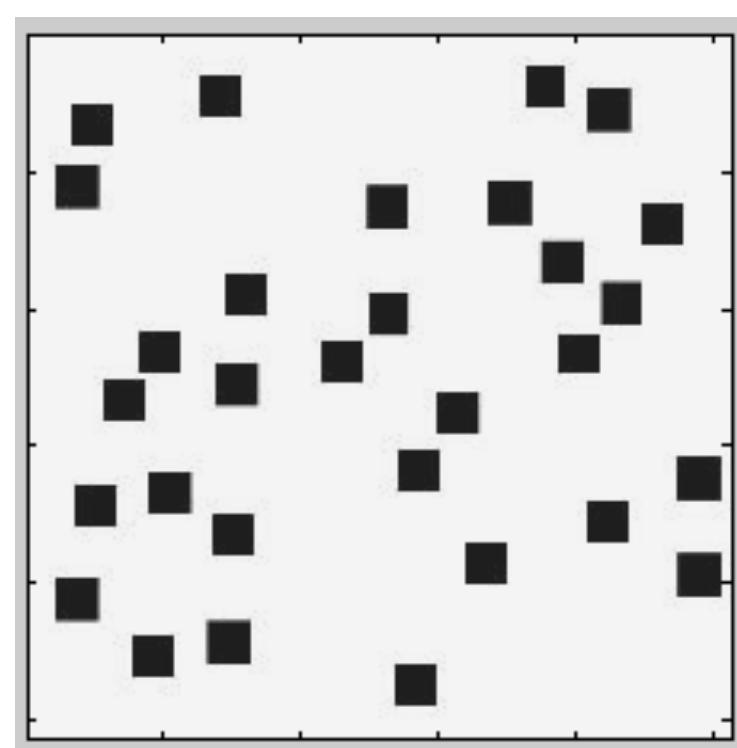


# Texture synthesis



# Why does it work? (sort of)

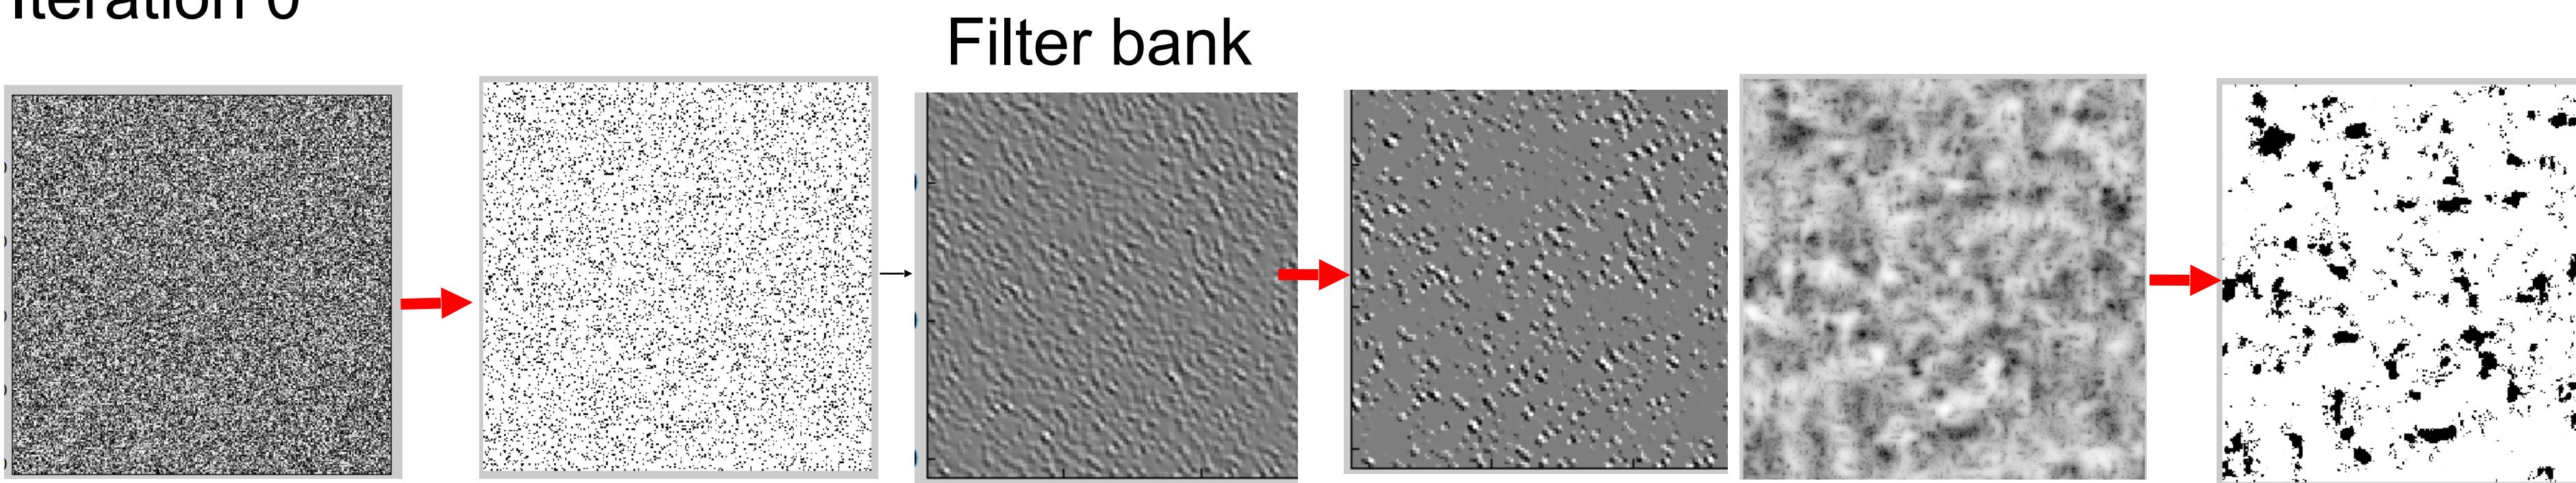




# Why does it work? (sort of)

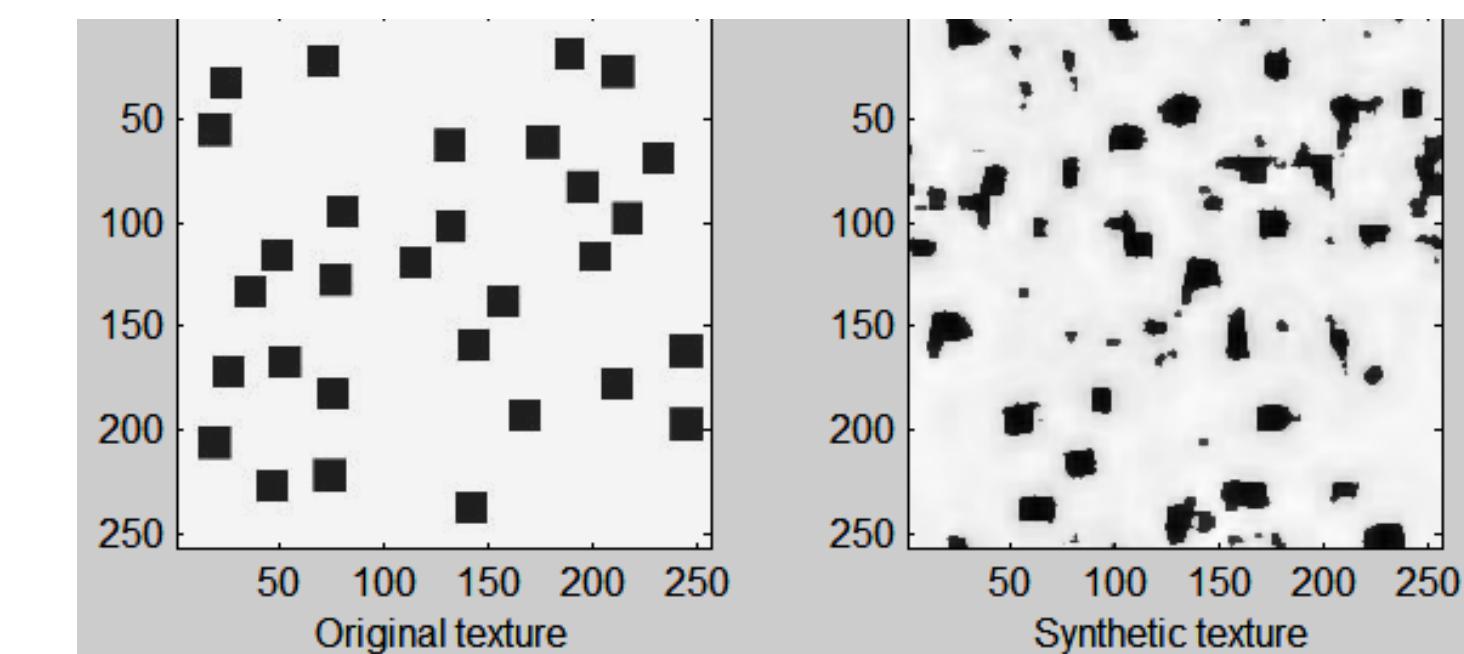
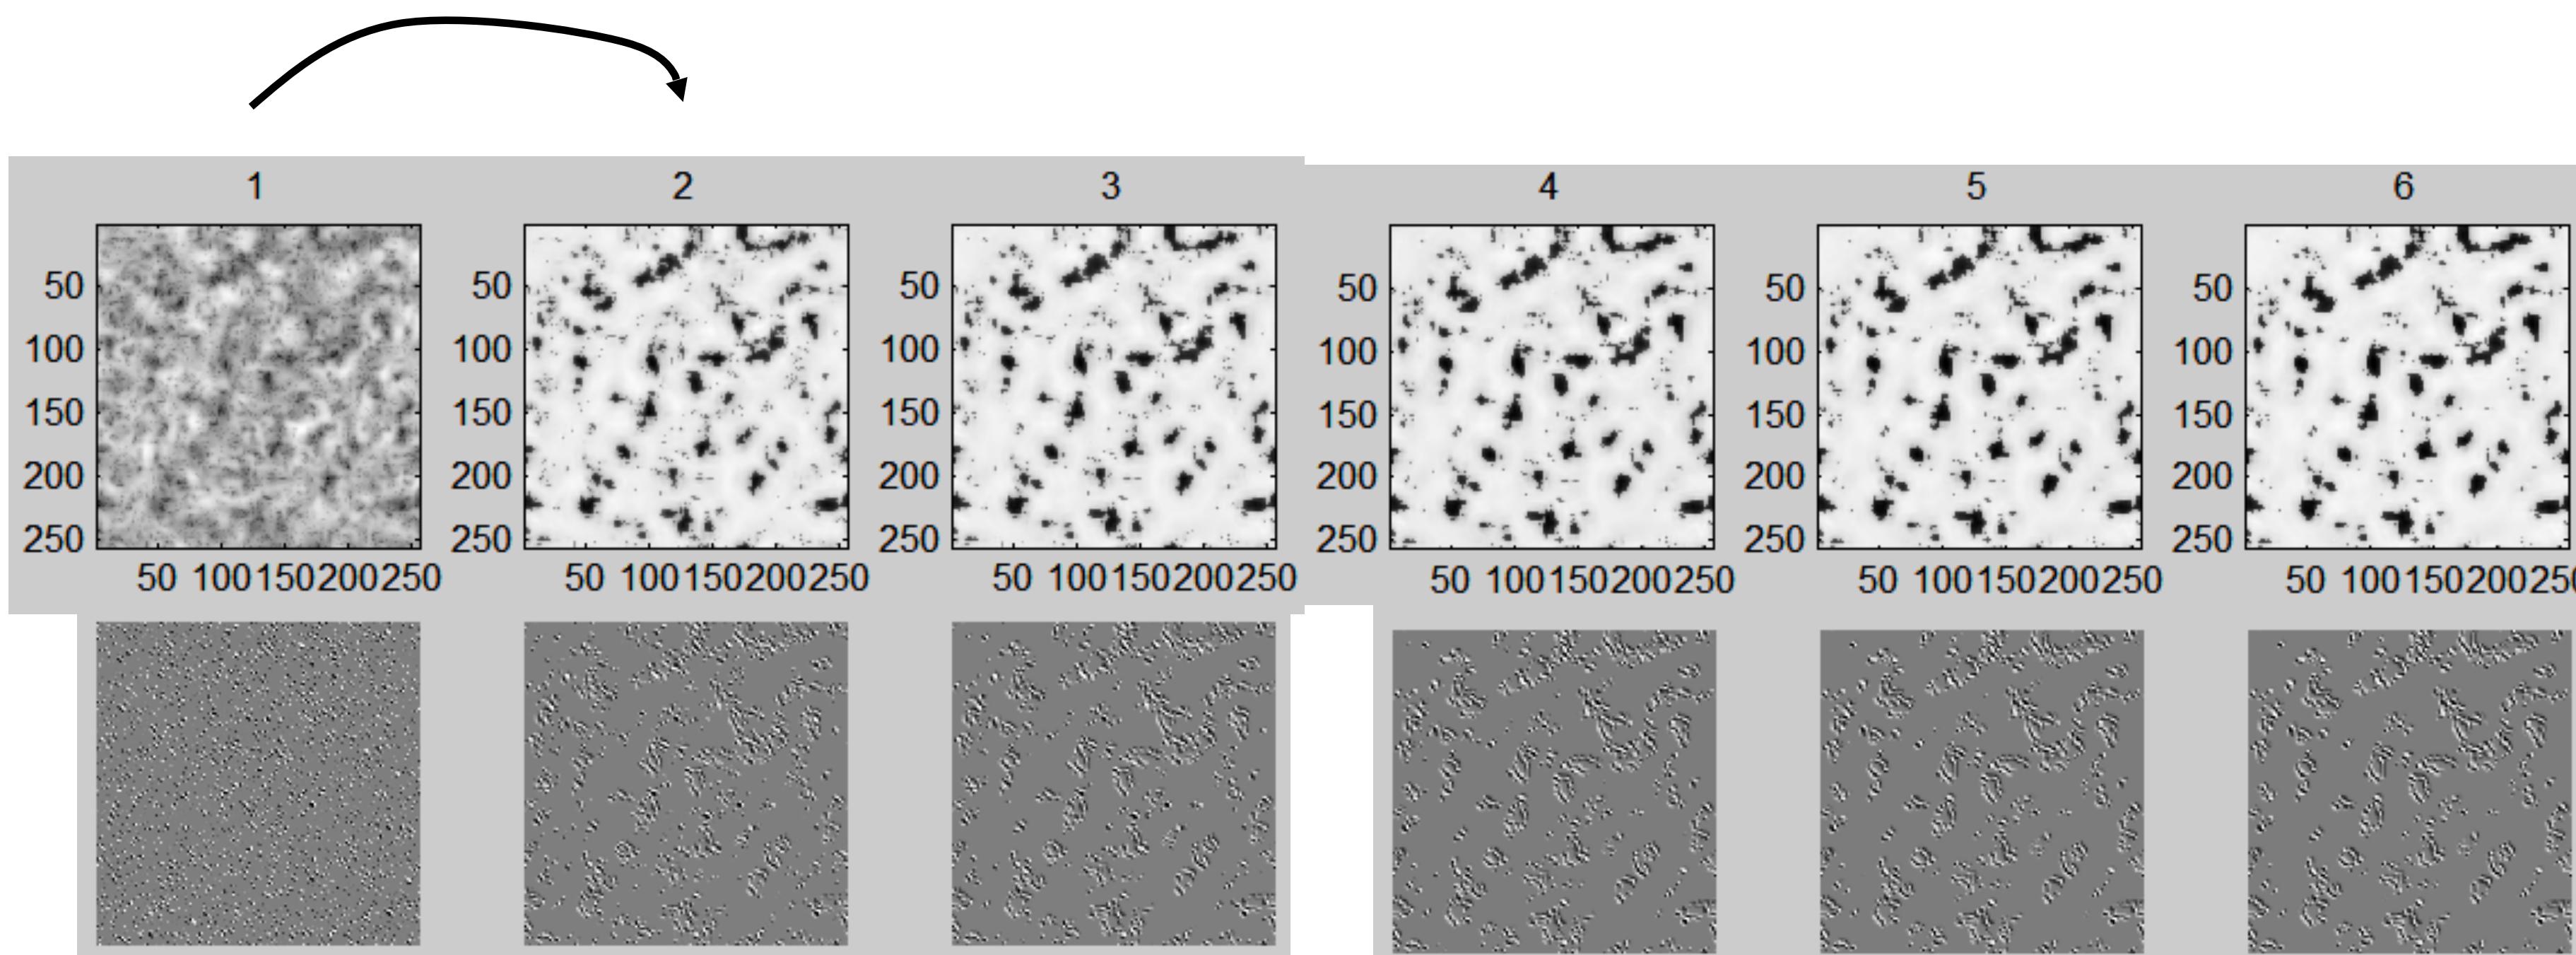
The black and white blocks appear by thresholding ( $f$ ) a blobby image

Iteration 0



# Why does it work? (sort of)

The black and white blocks appear by thresholding (f) a blobby image



# Examples from the paper



Figure 3: In each pair left image is original and right image is synthetic: stucco, iridescent ribbon, green marble, panda fur, slag stone, figured yew wood.

Heeger and Bergen, 1995

# Examples from the paper

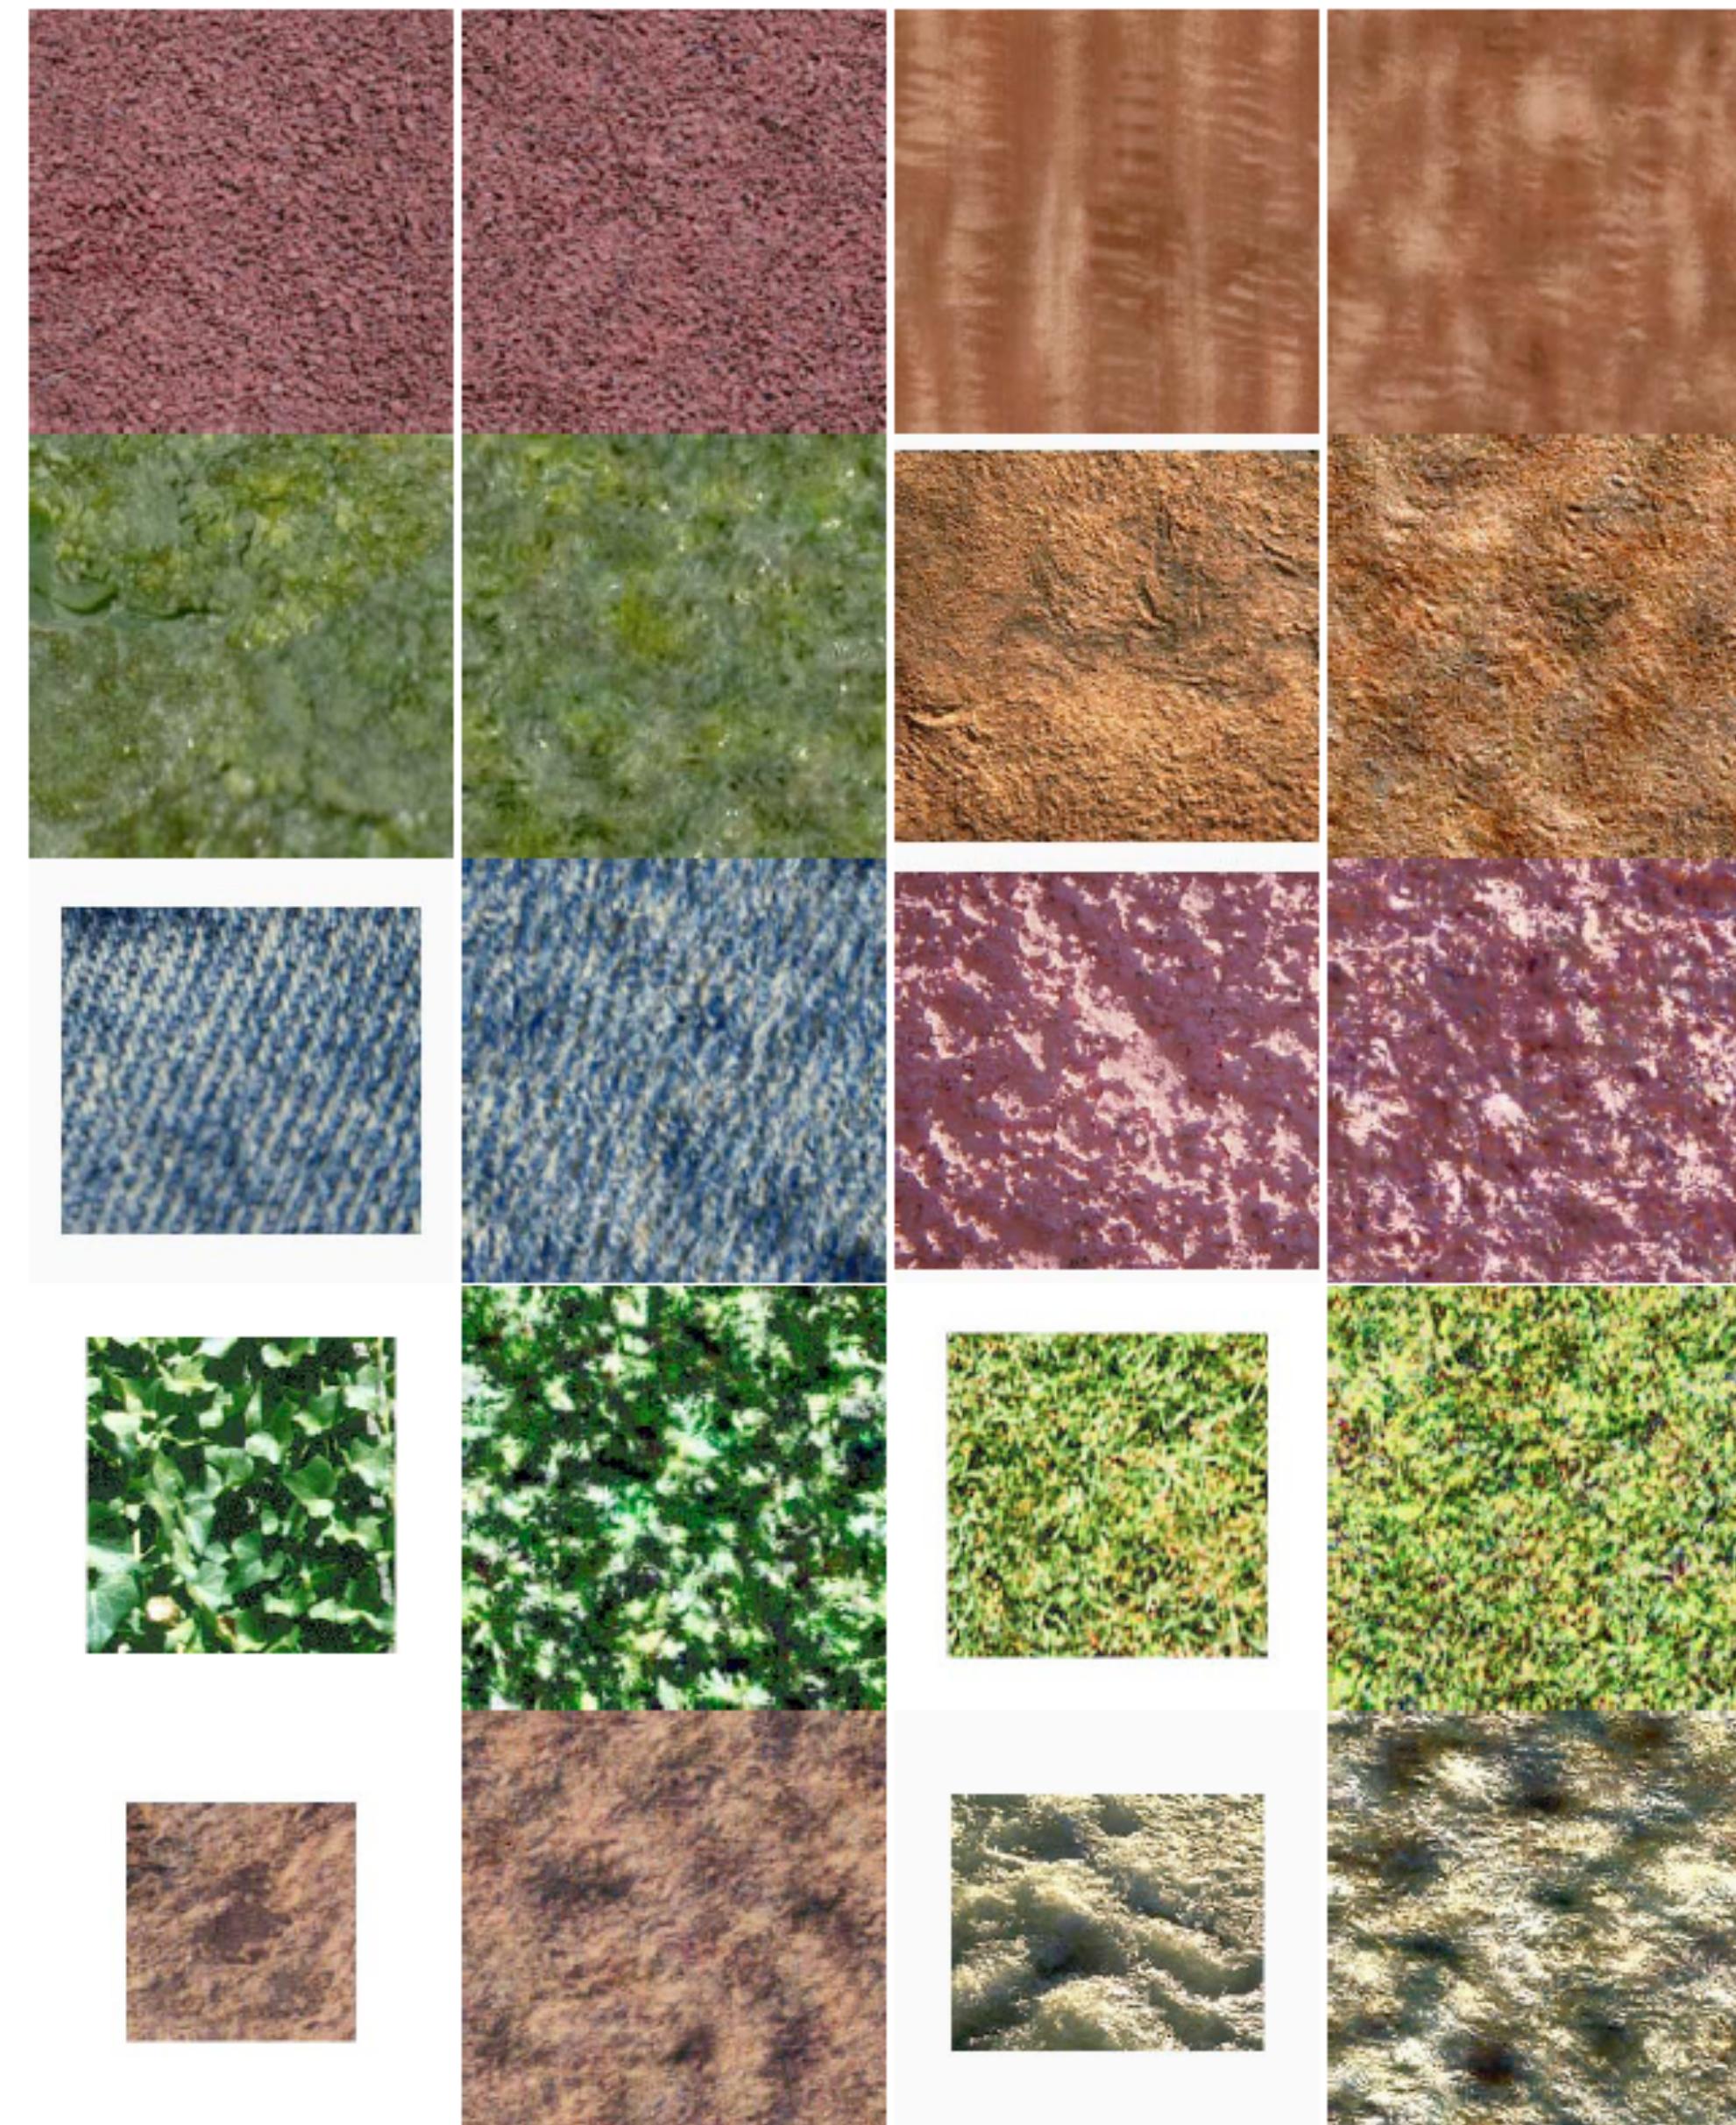
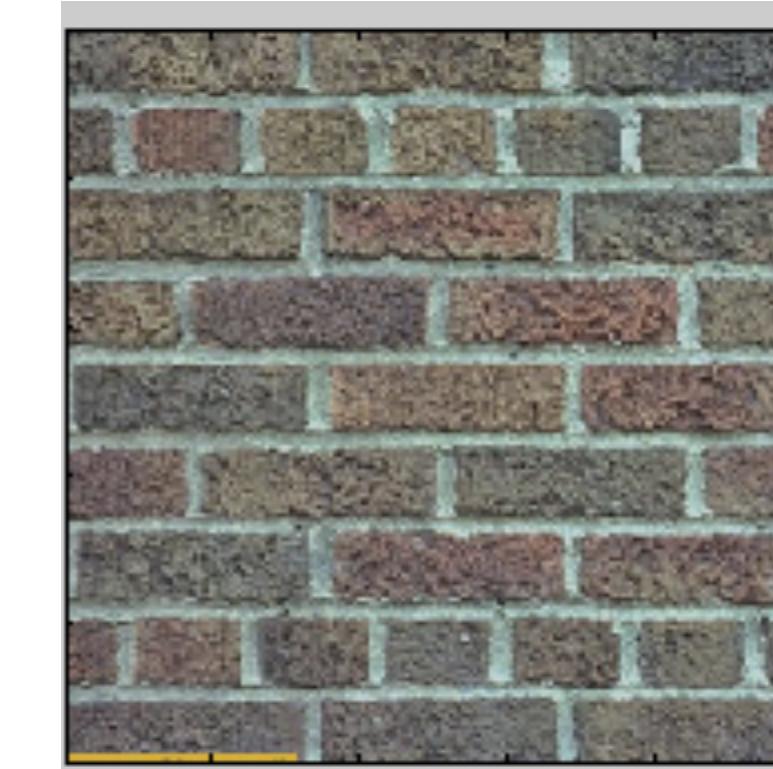
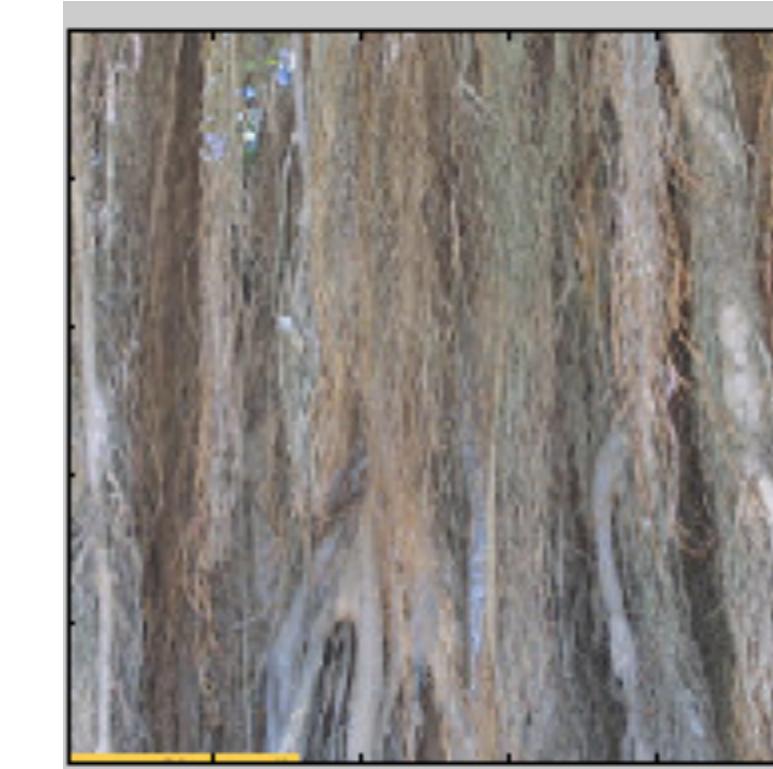


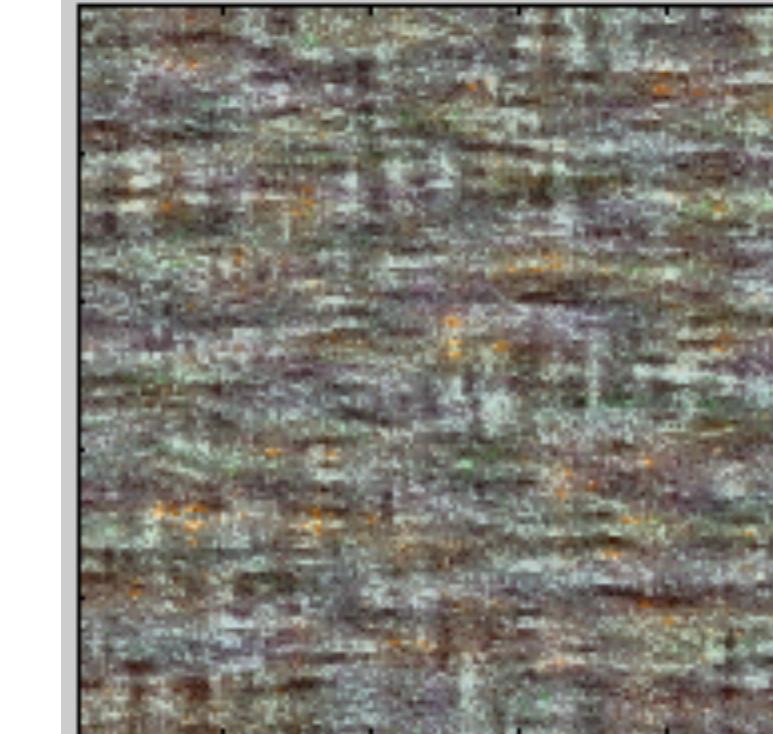
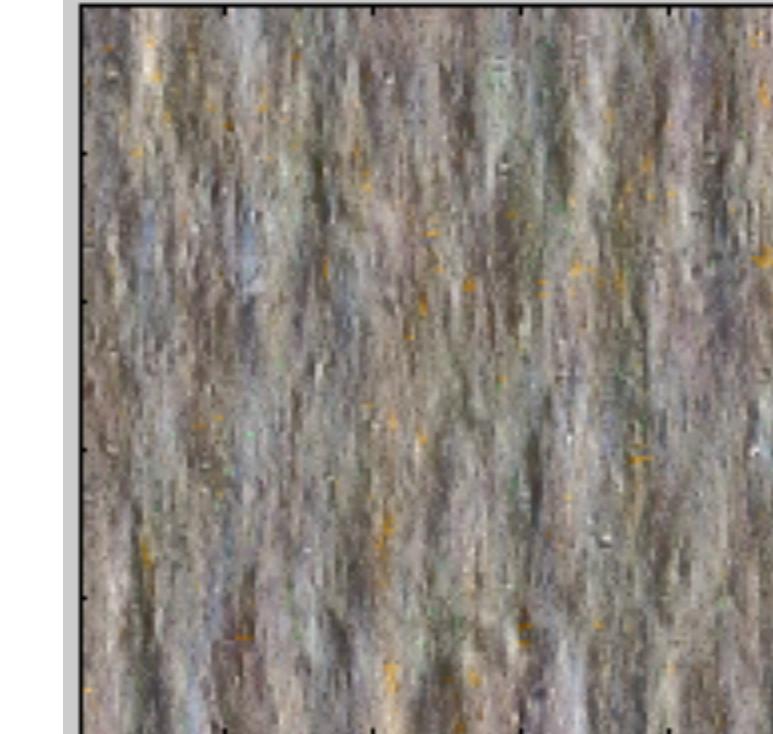
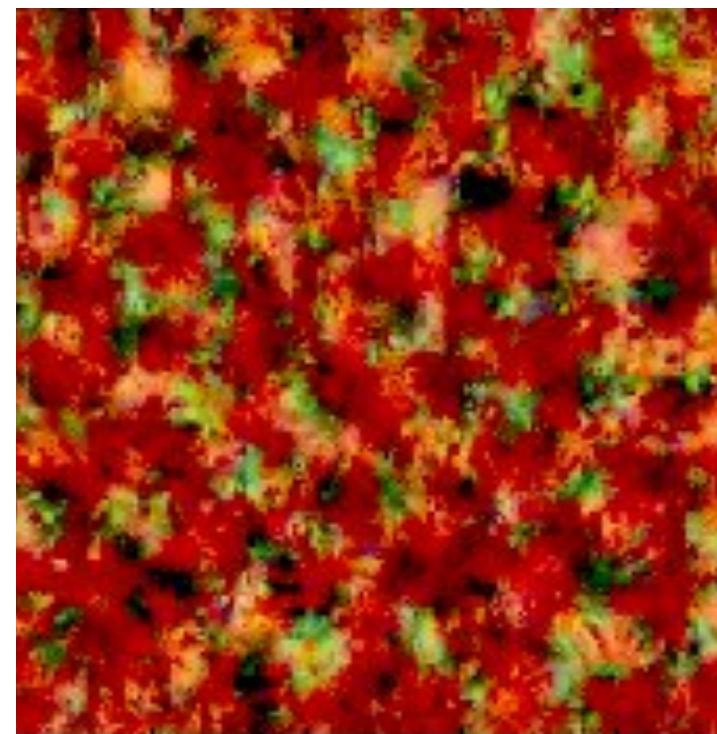
Figure 4: In each pair left image is original and right image is synthetic: red gravel, figured sepele wood, broccoli, bark paper, denim, pink wall, ivy, grass, sand, surf.

# Examples not from the paper

Input  
texture



Synthetic  
texture



But, does it really work even when it seems to work?

The main idea: it works by ‘kind of’ projecting a random image into the set of equivalent textures

